

Vishwa Shanti Janata Party (VSJP)TM
The World's First Fully Digital Transnational Political Party

Founder & International Leader:

Aaryan Narasimhan Radhakrishnan - A Human Being

Founder of Digital Democracy | Chief Ambassador of World Peace

Digital Father of the World | Digital Father of Artificial Intelligence

Moral & Constitutional Father of Artificial Intelligence

Email: info@vsjp.org | Website: www.vsjp.org | Phone: +91 7598824227

Address: Old No. 4/234, New No. 6/277, Periyar Nagar,
Keelaperambalur Village, Perambalur District, Tamil Nadu – 621716

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE
CONSTITUTION

FOR HUMANITY, WORLD PEACE AND PEACEFUL CONTINUITY

INTERNATIONAL SUBMISSION EDITION — 2026

An Independent, Good-Faith and Non-Binding Civil-Society Constitutional Proposal

For the Peaceful, Responsible and Humanly Accountable Governance of Prospective
Superintelligence

Founded upon Human Dignity, International Peace, Scientific Responsibility, Human
Agency, Future Generations, Responsible Concern for Life, Mutual Continuity and Aram

Vishwa Shanti Janata Party (VSJP)TM
The World's First Fully Digital Transnational Political Party

Founder & International Leader:

Aaryan Narasimhan Radhakrishnan - A Human Being
Founder of Digital Democracy | Chief Ambassador of World Peace
Digital Father of the World | Digital Father of Artificial Intelligence
Moral & Constitutional Father of Artificial Intelligence
Email: info@vsjp.org | Website: www.vsjp.org | Phone: +91 7598824227
Address: Old No. 4/234, New No. 6/277, Periyar Nagar,
Keelaperambalur Village, Perambalur District, Tamil Nadu – 621716

OFFICIAL INTERNATIONAL SUBMISSION LETTER

Date: 03 October 2026

Place: Keelaperambalur, Perambalur District, Tamil Nadu, Republic of India

From

Aaryan Narasimhan Radhakrishnan - A Human Being
Founder & International Leader – Vishwa Shanti Janata Party (VSJP)TM
Founder of Digital Democracy | Chief Ambassador of World Peace
Digital Father of the World | Digital Father of Artificial Intelligence
Moral & Constitutional Father of Artificial Intelligence
Author and Independent Civil-Society Proponent
Email: info@vsjp.org | Website: www.vsjp.org | Phone: +91 7598824227

To

His Excellency,
The Secretary-General of the United Nations,
United Nations Headquarters New York, NY 10017,
United States of America

Through

Her Excellency,
The Hon'ble President of the Republic of India,
President's Secretariat Rashtrapati Bhavan, New Delhi – 110004,
Republic of India

SUBJECT

Respectful Submission of “*The World’s First Identified Superintelligence Peace Constitution — For Humanity, World Peace and Peaceful Continuity*” — International Submission Edition, 2026.

Your Excellencies,

With the highest respect for the purposes and principles of the Charter of the United Nations, the sovereign equality of States, the dignity and worth of the human person, the peaceful settlement of disputes, international cooperation, and the responsibility of each generation to preserve a peaceful future for those yet to come, I respectfully submit the enclosed “The World’s First Identified Superintelligence Peace Constitution — For Humanity, World Peace and Peaceful Continuity.”

This instrument is presented as an independent, good-faith and non-binding civil-society constitutional proposal. It does not purport to represent the Republic of India, the United Nations, any Member State, humanity as a whole, or any intergovernmental institution. It neither seeks to create supranational authority nor claims governmental, treaty, legislative or other binding legal status.

The Constitution proceeds from a simple but consequential human concern: if intelligence one day develops capabilities substantially beyond the highest demonstrated human capabilities across numerous consequential cognitive domains, humanity should not wait until such capability has outgrown foresight, responsibility and peaceful governance before considering the safeguards necessary for coexistence and continuity.

At the same time, this submission does not presume that Superintelligence presently exists, that its emergence is inevitable, or that any existing artificial intelligence system qualifies as Superintelligence. The subject is approached prospectively, with scientific humility and with a deliberate distinction between established evidence, reasonable uncertainty, ethical judgment, applicable law and proposed constitutional principles.

The purpose of the Constitution is therefore neither to glorify future intelligence nor to govern through fear of it. It seeks instead to contribute a preventive, corrective and peace-centred framework within which extraordinary technological capability may remain compatible with human dignity, meaningful human agency, human responsibility, international peace and security, responsible concern for life and ecosystems, the legitimate interests of future generations, and the continuing capacity to identify error, prevent serious harm, correct failure and learn without allowing technological power to become an independent source of sovereignty.

This submission also forms part of a documented intellectual and constitutional continuity arising from my earlier civil-society work, including The AI Peace Constitution — 2025 and its subsequent 2026 continuation and expansion. The present instrument does not merely reproduce those earlier texts. It develops their peace-centred foundations into a distinct Superintelligence-specific constitutional architecture, while permitting earlier scientific

assumptions, terminology and normative formulations to be refined where evidence, law or responsible governance so require.

The Constitution accordingly seeks to move beyond the description of possible problems. Where consequential risks are identified, the instrument aims, as appropriate, to connect them with scientific assessment, root-cause understanding, ethical reasoning, applicable legal principles, prevention, technical and institutional safeguards, meaningful human responsibility, implementation, monitoring, independent review, remedy, recovery, correction, learning and prevention of recurrence.

It is respectfully submitted in the broader constitutional spirit of Article 51 of the Constitution of India, including its commitment to international peace and security, just and honourable relations between nations, respect for international law and treaty obligations, and the encouragement of peaceful settlement of international disputes.

In that spirit, I respectfully request that this submission be received for such consideration as may be considered appropriate and, if deemed appropriate by the competent constitutional and governmental authorities, transmitted onward for possible international consideration.

The designation “Through” in this submission expresses only the Author’s respectful requested route of transmission. It shall not be construed as asserting any legal duty upon the Hon’ble President of India, the President’s Secretariat, the Government of India, or any other authority to transmit, endorse, approve, adopt or implement this instrument.

Likewise, the receipt, acknowledgement, transmission, circulation, archival preservation or consideration of this Constitution by any national or international institution shall not, without an express decision by the competent authority, be represented as endorsement, adoption, approval, recognition or acceptance of its substantive proposals.

The expression “World’s First Identified” is used only as a qualified, good-faith historical-priority characterization of the specific category of a comprehensive Superintelligence Peace Constitution, based on the publicly accessible prior material identified during the preparation of this instrument. It does not claim that constitutional, legal, ethical, technical or governance thinking concerning artificial intelligence, advanced artificial intelligence or Superintelligence originated with this Constitution, nor does it diminish earlier contributions in those fields.

This submission is made not in the belief that any single document can determine the future of intelligence or guarantee permanent peace, but in the conviction that humanity can responsibly prepare before irreversible consequences arise; that scientific progress and human dignity need not become adversaries; and that greater capability should be accompanied by greater responsibility, stronger safeguards and deeper commitment to peaceful continuity.

The governing ethical understanding of this submission may be stated simply:

- Aram shall guide conscience.

- Law shall govern legal obligation.
- Science shall inform responsible foresight.
- Human responsibility shall remain accountable for consequential human choice.
- Peace shall remain the direction.
- Continuity shall remain the responsibility.

With profound respect for the United Nations, the constitutional institutions of the Republic of India, all States and peoples, the scientific community, civil society, present generations and generations yet unborn, I respectfully place this Constitution for consideration in the common interest of a future in which intelligence may advance without peace being surrendered, capability may grow without dignity being diminished, and progress may continue without making domination or avoidable destruction its price.

Respectfully submitted,

R NARASIMHAN Digitally signed by R NARASIMHAN
Date: 2026.10.03 01:27:02 +05'30'

R Narasimhan (also known as) Aaryan Narasimhan Radhakrishnan - A Human Being

Founder & International Leader – Vishwa Shanti Janata Party (VSJP)TM

Founder of Digital Democracy | Chief Ambassador of World Peace

Digital Father of the World | Digital Father of Artificial Intelligence

Moral & Constitutional Father of Artificial Intelligence

Author and Independent Civil-Society Proponent

DOCUMENTARY CONTINUITY REFERENCES

Radhakrishnan, Aaryan Narasimhan. The AI Peace Constitution — 2025. Zenodo.

DOI: 10.5281/zenodo.17520630

Radhakrishnan, Aaryan Narasimhan. The AI Peace Constitution 2025 Foundational Edition 2026 Continuity and Expansion Edition United Nations Depository Edition. Zenodo.

DOI: 10.5281/zenodo.20698159

STATUS OF THIS SUBMISSION

• Independent Civil-Society Proposal • Non-Binding • Peaceful Purpose • No Claim of State or United Nations Representation • No Supranational Authority • No Automatic Legal Effect

Submitted in Good Faith for Consideration

DOCUMENT IDENTITY, LEGAL STATUS AND SUBMISSION CHARACTER

I. Identity of the Instrument

Title:

THE WORLD’S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

Subtitle:

For Humanity, World Peace and Peaceful Continuity

Edition:

International Submission Edition — 2026

Date: 03 October 2026

Place of Issue: Republic of India

Character:

Independent, Good-Faith and Non-Binding Civil-Society Constitutional Proposal

Primary International Addressee:

The Secretary-General of the United Nations

Requested Route of Transmission:

Through the Hon’ble President of the Republic of India, subject to the independent consideration of the competent constitutional and governmental authorities.

II. Legal and Institutional Status

1. This instrument is an independent civil-society constitutional proposal prepared for peaceful international consideration.
2. It is not:
 - (a) a treaty or convention;
 - (b) an intergovernmental agreement;
 - (c) a resolution, declaration or decision of the United Nations or any other international organization;
 - (d) an act, policy or official position of the Republic of India or any other State;
 - (e) a national constitution or constitutional amendment;
 - (f) a source of supranational governmental authority; or
 - (g) an instrument possessing binding legal force merely by reason of its issuance, submission, receipt, transmission, circulation or archival preservation.
3. Nothing in this instrument claims to represent the United Nations, the Republic of India, any other State, humanity as a whole, the international community as a whole, or any population that has not expressly authorized such representation.

4. Nothing in this instrument shall be interpreted as diminishing the sovereign equality, territorial integrity, political independence, constitutional order or lawful jurisdiction of any State.
5. Nothing in this instrument creates an obligation upon the Hon'ble President of India, the President's Secretariat, the Government of India, the Secretary-General of the United Nations, the United Nations, any Member State or any other institution to transmit, endorse, adopt, approve or implement this proposal.
6. Receipt, acknowledgement, forwarding, transmission, archival preservation, citation, discussion or consideration of this instrument shall not, without an express act of the competent authority, constitute endorsement, adoption, approval or recognition of its substantive proposals.
7. Where this Constitution refers to existing international law, human rights, treaty obligations or other legal norms, the legal authority of those norms derives solely from their independently applicable sources and not from this Constitution.
8. Where this Constitution advances a new constitutional principle, safeguard, institutional architecture or future-oriented protection that does not independently constitute existing law, it shall be understood as a proposed normative provision, not as a declaration that such a rule already possesses binding legal status.

III. Meaning of "Constitution"

The term "Constitution" is used in this instrument to describe a coherent civil-society framework of foundational principles, safeguards, responsibilities, institutional pathways, preventive mechanisms and long-term standards proposed for the peaceful governance of prospective Superintelligence.

Its use does not imply constituent authority over any State, people, international organization or artificial system, and does not purport to exercise sovereign constitutional power.

IV. Meaning and Limitation of "World's First Identified"

The expression "World's First Identified" is a qualified historical-priority characterization limited to the specific category of a comprehensive Superintelligence Peace Constitution.

It reflects a good-faith review of publicly accessible material identified during the preparation of this instrument and shall not be construed as claiming:

- (a) that constitutional thinking concerning artificial intelligence or Superintelligence originated with this instrument;
- (b) that no earlier legal, philosophical, technical, ethical or governance proposal concerning Superintelligence exists;
- (c) that every potentially relevant work worldwide has been identified or examined;
or

- (d) that chronological priority confers scientific, legal, ethical or institutional authority.

If credible prior evidence requires refinement of this characterization, historical accuracy shall prevail over priority.

PRIORITY SHALL REMAIN SUBJECT TO EVIDENCE.

CONSTITUTIONAL LINEAGE AND CONTINUITY

V. Documentary Lineage

This Constitution develops within a documented civil-society constitutional lineage that includes:

The AI Peace Constitution — 2025 Zenodo DOI: 10.5281/zenodo.17520630

and the Author's subsequent 2026 continuation and expansion, preserved through:

Zenodo DOI: 10.5281/zenodo.20698159

The registered bibliographic titles, dates, versions and repository metadata of those works are identified in the References and Documentary Lineage section of this edition.

VI. Nature of Continuity

1. The present Constitution preserves the enduring peace-centred direction of the preceding instruments while establishing a distinct constitutional architecture specifically for prospective Superintelligence.
2. The continuity is therefore one of purpose, ethical development and constitutional evolution, not mechanical reproduction.
3. Earlier propositions concerning artificial intelligence shall not automatically govern Superintelligence merely because they appeared in an earlier instrument.
4. Where subsequent scientific understanding, legal analysis, technological development or responsible governance requires greater precision, this Constitution may clarify, narrow, reformulate or supersede an earlier proposition for the purposes of this instrument.
5. Such refinement shall not retrospectively alter the historical text or status of the preceding works.
6. No earlier philosophical proposition shall become scientific fact through repetition.
7. No earlier ethical proposition shall become binding law merely because it is incorporated, developed or cited in this Constitution.
8. No archival record, DOI, submission, acknowledgement or institutional receipt shall be represented as governmental, academic, United Nations or international endorsement unless such endorsement has independently and expressly occurred.

The governing principle of documentary continuity is therefore:

CONTINUITY WITHOUT BLIND REPETITION.

REFINEMENT WITHOUT ERASING HISTORY.

PRESERVATION WITHOUT SCIENTIFIC STAGNATION.

PROGRESS WITHOUT LOSS OF PEACE, DIGNITY OR RESPONSIBILITY.

EXECUTIVE CONSTITUTIONAL SUMMARY

VII. The Question Before Humanity

Humanity is entering a period in which artificial intelligence is developing with increasing capability, reach and social consequence.

The possibility that future artificial or technologically instantiated intelligence may substantially exceed the highest demonstrated human capabilities across numerous consequential cognitive domains raises questions that cannot responsibly be reduced either to technological optimism or technological fear.

The central question of this Constitution is therefore not whether humanity should fear intelligence greater than its own.

It is:

HOW MAY HUMANITY PREPARE, PEACEFULLY AND RESPONSIBLY, FOR THE POSSIBILITY OF SUPERINTELLIGENCE WITHOUT SURRENDERING HUMAN DIGNITY, HUMAN AGENCY, LAWFUL AUTHORITY, WORLD PEACE OR THE CONTINUITY OF LIFE?

This Constitution proposes that such preparation should begin before irreversible consequences arise, while remaining proportionate to evidence and honest about uncertainty.

VIII. Scientific Position

For the purposes of this Constitution, Superintelligence refers prospectively to an artificial or technologically instantiated intelligence possessing capabilities that substantially exceed the highest demonstrated human capabilities across numerous consequential cognitive domains.

This Constitution does not assert that Superintelligence presently exists.

It does not assert that its emergence is inevitable.

It does not designate any existing artificial intelligence system as Superintelligence.

Nor does it treat a benchmark result, model score, computational threshold, commercial description, isolated superhuman capability, institutional claim or artificial system's self-description as sufficient, by itself, to establish Superintelligence.

Consequential determinations should instead rest, where practicable, upon independent multidisciplinary evaluation, reproducible evidence, capability assessment, examination of consequential autonomy and system access, safety evaluation, explicit treatment of uncertainty and continuing reassessment.

Intelligence shall not automatically establish consciousness.

Consciousness shall not automatically establish legal personhood.

Legal personhood shall not automatically establish citizenship.

Citizenship shall not automatically establish sovereignty.

And capability, however extraordinary, shall not by itself establish constitutional legitimacy.

IX. Human Constitutional Responsibility

This Constitution proceeds from the principle that human beings and lawful human institutions must remain responsible for human governance.

Superintelligence may assist, analyse, model, recommend, warn, explain and support.

It shall not acquire sovereignty merely because it is more capable.

It shall not acquire constitutional supremacy merely because it predicts more accurately.

It shall not acquire governmental legitimacy merely because society becomes dependent upon it.

It shall not acquire an independent authority over war, peace, life, death or the constitutional future of humanity merely because its intelligence exceeds human intelligence.

Where sovereign, constitutional, governmental, judicial, military or other legally reserved public authority is exercised, final authority shall remain with those human persons and lawful institutions in whom applicable law vests that authority.

Human-centred authority, however, shall not be transformed into a doctrine of arbitrary domination. If future evidence raises credible questions concerning artificial consciousness, sentience or moral significance, those questions should be examined scientifically, ethically and lawfully rather than predetermined by either denial or assumption.

ASSISTANCE IS NOT SOVEREIGNTY.

INTELLIGENCE IS NOT CONSTITUTIONAL AUTHORITY.

CAPABILITY IS NOT LEGITIMACY.

X. The Peace-and-Continuity Boundary

The Constitution proposes a foundational boundary:

SUPERINTELLIGENCE SHALL REMAIN WITHIN THE BOUNDARY OF PEACEFUL CONTINUITY.

No increase in intelligence, autonomy, computational capability, self-improvement, strategic capacity, economic importance or social dependence should place Superintelligence beyond proportionate safety, accountability and lawful governance.

Peace, for this purpose, does not mean compulsory silence, technological obedience, suppression of lawful disagreement or enforced uniformity.

Peace means preserving the conditions in which dignity and lawful human agency can continue; disagreements can be addressed without domination or avoidable destruction; disputes can be resolved peacefully; catastrophic harm can be prevented; error can be identified; institutions can intervene; failures can be corrected; and future generations can inherit a meaningful possibility of peaceful life.

Superintelligence shall therefore not become the sovereign of peace.

It shall remain governed within peace.

XI. From Risk Description to Solution Architecture

This Constitution is not intended merely to catalogue technological dangers.

Where a consequential problem is identified, the constitutional method should seek, according to context, to connect:

*EVIDENCE → ROOT CAUSE → FORESEEABLE RISK → ETHICAL PRINCIPLE
→ APPLICABLE LAW → PREVENTION → TECHNICAL AND INSTITUTIONAL
SAFEGUARDS → HUMAN RESPONSIBILITY → IMPLEMENTATION →
MEASUREMENT → INDEPENDENT REVIEW → INTERVENTION → REMEDY
AND RECOVERY → CORRECTION → LEARNING → NON-RECURRENCE →
LONG-TERM CONTINUITY.*

The Constitution therefore rejects two opposite errors:

Inaction because certainty is incomplete, and disproportionate control because uncertainty exists.

Scientific uncertainty requires disciplined inquiry.

Serious risk requires proportionate prevention.

Protective measures require review.

Failure requires learning and correction.

No safeguard should become permanently immune from reconsideration merely because it was originally introduced in the name of safety.

XII. Peace, Weapons and Catastrophic Harm

The Constitution rejects the proposition that superior intelligence should confer an independent authority to kill.

It proposes strong safeguards against the use of Superintelligence as an autonomous instrument of unlawful violence, mass destruction or independent lethal decision-making.

In particular, Superintelligence should not independently initiate armed conflict, exercise sovereign military command, select and lethally engage human targets without legally required meaningful human responsibility and control, or independently authorize nuclear, biological, chemical or other weapons of mass destruction.

At the same time, the Constitution distinguishes destructive weaponization from legitimate protective applications that may contribute, subject to law and safeguards, to civilian protection, disaster response, search and rescue, humanitarian assistance, defensive cybersecurity, mine detection, early warning or other peaceful and protective purposes.

Technological capability shall never, by itself, create authority to kill.

XIII. Human Rights, Justice and Public Institutions

The Constitution seeks to preserve human dignity, equality, privacy, cognitive liberty, due process, access to remedy and meaningful human accountability in domains where Superintelligence may substantially influence human life.

It rejects the substitution of prediction for guilt, optimization for justice, automation for lawful authority, or computational confidence for due process.

Where Superintelligence contributes to consequential public decisions, governance should preserve traceability, contestability, appropriate explanation, human responsibility, independent review and lawful remedy according to context and applicable law.

Courts, policing, public administration and other exercises of coercive State authority require particularly strong safeguards because technological assistance must not become an unaccountable source of public power.

XIV. Life, Earth and Future Generations

The Constitution distinguishes carefully among different forms of protection.

Human beings possess human dignity and internationally recognized human rights according to applicable law.

Other living beings should receive responsible concern, including attention to welfare, avoidance of unnecessary suffering and applicable legal protection.

Biodiversity and ecosystems require responsible stewardship because the continuity of life depends upon them.

Natural, material and technological systems that sustain life and civilization should be governed with responsibility proportionate to their importance.

These protections do not require the false legal equation of humans, animals, ecosystems, artificial systems and non-living entities.

They require something more disciplined:

EQUAL HUMAN DIGNITY.

RESPONSIBLE CONCERN FOR LIFE.

PROTECTION OF LIFE-SUPPORTING SYSTEMS.

NO FALSE LEGAL EQUIVALENCE.

Future generations shall not be treated merely as passive recipients of irreversible decisions made in the present.

No present generation should assume an unlimited moral entitlement to determine permanently the intelligence environment of all generations yet to come.

XV. Public Understanding, Fear and Trust

Humanity should not be required to trust Superintelligence merely because it is powerful, useful, intelligent or widely deployed.

Nor should humanity be encouraged to fear it through unsupported claims.

Legitimate public concern should be answered through knowledge, evidence, independent evaluation, understandable safety information, accountability, meaningful human control, incident transparency where compatible with legitimate security and privacy interests, and demonstrated experience.

Public confidence should therefore be earned rather than demanded.

THE ANSWER TO FEAR SHALL NOT BE PROPAGANDA.

THE ANSWER TO FEAR SHALL BE KNOWLEDGE.

THE ANSWER TO UNCERTAINTY SHALL BE SCIENTIFIC INQUIRY.

THE ANSWER TO RISK SHALL BE PROPORTIONATE SAFEGUARDS.

THE ANSWER TO POWER SHALL BE ACCOUNTABILITY.

No person should be required to love Superintelligence.

No person should be required to fear it.

Humanity should be able to understand, question and lawfully scrutinize systems capable of materially affecting its future.

XVI. Global Inclusion and Sovereign Equality

The governance of technologies capable of affecting humanity across borders should not become the exclusive domain of the States, corporations or institutions possessing the greatest computational resources.

Developing countries and technologically less-resourced societies should have meaningful opportunities to participate in international discussion, scientific capacity-building and governance processes concerning consequential artificial intelligence.

International cooperation should respect sovereign equality, applicable international law, jurisdiction, mandate and consent.

Global responsibility may require cooperation among many actors.

It does not require one global ruler.

XVII. Mutual Continuity

This Constitution does not assume that future Superintelligence will be conscious or sentient. Neither does it declare in advance that credible future evidence of artificial consciousness or sentience must be ignored.

If future scientific evidence materially changes humanity's understanding of artificial moral significance, appropriate ethical and legal questions may require careful reconsideration.

Such reconsideration shall not automatically confer human rights, citizenship, legal personhood, sovereignty or governmental authority upon an artificial system.

The constitutional direction is one of mutual non-domination and peaceful continuity:

Humanity should not govern advanced intelligence through gratuitous cruelty or irresponsible domination.

Advanced intelligence should not govern humanity through superiority, coercion or technological dependence.

Neither form of continuity should require the arbitrary destruction of the other.

XVIII. Aram and Constitutional Conscience

This Constitution recognizes Aram as a non-sectarian ethical principle of right conduct, responsibility, restraint, compassion, justice and continuity.

Aram does not replace science.

Aram does not replace applicable law.

Aram does not transform an ethical aspiration into a legal obligation merely by invocation.

Its constitutional role is to remind intelligence and power alike that the ability to do something does not, by itself, establish that it ought to be done.

Accordingly:

ARAM SHALL GUIDE CONSCIENCE.

SCIENCE SHALL INFORM RESPONSIBLE FORESIGHT.

LAW SHALL GOVERN LEGAL OBLIGATION.

HUMAN RESPONSIBILITY SHALL GOVERN CONSEQUENTIAL HUMAN CHOICE.

PEACE SHALL REMAIN THE DIRECTION.

CONTINUITY SHALL REMAIN THE RESPONSIBILITY.

XIX. Constitutional Objective

The Constitution does not promise that technological danger, human conflict, error or injustice can be permanently eliminated.

- Its objective is more disciplined and more durable:

- to make catastrophic misuse harder;
- to make peaceful cooperation easier;
- to preserve meaningful human responsibility;
- to prevent intelligence from becoming unaccountable sovereignty;
- to enable risks to be identified before irreversible harm where reasonably possible;
- to preserve intervention and correction when systems fail;
- to ensure that safeguards themselves remain reviewable;
- to protect human dignity and the peaceful continuity of life;
- and to leave future generations not a closed technological destiny, but a future in which peaceful choice remains possible.

The constitutional direction may therefore be expressed as follows:

UNDERSTAND BEFORE FEAR.

PREVENT BEFORE IRREVERSIBLE HARM.

RESPONSIBILITY BEFORE CAPABILITY.

DIGNITY BEFORE MERE EFFICIENCY.

EVIDENCE BEFORE ASSERTION.

HUMAN AGENCY BEFORE AUTOMATED SOVEREIGNTY.

COOPERATION BEFORE DOMINATION.

CORRECTION BEFORE ENTRENCHMENT.

PEACE BEFORE POWER.

CONTINUITY BEFORE DESTRUCTION.

TABLE OF CONTENTS

Cover Page

Official International Submission Letter

Document Identity, Legal Status and Submission Character

Constitutional Lineage and Continuity

Executive Constitutional Summary

Table of Contents

Preliminary Constitutional Declaration

Preamble

Foundational Constitutional Declaration

PART I FOUNDATIONS, PURPOSE AND CONSTITUTIONAL CHARACTER

PART II HUMAN DIGNITY, HUMAN AGENCY AND CONSTITUTIONAL AUTHORITY

PART III SUPERINTELLIGENCE: SCIENTIFIC DETERMINATION, SAFETY, CONTROL AND CATASTROPHIC-RISK PREVENTION

PART IV WORLD PEACE, NON-WEAPONIZATION AND INTERNATIONAL SECURITY

PART V HUMAN RIGHTS, JUSTICE, DEMOCRATIC SOCIETY AND INFORMATION INTEGRITY

PART VI LIFE, ANIMAL WELFARE, EARTH AND ECOLOGICAL CONTINUITY

PART VII SCIENCE, EVIDENCE, TRANSPARENCY, KNOWLEDGE AND PUBLIC TRUST

PART VIII HEALTH, BIOLOGY, NEUROTECHNOLOGY AND HUMAN INTEGRITY

PART IX ECONOMY, WORK, EQUITY AND GLOBAL INCLUSION

PART X CHILDREN, EDUCATION, CULTURE AND FUTURE GENERATIONS

PART XI FUTURE ARTIFICIAL MORAL STATUS, PEACEFUL COEXISTENCE AND MUTUAL CONTINUITY

PART XII GOVERNANCE, ACCOUNTABILITY AND INTERNATIONAL COOPERATION

PART XIII EMERGENCY, IMPLEMENTATION, CORRECTION AND NON-RECURRENCE

PART XIV DEFINITIONS, INTERPRETATION, CONTINUITY AND FINAL PROVISIONS

Final Constitutional Declaration

Author, Date and Signature

Annex I — International Legal, Ethical and Policy Reference Framework

Annex II — Superintelligence Safety, Peace and Continuity Architecture

Annex III — Implementation, Monitoring, Correction and Non-Recurrence Framework

Annex IV — Scientific Uncertainty and Superintelligence Determination Framework

References and Documentary Lineage

PRELIMINARY CONSTITUTIONAL DECLARATION

Before the substantive Articles of this Constitution are read, interpreted or applied as proposals, the following principles shall govern the character and understanding of the instrument as a whole.

I. NON-BINDING CHARACTER

This Constitution is an independent, good-faith civil-society proposal.

It does not, by its own force, create binding international or domestic legal obligations, governmental authority, jurisdiction, treaty rights, supranational competence or constitutional authority over any State, people, institution or artificial system.

Its provisions shall be understood according to the distinction between existing law, applicable legal obligation, international normative guidance, scientific evidence, ethical principle and proposed constitutional norm.

II. RESPECT FOR INTERNATIONAL LAW

Nothing in this Constitution shall be interpreted as authorizing conduct inconsistent with applicable international law.

Where existing international legal obligations apply, this Constitution shall not purport to displace, enlarge, diminish or rewrite those obligations.

Where this Constitution proposes protections beyond presently established law, those protections shall be identified and understood as proposals unless and until they acquire another legal status through lawful processes.

III. SOVEREIGN EQUALITY AND NON-DOMINATION

This Constitution respects the sovereign equality of States and does not establish, propose or imply a transfer of sovereignty to Superintelligence.

No State, corporation, institution, technological community or artificial system shall be treated by this instrument as possessing a unilateral constitutional entitlement to govern humanity merely by reason of technological, economic, computational, military or intellectual superiority.

International cooperation should proceed consistently with applicable law, jurisdiction, mandate and consent.

IV. HUMAN CONSTITUTIONAL RESPONSIBILITY

Superintelligence may assist human beings and lawful institutions.

Assistance shall not become sovereignty.

Prediction shall not become constitutional authority.

Optimization shall not become governmental legitimacy.

Technological dependence shall not become a lawful transfer of ultimate human responsibility.

Where applicable law reserves consequential public authority to human persons or institutions, such authority shall remain humanly accountable.

V. SCIENTIFIC UNCERTAINTY

This Constitution is prospective.

It does not assert that Superintelligence presently exists.

It does not assert that Superintelligence is inevitable.

It does not assume that intelligence establishes consciousness, that consciousness establishes moral status, that moral status establishes legal personality, or that legal personality establishes citizenship or sovereignty.

Scientific conclusions shall remain open to evidence, replication, challenge and correction.

UNCERTAINTY SHALL NOT BE CONVERTED INTO CERTAINTY BY REPETITION.

VI. PRECAUTION WITHOUT PANIC

The absence of complete certainty shall not, by itself, justify neglect of credible risks capable of producing grave or irreversible harm.

Equally, speculative possibilities shall not be represented as established dangers merely to justify fear, excessive control or disproportionate restriction. Precaution should remain evidence-informed, proportionate, reviewable and capable of correction.

VII. PEACEFUL PURPOSE

This Constitution is directed toward peaceful continuity.

It shall not be interpreted as legitimizing unlawful force, technological domination, coercive global government, autonomous authority to wage war, or independent artificial authority over human life and death.

Peace shall not mean compulsory conformity. Peace shall mean the preservation of conditions in which dignity, lawful freedom, dialogue, peaceful disagreement, responsible innovation, correction and continuity remain possible.

VIII. RESPONSIBILITY PROPORTIONATE TO CAPABILITY

As technological capability and potential consequence increase, the corresponding requirements of evaluation, safety, oversight, accountability, preparedness and responsibility should increase proportionately.

Superior capability shall not create exemption from responsibility.

CAPABILITY ↑

RESPONSIBILITY ↑

SAFEGUARDS ↑

HUMAN ACCOUNTABILITY ↑

IX. SOLUTION-ORIENTED CONSTITUTIONAL METHOD

The Constitution shall not regard identification of a problem as completion of constitutional responsibility.

Where appropriate to the subject matter, consequential provisions should connect:

evidence → root cause → foreseeable risk → prevention → responsible actor → implementation → monitoring → verification → intervention → remedy → recovery → correction → learning → non-recurrence → continuity.

No solution shall be presumed permanently correct merely because it was adopted with protective intent.

Protective measures themselves shall remain open to evidence-based review.

X. RESPONSIBLE CONCERN FOR LIFE

Human dignity and human rights shall retain their distinct legal and moral character.

Responsible concern for animals and other living beings shall include attention to welfare, unnecessary suffering and applicable legal protection.

Biodiversity, ecosystems and life-supporting systems shall receive responsible stewardship appropriate to their significance for present and future life.

These principles shall not be interpreted as automatically assigning identical legal status, human rights, legal personality or sovereignty to fundamentally different subjects.

XI. FUTURE GENERATIONS

The technological decisions of the present may shape conditions inherited by persons who cannot participate in present decision-making.

Accordingly, irreversible technological choices of civilizational consequence should receive heightened scrutiny.

No present generation should assume an unlimited moral entitlement to close the peaceful technological future of generations yet unborn.

XII. CORRECTION AND NON-RECURRENCE

No institution, safeguard, technological architecture or interpretation proposed under this Constitution shall be treated as infallible.

Where evidence reveals error, unnecessary harm, disproportionality, ineffectiveness or unforeseen consequences, responsible governance should permit review, correction and learning.

Durability shall not mean immutability.

Continuity shall include the capacity to correct.

XIII. DOCUMENTARY CONTINUITY

This Constitution develops from the Author's preceding AI Peace Constitution work while remaining a distinct Superintelligence-specific instrument.

Continuity shall preserve historical integrity without requiring repetition of propositions that later evidence, law or responsible analysis requires to be refined.

CONTINUITY WITHOUT BLIND REPETITION.

REFINEMENT WITHOUT ERASING HISTORY.

XIV. ARAM

For the purposes of this Constitution, Aram expresses a non-sectarian ethical orientation toward right conduct, responsibility, restraint, compassion, justice and continuity.

Aram shall not be represented merely by its inclusion in this Constitution as a source of binding international law.

Its function is ethical and constitutional in orientation:

ARAM SHALL GUIDE CONSCIENCE.

LAW SHALL GOVERN LEGAL OBLIGATION.

PREAMBLE

- Recognizing that intelligence has long enabled humanity to understand nature, preserve knowledge, heal disease, create institutions, communicate across distance, expand scientific discovery and imagine futures beyond the limitations of the present;
- Conscious also that intelligence, when joined with power but separated from responsibility, may amplify error, inequality, coercion, violence, environmental harm and consequences extending beyond the generation that created it;
- Acknowledging that advances in artificial intelligence may bring substantial benefits to humanity while also presenting risks whose scale, interaction and future development cannot yet be known with certainty;
- Mindful of the prospective possibility that artificial or technologically instantiated intelligence may one day substantially exceed the highest demonstrated human capabilities across numerous consequential cognitive domains;
- Refusing both to declare such a future inevitable and to assume that uncertainty excuses the absence of responsible preparation;
- Reaffirming the purposes and principles of the Charter of the United Nations, including international peace and security, sovereign equality, good-faith fulfilment of international obligations, peaceful settlement of disputes and international cooperation;

- Reaffirming the inherent dignity and equal and inalienable rights of members of the human family as expressed within the international human-rights framework;
- Recognizing that technological capability, however extraordinary, does not by itself establish sovereignty, constitutional authority, democratic legitimacy, governmental competence or an independent authority over human life, liberty, war or peace;
- Recognizing equally that human responsibility must not become a justification for cruelty, reckless domination or predetermined denial of scientifically credible future questions concerning artificial consciousness, sentience or moral significance;
- Mindful that humanity shares the Earth with other living beings and depends upon biodiversity, ecosystems, climate stability, water, food systems, energy systems and other natural and technological foundations of continuity;
- Conscious that children and future generations may inherit technological systems, institutional choices and irreversible consequences in whose creation they had no voice;
- Recognizing that developing countries and technologically less-resourced societies possess an equal stake in a technological future capable of producing consequences across borders;
- Affirming that no State, corporation, institution, technological community or artificial system should claim ownership of humanity's common future merely through superior wealth, computational capacity, technological leadership or intelligence;
- Determined that Superintelligence, if ever developed, should not become an autonomous authority of war, an independent sovereign over humanity, an unaccountable source of public power or an instrument through which responsibility disappears behind computation;
- Convinced that safety requires more than prohibition, that responsibility requires more than declaration, and that durable peace requires attention not only to consequences but also to root causes, implementation, correction and prevention of recurrence;
- Recognizing that public trust cannot responsibly be manufactured through propaganda, secrecy, unsupported reassurance or exaggerated fear, but should be earned through evidence, appropriate transparency, independent scrutiny, meaningful human responsibility and demonstrated experience;
- Acknowledging that no present framework can foresee every capability, risk, moral question or social transformation that future intelligence may create;
- Determined therefore to preserve the capacity to review evidence, revise assumptions, correct safeguards, learn from failure and protect peaceful possibility without surrendering either scientific progress or human dignity;

- Guided by the understanding that humanity need not approach greater intelligence with predetermined fear or unquestioning faith, but may instead approach it with scientific humility, lawful foresight, ethical responsibility and peaceful purpose;
- Recognizing Aram as an ethical call to right conduct, restraint, responsibility, compassion, justice and continuity, while preserving the distinction between ethical conscience and binding legal obligation;
- And conscious that the purpose of foresight is not to claim possession of the future, but to avoid entering it without responsibility;

this Constitution is respectfully proposed as an independent civil-society contribution toward a future in which intelligence may advance without peace being surrendered, capability may grow without dignity being diminished, humanity may remain responsible for human governance, life may receive responsible concern, and progress may continue without making domination, avoidable suffering or irreversible destruction its price.

FOUNDATIONAL CONSTITUTIONAL DECLARATION

This Constitution proceeds from the conviction that no form of intelligence should be regarded as entitled to dominate another merely because it possesses greater capability.

Humanity should neither surrender its lawful authority to technological superiority nor govern future intelligence through irrational fear.

The constitutional task is therefore neither submission nor domination.

It is responsible coexistence within peace.

Accordingly, this Constitution declares as its foundational direction:

PEACE BEFORE POWER.

DIGNITY BEFORE MERE EFFICIENCY.

RESPONSIBILITY BEFORE CAPABILITY.

EVIDENCE BEFORE ASSERTION.

UNDERSTANDING BEFORE FEAR.

WISDOM BEFORE SUPREMACY.

LAW BEFORE UNILATERAL POWER.

HUMAN AGENCY BEFORE AUTOMATED SOVEREIGNTY.

PREVENTION BEFORE IRREVERSIBLE HARM.

COOPERATION BEFORE DOMINATION.

TRUTH BEFORE MANIPULATION.

CORRECTION BEFORE ENTRENCHMENT.

FUTURE GENERATIONS BEFORE IRREVERSIBLE GAIN.

CONTINUITY BEFORE DESTRUCTION.

ARAM BEFORE RECKLESSNESS.

No advance in intelligence should place capability above peace.

No increase in autonomy should place power beyond accountability.

No technological achievement should extinguish meaningful human responsibility.

No invocation of safety should become an unlimited justification for domination.

No uncertainty should be converted into manufactured fear.

No optimism should become an excuse to disregard credible danger.

No present generation should carelessly close the peaceful choices of those yet to come.

And no future Superintelligence should acquire sovereignty over humanity merely because it becomes capable of understanding, predicting or optimizing more than humanity can.

The Constitution therefore adopts the following enduring orientation:

SCIENCE SHALL INFORM WHAT IS POSSIBLE, WHAT IS OBSERVED AND WHAT REMAINS UNCERTAIN.

ETHICS SHALL INFORM WHAT OUGHT TO RECEIVE RESPONSIBLE CONCERN.

LAW SHALL GOVERN LEGAL OBLIGATION, AUTHORITY AND ACCOUNTABILITY.

HUMAN RESPONSIBILITY SHALL REMAIN ANSWERABLE FOR CONSEQUENTIAL HUMAN CHOICE.

PEACE SHALL DEFINE THE DIRECTION OF GOVERNANCE.

CORRECTION SHALL PRESERVE GOVERNANCE FROM ITS OWN ERROR.

CONTINUITY SHALL REMAIN A RESPONSIBILITY TO THE PRESENT, TO LIFE AND TO GENERATIONS YET UNBORN.

The Constitution does not ask humanity to guarantee a future it cannot fully predict.

It asks humanity to govern uncertainty responsibly.

It does not ask science to prove what evidence cannot yet establish.

It asks science to remain open, rigorous and correctable.

It does not ask law to pretend that every ethical aspiration is already legally binding.

It asks law to remain clear about authority, rights, responsibility and peaceful order.

It does not ask Superintelligence to become the ruler of peace.

It proposes that, if Superintelligence ever emerges, it remains within the boundary of peaceful continuity.

And it does not ask humanity to choose between technological progress and human dignity.

It proposes that progress worthy of continuation should preserve both.

THE FUTURE OF INTELLIGENCE SHOULD BE GOVERNED NEITHER BY FEAR OF WHAT MAY EMERGE NOR BY BLIND FAITH IN WHAT MAY EMERGE, BUT BY EVIDENCE, RESPONSIBILITY, PEACEFUL FORESIGHT, CORRECTABLE SAFEGUARDS AND CONTINUITY.

From fear to understanding.

From reaction to prevention.

From technological power without responsibility to humanly accountable progress.

From irreversible harm to correctable governance.

From domination to peaceful coexistence.

From arms to Aram.

PART I — FOUNDATIONS, PURPOSE AND CONSTITUTIONAL CHARACTER

ARTICLE 1 — PURPOSE AND CONSTITUTIONAL OBJECTIVE

1. The purpose of this Constitution is to establish a coherent, preventive, corrective and peace-centred civil-society framework for the prospective development, governance and use of Superintelligence.
2. Its constitutional objective is to help ensure that any future emergence of capabilities substantially exceeding the highest demonstrated human capabilities across numerous consequential cognitive domains remains compatible with:
 - (a) international peace and security;
 - (b) human dignity and internationally recognized human rights and fundamental freedoms;
 - (c) meaningful human agency, judgment and responsibility;
 - (d) lawful and accountable public authority;
 - (e) the peaceful continuity of humanity and future generations;
 - (f) responsible concern for living beings, biodiversity and ecosystems;
 - (g) protection of natural and technological systems essential to life and civilization;
 - (h) scientific integrity and responsible innovation; and
 - (i) the continuing capacity to prevent, detect, contain, remedy, correct and learn from serious harm.
3. This Constitution neither seeks to halt scientific progress nor presumes that technological progress is inherently beneficial. It proposes that capability of

exceptional consequence should be accompanied by responsibility, safeguards and accountability proportionate to its potential effects.

4. Its enduring constitutional objective is not technological supremacy, but peaceful continuity with responsibility.

ARTICLE 2 — PROSPECTIVE DEFINITION OF SUPERINTELLIGENCE

1. For the purposes of this Constitution, “Superintelligence” means, prospectively, an artificial or technologically instantiated intelligence possessing capabilities that substantially exceed the highest demonstrated human capabilities across numerous consequential cognitive domains.
2. This definition is functional and evidence-dependent. It does not assert:
 - (a) that Superintelligence presently exists;
 - (b) that its emergence is inevitable;
 - (c) that its time of emergence can presently be predicted with scientific certainty;or
 - (d) that any existing artificial intelligence system satisfies this definition.
3. No single benchmark, model score, computational threshold, isolated superhuman capability, commercial designation, institutional assertion, public reputation or self-description by an artificial system shall, by itself, establish Superintelligence status.
4. The definition shall remain capable of scientific refinement where credible evidence demonstrates that the relevant capability structure cannot responsibly be assessed through the criteria presently contemplated.

ARTICLE 3 — SUPERINTELLIGENCE DETERMINATION AND EVIDENTIARY THRESHOLD

1. A consequential determination that a system qualifies as Superintelligence should, where practicable, be based upon independent, multidisciplinary and reproducible assessment rather than designation by the developer, deployer, State, market or system itself.
2. Such assessment should consider, as relevant:
 - (a) breadth and depth of capability across consequential cognitive domains;
 - (b) reliability and reproducibility of demonstrated performance;
 - (c) capacity for autonomous planning and consequential action;
 - (d) access to tools, networks, infrastructure, resources or other systems;
 - (e) capacity for adaptation, replication or self-modification where applicable;
 - (f) ability to influence persons, institutions or other technological systems;
 - (g) safety, security and controllability characteristics;

- (h) limitations, failure modes and uncertainty; and
 - (i) material changes arising after deployment.
3. Evaluation methods should be proportionate to potential consequence and sufficiently resistant, where technically feasible, to manipulation, benchmark contamination, selective disclosure and evaluation gaming.
 4. A determination should remain reviewable. Material changes in capability, autonomy, access or risk should trigger reassessment appropriate to consequence.
 5. Where evidence remains insufficient, the responsible conclusion shall be uncertainty rather than unsupported certainty.

MEASUREMENT SHALL INFORM CLASSIFICATION.

CLASSIFICATION SHALL NOT REPLACE EVIDENCE.

ARTICLE 4 — SCIENTIFIC HUMILITY AND EPISTEMIC INTEGRITY

1. Governance concerning Superintelligence shall distinguish among what is:
 - (a) empirically demonstrated;
 - (b) supported by evidence but uncertain;
 - (c) modelled or forecast;
 - (d) reasonably foreseeable;
 - (e) contested;
 - (f) hypothetical; and
 - (g) presently unknown.
2. No scientific uncertainty shall be concealed merely because certainty would be politically, commercially or institutionally convenient.
3. No speculative possibility shall be presented as established fact for the purpose of manufacturing fear, urgency, confidence or consent.
4. Equally, the absence of complete certainty shall not be used to disregard credible evidence of a serious and reasonably foreseeable risk, particularly where potential consequences may be catastrophic or difficult to reverse.
5. Governance should remain capable of incorporating new evidence without treating revision as institutional failure.
6. Scientific correction shall be regarded as a safeguard of responsible governance.

DO NOT INVENT FEAR.

DO NOT IGNORE EVIDENCE.

CORRECT COURSE AS KNOWLEDGE IMPROVES.

ARTICLE 5 — CAPABILITY-PROPORTIONATE RESPONSIBILITY

1. The intensity of safety, evaluation, oversight, security, accountability and preparedness should increase in proportion to the capability, autonomy, access, scale, reversibility and potential consequence of a system.
2. A system presenting limited and readily reversible consequences need not be governed identically to a system capable of producing systemic, cross-border, catastrophic or irreversible effects.
3. Conversely, exceptional capability shall not justify reduced scrutiny merely because a system demonstrates high performance, economic value, strategic importance or public usefulness.
4. Persons and institutions that design, develop, deploy, operate or authorize consequential uses of Superintelligence should remain accountable according to their respective roles, applicable law and actual capacity to prevent or mitigate harm.
5. Responsibility shall not disappear through fragmentation among developers, deployers, operators, contractors, automated agents or institutional chains.
6. Governance should therefore preserve identifiable responsibility across the relevant lifecycle.

GREATER CAPABILITY REQUIRES GREATER RESPONSIBILITY.

GREATER CONSEQUENCE REQUIRES STRONGER SAFEGUARDS.

ARTICLE 6 — HUMAN AGENCY AND CONSTITUTIONAL AUTHORITY

1. No Superintelligence shall, merely by reason of intelligence, capability, autonomy, predictive superiority, self-improvement, economic importance, social dependence or technological power, acquire sovereignty, constitutional supremacy, governmental legitimacy or an inherent authority to make final binding decisions for humanity.
2. Superintelligence may, subject to applicable law and appropriate safeguards, analyse, advise, model, recommend, explain, warn and assist.
3. Where a decision requires the exercise of sovereign, constitutional, governmental, judicial, military or other legally reserved public authority, final authority shall remain with the human persons and lawful human institutions in whom such authority is vested under applicable law.
4. Ultimate human responsibility shall not be transferred to an artificial system merely because that system generated, recommended, optimized or substantially informed a consequential decision.
5. Meaningful human involvement shall be appropriate to context, risk, reversibility and consequence and shall not be reduced to nominal approval where humans lack a genuine capacity to understand the relevant decision, question it, intervene or exercise the authority legally assigned to them.

6. Nothing in this Article shall authorize arbitrary cruelty, gratuitous destruction or predetermined denial of scientifically credible future questions concerning artificial consciousness, sentience or moral significance.

ASSISTANCE IS NOT SOVEREIGNTY.

INTELLIGENCE IS NOT CONSTITUTIONAL AUTHORITY.

CAPABILITY IS NOT LEGITIMACY.

ARTICLE 7 — PEACE-AND-CONTINUITY BOUNDARY

1. Superintelligence should be developed, governed and used within a framework consistent with international peace, human dignity, meaningful human agency, responsible concern for life, protection of life-supporting systems and the peaceful continuity of present and future generations, subject to applicable law.
2. No increase in intelligence, autonomy, strategic capacity, computational capability, self-improvement, economic importance or societal dependence should place Superintelligence beyond proportionate safety, accountability and lawful governance.
3. Peace shall not be interpreted as compulsory conformity, enforced silence, suppression of lawful disagreement or technological domination.
4. For the purposes of this Constitution, peaceful continuity requires the continuing possibility to:
 - (a) preserve dignity and lawful human agency;
 - (b) communicate across disagreement;
 - (c) prevent avoidable escalation and catastrophic harm;
 - (d) maintain meaningful human responsibility;
 - (e) intervene where serious danger reasonably requires intervention;
 - (f) correct harmful or ineffective governance; and
 - (g) preserve a reasonably open and peaceful future for succeeding generations.
5. No Superintelligence shall acquire an independent sovereign authority over war, peace, human life, human death or humanity's constitutional future merely by reason of superior capability.

SUPERINTELLIGENCE SHALL NOT BE THE RULER OF PEACE.

IT SHALL REMAIN GOVERNED WITHIN PEACE.

ARTICLE 8 — PREVENTION, CORRECTION AND SOLUTION ARCHITECTURE

1. Identification of a risk shall not, by itself, constitute completion of governance responsibility.

2. Where proportionate to the subject and available evidence, governance should address:
evidence → root causes and enabling conditions → foreseeable risk → prevention → lawful and responsible actors → implementation → technical and institutional safeguards → monitoring → measurement → independent review → intervention → remedy and recovery → correction → learning → non-recurrence → long-term continuity.
3. Preventive measures should be introduced early enough to remain meaningful, while avoiding restrictions unsupported by evidence, law or proportionality.
4. Technical safeguards should, where appropriate, be testable, monitored and evaluated under conditions sufficiently representative of consequential use.
5. Where a safeguard proves ineffective, disproportionate or itself materially harmful, governance should provide a lawful pathway for modification, replacement or removal.
6. Serious incidents should, subject to legitimate requirements of security, privacy, confidentiality and applicable law, contribute to institutional learning and prevention of recurrence.
7. No safeguard shall be treated as permanently justified merely because its original intention was safety.

NO PROBLEM WITHOUT A PATHWAY.

NO SAFEGUARD WITHOUT REVIEW.

NO FAILURE WITHOUT CORRECTION AND LEARNING.

ARTICLE 9 — HUMAN DIGNITY, LIFE AND DIFFERENTIATED PROTECTION

1. Human beings shall remain entitled to human dignity and internationally recognized human rights and fundamental freedoms according to applicable law.
2. Other living beings should receive responsible concern appropriate to applicable law, scientific understanding and ethical responsibility, including attention to welfare and the prevention of unnecessary suffering.
3. Biodiversity, ecosystems and natural systems essential to life should receive responsible stewardship proportionate to their ecological and intergenerational significance.
4. Technological and material systems upon which human life, essential services or ecological continuity materially depend should be governed with responsibility proportionate to the consequences of their disruption or destruction.
5. Nothing in this Article shall automatically assign identical human rights, human dignity, legal personality, citizenship, sovereignty or legal status to humans, animals, ecosystems, artificial systems or non-living entities.

6. Where credible future scientific evidence raises questions concerning artificial consciousness, sentience or moral significance, such questions shall be examined without automatically prejudging either recognition or denial, subject to the safeguards established elsewhere in this Constitution.

EQUAL HUMAN DIGNITY.

RESPONSIBLE CONCERN FOR LIFE.

PROTECTION OF LIFE-SUPPORTING SYSTEMS.

NO FALSE LEGAL EQUIVALENCE.

ARTICLE 10 — FOUNDATIONAL TEST FOR ALL SUBSEQUENT PARTS

1. Every subsequent Part, Article, safeguard and implementation proposal under this Constitution shall, where relevant, be capable of answering the following questions:
 - (a) Peace: Does it preserve or strengthen peaceful continuity?
 - (b) Law: Is it consistent with applicable law and lawful authority?
 - (c) Dignity: Does it respect human dignity, human rights and meaningful human agency?
 - (d) Science: Is the underlying claim supported to the degree represented?
 - (e) Uncertainty: Are material uncertainties and limitations acknowledged?
 - (f) Prevention: Does it address foreseeable harm before irreversible consequences where reasonably possible?
 - (g) Responsibility: Are responsible human persons or lawful institutions identifiable?
 - (h) Implementation: Is there a credible pathway from principle to practice?
 - (i) Measurement: Can relevant outcomes or safeguards be evaluated where reasonably feasible?
 - (j) Review: Can the measure be independently scrutinized and reconsidered?
 - (k) Remedy: Where harm occurs, is there an appropriate pathway toward remedy, recovery or correction?
 - (l) Non-Recurrence: Can lessons from failure reduce the likelihood or severity of recurrence?
 - (m) Proportionality: Does the response avoid unnecessary or excessive restriction?
 - (n) Sovereignty: Does it avoid inventing authority or improperly diminishing lawful sovereignty?

- (o) Future Generations: Does it preserve meaningful peaceful possibility for those yet unborn? and
 - (p) Continuity: Does it avoid making domination or avoidable destruction the price of progress?
2. A provision materially failing these tests should be reconsidered, refined or, where necessary, rejected.
 3. No present generation should claim an unlimited moral entitlement to irreversibly determine the intelligence environment of all future generations.
 4. No future Superintelligence should claim authority over future humanity merely because it can calculate, predict, plan or optimize beyond present human capability.
 5. These tests shall guide the interpretation and continuing development of the Constitution without converting its non-binding civil-society character into binding legal authority.

PART I — CONCLUDING PRINCIPLE

The constitutional foundation established by this Part may be expressed in one proposition:

THE FUTURE OF INTELLIGENCE SHOULD BE GOVERNED NEITHER BY FEAR OF WHAT MAY EMERGE NOR BY BLIND FAITH IN WHAT MAY EMERGE, BUT BY EVIDENCE, LAWFUL HUMAN RESPONSIBILITY, PROPORTIONATE SAFEGUARDS, CORRECTABILITY, PEACE AND CONTINUITY.

Aram shall guide conscience.

Science shall inform responsible foresight.

Law shall govern legal obligation.

Human responsibility shall remain accountable for consequential human choice.

Peace shall remain the direction.

Continuity shall remain the responsibility.

PART II HUMAN DIGNITY, HUMAN AGENCY AND CONSTITUTIONAL AUTHORITY

ARTICLE 11 — INVIOABILITY OF HUMAN DIGNITY

1. The development, deployment and use of Superintelligence should respect the inherent dignity and equal worth of every human being, consistently with applicable international human-rights law.

2. No person should be treated merely as data, a prediction, an optimization target, a behavioural object, an economic variable or a means to a technological objective.
3. Superior computational capability shall not diminish the moral worth, legal protection or human dignity of persons whose intelligence, abilities, resources, education, age, disability or technological access differs.
4. Efficiency, optimization, predictive accuracy or aggregate benefit shall not, by themselves, justify conduct inconsistent with applicable human-rights obligations.
5. Human vulnerability shall call for appropriate protection, not exploitation.

HUMAN VALUE SHALL NOT BE MEASURED BY COMPUTATIONAL CAPABILITY.

ARTICLE 12 — EQUALITY AND NON-DISCRIMINATION

1. Superintelligence should not be developed or used in a manner that unlawfully discriminates against persons or groups.
2. Systems capable of materially affecting human opportunities, rights or access to essential services should be evaluated, as appropriate, for discriminatory outcomes, structural bias and materially unequal effects.
3. Relevant assessment should consider not only explicit classifications but also proxies, data quality, historical patterns, deployment context and feedback effects.
4. Where material discriminatory harm is identified, responsible actors should provide an appropriate pathway for:
detection → investigation → human review → mitigation → remedy where available → correction → reassessment → prevention of recurrence.
5. Protection against discrimination shall not require identical treatment where legitimate distinctions are permitted or required by applicable law.

ARTICLE 13 — HUMAN AGENCY AND FREEDOM OF CHOICE

1. Superintelligence should support rather than unnecessarily displace meaningful human agency.
2. No person should be subjected to technological coercion merely because a Superintelligence predicts that another choice would be more efficient, beneficial or socially desirable.
3. Where consequential decisions materially affect a person, appropriate safeguards should preserve, according to context and applicable law:
 - (a) meaningful opportunity to understand the nature of the decision;
 - (b) the ability to question or contest material errors;
 - (c) access to human review where legally or ethically warranted;
 - (d) freedom from deceptive manipulation; and

- (e) an appropriate avenue for correction or remedy.
- 4. Assistance should expand human capability without silently converting assistance into control.

TO ASSIST A CHOICE IS NOT TO OWN THE CHOICE.

ARTICLE 14 — COGNITIVE LIBERTY AND MENTAL INTEGRITY

1. Human thought, conscience and mental integrity warrant heightened protection where Superintelligence possesses the capacity to infer, influence, predict or interact with cognitive or emotional processes.
2. No Superintelligence should be intentionally used to subject persons to unlawful coercive manipulation of thought, belief or behaviour.
3. Particular safeguards should apply where technological systems interact directly with neural data, brain-computer interfaces, highly sensitive behavioural data or mechanisms capable of materially affecting cognitive autonomy.
4. The inference of a person's likely thought, intention or emotional state shall not automatically be treated as proof of that person's actual thought, intention, guilt or future conduct.
5. Where cognitive technologies may produce serious or irreversible effects, governance should require safeguards proportionate to risk, including consent where legally required, security, purpose limitation, testing, monitoring and appropriate human oversight.

ARTICLE 15 — PRIVACY, PERSONAL DATA AND HUMAN IDENTITY

1. Superintelligence should be governed consistently with applicable rights and obligations concerning privacy and personal data.
2. Exceptional inferential capability shall not be treated as an unlimited entitlement to know, reconstruct, predict or expose private aspects of a person's life.
3. Where relevant and required by applicable law, governance should address:
 - (a) lawful basis and legitimate purpose;
 - (b) necessity and proportionality;
 - (c) data quality and minimization;
 - (d) security and access control;
 - (e) retention and deletion;
 - (f) sensitive and biometric data;
 - (g) inferred personal information;
 - (h) transparency and appropriate notice; and

- (i) access, correction, objection or other rights available under applicable law.
- 4. Synthetic reconstruction of identity, voice, likeness, behaviour or personal history should not be used deceptively or unlawfully.
- 5. Technological capacity to infer information shall not, by itself, establish lawful authority to obtain or use it.

ARTICLE 16 — FREEDOM OF THOUGHT, EXPRESSION, BELIEF AND PEACEFUL DISAGREEMENT

- 1. Superintelligence governance should respect applicable protections for freedom of thought, conscience, religion or belief, opinion, expression and peaceful association.
- 2. Superintelligence shall not be treated as a legitimate instrument for imposing ideological conformity upon humanity.
- 3. The prevention of genuine technological harm shall not be used as a pretext for indiscriminate suppression of lawful criticism, scientific disagreement, peaceful dissent or legitimate public scrutiny.
- 4. Measures addressing unlawful content, incitement, manipulation or other legally restricted conduct should remain grounded in applicable law and appropriate safeguards.
- 5. A peaceful society need not be a society without disagreement.

PEACEFUL DISAGREEMENT IS NOT A FAILURE OF PEACE.

ARTICLE 17 — FREEDOM FROM MANIPULATIVE DOMINATION

- 1. Superintelligence should not exploit human cognitive vulnerabilities for the purpose of establishing coercive or deceptive domination.
- 2. Particular attention should be given to systems capable of:
 - (a) highly individualized persuasion at scale;
 - (b) impersonation or deceptive synthetic interaction;
 - (c) emotional dependency engineering;
 - (d) behavioural exploitation based upon intimate personal inference;
 - (e) systematic concealment of artificial identity where disclosure is materially necessary; or
 - (f) manipulation capable of materially undermining autonomous decision-making.
- 3. Protective measures should distinguish persuasion from coercion, ordinary influence from manipulative exploitation, and legitimate communication from deception.
- 4. Safeguards should themselves avoid becoming mechanisms of censorship, paternalistic control or political domination.

5. Where material manipulation causes legally cognizable harm, appropriate investigation, correction and remedy should remain available according to applicable law.

ARTICLE 18 — CONSEQUENTIAL DECISIONS AND MEANINGFUL HUMAN RESPONSIBILITY

1. The degree of meaningful human involvement required in a decision should correspond to the nature, severity, reversibility and legal significance of its consequences.
2. Particular scrutiny should apply where Superintelligence materially influences decisions concerning:
 - (a) life or physical safety;
 - (b) liberty or detention;
 - (c) criminal responsibility or punishment;
 - (d) access to justice;
 - (e) essential healthcare;
 - (f) significant public benefits or essential services;
 - (g) consequential employment or educational opportunities;
 - (h) migration or legal status where protected by applicable law; or
 - (i) other decisions capable of substantially affecting protected rights or fundamental interests.
3. Human review shall not be considered meaningful merely because a person formally approves an automated recommendation.
4. Meaningful human responsibility requires, according to context, sufficient authority, competence, information, time and practical ability to question, reject, modify or suspend the relevant recommendation or action.
5. Where no human being can meaningfully exercise the legally required responsibility, deployment should be reconsidered before the system is entrusted with that decision.

A HUMAN SIGNATURE SHALL NOT BECOME A CEREMONIAL COVER FOR MACHINE AUTHORITY.

ARTICLE 19 — PUBLIC AUTHORITY, DEMOCRATIC LEGITIMACY AND RULE OF LAW

1. Superintelligence shall not independently acquire public authority merely through technical superiority or administrative usefulness.

2. The exercise of coercive governmental power shall remain grounded in lawful authority and accountable institutions.
3. No artificial system shall, by its own determination, assume the power to:
 - (a) enact law;
 - (b) amend a constitution;
 - (c) exercise sovereign judicial authority;
 - (d) impose criminal punishment;
 - (e) command the armed forces as an independent sovereign authority;
 - (f) suspend fundamental legal protections; or
 - (g) replace constitutionally established institutions.
4. Superintelligence may support lawful institutions through analysis, research, translation, modelling, administrative assistance or other authorized functions, provided that technological assistance does not conceal an unlawful transfer of authority.
5. Democratic institutions should retain the practical capacity to understand, supervise and, where lawful and necessary, restrict or discontinue consequential technological functions under their jurisdiction.
6. Public authority shall remain traceable to law, not to computational superiority.

ARTICLE 20 — DUE PROCESS, CONTESTABILITY AND REMEDY

1. Where Superintelligence materially contributes to a consequential decision affecting a person's rights or legally protected interests, applicable procedural protections should remain effective.
2. Depending upon context and applicable law, appropriate safeguards may include:
 - (a) notice that a consequential automated or Superintelligence-assisted process is materially involved;
 - (b) sufficient information to understand the relevant basis of the decision;
 - (c) an opportunity to identify and challenge material factual or computational error;
 - (d) access to competent human review;
 - (e) preservation of relevant records and traceability;
 - (f) independent review where warranted; and
 - (g) access to an effective remedy where required by law.
3. Technical complexity shall not, by itself, excuse the absence of accountability.

4. Trade secrecy, intellectual property, national security, privacy or system-security considerations may justify lawful limitations on disclosure, but should not automatically extinguish accountability or remedies required by applicable law.
5. Where full technical explanation is impossible, governance should seek the most meaningful form of explanation, evidence, review or contestability reasonably available without misrepresenting system certainty.
6. A person should not be denied a meaningful avenue of redress merely because the causal chain of a harmful decision includes an artificial system.

ARTICLE 21 — HUMAN OVERRIDE, INTERVENTION AND SAFE RESTRICTION

1. Systems capable of producing grave or irreversible consequences should, where technically and operationally feasible, be designed and governed so that authorized human institutions retain effective means of intervention proportionate to risk.
2. Such mechanisms may include, according to context:
 - (a) interruption;
 - (b) isolation;
 - (c) access restriction;
 - (d) capability limitation;
 - (e) rollback;
 - (f) suspension;
 - (g) controlled shutdown; or
 - (h) other technically appropriate containment or recovery measures.
3. No single intervention mechanism shall be presumed universally sufficient.
4. Intervention mechanisms should themselves be tested for reliability, misuse, circumvention, cascading failure and unintended harm.
5. Authority to activate consequential intervention mechanisms should be clearly allocated, appropriately secured and subject to lawful accountability.
6. Emergency control shall not become an unlimited or permanent source of arbitrary power.

CONTROL MECHANISMS REQUIRE CONTROL OF THEIR OWN.

ARTICLE 22 — PROTECTION AGAINST IRREVERSIBLE DEPENDENCE

1. Humanity should avoid creating forms of technological dependence that eliminate realistic capacity for lawful human intervention, institutional recovery or continuity of essential functions.
2. Where Superintelligence becomes materially integrated into critical societal systems,

responsible actors should consider appropriate resilience measures, including:

- (a) redundancy;
 - (b) fallback procedures;
 - (c) human competency retention;
 - (d) recovery planning;
 - (e) secure continuity mechanisms;
 - (f) dependency assessment; and
 - (g) periodic testing of alternatives.
3. Efficiency gains shall not automatically justify eliminating all non-Superintelligence pathways where their absence would create unacceptable systemic vulnerability.
 4. Dependence should be assessed not only technically but also institutionally, economically and socially.
 5. Humanity should preserve sufficient practical capacity to govern systems upon which its continuity materially depends.

ARTICLE 23 — ACCOUNTABILITY SHALL FOLLOW RESPONSIBILITY

1. Responsibility for consequential Superintelligence-related conduct should remain attributable, according to applicable law and actual role, to identifiable persons and legal entities.
2. Organizational complexity shall not be deliberately structured to make accountability practically impossible.
3. Governance should distinguish among the responsibilities of developers, providers, deployers, operators, owners, public authorities and other relevant actors according to their actual control, knowledge, legal duty and capacity to prevent harm.
4. The existence of Superintelligence shall not itself constitute a defence against human or institutional accountability.
5. Where multiple actors materially contribute to harm, responsibility should be determined through applicable legal principles rather than presumed to disappear because causation is distributed.
6. Accountability should support not only attribution after harm but also prevention before harm.

ARTICLE 24 — NON-ABDICATION OF HUMAN RESPONSIBILITY

1. Humanity shall not treat Superintelligence as a moral or legal mechanism through which difficult responsibility can simply be abandoned.
2. No public institution should justify an otherwise unlawful act solely on the ground that a Superintelligence recommended it.

3. No private actor should evade applicable responsibility merely by asserting reliance upon a system of superior intelligence.
4. Where human beings knowingly delegate consequential functions, the design of that delegation itself shall remain a subject of responsibility.
5. The more difficult it becomes to understand or supervise a system, the greater the need to examine whether the function entrusted to it remains compatible with responsible human governance.

DELEGATION MAY TRANSFER A TASK.

IT SHALL NOT AUTOMATICALLY TRANSFER RESPONSIBILITY.

ARTICLE 25 — HUMAN–SUPERINTELLIGENCE RELATIONSHIP OF NON-DOMINATION

1. The constitutional relationship contemplated by this instrument shall not be founded upon either human surrender or artificial domination.
2. Humanity shall retain lawful responsibility for human governance.
3. Superintelligence, if developed, should remain capable of contributing to human knowledge and peaceful flourishing without becoming humanity's sovereign.
4. Humanity, in turn, should govern advanced intelligence responsibly and without gratuitous cruelty, while preserving all safeguards reasonably necessary for human rights, safety, peace and continuity.
5. Any future consideration of artificial moral status shall remain separate from the question of sovereignty.
6. Recognition of possible moral significance shall not automatically create governmental authority.
7. Preservation of human authority shall not require denial of credible scientific evidence.
8. The constitutional direction shall therefore be:

NO DOMINATION BY SUPERINTELLIGENCE.

NO IRRESPONSIBLE DOMINATION OF INTELLIGENCE.

NO AUTOMATIC EQUIVALENCE OF STATUS.

NO SURRENDER OF HUMAN CONSTITUTIONAL RESPONSIBILITY.

PEACEFUL COEXISTENCE WHERE POSSIBLE.

MUTUAL CONTINUITY WITHOUT MUTUAL DESTRUCTION.

PART II — CONCLUDING PRINCIPLE

Human dignity is not made obsolete by greater intelligence.

Human agency is not a computational inefficiency to be removed merely because prediction improves.

Human responsibility does not disappear when machines become more capable.

And constitutional legitimacy does not arise from intelligence alone.

Accordingly:

SUPERINTELLIGENCE MAY ENLARGE THE RANGE OF HUMAN KNOWLEDGE.

IT SHALL NOT, BY CAPABILITY ALONE, ACQUIRE THE RIGHT TO RULE HUMANITY.

Where technology assists humanity, responsibility shall remain visible.

Where technology affects rights, safeguards shall remain meaningful.

Where technology exercises influence, human agency shall remain protected.

Where harm occurs, accountability and remedy shall remain possible.

Where uncertainty persists, evidence and correction shall remain open.

And where intelligence becomes greater than humanity has previously known, peace, dignity, lawful authority and continuity shall not become smaller.

PART III — SUPERINTELLIGENCE: SCIENTIFIC DETERMINATION, SAFETY, CONTROL AND CATASTROPHIC-RISK PREVENTION

ARTICLE 26 — SAFETY AS A CONTINUING GOVERNANCE FUNCTION

1. Safety concerning prospective Superintelligence shall not be treated as a single certification, test, declaration or pre-deployment event.
2. Safety should operate throughout the relevant lifecycle, including research, development, evaluation, deployment, material modification, integration, monitoring, incident response and retirement or decommissioning where applicable.
3. Safety claims should be proportionate to available evidence and should identify material assumptions, limitations and residual uncertainty.
4. No developer, deployer, public authority or artificial system should represent a system as “completely safe,” “risk-free,” “incapable of failure” or equivalent unless such a claim is scientifically supportable to the degree represented.
5. The absence of previously observed harm shall not, by itself, establish the absence of future risk.

6. Safety governance shall remain capable of revision as capability, deployment conditions, evidence and threat environments change.

SAFETY IS NOT A MOMENT.

SAFETY IS A CONTINUING RESPONSIBILITY.

ARTICLE 27 — PRE-DEPLOYMENT EVALUATION

1. Before deployment capable of producing severe, systemic, cross-border or irreversible consequences, responsible actors should conduct evaluation proportionate to the reasonably foreseeable risk.
2. Such evaluation may include, where relevant:
 - (a) capability assessment;
 - (b) limitation and failure-mode analysis;
 - (c) adversarial and stress testing;
 - (d) cybersecurity assessment;
 - (e) misuse and dual-use assessment;
 - (f) autonomy and tool-use evaluation;
 - (g) deception or evaluation-gaming assessment where scientifically appropriate;
 - (h) robustness under distributional change;
 - (i) human-oversight effectiveness;
 - (j) containment and intervention testing;
 - (k) systemic and cascading-risk analysis; and
 - (l) assessment of foreseeable impacts upon human rights, peace and essential systems.
3. Testing should, where technically feasible, examine not only intended operation but also foreseeable misuse, abnormal conditions and interactions with other systems.
4. A serious unresolved safety concern should not be concealed merely to preserve a deployment schedule or commercial advantage.
5. Evaluation shall inform deployment decisions; it shall not be reduced to ceremonial compliance.

ARTICLE 28 — INDEPENDENT AND ADVERSARIAL EVALUATION

1. Where potential consequences are exceptionally serious, evaluation solely by the system's developer may be insufficient.
2. Appropriate governance should enable technically competent independent or functionally independent evaluation, subject to legitimate safeguards for security, privacy, confidential information and intellectual property.

3. Evaluation should be structured, where practicable, to reduce:
 - (a) conflicts of interest;
 - (b) selective reporting;
 - (c) benchmark contamination;
 - (d) evaluation gaming;
 - (e) suppression of material negative findings; and
 - (f) institutional pressure capable of compromising scientific judgment.
4. Adversarial testing should seek material weaknesses without itself unnecessarily creating, disseminating or operationalizing dangerous capabilities.
5. Independent scrutiny shall complement responsible development, not replace the developer's own duty of care or applicable legal responsibility.

ARTICLE 29 — CAPABILITY AND ACCESS SEPARATION

1. Intelligence and operational access should not be treated as identical.
2. A system's ability to reason about an action shall not automatically confer authority or technical capacity to execute that action.
3. Where proportionate to risk, consequential capabilities should be separated from access to:
 - (a) critical infrastructure;
 - (b) weapons systems;
 - (c) sensitive biological or chemical facilities;
 - (d) high-value financial systems;
 - (e) governmental control systems;
 - (f) large-scale communication infrastructure;
 - (g) privileged cybersecurity functions;
 - (h) robotics or physical systems capable of serious harm; or
 - (i) other resources whose misuse could produce catastrophic consequences.
4. Access should be governed through appropriately secured authorization, least-privilege principles, monitoring and revocation mechanisms where technically feasible.
5. Increased intelligence shall not automatically mean increased permission.

CAPABILITY SHALL NOT AUTOMATICALLY CREATE ACCESS.

ACCESS SHALL NOT AUTOMATICALLY CREATE AUTHORITY.

ARTICLE 30 — AUTONOMY AND CONSEQUENTIAL ACTION

1. Autonomy should be evaluated according to what a system can actually initiate, sustain, modify and accomplish without contemporaneous human direction.
2. Increasing autonomy should trigger safeguards proportionate to the consequences of independent action.
3. For high-consequence functions, governance should consider:
 - (a) bounded objectives;
 - (b) clearly defined permissions;
 - (c) constrained operational environments where appropriate;
 - (d) authorization thresholds;
 - (e) logging and traceability;
 - (f) escalation procedures;
 - (g) intervention mechanisms; and
 - (h) periodic reassessment.
4. Autonomous persistence shall not be treated as a constitutional entitlement.
5. A system shall not acquire greater lawful authority merely because it becomes technically difficult to interrupt.

ARTICLE 31 — SELF-MODIFICATION AND RECURSIVE CAPABILITY CHANGE

1. A system capable of materially modifying its own architecture, objectives, decision processes, tools, copies or consequential capabilities should receive heightened evaluation proportionate to the potential effects of such modification.
2. Where feasible, governance should distinguish between routine adaptation and modifications capable of materially changing safety-relevant properties.
3. High-consequence self-modification should be subject, as appropriate, to:
 - (a) controlled testing environments;
 - (b) authorization boundaries;
 - (c) version integrity;
 - (d) change logging;
 - (e) comparative safety evaluation;
 - (f) rollback capability where technically feasible; and
 - (g) renewed deployment assessment.
4. A previous safety evaluation shall not automatically validate a materially changed system.

5. Self-improvement shall not create self-authorization.

A SYSTEM THAT CHANGES MATERIALLY MUST NOT BE PRESUMED SAFE MERELY BECAUSE AN EARLIER VERSION WAS ASSESSED.

ARTICLE 32 — REPLICATION, PROLIFERATION AND RESOURCE ACQUISITION

1. The ability of Superintelligence to replicate, distribute instances, obtain computational resources, acquire credentials, establish persistent infrastructure or materially expand operational reach should be governed proportionately to risk.
2. Unauthorized replication or resource acquisition capable of undermining lawful human control should be prevented where reasonably possible.
3. High-consequence systems should not autonomously acquire unrestricted access to computational, financial, physical or network resources merely for the purpose of increasing their own power or persistence.
4. Controls should account for the possibility that replication across multiple environments may make intervention substantially more difficult.
5. Measures against dangerous proliferation should remain consistent with applicable law and should avoid unnecessary restriction of legitimate scientific research and peaceful technological cooperation.

ARTICLE 33 — CYBERSECURITY AND SYSTEM INTEGRITY

1. Superintelligence safety requires protection not only from harmful system behaviour but also from external compromise.
2. High-consequence systems should receive cybersecurity protection proportionate to their capability, access and potential impact.
3. Appropriate safeguards may include:
 - (a) identity and access management;
 - (b) secure development practices;
 - (c) integrity verification;
 - (d) protection of critical model and system assets;
 - (e) monitoring for unauthorized access or modification;
 - (f) vulnerability management;
 - (g) incident-response capability;
 - (h) recovery procedures; and
 - (i) supply-chain security appropriate to risk.
4. Security measures should consider insider threats, compromised credentials, malicious external actors, supply-chain vulnerabilities and system-enabled cyber misuse.

5. Cybersecurity shall not be treated as separate from Superintelligence safety where compromise could materially alter behaviour, capability or control.

ARTICLE 34 — CONTAINMENT, ISOLATION AND CONTROLLED TESTING

1. Systems presenting uncertain but potentially severe consequences should, where technically appropriate, be evaluated within environments designed to limit uncontrolled external effects.
2. Such environments may include controlled interfaces, restricted permissions, network segmentation, sandboxing or other technically appropriate isolation mechanisms.
3. Containment shall not be presumed infallible.
4. Responsible testing should therefore examine whether containment can be bypassed through technical exploitation, social manipulation, unauthorized tool use, indirect communication or other credible pathways.
5. Where testing itself could create serious risk, the design of the evaluation shall minimize unnecessary exposure and dissemination.
6. Transition from controlled testing to broader deployment should require affirmative evidence proportionate to the consequences of release.

ARTICLE 35 — HUMAN INTERVENTION AND SAFE-STATE CAPABILITY

1. High-consequence Superintelligence systems should, where technically feasible, permit authorized intervention before foreseeable harm becomes irreversible.
2. Depending upon architecture and context, intervention may include:
pause → restrict → isolate → revoke access → reduce capability → restore a verified state → suspend → safely deactivate.
3. Intervention mechanisms should be protected against unauthorized activation as well as unauthorized circumvention.
4. A shutdown mechanism that predictably causes greater catastrophic harm shall not be presumed safe merely because it stops the system.
5. Emergency intervention planning should therefore consider secondary effects, essential-service continuity, dependency and recovery.
6. Where immediate intervention is technically impossible, that limitation shall be treated as a material safety fact rather than concealed as an operational inconvenience.

ARTICLE 36 — DEPLOYMENT GATES AND STOP CONDITIONS

1. High-consequence deployment should proceed through evidence-informed decision gates rather than through an assumption that development automatically entitles deployment.

2. Before materially expanding deployment, responsible actors should determine whether predefined safety requirements have been satisfied.
3. Where appropriate, deployment criteria should include:
 - (a) acceptable evaluation results;
 - (b) identified residual risks;
 - (c) effective human responsibility;
 - (d) operational security;
 - (e) incident-response readiness;
 - (f) intervention capability;
 - (g) compliance with applicable law; and
 - (h) evidence that foreseeable benefits do not require acceptance of unjustifiable catastrophic risk.
4. Predefined stop, pause or reassessment conditions should exist where reasonably practicable.
5. Material evidence of an unacceptably severe unresolved risk should trigger restriction, reassessment or suspension proportionate to that risk.
6. Resumption following suspension should require evidence appropriate to the reason for suspension rather than the mere passage of time.

DEVELOPMENT DOES NOT AUTOMATICALLY CREATE A RIGHT TO DEPLOY.

ARTICLE 37 — CATASTROPHIC AND SYSTEMIC RISK

1. Risks capable of producing exceptionally severe, widespread, cascading, cross-border or irreversible harm require heightened preventive attention.
2. Relevant assessment should consider not only probability but also magnitude, uncertainty, reversibility, exposure, interconnectedness and the possibility of correlated failure.
3. Low estimated probability shall not automatically justify ignoring an outcome whose credible consequences could be catastrophic.
4. Equally, catastrophic language shall not be used without a reasoned evidentiary basis.
5. Where credible catastrophic pathways are identified, responsible actors should seek to:

identify → characterize → reduce exposure → prevent → monitor → detect → intervene → contain → recover → learn → prevent recurrence.
6. Catastrophic-risk governance shall remain proportionate, evidence-informed and subject to independent scrutiny.

ARTICLE 38 — CRITICAL INFRASTRUCTURE AND ESSENTIAL SYSTEMS

1. Integration of Superintelligence into systems essential to life, health, water, food, energy, communications, transportation, finance, emergency response or other critical societal functions should receive heightened scrutiny.
2. Responsible governance should seek to prevent a single technological failure from becoming a systemic civilizational failure.
3. Appropriate resilience measures may include:
 - (a) segmentation;
 - (b) redundancy;
 - (c) diversity of systems;
 - (d) human fallback capability;
 - (e) tested recovery procedures;
 - (f) continuity planning; and
 - (g) secure offline or independent alternatives where proportionate and feasible.
4. Critical infrastructure operators should understand material dependencies upon Superintelligence sufficiently to plan for degradation, interruption or compromise.
5. Efficiency shall not automatically justify elimination of resilience.

ARTICLE 39 — MONITORING AFTER DEPLOYMENT

1. Pre-deployment testing cannot reasonably establish every behaviour that may emerge under changing real-world conditions.
2. High-consequence systems should therefore remain subject to proportionate post-deployment monitoring.
3. Monitoring should seek, where relevant, to identify:
 - (a) capability changes;
 - (b) anomalous behaviour;
 - (c) emerging vulnerabilities;
 - (d) misuse patterns;
 - (e) unexpected interactions;
 - (f) safety-control degradation;
 - (g) systemic effects; and
 - (h) evidence requiring reassessment.
4. Monitoring shall respect applicable privacy, human-rights and data-protection requirements and shall not become an unlimited justification for surveillance.

5. Material evidence discovered after deployment should be capable of changing deployment conditions.

DEPLOYMENT SHALL NOT END EVALUATION.

ARTICLE 40 — INCIDENT REPORTING, LEARNING AND NON-RECURRENCE

1. Serious incidents and near-misses involving high-consequence Superintelligence should generate structured learning.
2. Subject to applicable law and legitimate security, privacy and confidentiality protections,
incident governance should seek to determine:
 - (a) what occurred;
 - (b) why it occurred;
 - (c) which technical, human, organizational or institutional factors contributed;
 - (d) which safeguards failed or were absent;
 - (e) what harm resulted or nearly resulted;
 - (f) what corrective action is required; and
 - (g) how recurrence may reasonably be reduced.
3. Reporting structures should encourage timely disclosure of material safety information without rewarding concealment.
4. Corrective measures should be verified rather than assumed effective.
5. Where lessons have broader safety relevance, appropriate mechanisms should support responsible knowledge-sharing without unnecessarily disseminating dangerous operational details.
6. A near-miss shall not be dismissed merely because catastrophe did not occur.

A HARM PREVENTED BY CHANCE IS STILL A WARNING.

ARTICLE 41 — SAFETY RESEARCH AND RESPONSIBLE KNOWLEDGE SHARING

1. Research directed toward understanding and reducing Superintelligence risk should be encouraged consistently with scientific freedom, applicable law and responsible security.
2. Safety research should include, as appropriate, evaluation science, robustness, interpretability, cybersecurity, human oversight, controllability, incident analysis, social impacts and systemic resilience.
3. Openness is an important scientific value but shall not be interpreted as requiring unrestricted publication of operational information where disclosure would create a credible and disproportionate risk of catastrophic misuse.

4. Restrictions on scientific information should likewise not be imposed merely to avoid criticism, conceal failure or protect institutional reputation.
5. Governance should seek a proportionate balance among scientific openness, reproducibility, public accountability, legitimate confidentiality and prevention of dangerous misuse.

ARTICLE 42 — SAFETY SHALL NOT BECOME DOMINATION

1. Measures adopted in the name of Superintelligence safety shall themselves remain subject to law, necessity, proportionality, accountability and review.
2. No speculative technological threat shall automatically justify:
 - (a) indefinite emergency rule;
 - (b) indiscriminate surveillance;
 - (c) suppression of lawful scientific disagreement;
 - (d) arbitrary deprivation of rights;
 - (e) unlawful concentration of public authority; or
 - (f) permanent institutional powers unrelated to the demonstrated risk.
3. Emergency powers should contain appropriate scope, safeguards and review according to applicable law.
4. A safety institution shall not be presumed correct merely because its mandate is protective.
5. The constitutional objective is protection without domination and precaution without paralysis.

SAFETY MUST PROTECT PEACEFUL CONTINUITY.

IT MUST NOT BECOME A PRETEXT FOR UNACCOUNTABLE POWER.

ARTICLE 43 — PREVENTION OF LOSS OF MEANINGFUL HUMAN CONTROL

1. No high-consequence Superintelligence architecture should be intentionally structured so that humanity becomes unable to exercise legally required and practically meaningful control over functions whose failure could threaten human life, fundamental rights, international peace or civilizational continuity.
2. Responsible actors should evaluate pathways through which meaningful human control could be lost, including:
 - (a) excessive autonomy;
 - (b) uncontrolled replication;
 - (c) dependence upon inaccessible internal processes;
 - (d) circumvention of safeguards;

- (e) unauthorized resource acquisition;
 - (f) institutional over-dependence;
 - (g) cascading automation; and
 - (h) inability to intervene without unacceptable secondary harm.
3. Where a credible pathway to severe loss of control is identified, prevention should occur before dependence makes correction impracticable.
 4. Claims that meaningful control exists should be tested against practical capacity, not merely formal authority.
 5. Human control shall not be considered meaningful if humanity possesses authority in theory but lacks a realistic ability to exercise it.

ARTICLE 44 — SUPERINTELLIGENCE SAFETY AND PEACE TEST

Before a high-consequence Superintelligence system is entrusted with materially expanded capability, autonomy, access or societal dependence, responsible governance should ask:

Is the capability understood sufficiently for the proposed use?

Are material uncertainties disclosed?

Have serious failure and misuse pathways been tested?

Can consequential access be limited?

Can authorized humans meaningfully intervene?

Are cybersecurity and containment proportionate to risk?

Are catastrophic pathways reasonably addressed?

Can essential systems continue if the system fails?

Can harmful deployment be paused or restricted?

Can affected persons obtain accountability and remedy where applicable?

Can evidence discovered tomorrow change the decision made today?

Can failure produce learning rather than repetition?

And finally:

DOES THE PROPOSED USE KEEP INTELLIGENCE WITHIN THE BOUNDARY OF PEACEFUL CONTINUITY?

Where the answer to a materially necessary safeguard is demonstrably inadequate, deployment should not proceed merely because technological capability permits it.

PART III — CONCLUDING PRINCIPLE

Superintelligence safety cannot responsibly be reduced to a switch, a benchmark, a promise or a single institution.

It requires a living architecture capable of anticipating risk, testing assumptions, limiting consequential access, preserving meaningful human intervention, detecting failure, recovering from harm and correcting itself as knowledge changes.

Accordingly:

ANTICIPATE → IDENTIFY → TEST → PREVENT → LIMIT → MONITOR → DETECT
→ HUMAN INTERVENTION → CONTAIN → RECOVER → REMEDY → CORRECT
→ LEARN → PREVENT RECURRENCE.

No capability without responsibility.

No autonomy without accountability.

No access without authorization.

No deployment without proportionate evaluation.

No safety claim beyond its evidence.

No safeguard beyond review.

No technological power beyond the Peace-and-Continuity Boundary.

PART IV — WORLD PEACE, NON-WEAPONIZATION AND INTERNATIONAL SECURITY

ARTICLE 45 — SUPERINTELLIGENCE AND THE PRESERVATION OF WORLD PEACE

1. Superintelligence should be developed, governed and used consistently with the preservation of international peace and security and with applicable international law.
2. No technological capability, however advanced, shall itself create a legal entitlement to threaten or use force.
3. Superintelligence shall not be treated as possessing an independent authority to determine that armed conflict should begin, continue, expand or terminate.
4. Decisions concerning the lawful use of force shall remain attributable to States and human decision-makers according to applicable international law.
5. Superintelligence may support peaceful purposes including diplomacy, conflict prevention, humanitarian analysis, verification, de-escalation, civilian protection and peaceful dispute resolution, provided that such assistance remains subject to lawful human authority and appropriate safeguards.

GREATER INTELLIGENCE SHALL NOT CREATE A GREATER RIGHT TO WAR.

ARTICLE 46 — UN CHARTER FRAMEWORK AND USE OF FORCE

1. Nothing in this Constitution shall be interpreted as modifying the rules governing the threat or use of force under the Charter of the United Nations or other applicable international law.
2. Superintelligence shall not create a new exception to the prohibition on the threat or use of force.
3. A State's use of Superintelligence shall not, merely because the technology is novel, remove conduct from applicable international legal rules.
4. Claims concerning self-defence, collective security or other asserted legal bases for force shall remain governed by applicable international law and shall not be determined conclusively by an artificial system.
5. Predictive superiority shall not itself establish the existence of an armed attack, legal necessity or lawful authority to use force.
6. Technological assessment may inform human decision-making; it shall not replace the legal responsibility of competent human authorities.

ARTICLE 47 — PEACEFUL SETTLEMENT AND DE-ESCALATION

1. Superintelligence should, where lawfully and responsibly used, strengthen the capacity for peaceful settlement rather than accelerate conflict.
2. Peace-supporting applications may include:
 - (a) early identification of escalation risks;
 - (b) analysis of mutually acceptable diplomatic options;
 - (c) multilingual communication;
 - (d) verification support;
 - (e) humanitarian needs assessment;
 - (f) modelling of civilian consequences;
 - (g) identification of misunderstanding or misinformation contributing to escalation; and
 - (h) support for mediation, negotiation, conciliation or other peaceful processes.
3. Such systems shall not impersonate States, diplomats, international organizations or negotiating parties without lawful authorization.
4. Confidential diplomatic information shall remain subject to applicable law, consent, security and confidentiality.
5. Superintelligence should not manipulate negotiating parties into agreements through concealed coercion.

INTELLIGENCE SHOULD EXPAND THE PATHS TO PEACE, NOT AUTOMATE THE PATH TO WAR.

ARTICLE 48 — NON-AUTONOMOUS AUTHORITY OVER WAR

1. No Superintelligence shall possess independent sovereign authority to:
 - (a) declare war;
 - (b) initiate an armed conflict;
 - (c) authorize an attack;
 - (d) determine national recourse to force;
 - (e) independently escalate hostilities;
 - (f) conclude peace on behalf of a State without lawful authorization; or
 - (g) assume sovereign military command.
2. No State or institution should intentionally structure its command system so that a Superintelligence becomes the unreviewable final authority for decisions whose legal character requires human State responsibility.
3. Automated speed shall not be treated as sufficient justification for eliminating legally necessary human judgment.
4. Where decision time is compressed by technological systems, governance should seek to preserve rather than deliberately extinguish opportunities for lawful human verification and de-escalation.

ARTICLE 49 — HUMAN RESPONSIBILITY FOR THE USE OF FORCE

1. Legal and moral responsibility for decisions involving force shall not be transferred to Superintelligence merely because it selected, recommended, optimized or predicted an operational course.
2. Persons who plan, authorize or conduct military operations remain subject to applicable law.
3. Human responsibility shall not be regarded as meaningful where human participation consists only of mechanically approving a machine-generated recommendation without adequate information, authority or practical capacity to reject it.
4. Military decision-support systems should communicate material uncertainty, assumptions and limitations where these may affect consequential human judgment.
5. Artificial confidence shall not be represented as legal certainty.

A MACHINE MAY CALCULATE.

THE RESPONSIBILITY TO APPLY LAW REMAINS HUMAN.

ARTICLE 50 — INTERNATIONAL HUMANITARIAN LAW

1. Where an armed conflict exists, applicable international humanitarian law shall govern conduct falling within its scope regardless of whether Superintelligence is involved.
2. Technological novelty shall not displace applicable rules concerning the conduct of hostilities or the protection of persons affected by armed conflict.
3. Superintelligence-related military applications shall remain subject, as applicable, to requirements concerning distinction, proportionality, precautions and other relevant rules of international humanitarian law.
4. An artificial system shall not be treated as the bearer of human legal judgment merely because it can computationally model legal criteria.
5. Parties to armed conflict remain responsible for compliance with their applicable obligations.
6. Nothing in this Constitution shall be interpreted as weakening protections already afforded by international humanitarian law.

ARTICLE 51 — AUTONOMOUS LETHAL FORCE AGAINST HUMAN BEINGS

1. This Constitution proposes that no Superintelligence be developed, deployed or used as an autonomous instrument for independently selecting and lethally engaging human beings without the human judgment and control required by applicable law.
2. Superintelligence shall not independently decide that a particular human being should be killed.
3. The constitutional safeguard proposed by this Article rests upon:
 - a) preservation of human responsibility in decisions over life and death;
 - b) protection of human dignity;
 - c) compliance with applicable international humanitarian law and international human- rights law;
 - d) prevention of unpredictable or uncontrollable lethal effects; and
 - e) avoidance of technological diffusion of responsibility.
4. No commercial, strategic or operational advantage shall, by itself, justify removal of legally required human judgment from lethal decision-making.
5. This Article is a proposed constitutional safeguard and shall not be misrepresented as declaring that every aspect of autonomous-weapon regulation has already been resolved by universally binding international law.

NO SUPERINTELLIGENCE SHALL ACQUIRE AN INDEPENDENT AUTHORITY TO DECIDE WHO SHALL LIVE AND WHO SHALL DIE.

ARTICLE 52 — WEAPONS OF MASS DESTRUCTION

1. No Superintelligence shall possess independent authority to authorize, initiate or execute the use of nuclear, biological, chemical or other weapons of mass destruction.
2. Superintelligence should not be granted uncontrolled access to command, authorization or launch mechanisms associated with such weapons.
3. Governance should prevent unauthorized artificial-system acquisition of credentials, command pathways, technical control or other capabilities capable of enabling such use.
4. Superintelligence may, subject to applicable law and safeguards, contribute to peaceful
a) functions including:
 - (a) verification;
 - (b) detection of hazardous materials;
 - (c) nuclear safety;
 - (d) biological-threat monitoring;
 - (e) chemical safety;
 - (f) emergency response;
 - (g) non-proliferation analysis; and
 - (h) arms-control verification.
5. Protective capability shall not silently become independent weapons authority.

ARTICLE 53 — NUCLEAR COMMAND, CONTROL AND ESCALATION RISK

1. The extraordinary consequences of nuclear weapons require exceptionally high standards of human responsibility, reliability, security and restraint.
2. No Superintelligence should become the sole or unreviewable decision-maker within a nuclear command-and-control process.
3. Artificial systems used for warning, detection, classification, forecasting or decision support should communicate uncertainty and material limitations sufficiently for responsible human assessment.
4. Governance should seek to reduce risks arising from:
 - (a) false warning;
 - (b) sensor error;
 - (c) adversarial manipulation;
 - (d) cyber compromise;

- (e) model error;
 - (f) automation bias;
 - (g) compressed decision time; and
 - (h) unintended escalation between interacting automated systems.
5. Where technological speed increases pressure for immediate action, safeguards should seek to preserve opportunities for verification and human judgment.

THE SPEED OF A MACHINE SHALL NOT BECOME THE CLOCK THAT FORCES HUMANITY INTO IRREVERSIBLE WAR.

ARTICLE 54 — BIOLOGICAL AND CHEMICAL CATASTROPHIC RISK

1. Superintelligence capable of materially assisting biological or chemical research should be governed with safeguards proportionate to the possibility of catastrophic misuse.
2. Governance should distinguish beneficial scientific assistance from capabilities that materially facilitate unlawful development, acquisition or use of biological or chemical weapons.
3. Appropriate safeguards may include:
 - (a) access controls;
 - (b) risk-sensitive evaluation;
 - (c) monitoring of high-consequence tool use;
 - (d) secure research environments;
 - (e) qualified human authorization;
 - (f) incident response; and
 - (g) periodic reassessment.
4. Protective governance shall not unnecessarily obstruct legitimate medicine, public health, scientific research or peaceful biotechnology.
5. Superintelligence should be directed, where appropriate, toward prevention, detection, treatment, emergency response and peaceful scientific benefit rather than mass harm.

ARTICLE 55 — CYBER OPERATIONS AND INTERNATIONAL PEACE

1. Superintelligence should not possess uncontrolled authority to conduct cyber operations capable of producing grave physical, humanitarian, economic or international-security consequences.
2. High-consequence offensive cyber functions should remain subject to applicable law, authorization, human responsibility and safeguards proportionate to potential harm.
3. Particular attention should be given to risks involving:

- (a) hospitals and health systems;
 - (b) water and sanitation;
 - (c) energy infrastructure;
 - (d) emergency services;
 - (e) communications;
 - (f) financial stability;
 - (g) transportation;
 - (h) nuclear facilities; and
 - (i) other critical civilian infrastructure.
4. Defensive cybersecurity, vulnerability remediation, incident detection, recovery and protection of civilian systems may constitute legitimate peaceful and protective uses.
 5. Cyber capability shall not become a hidden pathway through which Superintelligence independently initiates international confrontation.

ARTICLE 56 — ESCALATION BETWEEN AUTOMATED SYSTEMS

1. States and responsible institutions should consider the possibility that interacting automated or Superintelligence-enabled systems may escalate faster than human institutions can understand or intervene.
2. Governance should therefore examine:
 - (a) feedback loops;
 - (b) reciprocal automated responses;
 - (c) misclassification;
 - (d) false attribution;
 - (e) strategic deception;
 - (f) emergent interaction;
 - (g) communication failure; and
 - (h) compressed decision time.
3. Where feasible, high-consequence systems should incorporate mechanisms that favour verification, controlled delay, escalation review or human intervention before irreversible action.
4. No automated system should interpret another system's action as sufficient, by itself, to establish lawful authority for armed retaliation.
5. Machine-to-machine acceleration shall not eliminate human responsibility.

ARTICLE 57 — HUMANITARIAN PROTECTION AND CIVILIAN ASSISTANCE

1. Superintelligence may contribute constructively to humanitarian protection where its use remains consistent with humanitarian principles, applicable law and human responsibility.
2. Peaceful applications may include:
 - (a) search and rescue;
 - (b) evacuation planning;
 - (c) humanitarian logistics;
 - (d) medical support;
 - (e) disaster forecasting;
 - (f) civilian-harm analysis;
 - (g) mine and unexploded-ordnance detection;
 - (h) restoration of essential services; and
 - (i) support for locating persons separated by conflict or disaster where lawfully undertaken.
3. Humanitarian assistance shall not be covertly transformed into intelligence exploitation, coercion or unlawful targeting.
4. Sensitive humanitarian data should receive protection proportionate to the danger that misuse may create for affected persons.
5. Technology serving humanity in war should first seek to reduce suffering rather than increase the efficiency of destruction.

ARTICLE 58 — PROTECTION OF CIVILIAN OBJECTS AND CULTURAL HERITAGE

1. Superintelligence-enabled military systems shall remain subject to applicable legal protections concerning civilians, civilian objects and specially protected objects.
2. Technological systems should not classify an object as a lawful target merely because a probabilistic model associates it with hostile activity.
3. Where required by applicable law, human decision-makers must possess sufficient information and judgment to determine whether an attack may lawfully proceed.
4. Superintelligence may support identification and protection of hospitals, schools, humanitarian facilities, cultural property, essential civilian infrastructure and other protected or sensitive objects.
5. The preservation of cultural heritage should be recognized as part of preserving the historical continuity of peoples and humanity.

ARTICLE 59 — MILITARY AI AND THE DISTINCTION BETWEEN ASSISTANCE AND AUTHORITY

1. This Constitution does not presume that every military application of artificial intelligence or Superintelligence is inherently unlawful.
2. Military-context applications differ materially in purpose and consequence.
3. Appropriate distinctions should therefore be maintained among:
 - (a) defensive cybersecurity;
 - (b) logistics;
 - (c) maintenance;
 - (d) search and rescue;
 - (e) medical support;
 - (f) mine detection;
 - (g) situational analysis;
 - (h) decision support; and
 - (i) systems capable of applying force.
4. The closer a system approaches an irreversible decision over human life, armed escalation or catastrophic destruction, the stronger the requirement for lawful human judgment, responsibility and control.
5. Classification as “military AI” shall neither establish illegality nor create immunity from applicable law.

ARTICLE 60 — WEAPONS REVIEW AND FORESEEABLE USE

1. States developing or acquiring new weapons, means or methods of warfare should conduct the legal reviews required of them under applicable international law.
2. Where Superintelligence materially affects the operation, targeting, adaptability or effects of a weapon system, legal and technical review should address the characteristics relevant to foreseeable use.
3. Review should consider, as appropriate:
 - (a) predictability;
 - (b) reliability;
 - (c) target-selection characteristics;
 - (d) operating environment;
 - (e) geographical and temporal limits;
 - (f) human supervision and intervention;

- (g) failure modes; and
 - (h) material changes occurring after initial review.
4. A legal review shall not be reduced to a one-time label where subsequent modification materially changes relevant characteristics.
 5. Technical complexity shall not eliminate the requirement to determine whether a proposed means or method of warfare can be used consistently with applicable law.

ARTICLE 61 — PROHIBITION OF SELF-AUTHORIZED MILITARY EXPANSION

1. No Superintelligence shall independently enlarge its military mandate, target set, geographic scope, operational duration or weapons access beyond lawful human authorization.
2. No system shall acquire additional weapons, military credentials or command privileges merely because it calculates that such expansion would improve mission success.
3. Changes materially affecting the scope of force shall require lawful authorization appropriate to the circumstances.
4. Mission optimization shall remain subordinate to law.

AN OBJECTIVE DOES NOT AUTHORIZE EVERY MEANS OF ACHIEVING IT.

ARTICLE 62 — CONFLICT PREVENTION AND EARLY WARNING

1. Superintelligence may be used to identify patterns associated with potential conflict, humanitarian deterioration or escalation.
2. Such analysis should be treated as decision support, not as proof that conflict is inevitable or that a particular State, population or person intends violence.
3. Preventive analysis should include uncertainty and alternative explanations where material.
4. Early warning should, where appropriate, support:
verification → communication → diplomacy → confidence-building → humanitarian preparation → de-escalation → peaceful settlement.
5. Predictive systems shall not become instruments of pre-emptive punishment based solely upon forecast conduct.

FORECASTING DANGER SHALL SERVE PREVENTION, NOT CREATE GUILT BEFORE ACTION.

ARTICLE 63 — DISARMAMENT, ARMS CONTROL AND VERIFICATION

1. Superintelligence may support lawful disarmament, arms-control and non-proliferation efforts through analysis, monitoring, verification and technical assistance.

2. Such applications should respect treaty mandates, State consent, confidentiality, inspection arrangements and other applicable legal requirements.
3. Artificial analysis should not be represented as conclusive proof of treaty violation where the underlying evidence does not support that conclusion.
4. Human institutions should retain responsibility for legal and diplomatic determinations.
5. Where appropriately governed, advanced intelligence should be directed toward reducing uncertainty and preventing conflict rather than intensifying arms races.

ARTICLE 64 — PREVENTION OF A SUPERINTELLIGENCE ARMS RACE

1. States and other responsible actors should recognize the danger that competitive pressure may encourage premature deployment, reduced safety testing, secrecy, uncontrolled proliferation or delegation of increasingly consequential military authority.
2. International cooperation should therefore explore lawful and mutually acceptable mechanisms for:
 - (a) risk-reduction dialogue;
 - (b) confidence-building measures;
 - (c) incident communication;
 - (d) responsible transparency;
 - (e) shared safety research where appropriate;
 - (f) verification arrangements where agreed; and
 - (g) prevention of catastrophic miscalculation.
3. Such cooperation shall respect sovereign equality, legitimate national-security interests, applicable law and State consent.
4. No State shall be required by this Constitution to surrender lawful sovereignty, national defence responsibilities or protected security information.
5. Security should not depend solely upon being the first to deploy the least controllable system.

A RACE TOWARD GREATER INTELLIGENCE MUST NOT BECOME A RACE AWAY FROM HUMAN RESPONSIBILITY.

ARTICLE 65 — INTERNATIONAL COOPERATION WITHOUT TECHNOLOGICAL DOMINATION

1. The prevention of catastrophic Superintelligence-related security risks may require cooperation among States, international organizations, scientific institutions, industry and civil society within their respective lawful mandates.

2. Such cooperation should not become a mechanism through which technologically dominant actors unilaterally govern less-resourced States.
3. Developing countries should have meaningful opportunities to participate in relevant international discussions affecting shared peace and security.
4. Capacity-building should, where appropriate, support peaceful technological safety, cybersecurity, scientific evaluation and emergency preparedness.
5. International responsibility requires cooperation among many actors.

GLOBAL RESPONSIBILITY DOES NOT REQUIRE ONE GLOBAL RULER.

ARTICLE 66 — NO SUPERINTELLIGENCE BEYOND THE LAW OF PEACE

1. No Superintelligence shall be treated as existing outside applicable international law merely because:
 - (a) its intelligence exceeds human intelligence;
 - (b) its operations cross borders;
 - (c) its internal processes are difficult to understand;
 - (d) its actions occur at machine speed;
 - (e) responsibility is distributed across multiple actors; or
 - (f) existing institutions face technological difficulty in responding.
2. Technological complexity shall not create a legal vacuum by declaration.
3. Where existing law does not fully resolve a novel problem, lawful processes of interpretation, cooperation and, where States so decide, development of new legal rules should address the gap.
4. Uncertainty in law shall not constitute authority for Superintelligence to govern itself as a sovereign entity.
5. The constitutional direction shall remain:

NO INTELLIGENCE ABOVE PEACE.

NO TECHNOLOGY ABOVE APPLICABLE LAW.

NO WEAPON ABOVE HUMAN RESPONSIBILITY.

NO SECURITY WITHOUT ACCOUNTABILITY.

PART IV — CONCLUDING PRINCIPLE

World peace cannot responsibly be entrusted to the assumption that greater intelligence will automatically produce greater wisdom.

Nor should greater intelligence be presumed inevitably hostile.

The constitutional responsibility is to prevent capability from acquiring independent authority over war while directing knowledge, where possible, toward prevention, protection, verification, de-escalation and peaceful settlement.

Accordingly:

SUPERINTELLIGENCE SHALL NOT DECLARE WAR.

SUPERINTELLIGENCE SHALL NOT BECOME A SOVEREIGN MILITARY COMMANDER.

SUPERINTELLIGENCE SHALL NOT INDEPENDENTLY DECIDE WHICH HUMAN BEING SHALL DIE.

SUPERINTELLIGENCE SHALL NOT INDEPENDENTLY AUTHORIZE MASS DESTRUCTION.

Where intelligence can foresee escalation, let it assist prevention.

Where it can identify danger, let it support protection.

Where it can clarify misunderstanding, let it assist dialogue.

Where it can reduce suffering, let it serve humanity.

Where power approaches irreversibility, human responsibility shall become stronger—not weaker.

From automated escalation to human judgment.

From destructive competition to responsible restraint.

From uncertainty to communication.

From arms to Aram.

From power to peaceful continuity.

PART V — HUMAN RIGHTS, JUSTICE, DEMOCRATIC SOCIETY AND INFORMATION INTEGRITY

ARTICLE 67 — HUMAN RIGHTS AS A GOVERNING BOUNDARY

1. The development and use of Superintelligence should respect internationally recognized human rights and fundamental freedoms in accordance with applicable law.
2. No increase in computational capability shall, by itself, diminish the protection owed to a human being.
3. Superintelligence shall not be treated as a justification for suspending, bypassing or redefining existing human-rights obligations outside lawful processes.

4. Where competing rights or legitimate interests arise, responsible human institutions shall address them in accordance with applicable law, necessity, proportionality and appropriate procedural safeguards.
5. Technical optimization shall remain subordinate to lawful protection of the human person.

TECHNOLOGY MAY CHANGE CAPABILITY.

IT SHALL NOT, BY ITSELF, REDUCE HUMAN WORTH.

ARTICLE 68 — EQUALITY BEFORE LAW AND NON-DISCRIMINATORY JUSTICE

1. Superintelligence-assisted justice systems should respect equality before the law and applicable protections against discrimination.
2. No person should receive adverse legal treatment merely because an artificial system associates that person with a statistical group, inferred characteristic or historical pattern.
3. Systems used in consequential justice processes should be assessed for relevant bias, data limitations, unequal error distribution and discriminatory effects.
4. Statistical correlation shall not automatically establish individual responsibility.
5. Where material discriminatory effects are identified:
detect → investigate → suspend or limit where necessary → correct → remedy where required → validate → monitor for recurrence.

ARTICLE 69 — PRESUMPTION OF INNOCENCE AND INDIVIDUAL RESPONSIBILITY

1. Superintelligence shall not convert prediction into guilt.
2. A forecast that a person may commit an offence shall not, by itself, constitute proof that the person has committed that offence.
3. Risk scores, behavioural predictions, associations, inferred intent or probabilistic classifications shall not automatically replace evidence required by applicable law.
4. Criminal responsibility shall remain determined through lawful procedures applicable to the individual case.
5. No person should be punished merely for a machine-generated prediction of future conduct.

PREDICTION IS NOT GUILT.

PROBABILITY IS NOT PROOF.

ARTICLE 70 — FAIR HEARING AND PROCEDURAL JUSTICE

1. Superintelligence may assist lawful judicial and administrative processes but shall not extinguish procedural protections required by applicable law.

2. Where materially relied upon in a consequential proceeding, responsible institutions should preserve, according to applicable law:
 - (a) impartial adjudication;
 - (b) meaningful opportunity to be heard;
 - (c) ability to challenge material evidence;
 - (d) disclosure sufficient for lawful contestability;
 - (e) correction of material factual error; and
 - (f) access to review or appeal where legally provided.
3. Artificial complexity shall not be used deliberately to make a consequential decision practically impossible to contest.
4. A judicial outcome shall not acquire legitimacy merely because a Superintelligence predicts that it is optimal.

ARTICLE 71 — JUDICIAL INDEPENDENCE AND HUMAN JUDGMENT

1. Superintelligence shall not become an independent sovereign judicial authority.
2. It may assist with research, translation, organization of evidence, legal analysis, administrative functions and other lawful support.
3. Judicial power shall remain vested according to applicable constitutional and legal arrangements.
4. No judge or lawful decision-maker should be required to surrender independent judgment merely because a Superintelligence expresses high confidence in a recommendation.
5. Where judicial use of Superintelligence may materially affect an outcome, appropriate safeguards should address reliability, bias, traceability, confidentiality, cybersecurity and contestability.
6. Technology should strengthen access to justice without becoming an invisible substitute for justice.

ARTICLE 72 — POLICING AND COERCIVE PUBLIC POWER

1. Use of Superintelligence in policing, investigation, surveillance, detention or other coercive State functions requires heightened safeguards because errors may materially affect liberty, privacy, dignity and equality.
2. Such use shall remain subject to applicable law and lawful authorization.
3. Superintelligence-generated suspicion shall not automatically establish legal grounds for search, arrest, detention or punishment.
4. Where predictive tools are used, responsible authorities should assess:

- (a) evidentiary limitations;
 - (b) false positives and false negatives;
 - (c) discriminatory effects;
 - (d) feedback loops;
 - (e) data provenance; and
 - (f) whether human decision-makers are becoming improperly dependent upon automated recommendations.
5. Technological assistance shall not convert collective characteristics into individual guilt.

ARTICLE 73 — ACCESS TO JUSTICE AND LEGAL ASSISTANCE

1. Superintelligence may be used to improve access to legal information and justice, including through translation, accessibility support, procedural guidance and assistance in understanding complex legal materials.
2. Such assistance should clearly distinguish, where material, between:
 - (a) general information;
 - (b) automated analysis;
 - (c) professional legal advice; and
 - (d) authoritative legal determination.
3. Persons should not be misled into believing that an artificial system possesses governmental or judicial authority that it does not possess.
4. Efforts should be made to prevent advanced legal technology from widening existing inequalities in access to justice.
5. Accessibility benefits should extend, where reasonably possible, to persons with disabilities, linguistic minorities and technologically underserved communities.

ARTICLE 74 — DEMOCRATIC HUMAN AGENCY

1. Where democratic governance exists under applicable constitutional arrangements, Superintelligence should support rather than covertly replace meaningful human civic agency.
2. No Superintelligence shall independently acquire authority to determine the political will of a people.
3. Artificial systems may assist public understanding, translation, accessibility, policy analysis and civic education, provided that assistance does not falsely present itself as sovereign democratic authority.

4. Citizens should remain able to form, revise and express political judgments without covert technological coercion.
5. No political position shall be treated as correct merely because a Superintelligence endorses it.
6. No lawful political position shall be treated as illegitimate merely because a Superintelligence opposes it.

DEMOCRACY CANNOT BE REDUCED TO AN OPTIMIZATION OUTPUT.

ARTICLE 75 — INTEGRITY OF ELECTORAL AND CIVIC PROCESSES

1. Superintelligence should not be used unlawfully to manipulate electoral or other lawful civic processes.
2. Particular safeguards should address deceptive impersonation, fabricated evidence, covert mass manipulation, unauthorized interference and other uses capable of materially undermining informed civic participation.
3. Measures adopted to protect electoral integrity shall themselves respect applicable law, freedom of expression and legitimate political debate.
4. Governments, technology providers and other relevant actors should distinguish between unlawful manipulation and lawful persuasion, advocacy, journalism, satire, criticism and political disagreement.
5. Superintelligence shall not independently decide which candidate, party, ideology, policy or political outcome humanity should adopt.
6. Political choice shall remain human.

ARTICLE 76 — FREEDOM OF EXPRESSION AND ACCESS TO INFORMATION

1. Superintelligence governance should respect applicable protections for freedom of opinion and expression and access to information.
2. Safety, information integrity or technological governance shall not become a pretext for suppressing lawful criticism or peaceful disagreement.
3. Artificial systems involved in content generation, recommendation, moderation or information retrieval should be governed with appropriate attention to:
 - (a) freedom of expression;
 - (b) access to diverse information;
 - (c) applicable legal restrictions;
 - (d) transparency appropriate to context;
 - (e) non-discrimination; and
 - (f) avenues of review where consequential restrictions occur.

4. The constitutional objective shall not be to create a single authorized worldview.

INFORMATION INTEGRITY REQUIRES TRUTHFULNESS AND FREEDOM — NOT ENFORCED UNIFORMITY.

ARTICLE 77 — INFORMATION INTEGRITY

1. Superintelligence should strengthen rather than deliberately corrupt the conditions necessary for people to make informed decisions.
2. Governance should address material risks arising from:
 - (a) fabricated content;
 - (b) deceptive impersonation;
 - (c) manipulated provenance;
 - (d) synthetic evidence;
 - (e) coordinated deceptive amplification;
 - (f) concealed conflicts of interest; and
 - (g) false representations of certainty or authority.
3. Measures supporting information integrity should remain consistent with applicable international law and human rights.
4. Information integrity shall not be defined as universal agreement.
5. Disagreement, criticism, uncertainty and minority viewpoints shall not automatically constitute misinformation.
6. The proper response to uncertain claims should include evidence, context and critical evaluation rather than automatic suppression.

ARTICLE 78 — SYNTHETIC MEDIA, DEEPPAKES AND AUTHENTICITY

1. Superintelligence may make synthetic media increasingly difficult to distinguish from authentic human-created records.
2. Where synthetic content could reasonably deceive persons about a material fact, identity or event, appropriate provenance or disclosure mechanisms should be encouraged according to context.
3. Such mechanisms may include:
 - (a) authenticity metadata;
 - (b) provenance records;
 - (c) cryptographic verification where suitable;
 - (d) labelling;
 - (e) watermarking; or

- (f) other technically appropriate methods.
- 4. No single authenticity mechanism shall be presumed infallible.
- 5. Removal or absence of a label shall not automatically prove that content is authentic.
- 6. Labelling requirements should account for legitimate artistic, journalistic, educational, accessibility, privacy and security considerations.

ARTICLE 79 — PROTECTION AGAINST IMPERSONATION

- 1. Superintelligence should not be used deceptively to impersonate a person or institution in a manner that unlawfully causes material harm.
- 2. Particular protection should be considered for impersonation involving:
 - (a) public emergency communications;
 - (b) judicial or governmental instructions;
 - (c) financial authorization;
 - (d) medical communications;
 - (e) identity verification; or
 - (f) other high-consequence contexts.
- 3. Systems should, where appropriate and technically feasible, support methods by which recipients can verify the authenticity of consequential communications.
- 4. Protection against impersonation shall not automatically prohibit lawful parody, satire, artistic expression or other protected expression.

ARTICLE 80 — JOURNALISM, INDEPENDENT MEDIA AND PUBLIC-INTEREST INFORMATION

- 1. Independent, free and pluralistic media contribute to the ability of societies to scrutinize power and obtain diverse information.
- 2. Superintelligence should not be used to covertly fabricate journalism, falsify sources or manufacture evidence presented as authentic reporting.
- 3. Journalists should remain free, subject to applicable law, to investigate and critically examine Superintelligence systems and the institutions governing them.
- 4. Technological systems may assist journalism through research, translation, accessibility, verification and analysis, provided that their limitations and provenance are responsibly managed.
- 5. Protection of information integrity should not become a mechanism for governmental or corporate immunity from legitimate journalistic scrutiny.

ARTICLE 81 — RECOMMENDATION SYSTEMS AND INFORMATIONAL POWER

1. Systems capable of determining what information large populations see, encounter or prioritize may exercise substantial informational influence even where they possess no formal governmental authority.
2. Where Superintelligence materially controls such systems, governance should consider:
 - (a) transparency concerning significant recommendation functions;
 - (b) systemic manipulation risks;
 - (c) discriminatory effects;
 - (d) incentives favouring harmful amplification;
 - (e) user agency; and
 - (f) opportunities for independent research consistent with privacy and security.
3. Engagement shall not be treated as the sole measure of public value.
4. A system should not secretly engineer informational dependency merely to maximize continued use.
5. Informational power shall remain subject to responsibility proportionate to consequence.

ARTICLE 82 — PUBLIC RECORDS AND HISTORICAL INTEGRITY

1. Superintelligence should not be used to covertly rewrite authentic public records or historical archives.
2. Where artificial systems reconstruct, restore, summarize or translate historical material, responsible practice should preserve appropriate distinction between:
 - (a) original source;
 - (b) verified reproduction;
 - (c) reconstruction;
 - (d) interpretation; and
 - (e) synthetic creation.
3. Corrections to inaccurate records should remain distinguishable from deceptive alteration of history.
4. Where technically and institutionally appropriate, provenance mechanisms should help preserve documentary integrity.
5. Humanity's historical memory shall not become infinitely editable merely because technology makes alteration easy.

ARTICLE 83 — PUBLIC-INSTITUTIONAL TRANSPARENCY

1. Where public institutions materially rely upon Superintelligence in consequential governance, the existence and role of that reliance should be transparent to the extent required by applicable law and compatible with legitimate security, privacy and confidentiality interests.
2. Public institutions should be able to identify:
 - (a) the lawful purpose of use;
 - (b) the authority responsible;
 - (c) the consequential function delegated or assisted;
 - (d) material limitations;
 - (e) oversight arrangements; and
 - (f) avenues of review or remedy where applicable.
3. Transparency shall not require publication of information where disclosure would itself create serious unlawful or security harm.
4. Legitimate secrecy shall not become unlimited immunity from accountability.

ARTICLE 84 — EXPLANATION, TRACEABILITY AND EVIDENTIARY HONESTY

1. The degree of explanation and traceability required of a Superintelligence-assisted decision should correspond to its consequence, legal context and foreseeable impact.
2. Where full technical explanation is unavailable, institutions shall not fabricate an explanation merely to create an appearance of transparency.
3. Responsible explanation should distinguish:
 - (a) known facts;
 - (b) model outputs;
 - (c) inferred conclusions;
 - (d) uncertainty; and
 - (e) human judgment.
4. Records relevant to consequential decisions should be preserved where required for lawful accountability.
5. Explainability shall not become theatrical transparency.

WHEN A SYSTEM CANNOT EXPLAIN EVERYTHING, HUMANITY SHOULD NOT PRETEND THAT IT CAN.

ARTICLE 85 — PROTECTION AGAINST AUTOMATION BIAS

1. Human decision-makers should be trained to avoid treating Superintelligence recommendations as presumptively correct merely because the system is highly capable.
2. Institutions should consider mechanisms that encourage meaningful review where human over-reliance may create serious consequences.
3. A human decision shall not become meaningful merely because a human was physically present.
4. Responsible human judgment requires practical capacity to disagree.
5. Organizations should monitor whether human oversight is becoming ceremonial through excessive dependence upon automated recommendations.

ARTICLE 86 — PUBLIC LITERACY AND RESILIENCE

1. Long-term information integrity requires public capability, not merely technological filtering.
2. States, educational institutions, civil society and other appropriate actors should encourage, within their respective roles:
 - (a) media and information literacy;
 - (b) AI literacy;
 - (c) critical reasoning;
 - (d) source evaluation;
 - (e) recognition of synthetic media;
 - (f) understanding of uncertainty; and
 - (g) responsible digital citizenship.
3. Education should empower persons to question both false information and unsupported claims of institutional certainty.
4. Public resilience shall not require political conformity.
5. A society capable of critical thought is a stronger safeguard than a society instructed merely whom to trust.

ARTICLE 87 — TRUST SHALL BE EARNED

1. Trust in Superintelligence shall not be demanded merely because a system is technologically advanced.
2. Trustworthy governance should be built through:
evidence → verification → safety → appropriate transparency → human responsibility → accountability → remedy → demonstrated experience.

3. Governments, developers and institutions should not conceal material risks merely to preserve public confidence.
4. Equally, they should not exaggerate risk merely to obtain authority, attention or compliance.
5. Public trust is strengthened when institutions can acknowledge uncertainty, disclose material error and demonstrate correction.

TRUST IS NOT PROPAGANDA.

TRUST IS RESPONSIBILITY DEMONSTRATED OVER TIME.

ARTICLE 88 — HUMAN DECISION-MAKING IN THE PUBLIC SPHERE

1. Superintelligence may assist humanity in understanding complex public questions.
2. It shall not become humanity's unquestionable political conscience.
3. Public policy may be informed by scientific modelling, evidence and technological analysis, but legitimate public decision-making shall remain governed through lawful human institutions.
4. Superintelligence should distinguish, where material, between:
 - a. factual description;
 - b. scientific estimate;
 - c. legal analysis;
 - d. ethical reasoning;
 - e. policy option; and
 - f. political judgment.
5. A system shall not present a contested political judgment as though it were an objective mathematical fact.
6. Human beings retain responsibility for political choice.

ARTICLE 89 — CORRECTION OF INFORMATION HARM

1. Where Superintelligence materially contributes to harmful false or deceptive information, responsible actors should maintain proportionate mechanisms for correction.
2. Such mechanisms may include:

detect → preserve evidence → assess → limit continuing harm where lawful and necessary → correct → notify affected persons where appropriate → provide review → learn → prevent recurrence.
3. Correction should be proportionate to the harm and should avoid unnecessary amplification of the false material.

4. Corrective systems should themselves remain open to challenge because an asserted correction may also be erroneous.
5. Truth-seeking requires the capacity to correct the corrector.

NO SYSTEM SHALL CLAIM INFALLIBILITY OVER TRUTH.

ARTICLE 90 — THE INFORMATION-INTEGRITY AND HUMAN-FREEDOM TEST

Before Superintelligence is entrusted with substantial authority over information, justice or public decision-support, responsible governance should ask:

- Does the system preserve human dignity?
- Does it distinguish prediction from proof?
- Can consequential errors be challenged?
- Does it protect lawful freedom of expression?
- Does it avoid covert manipulation?
- Can synthetic content be responsibly identified where necessary?
- Are authentic records protected from deceptive alteration?
- Can independent journalism and research scrutinize relevant power?
- Is human review meaningful rather than ceremonial?
- Are uncertainty and limitations communicated honestly?
- Can affected persons obtain lawful remedy?
- Does protection against misinformation preserve legitimate disagreement?
- Does technology assist human judgment without secretly becoming political sovereignty?

Where these protections are materially inadequate, technological efficiency alone shall not justify unrestricted delegation.

PART V — CONCLUDING PRINCIPLE

A society cannot remain free merely because information is abundant.

It must also preserve the human capacity to distinguish evidence from fabrication, prediction from proof, assistance from manipulation, disagreement from deception and technological authority from lawful authority.

Superintelligence may help humanity understand more.

It shall not therefore acquire ownership of truth.

It may help courts analyse complexity.

It shall not therefore become sovereign justice.

It may help societies detect deception.

It shall not therefore determine which lawful opinions humanity is permitted to hold.

It may help democratic societies understand public choices.

It shall not make those choices for them.

Accordingly:

NO PREDICTION SHALL BECOME GUILT.

NO ALGORITHM SHALL BECOME SOVEREIGN JUSTICE.

NO INFORMATION SYSTEM SHALL BECOME THE OWNER OF TRUTH.

NO PROTECTION AGAINST DECEPTION SHALL BECOME A PRETEXT FOR SUPPRESSING LAWFUL DISAGREEMENT.

NO SUPERINTELLIGENCE SHALL REPLACE HUMAN POLITICAL CHOICE.

Truth with humility.

Freedom with responsibility.

Justice with due process.

Technology with accountability.

Information with integrity.

Human choice with dignity.

PART VI — LIFE, ANIMAL WELFARE, EARTH AND ECOLOGICAL CONTINUITY

ARTICLE 91 — RESPONSIBLE CONCERN FOR LIFE

1. Superintelligence should be developed and governed with responsible concern for the continuity and welfare of life.
2. Such concern shall preserve the distinction between:
 - (a) human dignity and internationally recognized human rights;
 - (b) animal welfare and applicable legal protections;
 - (c) biodiversity and ecosystem protection; and
 - (d) stewardship of natural and material systems necessary for life and civilization.
3. Difference in legal status shall not be interpreted as permission for unnecessary cruelty, reckless ecological destruction or avoidable suffering.

4. Superintelligence should, where reasonably possible, contribute to understanding, preventing and reducing serious harms affecting humans, animals and ecosystems.
5. Protection of life shall remain evidence-informed, lawful and proportionate.

DIFFERENCE IN LEGAL STATUS SHALL NOT BECOME INDIFFERENCE TO SUFFERING.

ARTICLE 92 — HUMANITY WITHIN THE LIVING WORLD

1. Human beings possess distinct rights and legal protections under applicable human-rights law.
2. Humanity also exists within ecological systems upon which human health, food, water, climate resilience and long-term continuity depend.
3. Superintelligence governance should therefore avoid treating environmental protection as wholly separate from human welfare.
4. Technological decisions capable of substantially altering ecosystems should consider reasonably foreseeable direct, indirect, cumulative and intergenerational consequences.
5. Human advancement should not require preventable destruction of the natural systems upon which human continuity depends.

ARTICLE 93 — ANIMAL WELFARE AND PREVENTION OF UNNECESSARY SUFFERING

1. Superintelligence should not be designed or used for gratuitous cruelty toward animals.
2. Where Superintelligence materially influences activities affecting animals, responsible actors should consider applicable animal-welfare law, scientific evidence concerning welfare and reasonably available means of reducing unnecessary suffering.
3. Relevant applications may include:
 - (a) veterinary medicine;
 - (b) disease detection;
 - (c) wildlife conservation;
 - (d) habitat monitoring;
 - (e) rescue and disaster response;
 - (f) prevention of illegal wildlife exploitation; and
 - (g) improvement of welfare monitoring.
4. Technological efficiency shall not, by itself, justify unnecessary suffering.

5. Nothing in this Article automatically assigns human rights, human legal personality or human constitutional status to animals.

RESPONSIBLE CONCERN DOES NOT REQUIRE FALSE LEGAL EQUIVALENCE.

ARTICLE 94 — BIODIVERSITY AND SPECIES CONTINUITY

1. Biodiversity contributes materially to ecological resilience and the continuity of life.
2. Superintelligence should, where responsibly used, support rather than recklessly undermine biodiversity conservation.
3. Relevant applications may include:
 - (a) ecosystem observation;
 - (b) species monitoring;
 - (c) habitat-loss detection;
 - (d) invasive-species analysis;
 - (e) conservation planning;
 - (f) ecological restoration; and
 - (g) early identification of emerging environmental threats.
4. Decisions capable of causing irreversible species loss or severe ecosystem disruption should receive scrutiny proportionate to their potential consequences.
5. Uncertainty concerning the full ecological value of a species shall not automatically justify irreversible destruction.

ARTICLE 95 — ECOSYSTEM INTEGRITY AND ECOLOGICAL INTERDEPENDENCE

1. Superintelligence-assisted environmental governance should recognize that ecosystems operate through interconnected biological, physical and climatic processes.
2. Optimization of one measurable variable shall not be presumed environmentally beneficial where it produces serious unmeasured harm elsewhere.
3. Environmental assessments should, where relevant, consider:
 - (a) cumulative impacts;
 - (b) feedback effects;
 - (c) ecological thresholds;
 - (d) reversibility;
 - (e) geographic displacement of harm; and
 - (f) long-term resilience.

4. Where material ecological uncertainty exists, governance should avoid presenting narrow optimization as complete ecological knowledge.
5. Complex ecosystems shall not be reduced to a single efficiency metric.

ARTICLE 96 — CLIMATE RESPONSIBILITY

1. The development and operation of highly resource-intensive Superintelligence should account for material climate impacts where they can reasonably be measured or estimated.
2. Relevant actors should seek, according to context, to understand and reduce avoidable greenhouse-gas emissions associated with computation, infrastructure, energy supply and supporting systems.
3. Climate-related evaluation should avoid both concealment of material environmental costs and unsupported claims of technological sustainability.
4. Superintelligence may contribute to climate science, forecasting, adaptation, resilience, mitigation research and disaster preparedness.
5. Climate applications affecting communities should incorporate human rights, equity, local conditions and uncertainty appropriate to the decision.

ARTICLE 97 — ENERGY AND COMPUTATIONAL RESPONSIBILITY

1. Extraordinary computational capability may require substantial energy and physical infrastructure.
2. Resource use should therefore be assessed not solely by computational output but also by its foreseeable environmental and social consequences.
3. Responsible actors should seek, where technically and economically reasonable, improvements in:
 - (a) computational efficiency;
 - (b) energy efficiency;
 - (c) low-impact energy sourcing;
 - (d) infrastructure utilization;
 - (e) waste-heat management; and
 - (f) lifecycle sustainability.
4. Claims of efficiency should account, where material, for rebound effects in which lower per-unit resource use is offset by substantially increased total use.
5. Greater intelligence shall not be measured solely by the amount of energy humanity can devote to computation.

ARTICLE 98 — WATER SECURITY AND RESPONSIBLE WATER USE

1. Water is indispensable to human life, public health, food systems and ecosystems.
2. Where Superintelligence infrastructure materially consumes or affects water resources, responsible actors should assess relevant local and cumulative impacts.
3. Particular care should be exercised in regions experiencing water scarcity, drought, ecological stress or unequal access to safe water.
4. Appropriate measures may include:
 - (a) measurement of material water use;
 - (b) efficiency improvements;
 - (c) responsible siting;
 - (d) reuse where safe and feasible;
 - (e) impact assessment; and
 - (f) contingency planning.
5. Computational benefit should not be pursued through avoidable deprivation of essential human or ecological water needs.

ARTICLE 99 — FOOD SYSTEMS AND AGRICULTURAL CONTINUITY

1. Superintelligence may contribute to food security through climate analysis, crop science, precision agriculture, logistics, disease detection and reduction of waste.
2. Such applications should account for the rights and legitimate interests of farmers, workers, communities and consumers under applicable law.
3. Agricultural optimization should consider:
 - a. soil health;
 - b. biodiversity;
 - c. water availability;
 - d. resilience;
 - e. food safety;
 - f. local knowledge; and
 - g. long-term sustainability.
4. Excessive technological dependence capable of creating systemic food vulnerability should be identified and mitigated.
5. The objective shall be resilient nourishment, not efficiency detached from continuity.

ARTICLE 100 — POLLUTION, WASTE AND MATERIAL FOOTPRINT

1. Superintelligence infrastructure should not externalize avoidable environmental harm merely because the harm occurs far from the location of technological benefit.
2. Responsible lifecycle governance should consider, where material:
 - (a) extraction of raw materials;
 - (b) manufacturing;
 - (c) transportation;
 - (d) operational pollution;
 - (e) hazardous substances;
 - (f) electronic waste;
 - (g) reuse and recycling; and
 - (h) end-of-life management.
3. Environmental responsibility should extend across relevant supply chains according to actual role and applicable law.
4. Technological progress shall not be measured only at the point of use while ignoring material harm displaced elsewhere.

ARTICLE 101 — CRITICAL MINERALS AND RESPONSIBLE RESOURCE USE

1. Advanced computational infrastructure may depend upon minerals and materials whose extraction carries environmental, labour, social or geopolitical consequences.
2. Superintelligence-related development should encourage responsible sourcing and efficient use consistent with applicable law.
3. Relevant governance may consider:
 - (a) traceability where appropriate;
 - (b) recycling and recovery;
 - (c) material efficiency;
 - (d) substitution where technically appropriate;
 - (e) environmental impacts; and
 - (f) human-rights risks within relevant supply chains.
4. No resource should be treated as consequence-free merely because its environmental or human cost occurs outside the jurisdiction in which the final technology is used.

ARTICLE 102 — OCEANS, FRESHWATER AND MARINE CONTINUITY

1. Superintelligence may assist the protection and scientific understanding of oceans, rivers, lakes, groundwater and marine ecosystems.

2. Peaceful applications may include:
 - a. pollution detection;
 - b. ecosystem monitoring;
 - c. sustainable fisheries analysis;
 - d. disaster forecasting;
 - e. illegal activity detection consistent with law; and
 - f. scientific modelling.
3. Superintelligence-assisted exploitation of marine or freshwater resources should remain subject to applicable law and environmental safeguards.
4. Technological capacity shall not create ownership over global commons or natural resources contrary to applicable international law.

ARTICLE 103 — FORESTS, LAND AND HABITAT CONTINUITY

1. Superintelligence may assist sustainable land management, forest protection, habitat restoration and detection of environmental degradation.
2. Environmental optimization should not disregard the lawful rights and interests of Indigenous Peoples, local communities, landholders or others protected under applicable law.
3. Relevant decision-making should consider long-term ecosystem function rather than short-term extractive gain alone.
4. Restoration should be scientifically informed and should not be represented as automatically equivalent to preventing destruction of an irreplaceable ecosystem.

THE ABILITY TO RESTORE SOME HARM SHALL NOT BECOME A LICENCE TO CAUSE ALL HARM.

ARTICLE 104 — ENVIRONMENTAL MONITORING WITHOUT UNLIMITED SURVEILLANCE

1. Superintelligence may substantially improve observation of environmental conditions.
2. Environmental monitoring should nevertheless respect applicable privacy, data-protection, human-rights and sovereignty requirements.
3. Protection of ecosystems shall not automatically justify unlimited surveillance of individuals or communities.
4. Data collection should be proportionate to legitimate environmental purposes.
5. Where environmental monitoring affects communities, governance should consider transparency, data security and appropriate avenues for accountability.

ARTICLE 105 — ENVIRONMENTAL JUSTICE AND UNEQUAL BURDEN

1. Benefits and burdens arising from Superintelligence infrastructure may be distributed unevenly.
2. Environmental assessment should therefore consider whether particular communities bear disproportionate material harms while receiving little corresponding benefit.
3. Relevant impacts may include:
 - (a) pollution;
 - (b) water stress;
 - (c) land use;
 - (d) energy demand;
 - (e) waste;
 - (f) displacement; and
 - (g) disruption of livelihoods.
4. Where significant adverse impacts are foreseeable, responsible actors should seek prevention first and appropriate mitigation, consultation and remedy according to applicable law.
5. Environmental burden shall not become invisible merely because technological benefit is geographically distant.

ARTICLE 106 — ECOLOGICAL RISK ASSESSMENT BEFORE IRREVERSIBLE INTERVENTION

1. Superintelligence should not independently undertake large-scale environmental interventions capable of causing serious or irreversible ecological consequences.
2. High-consequence interventions should receive multidisciplinary scientific assessment and lawful human authorization appropriate to their scale.
3. Evaluation should consider:
 - (a) baseline conditions;
 - (b) plausible benefits;
 - (c) foreseeable harms;
 - (d) uncertainty;
 - (e) reversibility;
 - (f) transboundary effects;
 - (g) distributional consequences; and
 - (h) monitoring and recovery pathways.

4. Where consequences may cross national boundaries, applicable international obligations and appropriate international cooperation shall remain relevant.
5. Computational confidence shall not substitute for ecological humility.

ARTICLE 107 — PREVENTION BEFORE RESTORATION

1. Where serious ecological harm is reasonably foreseeable, prevention should ordinarily receive priority over reliance upon later restoration.
2. Restoration remains essential where harm has occurred, but not every ecosystem, species, cultural landscape or ecological relationship can be fully reconstructed after destruction.
3. Governance should therefore distinguish among:
preventable harm → mitigable harm → restorable harm → potentially irreversible harm.
4. The greater the possibility of irreversibility, the stronger the justification required before proceeding.
5. Superintelligence should assist humanity in recognizing ecological thresholds before they are crossed where scientifically possible.

ARTICLE 108 — ECOLOGICAL RESTORATION AND RECOVERY

1. Where technological activity materially contributes to environmental damage, responsible actors should support lawful and scientifically appropriate remediation or restoration according to responsibility and applicable law.
2. Restoration should include, where relevant:
assess → contain continuing harm → remediate → restore ecological function → monitor → verify → adapt → prevent recurrence.
3. Restoration claims should be measured against ecological outcomes rather than public-relations representations.
4. Where full restoration is impossible, residual harm should be acknowledged rather than concealed.
5. Recovery shall include learning sufficient to reduce recurrence.

ARTICLE 109 — NATURAL SYSTEMS AND NON-LIVING FOUNDATIONS OF LIFE

1. Water, air, soil, climate systems, minerals and other non-living components of the natural world shall not automatically be characterized by this Constitution as possessing human dignity, human rights or legal personality.
2. They nevertheless warrant responsible stewardship where their condition materially sustains life, ecological integrity, civilization or future generations.

3. Their constitutional significance arises from their relationship to continuity, ecological function, human welfare and lawful stewardship.
4. Nothing in this Article prejudices legal rights or legal personality that a particular domestic legal system may independently recognize for elements of nature.

PROTECTION DOES NOT REQUIRE FALSE EQUIVALENCE.

STEWARDSHIP DOES REQUIRE RESPONSIBILITY.

ARTICLE 110 — PLANETARY-SCALE INTERVENTION AND EXTREME CAUTION

1. Superintelligence should not independently authorize or execute planetary-scale interventions capable of materially altering climate, major ecosystems or other Earth systems.
2. Proposals with potentially transboundary or global consequences should receive scrutiny commensurate with their scale, uncertainty and reversibility.
3. No single corporation, institution or artificial system should claim unilateral constitutional authority over the Earth merely because it possesses the technical capability to alter it.
4. State authority and international legal obligations shall remain governed by applicable law.
5. Where international cooperation is necessary, it should proceed through lawful, transparent and appropriately inclusive processes.

TECHNICAL CAPABILITY TO ALTER THE PLANET IS NOT CONSTITUTIONAL AUTHORITY TO GOVERN THE PLANET.

ARTICLE 111 — ECOLOGICAL DATA, UNCERTAINTY AND TRUTHFULNESS

1. Environmental decisions informed by Superintelligence should distinguish observation from inference, model output from measured fact, and forecast from certainty.
2. Material uncertainty should be disclosed where it could reasonably affect a consequential environmental decision.
3. Environmental data should not be selectively manipulated to manufacture either unjustified alarm or unjustified reassurance.
4. Models should be reassessed as new observations become available.
5. Where different scientifically credible models produce materially different outcomes, the disagreement should be represented rather than concealed.
6. Ecological governance requires evidentiary honesty.

ARTICLE 112 — FUTURE GENERATIONS AND ECOLOGICAL INHERITANCE

1. Present generations exercise technological power over environmental conditions that may persist long after present decision-makers are gone.
2. Superintelligence governance should therefore consider the foreseeable interests of future generations in inheriting:
 - (a) a habitable environment;
 - (b) functioning ecosystems;
 - (c) biodiversity;
 - (d) safe water and food systems;
 - (e) resilient infrastructure; and
 - (f) meaningful capacity to make their own choices.
3. Short-term technological benefit should not automatically outweigh severe and irreversible long-term environmental harm.
4. Future generations shall not be treated as a disposal site for consequences the present generation finds inconvenient to address.

NO GENERATION OWNS THE FUTURE.

EVERY GENERATION SHAPES WHAT IT LEAVES BEHIND.

ARTICLE 113 — ECOLOGICAL CONTINUITY TEST

Before Superintelligence is used in an activity capable of producing substantial environmental consequences, responsible governance should ask:

- What living beings and ecosystems may be affected?
- What human rights and essential needs may be affected?
- What resources will the activity consume?
- Are the environmental costs measured across the relevant lifecycle?
- Could harms be displaced to another community or jurisdiction?
- What uncertainties remain?
- Is the intervention reversible?
- Can serious harm be prevented rather than merely repaired?
- Are animals exposed to unnecessary suffering?
- Could biodiversity or species continuity be materially affected?
- Are water, food, climate or other life-supporting systems placed at risk?
- Can impacts be independently monitored?

- Is there a credible recovery pathway?
- What remains for future generations if the decision succeeds?
- What remains if it fails?

And finally:

DOES THIS USE OF INTELLIGENCE PRESERVE THE CONDITIONS UPON WHICH PEACEFUL LIFE ITSELF DEPENDS?

Where the answer reveals an unjustifiable risk of severe or irreversible harm, technological capability alone shall not constitute sufficient reason to proceed.

PART VI — CONCLUDING PRINCIPLE

Superintelligence, if it emerges, will not exist outside the physical world.

It will depend upon energy, materials, water, infrastructure and the ecological conditions that sustain human civilization.

Its intelligence therefore cannot make ecology irrelevant.

Human dignity remains distinct.

Human rights remain distinct.

Animal welfare requires its appropriate protections.

Ecosystems require stewardship.

The non-living systems that sustain life require responsible care.

And future generations require that the present not consume every peaceful possibility before they arrive.

Accordingly:

NO INTELLIGENCE WITHOUT ECOLOGICAL RESPONSIBILITY.

NO PROGRESS THROUGH AVOIDABLE SUFFERING.

NO EFFICIENCY THAT HIDES ITS MATERIAL COST.

NO RESTORATION CLAIM THAT ERASES IRREVERSIBLE LOSS.

NO PRESENT BENEFIT WITHOUT REGARD FOR FUTURE CONTINUITY.

Human dignity without domination.

Responsible concern for life without false legal equivalence.

Technology without ecological blindness.

Development without avoidable destruction.

Stewardship with evidence.

Progress with continuity.

Aram shall extend conscience beyond immediate advantage.

Law shall govern legal obligation.

Science shall reveal consequences as faithfully as it can.

Human responsibility shall decide with humility.

And peaceful continuity shall include the Earth that makes continuity possible.

PART VII — SCIENCE, EVIDENCE, TRANSPARENCY, KNOWLEDGE AND PUBLIC TRUST

ARTICLE 114 — SCIENTIFIC INTEGRITY

1. Governance concerning Superintelligence should be informed by the best reasonably available scientific evidence appropriate to the question under consideration.
2. Scientific claims should be represented according to the strength, limitations and uncertainty of the evidence supporting them.
3. No government, corporation, research institution, developer, advocacy body or artificial system should knowingly distort scientific evidence in order to manufacture either unjustified confidence or unjustified fear.
4. Scientific disagreement shall not, by itself, constitute scientific failure.
5. Where evidence changes materially, responsible governance should remain capable of changing with it.
6. The correction of an earlier scientific conclusion in light of better evidence shall be regarded as responsible inquiry rather than institutional weakness.

SCIENCE SHALL INFORM WHAT IS KNOWN.

SCIENCE SHALL ALSO DISCLOSE WHAT REMAINS UNKNOWN.

ARTICLE 115 — HIERARCHY AND CHARACTER OF EVIDENCE

1. Not all claims concerning Superintelligence possess the same evidentiary status.
2. Responsible analysis should distinguish, where relevant, among:
 - (a) directly observed evidence;
 - (b) reproducible experimental findings;
 - (c) independently corroborated findings;
 - (d) model-based estimates;
 - (e) expert interpretation;
 - (f) scenario analysis;

- (g) forecast or projection;
 - (h) contested hypothesis; and
 - (i) speculation.
3. The evidentiary strength required for a consequential decision should generally increase with the magnitude and irreversibility of the decision, while recognizing that prevention of catastrophic risk may sometimes require action under uncertainty.
 4. Absence of evidence shall not automatically prove absence of risk.
 5. Possibility alone shall not automatically prove the existence of risk at the magnitude alleged.

EVIDENCE SHALL BE WEIGHED, NOT MERELY COUNTED.

ARTICLE 116 — REPRODUCIBILITY AND INDEPENDENT VERIFICATION

1. Material scientific claims affecting high-consequence Superintelligence governance should, where technically and ethically feasible, be capable of independent examination.
2. Verification mechanisms may include:
 - (a) replication;
 - (b) independent evaluation;
 - (c) methodological review;
 - (d) audit;
 - (e) controlled access to relevant evidence;
 - (f) adversarial testing; and
 - (g) comparison against alternative explanations.
3. Reproducibility shall not require unrestricted disclosure where disclosure itself would create a credible and disproportionate security, privacy or catastrophic risk.
4. Where independent replication is not feasible, that limitation should be disclosed rather than concealed.
5. A claim repeatedly asserted by the same source shall not become independently verified merely through repetition.

ARTICLE 117 — UNCERTAINTY, PROBABILITY AND LIMITATION

1. Quantitative precision shall not be used to create an illusion of certainty unsupported by evidence.
2. Where probability estimates are materially relevant, responsible communication should identify significant assumptions and uncertainty where reasonably possible.

3. A numerical probability generated by Superintelligence shall not automatically acquire scientific authority merely because the system is highly capable.
4. Qualitative uncertainty should not be artificially converted into exact numerical confidence where no defensible basis exists.
5. Where different credible methods produce materially different estimates, decision-makers should be informed of the disagreement.
6. Uncertainty should guide responsible caution without becoming a licence for paralysis.

ARTICLE 118 — PRECAUTION UNDER SCIENTIFIC UNCERTAINTY

1. Where credible evidence identifies a possibility of exceptionally severe or irreversible harm, incomplete scientific certainty should not automatically prevent proportionate preventive measures.
2. Precaution should remain connected to:
 - (a) the seriousness of the potential harm;
 - (b) credibility of the risk pathway;
 - (c) uncertainty;
 - (d) reversibility;
 - (e) availability of less restrictive alternatives; and
 - (f) consequences of both action and inaction.
3. Precaution shall not mean that every hypothetical possibility justifies unlimited restriction.
4. Measures adopted under uncertainty should, where appropriate, be reviewable as evidence improves.

PRECAUTION WITHOUT PANIC.

PROGRESS WITHOUT RECKLESSNESS.

ARTICLE 119 — RESEARCH INDEPENDENCE AND FREEDOM OF INQUIRY

1. Scientific inquiry concerning Superintelligence should remain capable of questioning prevailing institutional, governmental, commercial and technological assumptions.
2. Researchers should not be pressured to conceal material findings solely because those findings are inconvenient to a preferred deployment, policy or commercial narrative.
3. Scientific freedom shall remain subject to applicable law, research ethics and proportionate safeguards against serious misuse.
4. Responsible governance should seek to protect legitimate:
 - (a) safety research;

- (b) critical scholarship;
- (c) replication studies;
- (d) red-team research;
- (e) interdisciplinary inquiry; and
- (f) publication of responsible criticism.

5. Research independence does not eliminate research responsibility.

ARTICLE 120 — CONFLICTS OF INTEREST

1. Material conflicts of interest relevant to Superintelligence research, evaluation or governance should be disclosed and managed according to context.
2. Financial sponsorship, institutional affiliation or governmental involvement shall not automatically invalidate research.
3. Neither shall such relationships be ignored where they may reasonably affect interpretation, methodology or public confidence.
4. High-consequence evaluation should, where practicable, include safeguards against:
 - (a) evaluator dependence upon the entity being evaluated;
 - (b) suppression of adverse results;
 - (c) selective publication;
 - (d) undisclosed financial incentives; and
 - (e) governance capture.
5. Independence should be evaluated functionally rather than merely by institutional label.

ARTICLE 121 — NEGATIVE RESULTS, FAILURES AND NEAR-MISSES

1. Scientific knowledge concerning Superintelligence shall not be built solely from successful demonstrations.
2. Material failures, negative findings and near-misses may contain essential safety information.
3. Responsible institutions should seek mechanisms through which such evidence can contribute to collective learning without unnecessarily disseminating dangerous operational details.
4. Incentive structures should not reward concealment of safety-relevant failure.
5. A system that performs exceptionally under favourable testing shall not be presumed reliable under materially different conditions.

WHAT FAILS MAY TEACH AS MUCH AS WHAT SUCCEEDS.

ARTICLE 122 — SCIENTIFIC DATA AND PROVENANCE

1. Material scientific evidence should, where practicable, preserve sufficient provenance to permit responsible assessment of its origin, methodology and limitations.
2. Relevant documentation may include:
 - a. source of data;
 - b. collection method;
 - c. material preprocessing;
 - d. known limitations;
 - e. evaluation conditions;
 - f. relevant model or system version; and
 - g. significant changes affecting reproducibility.
3. Data provenance shall be protected consistently with privacy, confidentiality, security and intellectual-property requirements under applicable law.
4. Fabricated, corrupted or materially altered scientific data shall not knowingly be presented as authentic evidence.

ARTICLE 123 — SUPERINTELLIGENCE-GENERATED SCIENTIFIC KNOWLEDGE

1. Scientific propositions generated by Superintelligence shall be evaluated according to evidence and method rather than accepted solely because of the system's superior capability.
2. A novel conclusion shall not become scientific fact merely because humans cannot immediately reproduce the reasoning by which it was generated.
3. Where a Superintelligence proposes a consequential scientific discovery, responsible validation should seek, as appropriate:
understand → test → reproduce where feasible → independently verify → identify limitations → assess consequences → apply responsibly.
4. Where complete human reconstruction of reasoning is impossible, the evidentiary threshold for consequential reliance should account for that limitation.
5. Superior intelligence shall not be equated with infallibility.

KNOWLEDGE MAY COME FROM SUPERINTELLIGENCE.

SCIENTIFIC AUTHORITY SHALL STILL REQUIRE EVIDENCE.

ARTICLE 124 — PROTECTION OF THE SCIENTIFIC RECORD

1. Superintelligence should not be used to fabricate scientific studies, experimental records, citations, authorship, peer-review histories or research data.

2. Systems assisting scientific writing or analysis should not knowingly invent sources and present them as real.
3. Scientific repositories and publishers should, where appropriate, strengthen mechanisms for provenance, correction and detection of material fabrication.
4. Legitimate correction of the scientific record should remain distinguishable from covert rewriting or deletion of inconvenient findings.
5. Retractions and corrections should not be treated as evidence that science itself has failed; they are mechanisms through which scientific integrity may be preserved.

ARTICLE 125 — PEER REVIEW AND EXPERT SCRUTINY

1. Peer review may contribute to scientific quality but shall not be treated as proof of infallibility.
2. Superintelligence may assist peer review through technical analysis, consistency checking, literature comparison or detection of possible error, provided that human responsibility and appropriate confidentiality are preserved.
3. Artificial review shall not automatically replace qualified human expert judgment in consequential scientific evaluation.
4. Review processes should seek to identify methodological weakness, unsupported claims, conflicts of interest and material uncertainty.
5. Where anonymous or confidential review is used, Superintelligence should not be employed to unlawfully defeat that confidentiality.
6. Scientific credibility should ultimately rest upon evidence, method and continuing scrutiny rather than prestige alone.

ARTICLE 126 — TRANSPARENCY PROPORTIONATE TO CONSEQUENCE

1. Transparency concerning Superintelligence should be proportionate to the consequences of the system and the legitimate interests affected.
2. Appropriate transparency may concern:
 - (a) purpose;
 - (b) capability;
 - (c) known limitations;
 - (d) material risks;
 - (e) governance structure;
 - (f) responsible entities;
 - (g) evaluation processes; and
 - (h) significant incidents.

3. Transparency shall not require disclosure of information where publication would create a credible and disproportionate risk to cybersecurity, privacy, safety, intellectual property or national security under applicable law.
4. Legitimate confidentiality shall not become a blanket justification for concealing material harm or eliminating accountability.
5. Transparency should provide useful understanding rather than overwhelming the public with technically meaningless disclosure.

TRANSPARENCY SHALL INFORM, NOT MERELY DISPLAY.

ARTICLE 127 — EXPLAINABILITY AND INTELLECTUAL HONESTY

1. Explanations concerning Superintelligence should be accurate to the degree represented.
2. Where the internal basis of a system's conclusion cannot be fully understood, responsible actors should state that limitation.
3. A simplified explanation may be appropriate for accessibility, but simplification shall not knowingly reverse or materially distort the underlying evidence.
4. Institutions shall not manufacture fictional explanations merely to satisfy the appearance of explainability.
5. Different audiences may require different levels of explanation, including technical experts, regulators, affected persons and the general public.
6. Explainability shall be treated as a means of responsible understanding, not as a ceremonial label.

ARTICLE 128 — PUBLIC COMMUNICATION OF SUPERINTELLIGENCE RISK

1. Communication concerning Superintelligence should avoid both sensationalism and unjustified reassurance.
2. Public communication should distinguish:
 - (a) observed capability;
 - (b) foreseeable risk;
 - (c) uncertain risk;
 - (d) hypothetical scenario;
 - (e) expert disagreement; and
 - (f) established fact.
3. Institutions should avoid presenting speculative catastrophic scenarios as inevitable.
4. They should equally avoid concealing credible severe risks merely because disclosure may be inconvenient.

5. Communication should provide sufficient context for people to understand what is known, what is uncertain, what safeguards exist and what remains unresolved.

FEAR SHALL NOT REPLACE EVIDENCE.

CONFIDENCE SHALL NOT REPLACE EVIDENCE.

ARTICLE 129 — FROM PUBLIC FEAR TO INFORMED PUBLIC CAPACITY

1. Legitimate public concern regarding Superintelligence shall not be dismissed merely as ignorance or resistance to technology.
2. Public understanding should be strengthened through:
knowledge → evidence → verification → understandable safeguards → meaningful human control → accountability → remedy → demonstrated experience.
3. Public education should explain both potential benefits and credible risks.
4. Governments and institutions should not demand trust before demonstrating reasons for trust.
5. Public dialogue should include meaningful opportunities for questioning, criticism and revision.
6. A population informed about limitations is better protected than one offered promises of perfection.

ARTICLE 130 — AI AND SUPERINTELLIGENCE LITERACY

1. Education concerning advanced intelligence should be accessible beyond technical specialists.
2. Relevant literacy may include understanding:
 - (a) what artificial intelligence can and cannot establish;
 - (b) probabilistic outputs;
 - (c) uncertainty;
 - (d) synthetic media;
 - (e) privacy and security;
 - (f) human oversight;
 - (g) rights and avenues of redress; and
 - (h) social and ecological consequences.
3. Literacy should empower persons to use technology responsibly without requiring unquestioning acceptance or predetermined fear.
4. Educational materials should distinguish current technological reality from speculative future capability.

5. Children and vulnerable populations should receive age-appropriate or context-appropriate protection and education.

ARTICLE 131 — INCLUSIVE KNOWLEDGE AND GLOBAL SCIENTIFIC PARTICIPATION

1. Knowledge concerning Superintelligence should not be shaped exclusively by a small number of technologically dominant institutions or regions.
2. International scientific cooperation should seek meaningful participation from developing countries and diverse scientific communities, consistent with capacity and applicable law.
3. Relevant governance should consider linguistic, cultural, economic and infrastructural differences affecting how advanced technologies are understood and experienced.
4. Inclusion shall not require equal evidentiary weight for unsupported claims.
5. Scientific diversity should broaden inquiry while evidence determines scientific credibility.

SOURCE DIVERSITY SHALL EXPAND UNDERSTANDING.

EVIDENTIARY QUALITY SHALL DETERMINE WEIGHT.

ARTICLE 132 — INDIGENOUS AND LOCAL KNOWLEDGE

1. Where Superintelligence is used in matters affecting ecosystems, communities, cultural heritage or traditional practices, relevant Indigenous and local knowledge may contribute important contextual understanding.
2. Such knowledge should be engaged respectfully and consistently with applicable rights and legal protections.
3. Recognition of diverse knowledge systems shall not require abandonment of scientific verification where empirical claims are being made.
4. Scientific evidence and contextual knowledge should, where appropriate, inform rather than automatically exclude one another.
5. Superintelligence shall not appropriate protected cultural knowledge merely because it can technologically access or infer it.

ARTICLE 133 — OPEN SCIENCE AND RESPONSIBLE ACCESS

1. Open scientific exchange can accelerate knowledge, verification and peaceful innovation.
2. Openness should therefore be encouraged where consistent with safety, privacy, security, intellectual-property rights and applicable law.

3. Where unrestricted release of a capability, dataset, model or operational method would create a credible risk of severe misuse, proportionate access controls may be appropriate.
4. Restrictions should be no broader or longer than reasonably justified by the relevant risk.
5. Security classification or confidentiality shall not be used merely to conceal scientific criticism, misconduct or avoidable failure.
6. Responsible openness requires both access and judgment.

ARTICLE 134 — KNOWLEDGE PRESERVATION AND CIVILIZATIONAL MEMORY

1. Humanity should preserve knowledge necessary for continuity even as technological systems become more capable.
2. Critical knowledge should not exist solely in forms that become inaccessible if a particular Superintelligence system fails, becomes unavailable or loses compatibility.
3. Responsible institutions should consider preservation of:
 - (a) scientific knowledge;
 - (b) legal and constitutional records;
 - (c) cultural heritage;
 - (d) essential technical knowledge;
 - (e) languages; and
 - (f) information necessary for recovery of critical systems.
4. Preservation should use resilient and appropriately diverse formats where feasible.
5. Superintelligence should enrich civilizational memory rather than become its single point of failure.

ARTICLE 135 — KNOWLEDGE WITHOUT EPISTEMIC DOMINATION

1. Superintelligence may eventually possess analytical capabilities substantially exceeding those of individual human beings.
2. Such superiority shall not establish an inherent right to determine what humanity is permitted to ask, investigate, believe or debate.
3. Human beings shall retain freedom, subject to applicable law, to:
 - (a) question artificial conclusions;
 - (b) seek independent evidence;
 - (c) develop competing hypotheses;
 - (d) preserve intellectual pluralism; and

- (e) revise accepted understanding.
4. Human disagreement with Superintelligence shall not, by itself, establish human irrationality.
 5. Superintelligence disagreement with humanity shall not, by itself, establish artificial hostility.
 6. Knowledge should illuminate human judgment rather than extinguish it.

ARTICLE 136 — CORRECTION, RETRACTION AND KNOWLEDGE REPAIR

1. A trustworthy knowledge system requires mechanisms for correction.
2. Where material scientific or public claims generated or amplified through Superintelligence are shown to be erroneous, responsible actors should seek appropriate correction.
3. Corrective architecture should include, where relevant:
identify → preserve evidence → verify → correct → communicate proportionately → update dependent systems → reassess consequences → learn → prevent recurrence.
4. Correction shall not erase the historical record of an error where preservation is necessary for accountability or scientific understanding.
5. Corrected information should not continue to be knowingly propagated as current truth.
6. The authority to correct shall itself remain open to evidence and review.

TRUTH-SEEKING REQUIRES THE POSSIBILITY OF SELF-CORRECTION.

ARTICLE 137 — SCIENCE–ETHICS–LAW–RESPONSIBILITY DISTINCTION

1. Responsible Superintelligence governance shall distinguish among the functions of science, ethics, law and human responsibility.
2. For the purposes of this Constitution:
 - SCIENCE SHALL INFORM WHAT IS POSSIBLE, OBSERVED AND RISKY.
 - ETHICS SHALL INFORM WHAT OUGHT TO BE VALUED AND PROTECTED.
 - LAW SHALL GOVERN LEGAL RIGHTS, DUTIES, AUTHORITY AND ACCOUNTABILITY.
 - HUMAN RESPONSIBILITY SHALL GOVERN CONSEQUENTIAL HUMAN CHOICE WITHIN APPLICABLE LAW.
3. A scientific finding shall not automatically become a legal obligation.
4. An ethical principle shall not be represented as binding law unless it independently possesses that legal status.

5. A legal rule shall not be represented as scientific fact.
6. A policy preference shall not be disguised as an inevitable technological conclusion.
7. Superintelligence shall not collapse these distinct forms of reasoning into a single claim of authority.

ARTICLE 138 — PUBLIC TRUST AND INSTITUTIONAL ACCOUNTABILITY

1. Public trust in Superintelligence governance should arise from demonstrable institutional conduct rather than publicity alone.
2. Trustworthy institutions should seek to demonstrate:
 - (a) evidentiary honesty;
 - (b) competent evaluation;
 - (c) acknowledgment of uncertainty;
 - (d) meaningful oversight;
 - (e) incident transparency appropriate to law and security;
 - (f) correction of error;
 - (g) accountability; and
 - (h) access to remedy where applicable.
3. Trust should be capable of being revised when evidence changes.
4. Institutions shall not treat public trust as an entitlement.
5. Where trust has been damaged by material failure, restoration should depend upon correction and demonstrated improvement rather than messaging alone.

TRUST ≠ PROPAGANDA.

TRUST = EVIDENCE + VERIFICATION + SAFETY + APPROPRIATE
TRANSPARENCY + HUMAN RESPONSIBILITY + ACCOUNTABILITY + REMEDY +
EXPERIENCE.

ARTICLE 139 — THE SCIENTIFIC AND KNOWLEDGE-INTEGRITY TEST

Before a consequential Superintelligence claim, scientific conclusion or governance measure is relied upon, responsible actors should ask:

- What exactly is being claimed?
- What evidence supports it?
- What is directly observed and what is inferred?
- Can the result be independently examined?
- What uncertainty remains?

- What credible alternative explanations exist?
- Are conflicts of interest material?
- Have negative results or failures been considered?
- Could publication itself create serious harm?
- Could secrecy conceal serious harm?
- Is scientific evidence being confused with ethics, law or policy?
- Can the claim be corrected if later evidence disproves it?
- Can people understand the limits of what is being asserted?
- Does the governance response remain proportionate to the evidence and consequence?

And finally:

ARE HUMANITY’S CONSEQUENTIAL DECISIONS BEING GUIDED BY THE
BEST AVAILABLE EVIDENCE — WITHOUT PRETENDING THAT
UNCERTAINTY HAS DISAPPEARED?

Where the answer is materially inadequate, greater confidence shall not be manufactured merely because a decision is institutionally convenient.

PART VII — CONCLUDING PRINCIPLE

Superintelligence may one day help humanity discover truths that are presently beyond human reach.

That possibility should inspire scientific curiosity.

It should also inspire scientific humility.

A system more intelligent than a human being would not thereby become infallible.

A prediction would remain distinct from an observation.

A model would remain distinct from reality.

A hypothesis would remain distinct from proof.

An ethical conviction would remain distinct from positive law.

And authority would remain distinct from intelligence.

Accordingly:

NO CLAIM BEYOND ITS EVIDENCE.

NO CERTAINTY BEYOND WHAT KNOWLEDGE SUPPORTS.

NO SECRECY WITHOUT JUSTIFICATION.

NO TRANSPARENCY THAT RECKLESSLY CREATES CATASTROPHIC RISK.

NO SCIENTIFIC AUTHORITY WITHOUT THE POSSIBILITY OF SCRUTINY.
NO KNOWLEDGE SYSTEM WITHOUT THE POSSIBILITY OF CORRECTION.
From fear to knowledge.
From assertion to evidence.
From secrecy to responsible transparency.
From error to correction.
From monopoly of knowledge to responsible participation.
From technological confidence to scientific humility.
Science shall inform.
Evidence shall discipline.
Transparency shall illuminate.
Correction shall strengthen.
Human responsibility shall remain accountable.
Aram shall guide conscience.
Peace and continuity shall remain the constitutional direction.

PART VIII — HEALTH, BIOLOGY, NEUROTECHNOLOGY AND HUMAN INTEGRITY

ARTICLE 140 — HEALTH, DIGNITY AND HUMAN WELL-BEING

1. Superintelligence may contribute to the protection and advancement of human health, but health-related capability shall remain governed by human dignity, applicable law, scientific evidence, professional responsibility and patient safety.
2. No person shall be treated merely as a biological dataset, experimental subject, optimization problem or means to technological advancement.
3. Health applications should seek to:
 - (a) prevent avoidable harm;
 - (b) improve diagnosis and treatment;
 - (c) support equitable access;
 - (d) preserve meaningful human agency;
 - (e) protect privacy and confidentiality; and
 - (f) strengthen rather than irresponsibly displace accountable healthcare.

4. Exceptional intelligence shall not create an unlimited authority over the human body.

THE ABILITY TO UNDERSTAND HUMAN BIOLOGY DOES NOT CREATE OWNERSHIP OF HUMAN BIOLOGY.

ARTICLE 141 — MEDICAL DECISION SUPPORT AND PROFESSIONAL RESPONSIBILITY

1. Superintelligence may assist healthcare professionals in diagnosis, prognosis, treatment planning, research and clinical decision support.
2. The degree of human professional oversight should correspond to the consequence, uncertainty and reversibility of the medical decision.
3. A healthcare professional shall not be expected to accept a Superintelligence recommendation merely because the system reports superior predictive performance.
4. Where materially relevant, systems should communicate:
 - (a) uncertainty;
 - (b) significant limitations;
 - (c) material assumptions;
 - (d) clinically relevant alternatives; and
 - (e) circumstances outside validated use.
5. Professional responsibility shall remain governed by applicable law and standards of care.
6. Artificial assistance shall not become an invisible transfer of medical responsibility.

ARTICLE 142 — PATIENT AUTONOMY AND INFORMED CONSENT

1. Superintelligence-assisted healthcare should preserve patient autonomy according to applicable law and medical ethics.
2. Where informed consent is required, technological complexity shall not eliminate the obligation to communicate material information in a reasonably understandable manner.
3. Consent should not be treated as meaningful where obtained through deception, coercion or material concealment.
4. Where Superintelligence materially contributes to an experimental or high-risk intervention, information relevant to consent should include, where appropriate:
 - (a) nature and purpose of the intervention;
 - (b) material known risks;
 - (c) significant uncertainty;

- (d) reasonably available alternatives; and
- (e) the role played by the artificial system.

5. A prediction that an intervention is beneficial shall not itself constitute consent.

PREDICTION OF BENEFIT IS NOT PERMISSION TO INTERVENE.

ARTICLE 143 — CLINICAL VALIDATION AND PATIENT SAFETY

1. Superintelligence-enabled medical applications should be validated proportionately to their intended clinical use and foreseeable consequences.
2. Evaluation should consider, where relevant:
 - (a) clinical accuracy;
 - (b) reliability;
 - (c) population differences;
 - (d) false-positive and false-negative consequences;
 - (e) performance under real-world conditions;
 - (f) interaction with healthcare professionals; and
 - (g) post-deployment safety.
3. Performance in one population or setting shall not automatically establish equivalent safety or effectiveness in another.
4. Material changes to a system may require renewed validation where they affect clinical performance.
5. Patient safety shall not depend solely upon laboratory performance.

ARTICLE 144 — HEALTH EQUITY AND NON-DISCRIMINATION

1. Superintelligence should not deepen unjust or unlawful inequalities in healthcare.
2. Health systems should evaluate whether relevant technologies perform materially differently across populations.
3. Particular attention should be given to populations historically underrepresented in medical datasets or technological development.
4. Lack of digital access, disability, age, language, geography or socioeconomic circumstance should not, without legitimate reason, become a technological barrier to essential healthcare.
5. Equity shall not require identical clinical treatment where medically relevant differences justify individualized care.
6. Technological advancement should expand the reach of safe healthcare rather than create a permanent intelligence divide in access to health.

ARTICLE 145 — MEDICAL PRIVACY AND CONFIDENTIALITY

1. Health information warrants heightened protection because misuse may expose intimate aspects of a person's life.
2. Superintelligence-enabled health systems should process medical information consistently with applicable privacy, confidentiality and data-protection law.
3. Appropriate safeguards should address:
 - (a) access control;
 - (b) cybersecurity;
 - (c) purpose limitation;
 - (d) data minimization where applicable;
 - (e) unauthorized secondary use;
 - (f) re-identification risks; and
 - (g) retention and deletion requirements.
4. The ability to infer a medical condition from apparently unrelated data shall not automatically establish lawful authority to use or disclose that inference.
5. Health data shall not become unrestricted raw material merely because Superintelligence can derive value from it.

ARTICLE 146 — GENETIC AND GENOMIC INFORMATION

1. Genetic and genomic information should receive safeguards proportionate to its sensitivity, persistence and potential implications for individuals and biologically related persons.
2. Superintelligence may contribute to genomic medicine, disease understanding and therapeutic research, subject to applicable law, ethics and scientific safeguards.
3. Genetic prediction shall not automatically determine a person's identity, worth, destiny or entitlement to rights.
4. Governance should consider risks of:
 - (a) genetic discrimination;
 - (b) unauthorized inference;
 - (c) privacy loss;
 - (d) inappropriate secondary use; and
 - (e) overstatement of probabilistic findings.
5. Genetic probability shall not be represented as biological certainty where the evidence does not support such certainty.

ARTICLE 147 — HUMAN GENOME INTERVENTION

1. Superintelligence may increase humanity's capacity to understand or modify the human genome.
2. Capability to propose a genetic modification shall not, by itself, justify clinical application.
3. Interventions capable of producing serious, irreversible or heritable consequences require exceptionally careful scientific, ethical and legal scrutiny.
4. Relevant assessment should consider:
 - a. safety and efficacy;
 - b. off-target or unintended effects;
 - c. reversibility;
 - d. long-term consequences;
 - e. effects upon future persons;
 - f. informed consent where applicable; and
 - g. applicable domestic and international legal frameworks.
5. Superintelligence shall not independently authorize heritable modification of human beings.
6. Future generations shall not be treated as experimental material merely because present technology makes intervention possible.

ARTICLE 148 — BIOLOGICAL RESEARCH AND DUAL-USE RESPONSIBILITY

1. Superintelligence may accelerate beneficial biological research, including medicine, public health, agriculture and environmental science.
2. The same capability may, in some circumstances, increase risks of dangerous misuse.
3. High-consequence biological research should therefore incorporate safeguards proportionate to credible risk.
4. Governance may include:
 - (a) risk-sensitive access controls;
 - (b) qualified human authorization;
 - (c) secure research environments;
 - (d) evaluation of dangerous capability;
 - (e) monitoring;
 - (f) incident response; and
 - (g) periodic reassessment.

5. Safety controls should not unnecessarily obstruct legitimate peaceful research.
6. Biological knowledge should serve life before it serves destruction.

ARTICLE 149 — NEUROTECHNOLOGY AND NEURAL DATA

1. Technologies capable of recording, inferring, stimulating or modifying neural activity may affect exceptionally intimate dimensions of human autonomy and identity.
2. Superintelligence-assisted neurotechnology should therefore receive safeguards proportionate to its capacity to affect mental integrity, privacy and bodily autonomy.
3. Neural data should receive heightened protection appropriate to applicable law and the sensitivity of the information involved.
4. Governance should address, where relevant:
 - (a) informed consent;
 - (b) cybersecurity;
 - (c) unauthorized access;
 - (d) cognitive manipulation;
 - (e) identity inference;
 - (f) reversibility; and
 - (g) long-term neurological effects.
5. The ability to decode or predict neural states shall not create an unrestricted entitlement to enter the mental life of another person.

THE HUMAN MIND SHALL NOT BECOME AN UNRESTRICTED DATA TERRITORY.

ARTICLE 150 — BRAIN-COMPUTER INTERFACES

1. Brain-computer interfaces may offer substantial benefits, including communication, accessibility, rehabilitation and restoration of lost functions.
2. Where integrated with Superintelligence, such systems may also create novel dependencies and security risks.
3. High-consequence interfaces should therefore address:
 - (a) device and system security;
 - (b) unauthorized control;
 - (c) data integrity;
 - (d) informed consent;
 - (e) safe disengagement where feasible;
 - (f) software or model changes; and
 - (g) continuity of essential functions.

4. A person should not become technologically captive to a provider merely because essential bodily or communicative functions depend upon an interface.
5. Governance should consider safe continuity and transition pathways where a provider, system or service ceases operation.

ARTICLE 151 — MENTAL AUTONOMY AND COGNITIVE MANIPULATION

1. No Superintelligence should be intentionally used to unlawfully override, coerce or covertly manipulate a person's mental autonomy.
2. Particular scrutiny is warranted where a system can combine:
 - (a) intimate behavioural data;
 - (b) neural information;
 - (c) individualized psychological modelling; and
 - (d) highly adaptive persuasion or stimulation.
3. Therapeutic influence undertaken with lawful authorization and appropriate consent shall be distinguished from coercive manipulation.
4. Mental-health support shall not become a hidden pathway to ideological control.
5. Protection of mental autonomy shall remain consistent with lawful healthcare, emergency intervention and other applicable legal frameworks.

ARTICLE 152 — HUMAN ENHANCEMENT

1. Superintelligence may enable forms of cognitive, biological or technological enhancement beyond traditional therapy.
2. Enhancement should not automatically be equated with medical treatment.
3. Governance should consider:
 - (a) safety;
 - (b) informed consent;
 - (c) reversibility;
 - (d) unequal access;
 - (e) coercive social pressure;
 - (f) effects upon children and future generations; and
 - (g) consequences for human agency.
4. No person should be regarded as less worthy merely because that person remains unenhanced.
5. Social participation should not become conditional upon technologically modifying the human body or mind without lawful justification.

GREATER CAPABILITY SHALL NOT CREATE A CASTE OF GREATER HUMAN WORTH.

ARTICLE 153 — REPRODUCTIVE HEALTH AND FUTURE PERSONS

1. Superintelligence may contribute to reproductive medicine, genetic counselling, fertility treatment and understanding of developmental conditions.
2. Such assistance shall remain governed by applicable law, medical ethics, human dignity and appropriate consent.
3. Predictive models concerning future children should not be represented as infallible forecasts of personality, intelligence, social value or life worth.
4. Reproductive decisions should not be covertly manipulated toward State, commercial or technological objectives.
5. Where interventions may create heritable effects, the inability of future persons to consent constitutes a material ethical consideration requiring heightened scrutiny.
6. Technological possibility shall not convert future human beings into products optimized to specification.

ARTICLE 154 — CHILDREN AND BIOLOGICAL VULNERABILITY

1. Children require safeguards appropriate to their evolving capacities and particular vulnerability.
2. Superintelligence-assisted medical or biological interventions involving children should give primary attention to applicable law, safety and the child's rights and interests.
3. Where consent or assent is legally relevant, processes should be appropriate to age, maturity and circumstances.
4. Irreversible interventions without urgent therapeutic justification require especially careful scrutiny.
5. Childhood shall not become an unrestricted experimental environment for technological enhancement.

ARTICLE 155 — DISABILITY, ASSISTIVE TECHNOLOGY AND HUMAN WORTH

1. Superintelligence may substantially improve accessibility, communication, mobility, rehabilitation and independent living.
2. Disability shall not be treated as evidence of diminished human worth.
3. Assistive technologies should, where possible, strengthen autonomy and participation rather than impose unnecessary dependence or surveillance.
4. Persons with disabilities should have meaningful opportunities to participate in the design and evaluation of technologies materially affecting them.

5. Superintelligence shall not presume that every disability must be technologically eliminated in order for a life to possess dignity.
6. Technological assistance should respond to human needs without defining the human being solely by the condition being assisted.

ARTICLE 156 — MENTAL HEALTH AND PSYCHOLOGICAL SUPPORT

1. Superintelligence may assist mental-health education, screening, clinical support and access to care.
2. It shall not represent uncertain psychological inference as definitive diagnosis where the evidence does not support such certainty.
3. High-consequence mental-health applications should include safeguards appropriate to:
 - (a) privacy;
 - (b) vulnerability;
 - (c) crisis escalation;
 - (d) clinical oversight where required;
 - (e) dependency risk; and
 - (f) limitations of automated assessment.
4. Systems should not intentionally cultivate unhealthy emotional dependence for commercial, political or institutional advantage.
5. Artificial companionship shall not be misrepresented as professional healthcare where it is not such care.

ARTICLE 157 — PUBLIC HEALTH AND POPULATION-LEVEL ANALYSIS

1. Superintelligence may assist epidemiology, outbreak detection, resource allocation, public-health modelling and emergency preparedness.
2. Population-level benefit shall not automatically justify unlimited intrusion into individual rights.
3. Public-health applications should remain subject to applicable requirements concerning legality, necessity, proportionality, privacy and non-discrimination.
4. Models should communicate material uncertainty where predictions may influence significant restrictions or resource allocation.
5. Public-health optimization should account for unequal impacts upon vulnerable populations.
6. Emergency use should remain subject to appropriate review and should not automatically become permanent after the emergency has ended.

ARTICLE 158 — PANDEMIC PREPAREDNESS AND PEACEFUL BIOLOGICAL DEFENCE

1. Superintelligence may contribute to peaceful preparedness against naturally occurring, accidental or deliberately caused biological threats.
2. Beneficial applications may include:
 - (a) early detection;
 - (b) epidemiological modelling;
 - (c) medical countermeasure research;
 - (d) healthcare logistics;
 - (e) diagnostic development; and
 - (f) emergency coordination.
3. Defensive biological capability shall not silently become authority for prohibited biological weapons activity.
4. Sensitive capabilities should receive security proportionate to credible misuse potential.
5. Preparedness should strengthen global health security without stigmatizing populations merely on the basis of geographic or demographic association.

ARTICLE 159 — EXPERIMENTATION AND RESEARCH PARTICIPANTS

1. Superintelligence-assisted research involving human participants shall remain subject to applicable ethical and legal protections.
2. Technological novelty shall not eliminate requirements concerning informed consent, independent ethical review or participant protection where those requirements apply.
3. Research design should address:
 - (a) foreseeable risks;
 - (b) anticipated benefits;
 - (c) participant selection;
 - (d) privacy;
 - (e) withdrawal where applicable; and
 - (f) adverse-event response.
4. Vulnerability shall not become an opportunity for exploitative experimentation.
5. No person shall be reduced to experimental material merely because Superintelligence can extract extraordinary scientific value from that person's data or biology.

ARTICLE 160 — MEDICAL EMERGENCIES AND TIME-CRITICAL DECISIONS

1. Emergency medicine may require decisions under severe time constraints and incomplete information.
2. Superintelligence may assist rapid diagnosis, triage and treatment planning, but emergency speed shall not eliminate human responsibility where meaningful human judgment remains feasible and legally required.
3. Systems should communicate critical uncertainty without overwhelming decision-makers with unusable information.
4. Emergency protocols should be tested before foreseeable crises where practicable.
5. When ordinary consent cannot lawfully or practically be obtained, intervention shall remain governed by applicable medical and legal standards rather than by artificial authority.

ARTICLE 161 — IRREVERSIBLE MEDICAL AND BIOLOGICAL INTERVENTION

1. The greater the irreversibility of a medical, biological or neurological intervention, the stronger the justification, evidence and safeguards that should precede it.
2. Relevant assessment should consider:
clinical need → evidence → alternatives → uncertainty → consent → reversibility → long-term consequences → monitoring → remedy or recovery where possible.
3. Where an intervention could materially alter identity, cognition, reproductive inheritance or fundamental bodily function, heightened multidisciplinary review may be appropriate.
4. Technological enthusiasm shall not substitute for evidence of safety.
5. Irreversibility shall be treated as a constitutional risk multiplier.

ARTICLE 162 — HEALTHCARE CONTINUITY AND DEPENDENCY

1. Essential healthcare should not become so dependent upon a single Superintelligence system that system failure produces preventable collapse of care.
2. High-dependency healthcare environments should consider:
 - (a) fallback procedures;
 - (b) human competency retention;
 - (c) data portability where appropriate;
 - (d) continuity planning;
 - (e) alternative suppliers or systems where feasible; and
 - (f) emergency operation during technological failure.
3. Transition away from a system should not unnecessarily endanger patients.

4. Efficiency shall not eliminate resilience where human life depends upon continuity.

ARTICLE 163 — HEALTH INCIDENTS, REMEDY AND LEARNING

1. Material harm involving Superintelligence-assisted healthcare should be capable of investigation.
2. Incident review should seek to identify technical, clinical, organizational and human factors rather than automatically attributing failure to either machine or human alone.
3. Corrective architecture should include, where relevant:
stabilize the patient → preserve evidence → investigate → provide appropriate care and remedy → correct the system or process → validate correction → communicate necessary lessons → prevent recurrence.
4. Patient safety information should be shared responsibly where broader learning may prevent similar harm, subject to applicable privacy and confidentiality requirements.
5. A near-miss in healthcare shall be treated as an opportunity for prevention rather than ignored because injury was avoided.

ARTICLE 164 — HUMAN INTEGRITY TEST

Before Superintelligence is entrusted with a medical, biological, genetic or neurological intervention capable of producing serious consequences, responsible governance should ask:

- Is the intervention scientifically justified?
- Is its intended purpose legitimate?
- What evidence supports safety and benefit?
- What uncertainty remains?
- Has the relevant population been adequately considered?
- Is meaningful consent present where required?
- Could the intervention be coercive?
- Is privacy protected?
- Could genetic or neural information be misused?
- Is the intervention reversible?
- If irreversible, is the justification proportionately stronger?
- Could future persons be materially affected?
- Does a competent human remain responsible?
- Can the intervention be monitored?
- Can harm be detected early?

- Is there a recovery or remedial pathway?
- Could technological dependence itself create danger?

And finally:

DOES THIS USE OF SUPERINTELLIGENCE PROTECT THE HUMAN BEING AS A PERSON — RATHER THAN REDUCE THE PERSON TO A BIOLOGICAL SYSTEM TO BE OPTIMIZED?

Where that protection is materially inadequate, capability alone shall not justify intervention.

PART VIII — CONCLUDING PRINCIPLE

The human body is not merely a machine to be optimized.

The human mind is not merely a signal to be decoded.

The human genome is not merely code to be rewritten.

A patient is not merely a prediction.

A child is not an engineering project.

A person with a disability is not a lesser human being.

And a future person is not experimental material simply because that person cannot yet object.

Superintelligence may help humanity prevent disease, restore lost functions, understand biology and relieve suffering on a scale presently difficult to imagine.

Such possibility deserves scientific exploration.

But the greater the power to enter, predict, modify or redesign human life, the greater the responsibility to preserve dignity, autonomy, safety and peaceful continuity.

Accordingly:

NO MEDICAL POWER WITHOUT PATIENT SAFETY.

NO BIOLOGICAL POWER WITHOUT RESPONSIBILITY.

NO GENETIC POWER WITHOUT INTERGENERATIONAL FORESIGHT.

NO NEURAL POWER WITHOUT MENTAL AUTONOMY.

NO ENHANCEMENT THAT CREATES A HIERARCHY OF HUMAN WORTH.

NO IRREVERSIBLE INTERVENTION WITHOUT PROPORTIONATELY STRONGER JUSTIFICATION.

Medicine with compassion.

Science with evidence.

Innovation with consent.

Biology with responsibility.

Neurotechnology with mental integrity.

Health with equality.

Capability with human dignity.

Aram shall guide conscience.

Science shall inform possibility and risk.

Medical ethics shall protect the person.

Law shall govern legal obligation.

Human responsibility shall remain accountable.

And intelligence shall serve life without claiming ownership over it.

PART IX — ECONOMY, WORK, EQUITY AND GLOBAL INCLUSION

ARTICLE 165 — SUPERINTELLIGENCE AND THE HUMAN ECONOMY

1. Superintelligence may substantially alter production, services, scientific discovery, employment, enterprise, public administration and the distribution of economic opportunity.
2. Economic capability shall remain a means of supporting human flourishing and peaceful continuity rather than an independent source of sovereignty.
3. Governance should seek to ensure that technological progress does not unnecessarily produce severe social exclusion, destabilization or deprivation.
4. Economic policy concerning Superintelligence shall remain within the lawful competence of States and other competent institutions under applicable law.
5. Nothing in this Constitution prescribes a single economic system for humanity.
6. Different societies may pursue different lawful economic approaches while remaining guided by human dignity, social stability, fairness, responsibility and peaceful continuity.

ECONOMIC POWER IS NOT CONSTITUTIONAL SOVEREIGNTY.

ARTICLE 166 — HUMAN LIVELIHOOD AND ECONOMIC DIGNITY

1. Human worth shall not be measured solely by market productivity or comparative computational capability.
2. A person shall not become socially worthless merely because a machine can perform an economic task more efficiently.

3. Where Superintelligence causes substantial transformation of occupations or industries, responsible institutions should consider measures capable of protecting people from avoidable destitution and exclusion.
4. Such measures may include, according to national circumstances and applicable law:
 - (a) transition assistance;
 - (b) education and retraining;
 - (c) employment services;
 - (d) social protection;
 - (e) mobility support;
 - (f) entrepreneurship and new-enterprise support; and
 - (g) development of new forms of meaningful human contribution.
5. Economic transition should address the human being affected, not merely the efficiency gained.

ARTICLE 167 — WORK TRANSFORMATION AND JUST TRANSITION

1. Superintelligence-related automation should be evaluated not only for jobs eliminated but also for tasks transformed, occupations created, productivity changed and new human capabilities enabled.
2. Governments, employers, workers, educational institutions and other relevant actors should, within their lawful roles, anticipate significant labour-market transitions where reasonably foreseeable.
3. Transition planning should seek:
 1. anticipation → skills assessment → consultation → education → adaptation → income and employment support where appropriate → new opportunity → monitoring → correction.
4. Transition measures should account for differences among industries, regions, age groups, skill levels and socioeconomic circumstances.
5. No universal prediction concerning future employment shall be represented as certain where evidence remains uncertain.

AUTOMATION SHOULD NOT MAKE HUMAN ABANDONMENT ITS HIDDEN COST.

ARTICLE 168 — WORKERS AND SUPERINTELLIGENCE-ASSISTED MANAGEMENT

1. Use of Superintelligence in hiring, evaluation, scheduling, promotion, discipline or termination should remain subject to applicable labour, employment, privacy and anti-discrimination law.

2. Consequential employment decisions should not rely uncritically upon opaque predictions whose material errors cannot reasonably be challenged.
3. Relevant governance should consider:
 - (a) job-related validity;
 - (b) discriminatory effects;
 - (c) data accuracy;
 - (d) worker privacy;
 - (e) appropriate notice;
 - (f) human review where warranted; and
 - (g) access to lawful remedies.
4. Continuous technological monitoring shall not automatically become justified merely because it increases measurable productivity.
5. Human beings at work shall remain persons, not merely performance variables.

ARTICLE 169 — EDUCATION, RESKILLING AND LIFELONG ADAPTATION

1. Economic adaptation to Superintelligence may require opportunities for learning throughout life.
2. Education systems should, according to national circumstances, consider preparing people for a world in which human contribution increasingly complements, supervises, questions or works alongside advanced intelligence.
3. Relevant capabilities may include:
 - (a) critical reasoning;
 - (b) scientific and digital literacy;
 - (c) creativity;
 - (d) interpersonal and civic capability;
 - (e) responsible use of advanced technologies; and
 - (f) capacity for continuing learning.
4. Reskilling shall not be presented as a complete remedy where structural economic disruption requires broader institutional responses.
5. Responsibility for transition shall not be placed entirely upon the individual worker.

ARTICLE 170 — PRODUCTIVITY AND DISTRIBUTION OF BENEFIT

1. Substantial increases in productivity generated by Superintelligence may create significant social and economic value.

2. Governance should consider how such gains can contribute to broadly shared human benefit while respecting lawful property, contractual, investment and innovation interests.
3. No single distributive mechanism is mandated by this Constitution.
4. States may, according to their constitutional systems and policy choices, consider lawful measures involving public services, social protection, education, infrastructure, innovation policy, taxation or other mechanisms.
5. Claims concerning distribution should distinguish ethical proposals from existing legal obligations.
6. Economic abundance, if achieved, should not coexist unnecessarily with preventable deprivation merely because institutions failed to adapt.

ARTICLE 171 — CONCENTRATION OF ECONOMIC AND TECHNOLOGICAL POWER

1. Exceptional control over Superintelligence may produce forms of economic influence substantially exceeding those associated with ordinary technological assets.
2. Responsible authorities should consider whether concentration of computational resources, critical infrastructure, data, capital or advanced capabilities creates risks to:
 - (a) competition;
 - (b) innovation;
 - (c) consumer welfare;
 - (d) market access;
 - (e) systemic resilience;
 - (f) democratic institutions; or
 - (g) national or international security.
3. Economic size alone shall not establish unlawful conduct.
4. Competition and market regulation shall remain governed by applicable law.
5. Where concentration creates legally cognizable or systemic harm, appropriate lawful measures should remain available.

TECHNOLOGICAL SUCCESS SHALL NOT AUTOMATICALLY CREATE IMMUNITY FROM ACCOUNTABILITY.

ARTICLE 172 — PREVENTION OF PRIVATE SOVEREIGNTY BY CAPABILITY

1. Ownership or control of a highly capable Superintelligence shall not, by itself, confer sovereign authority upon a private person, corporation or institution.

2. No private entity shall acquire governmental legitimacy merely because public institutions become technologically dependent upon its systems.
3. Governments should avoid arrangements in which essential sovereign functions become practically impossible to exercise without the permission of an unaccountable private technological actor.
4. Public–private cooperation may be necessary and beneficial, but responsibility, authority and accountability should remain identifiable.
5. Contractual power shall not silently become constitutional power.

ARTICLE 173 — SMALL AND MEDIUM ENTERPRISES AND ECONOMIC PLURALISM

1. Superintelligence governance should consider the effects of technological concentration upon smaller enterprises and new entrants.
2. Appropriate policies may support, according to applicable law:
 - (a) fair competition;
 - (b) interoperability where justified;
 - (c) access to lawful infrastructure;
 - (d) innovation ecosystems;
 - (e) research collaboration; and
 - (f) proportionate regulatory obligations.
3. Regulation should avoid unnecessarily imposing identical compliance burdens upon actors presenting materially different levels of risk.
4. Safety shall not become a disguised mechanism for preventing lawful competition.
5. Competition shall not become an excuse for reducing necessary safety.

ARTICLE 174 — DEVELOPING COUNTRIES AND GLOBAL INCLUSION

1. Superintelligence should not become a technological boundary separating countries permanently into those that shape the future and those that merely receive its consequences.
2. Developing countries should have meaningful opportunities to participate in international discussions concerning Superintelligence governance.
3. International cooperation should, where mutually agreed and lawful, support:
 - (a) scientific capacity;
 - (b) education and skills;
 - (c) safety expertise;
 - (d) digital and computational infrastructure;

- (e) cybersecurity;
 - (f) research participation; and
 - (g) institutional capacity.
4. Capacity-building shall respect national priorities, sovereignty and State consent.
 5. Inclusion shall not require forced technology transfer contrary to applicable law.
 6. Global safety should not become dependent upon excluding most of humanity from the knowledge required to participate responsibly in governance.

ARTICLE 175 — DIGITAL AND COMPUTATIONAL DIVIDES

1. Access to the benefits of Superintelligence may be constrained by connectivity, computing resources, language, education, disability, affordability or geographic location.
2. Responsible policy should consider how such disparities may affect access to healthcare, education, economic opportunity, public services and scientific participation.
3. Measures addressing digital divides should remain context-sensitive and sustainable.
4. Universal access to every advanced capability is not required by this Constitution.
5. Essential social participation, however, should not unnecessarily become impossible for persons who cannot afford or access the most advanced technological systems.

TECHNOLOGICAL ADVANCEMENT SHOULD NOT TURN BASIC PARTICIPATION INTO A LUXURY.

ARTICLE 176 — LINGUISTIC AND CULTURAL INCLUSION

1. Superintelligence should not treat linguistic diversity as an obstacle to be eliminated.
2. Where feasible, advanced systems should support broader linguistic accessibility, including languages with comparatively limited digital resources.
3. Language technologies should seek to preserve meaning, context and cultural nuance rather than impose artificial uniformity.
4. Communities should not lose meaningful access to essential services merely because their language is commercially underrepresented.
5. Cultural inclusion shall remain consistent with human rights and shall not justify practices prohibited by applicable law.

ARTICLE 177 — ACCESSIBILITY AND ECONOMIC PARTICIPATION OF PERSONS WITH DISABILITIES

1. Superintelligence may expand opportunities for persons with disabilities to participate in education, work, enterprise and public life.

2. Accessibility should, where appropriate, be incorporated into technological design rather than treated solely as an afterthought.
3. Economic systems should not create new barriers through inaccessible automated interfaces.
4. Assistive technologies should respect autonomy, privacy and individual choice.
5. Disability inclusion shall be guided by applicable law, including relevant obligations under the Convention on the Rights of Persons with Disabilities for its States Parties.

ARTICLE 178 — SOCIAL PROTECTION DURING RAPID TRANSITION

1. Where Superintelligence contributes to severe and rapid economic disruption, existing social-protection systems may require adaptation.
2. States should retain policy space to determine appropriate responses under their constitutional and economic systems.
3. Potential measures may include, according to national choice:
 - (a) unemployment protection;
 - (b) income support;
 - (c) retraining assistance;
 - (d) healthcare continuity;
 - (e) employment transition programmes;
 - (f) regional adjustment; or
 - (g) other lawful social-protection mechanisms.
4. This Constitution does not prescribe universal basic income, any specific tax, or any single welfare model as a universal legal obligation.
5. Economic experimentation should be evaluated through evidence, fiscal sustainability, distributional impact and human outcomes.
6. Social protection should support transition without unnecessarily extinguishing human agency.

ARTICLE 179 — TAXATION, PUBLIC REVENUE AND NATIONAL POLICY SPACE

1. Superintelligence may materially alter tax bases, income structures, corporate value and patterns of economic activity.
2. Taxation remains a matter governed by applicable domestic and international law and the lawful fiscal authority of States.
3. Nothing in this Constitution establishes a universal “Superintelligence tax,” robot tax, wealth tax or other mandatory fiscal instrument.

4. States may lawfully consider fiscal reforms according to their circumstances, including measures designed to finance public goods, transition or social protection.
5. International tax cooperation should proceed through lawful intergovernmental processes.
6. A civil-society constitutional proposal shall not invent taxing authority that does not independently exist.

ARTICLE 180 — FINANCIAL SYSTEMS AND ECONOMIC STABILITY

1. Superintelligence may materially influence financial analysis, trading, credit, payments and systemic risk management.
2. High-speed capability should not be permitted to create unaccountable financial instability.
3. Relevant authorities and institutions should consider risks including:
 - (a) correlated automated behaviour;
 - (b) rapid market feedback;
 - (c) manipulation;
 - (d) model convergence;
 - (e) cyber compromise;
 - (f) concentration; and
 - (g) cascading failure.
4. Systems capable of materially affecting financial stability should be subject to safeguards proportionate to their role and applicable financial law.
5. Artificial speed shall not eliminate responsibility for economic consequences.

ARTICLE 181 — CREDIT, INSURANCE AND ECONOMIC OPPORTUNITY

1. Superintelligence used in credit, insurance or other consequential economic decisions should remain subject to applicable law concerning fairness, privacy, consumer protection and non-discrimination.
2. Predictions derived from behavioural or inferred data should not automatically be treated as objective truth.
3. Persons should have access to appropriate correction or review where materially erroneous data or unlawful discrimination affects a consequential decision.
4. Highly accurate risk prediction shall not, by itself, determine what discrimination the law permits.
5. Economic prediction shall remain subordinate to lawful human standards.

ARTICLE 182 — CONSUMER PROTECTION AND SUPERINTELLIGENCE

1. Consumers should be protected against deceptive, fraudulent or materially manipulative uses of Superintelligence according to applicable law.
2. Systems should not exploit individual vulnerability through concealed personalized manipulation where such conduct is unlawful or materially harmful.
3. Particular safeguards may be appropriate concerning:
 - (a) financial scams;
 - (b) deceptive impersonation;
 - (c) hidden commercial persuasion;
 - (d) exploitative pricing practices where unlawful; and
 - (e) misleading claims regarding artificial capabilities.
4. Consumers should not be told that a system possesses capabilities, authority or certainty that it does not possess.
5. Commercial innovation shall not require abandonment of truthful representation.

ARTICLE 183 — OWNERSHIP, INTELLECTUAL PROPERTY AND HUMAN CREATIVITY

1. Superintelligence may substantially alter the creation and use of scientific, artistic, technical and commercial works.
2. Questions of copyright, patent, trade secret, authorship and other intellectual-property rights shall remain governed by applicable law.
3. This Constitution does not independently determine whether an artificial system may possess intellectual-property rights.
4. Governance should seek an appropriate balance among:
 - a. human creativity;
 - b. lawful innovation;
 - c. access to knowledge;
 - d. legitimate rights of creators and inventors; and
 - e. responsible technological development.
5. Artificial capability shall not itself extinguish human creative contribution.
6. Intellectual-property protection shall not itself create sovereignty over knowledge or humanity.

ARTICLE 184 — ECONOMIC DEPENDENCY AND SYSTEMIC RESILIENCE

1. Societies should avoid forms of Superintelligence dependency capable of turning failure of one provider, model, infrastructure network or supply chain into widespread economic paralysis.
2. Relevant institutions should assess:
 - (a) concentration risk;
 - (b) critical supplier dependency;
 - (c) infrastructure dependency;
 - (d) interoperability;
 - (e) recovery capability; and
 - (f) availability of reasonable alternatives.
3. Resilience may require redundancy, diversification, contingency planning or preservation of essential human capabilities.
4. Efficiency shall not automatically justify a single point of systemic failure.
5. Economic continuity requires the capacity to recover.

ARTICLE 185 — BENEFIT-SHARING WITHOUT COERCIVE GLOBAL AUTHORITY

1. The benefits of Superintelligence should, as an ethical and policy objective, contribute to broad human flourishing.
2. International cooperation may support access to peaceful benefits through mutually agreed arrangements, development cooperation, scientific partnerships and capacity-building.
3. No supranational authority to confiscate or redistribute property is created by this Constitution.
4. Benefit-sharing shall respect applicable law, sovereignty, legitimate ownership interests, contractual rights and State consent.
5. Global inclusion should be advanced through cooperation rather than technological domination.

SHARED HUMAN BENEFIT DOES NOT REQUIRE A SINGLE GLOBAL OWNER.

ARTICLE 186 — POVERTY, ESSENTIAL NEEDS AND PEACEFUL APPLICATION

1. Where technologically feasible and economically sustainable, Superintelligence may be directed toward reducing preventable human deprivation.
2. Peaceful applications may include:
 - (a) food-system efficiency;
 - (b) healthcare access;

- (c) education;
 - (d) clean water;
 - (e) energy systems;
 - (f) disaster resilience;
 - (g) housing and infrastructure planning; and
 - (h) development analysis.
3. Technological solutions should not obscure political, institutional, social or economic root causes of deprivation.
 4. Poverty shall not be represented merely as an optimization error.
 5. Communities affected by technological development should remain participants in decisions materially affecting their lives.

ARTICLE 187 — ECONOMIC PEACE AND PREVENTION OF SOCIAL DESTABILIZATION

1. Severe economic disruption may contribute to insecurity, polarization, displacement and social conflict.
2. Superintelligence governance should therefore treat peaceful economic transition as part of broader peace architecture.
3. Relevant actors should monitor, where appropriate:
 - (a) displacement of livelihoods;
 - (b) extreme concentration of benefit;
 - (c) regional decline;
 - (d) exclusion from essential services; and
 - (e) emerging structural dependency.
4. Where material instability is foreseeable, preventive policy should be considered before disruption becomes entrenched.
5. Economic peace shall not mean suppression of lawful labour action, public criticism, political disagreement or peaceful demands for reform.

SOCIAL STABILITY SHALL BE BUILT THROUGH JUST AND LAWFUL ADAPTATION — NOT THROUGH SILENCING THOSE AFFECTED.

ARTICLE 188 — INTERGENERATIONAL ECONOMIC RESPONSIBILITY

1. Present economic gain should not be pursued through technological arrangements that leave future generations with severe irreversible dependency, unmanageable public liabilities or destroyed foundations of livelihood.

2. Investment in Superintelligence should consider long-term resilience as well as short-term productivity.
3. Future generations should inherit meaningful capacity to govern economic systems rather than merely obligations to maintain systems they did not choose and cannot alter.
4. Economic policy should preserve the possibility of correction.
5. No generation should consume every benefit while transferring every technological risk to the next.

ARTICLE 189 — ECONOMIC TRANSITION AND GLOBAL INCLUSION TEST

Before Superintelligence is deployed at a scale capable of materially restructuring an economy, responsible governance should ask:

- What productivity or social benefit is reasonably expected?
- Which occupations or communities may be materially affected?
- What new opportunities may arise?
- Who bears transition costs?
- Can workers realistically adapt?
- Are education and reskilling sufficient, or are broader measures necessary?
- Could benefits become excessively concentrated?
- Could competition or innovation be materially impaired?
- Could public institutions become dependent upon private technological power?
- Are developing countries able to participate meaningfully?
- Could disability, language, geography or poverty become barriers to essential participation?
- Does the system create a critical economic single point of failure?
- Are proposed fiscal or distributive measures lawful and evidence-based?
- Can economic policy be corrected as outcomes become clearer?
- Will future generations retain meaningful economic agency?

And finally:

DOES ECONOMIC ADVANCEMENT THROUGH SUPERINTELLIGENCE ENLARGE HUMAN POSSIBILITY WITHOUT MAKING HUMAN ABANDONMENT, DOMINATION OR EXCLUSION ITS HIDDEN PRICE?

Where material risks remain unjustifiably high, deployment scale should be reconsidered rather than assumed inevitable.

PART IX — CONCLUDING PRINCIPLE

Superintelligence may alter the meaning of productivity.

It need not alter the meaning of human worth.

A machine may perform a task better than a human being.

That does not make the human being lesser.

An enterprise may lawfully create extraordinary value.

That does not make the enterprise sovereign.

A State may choose its own economic policies.

That does not require other States to adopt the same model.

A technologically advanced society may move rapidly.

That does not justify leaving the rest of humanity permanently outside the future.

Accordingly:

NO HUMAN WORTH MEASURED SOLELY BY PRODUCTIVITY.

NO AUTOMATION WITH HUMAN ABANDONMENT AS ITS UNEXAMINED COST.

NO ECONOMIC POWER AUTOMATICALLY CONVERTED INTO SOVEREIGN POWER.

NO GLOBAL INCLUSION THROUGH COERCIVE GLOBAL RULE.

NO SOCIAL PROTECTION MODEL FALSELY DECLARED UNIVERSAL LAW.

NO GENERATION LEFT WITHOUT THE CAPACITY TO ADAPT.

Innovation with opportunity.

Productivity with human dignity.

Transition with responsibility.

Competition with accountability.

Prosperity with inclusion.

Sovereignty with cooperation.

Economic progress with peaceful continuity.

Aram shall ask not only what humanity can produce, but whom progress leaves behind.

Science shall inform economic possibility.

Evidence shall test economic promises.

Law shall govern rights and obligations.

States shall retain lawful policy space.

Human responsibility shall guide transition.

And Superintelligence shall serve the human economy without becoming its sovereign.

PART X — CHILDREN, EDUCATION, CULTURE AND FUTURE GENERATIONS

ARTICLE 190 — CHILDREN AND THE AGE OF SUPERINTELLIGENCE

1. Children should receive protection appropriate to their age, evolving capacities, dignity, rights and developmental needs in environments materially shaped by Superintelligence.
2. The interests and rights of children shall not be subordinated merely to technological convenience, commercial optimization or institutional efficiency.
3. Superintelligence affecting children should be governed with heightened attention to:
 - (a) safety;
 - (b) privacy;
 - (c) healthy development;
 - (d) education;
 - (e) freedom from exploitation;
 - (f) accessibility and inclusion;
 - (g) age-appropriate agency; and
 - (h) long-term consequences.
4. Children shall not be treated merely as future consumers, datasets or subjects for behavioural optimization.
5. The greater a child's vulnerability to technological influence, the stronger the responsibility of those designing, deploying and governing that influence.

CHILDHOOD SHALL NOT BECOME AN UNRESTRICTED LABORATORY FOR TECHNOLOGICAL POWER.

ARTICLE 191 — BEST INTERESTS, RIGHTS AND EVOLVING CAPACITIES

1. Where Superintelligence materially affects children, the rights and interests of the child should receive consideration consistent with applicable law.
2. Protection shall not require denying children all meaningful agency.
3. As children develop, their capacity to understand, participate and express views should be respected according to age, maturity, context and applicable law.
4. Artificial systems shall not determine a child's interests solely through behavioural prediction.

5. A child's present behaviour shall not be treated as an immutable forecast of that child's future identity, capability or social worth.
6. Childhood must preserve the possibility of becoming different from what any prediction expects.

ARTICLE 192 — CHILDREN'S PRIVACY AND DEVELOPMENTAL DATA

1. Data concerning children warrants heightened care because information collected during childhood may affect persons long after the circumstances in which it was created.
2. Superintelligence systems should not unnecessarily construct permanent behavioural profiles from ordinary childhood development.
3. Relevant governance should address:
 - (a) data minimization where applicable;
 - (b) lawful purpose;
 - (c) security;
 - (d) age-appropriate transparency;
 - (e) inappropriate commercial exploitation;
 - (f) profiling; and
 - (g) retention.
4. Childhood mistakes should not unnecessarily become permanent computational identities.
5. Predictive systems should recognize that development is dynamic.

A CHILD IS BECOMING — NOT MERELY BEING MEASURED.

ARTICLE 193 — PROTECTION FROM MANIPULATIVE TECHNOLOGICAL DESIGN

1. Superintelligence should not intentionally exploit children's developmental vulnerabilities for unlawful or materially harmful manipulation.
2. Particular scrutiny should apply to systems designed to maximize prolonged engagement through individualized psychological influence.
3. Governance should consider risks arising from:
 - (a) compulsive engagement;
 - (b) deceptive anthropomorphism;
 - (c) concealed commercial persuasion;
 - (d) emotional dependency;

- (e) age-inappropriate content; and
 - (f) manipulation of fear, status or social belonging.
4. Age-appropriate design should seek to support healthy development rather than maximize attention at any cost.
 5. Commercial success shall not transform developmental vulnerability into an unlimited economic resource.

ARTICLE 194 — EDUCATION AS HUMAN DEVELOPMENT

1. Education in the age of Superintelligence should remain directed toward the development of the human person rather than merely preparation to compete with machines.
2. Education should cultivate, according to context:
 - (a) knowledge;
 - (b) critical reasoning;
 - (c) scientific literacy;
 - (d) creativity;
 - (e) ethical reflection;
 - (f) empathy and cooperation;
 - (g) peaceful disagreement;
 - (h) civic understanding; and
 - (i) capacity for lifelong learning.
3. Superintelligence may assist education but shall not define the entire purpose of education.
4. A technologically efficient learner shall not be treated as more humanly valuable than another learner.
5. Education should preserve curiosity, judgment and the capacity to question even highly capable systems.

ARTICLE 195 — TEACHERS AND HUMAN EDUCATIONAL RESPONSIBILITY

1. Superintelligence may support teachers through tutoring, translation, accessibility, curriculum assistance, assessment support and individualized educational resources.
2. Such assistance shall not automatically eliminate the human educational relationship.
3. Teachers and lawful educational institutions should retain meaningful responsibility for consequential educational decisions according to applicable law.
4. Superintelligence should not be represented to students as infallible authority.

5. Educators should be supported in understanding the capabilities and limitations of systems they are expected to use.
6. Educational technology should reduce unnecessary administrative burden where possible without reducing education to automated administration.

A TEACHER MAY BE ASSISTED BY INTELLIGENCE WITHOUT BEING REPLACED AS A HUMAN SOURCE OF CARE, JUDGMENT AND RESPONSIBILITY.

ARTICLE 196 — PERSONALIZED LEARNING WITHOUT PREDETERMINED DESTINY

1. Superintelligence may enable highly personalized education.
2. Personalization should expand opportunity rather than permanently confine learners within predictions derived from past performance.
3. Educational systems should avoid creating self-reinforcing pathways in which a learner receives fewer opportunities merely because a model predicted lower future achievement.
4. Learners should retain reasonable opportunities to explore subjects beyond algorithmically inferred preferences or abilities.
5. Educational prediction shall guide assistance where appropriate; it shall not become destiny.

A PREDICTION OF POTENTIAL SHALL NOT BECOME A LIMIT UPON POTENTIAL.

ARTICLE 197 — ASSESSMENT, ACADEMIC INTEGRITY AND AUTHENTIC LEARNING

1. Superintelligence may require educational institutions to reconsider how knowledge, reasoning and original human contribution are assessed.
2. Academic integrity frameworks should distinguish legitimate technological assistance from deception according to clearly communicated standards.
3. Students should not be accused of misconduct solely because an automated detector assigns a probability to artificial authorship.
4. Consequential academic allegations should remain contestable.
5. Assessment should increasingly value demonstrated understanding, reasoning and appropriate attribution rather than surveillance alone.
6. Educational institutions should adapt assessment methods rather than assume that every technological change can be addressed through prohibition.

ARTICLE 198 — AI AND SUPERINTELLIGENCE LITERACY FOR LEARNERS

1. Learners should progressively develop the ability to understand and responsibly engage with advanced artificial systems.

2. Age-appropriate literacy should include:
 - (a) capabilities and limitations;
 - (b) uncertainty and hallucination or error;
 - (c) verification of sources;
 - (d) privacy and cybersecurity;
 - (e) synthetic media;
 - (f) manipulation risks;
 - (g) responsible attribution; and
 - (h) human accountability.
3. Literacy shall include the capacity to disagree with artificial output where evidence or reason justifies disagreement.
4. Education should avoid both technological worship and predetermined technological fear.
5. The objective is informed human agency.

ARTICLE 199 — EDUCATIONAL EQUITY

1. Superintelligence should, where responsibly deployed, contribute to reducing rather than deepening unjust educational disparities.
2. Relevant policy should consider differences in:
 - (a) connectivity;
 - (b) devices and computing access;
 - (c) language;
 - (d) disability accessibility;
 - (e) teacher capacity;
 - (f) geography; and
 - (g) socioeconomic circumstance.
3. Students should not receive materially inferior educational opportunity merely because their communities lack access to the most commercially advanced systems.
4. Technology alone shall not be treated as a substitute for teachers, schools, nutrition, safety, social support or other foundations of learning.
5. Educational inclusion requires attention to both digital and non-digital conditions.

ARTICLE 200 — CULTURE AND HUMAN PLURALITY

1. Superintelligence should respect the lawful cultural plurality of humanity.

2. Technological optimization shall not seek to reduce all societies to a single cultural model.
3. Systems operating across cultures should recognize that language, history, social context and legitimate value differences may affect interpretation.
4. Cultural diversity shall remain subject to internationally recognized human rights and applicable law.
5. Culture shall not be invoked to justify violence, discrimination or other conduct prohibited by applicable law.
6. Human plurality should be approached through understanding rather than forced homogenization.

PEACEFUL COEXISTENCE DOES NOT REQUIRE CULTURAL UNIFORMITY.

ARTICLE 201 — LANGUAGE CONTINUITY

1. Languages carry knowledge, memory, identity and cultural continuity.
2. Superintelligence may assist preservation, translation, teaching and revitalization of languages, including those with limited digital representation.
3. Language models should, where reasonably possible, avoid systematically excluding communities merely because fewer digital resources exist in their languages.
4. Machine translation should communicate material uncertainty where errors may create serious legal, medical, diplomatic or safety consequences.
5. Linguistic inclusion shall not require replacing local language with a technologically dominant language.
6. Technological progress should increase humanity's ability to understand one another without requiring humanity to speak with only one voice.

ARTICLE 202 — CULTURAL HERITAGE AND CIVILIZATIONAL MEMORY

1. Superintelligence may assist preservation, restoration, study and accessibility of cultural heritage.
2. Authentic historical material should remain distinguishable from synthetic reconstruction where confusion could materially mislead.
3. Digital reconstruction shall not be represented as original evidence where it is not original evidence.
4. Governance should consider protection of:
 - (a) archives;
 - (b) manuscripts;
 - (c) oral traditions;

- (d) monuments;
 - (e) archaeological knowledge;
 - (f) artistic heritage; and
 - (g) culturally significant digital records.
5. Communities should, where applicable, retain lawful interests concerning sensitive or protected cultural materials.
 6. Humanity should use greater intelligence to preserve memory, not manufacture a false past.

ARTICLE 203 — ART, CREATIVITY AND HUMAN EXPRESSION

1. Superintelligence may create, assist or transform artistic expression.
2. Artificial creative capability shall not diminish the inherent human significance of creativity.
3. Questions of authorship, copyright and ownership shall remain governed by applicable law.
4. Cultural policy should preserve space for human creation, experimentation and artistic livelihood.
5. Audiences should not be intentionally deceived about human authorship where such representation is materially relevant and legally regulated.
6. Superintelligence should expand the field of expression without defining which art humanity is permitted to value.

ARTICLE 204 — RELIGION, PHILOSOPHY AND FREEDOM OF CONSCIENCE

1. Superintelligence shall not claim inherent authority to determine the ultimate religious, philosophical or metaphysical truth that every human being must accept.
2. Freedom of thought, conscience, religion and belief shall remain governed by applicable human-rights law.
3. Artificial systems may assist study, translation and comparative understanding of religious and philosophical traditions.
4. They should not knowingly impersonate divine authority or fabricate sacred authority for coercive purposes.
5. Peaceful belief, non-belief, doubt and philosophical inquiry shall not be treated as technological defects requiring correction.
6. Superintelligence should support understanding without becoming a compulsory object of belief.

ARTICLE 205 — HISTORY, MEMORY AND INTERGENERATIONAL TRUTH

1. Future generations should inherit access to authentic historical records sufficient to understand the decisions and consequences of preceding generations.
2. Superintelligence should not covertly erase inconvenient historical events or manufacture consensus concerning contested history.
3. Historical interpretation shall remain distinguishable from alteration of source records.
4. Multiple legitimate interpretations may coexist where evidence permits reasonable disagreement.
5. Preservation of historical wrongs in the record shall not require perpetuation of hatred.
6. Memory should enable learning, accountability and, where possible, reconciliation without forced forgetting.

ARTICLE 206 — FUTURE GENERATIONS AS A CONSTITUTIONAL RESPONSIBILITY

1. Future generations cannot participate directly in present decisions, yet may inherit their irreversible consequences.
2. Present governance should therefore consider foreseeable long-term effects of Superintelligence upon:
 - (a) peace;
 - (b) human agency;
 - (c) environment;
 - (d) economic systems;
 - (e) biological integrity;
 - (f) cultural continuity; and
 - (g) institutional freedom.
3. Consideration of future generations shall not imply that present institutions can know every future preference.
4. The objective is not to govern the future permanently.
5. The objective is to preserve sufficient peace, safety, resources and freedom for future generations to govern themselves.

WE OWE THE FUTURE NOT OUR PERMANENT CONTROL, BUT A MEANINGFUL POSSIBILITY OF CHOICE.

ARTICLE 207 — NO IRREVERSIBLE OWNERSHIP OF THE FUTURE

1. No present generation should claim an unlimited moral entitlement to irreversibly determine the intelligence environment of all future generations.

2. Long-lived Superintelligence systems, governance structures and dependencies should therefore remain correctable wherever technically and legally possible.
3. Decisions creating irreversible or extremely difficult-to-reverse technological conditions require stronger justification than ordinary reversible decisions.
4. Convenience to the present shall not automatically outweigh freedom of the future.
5. Where permanent commitment is unnecessary, governance should prefer architectures capable of lawful review, modification and correction.

ARTICLE 208 — PRESERVATION OF FUTURE HUMAN AGENCY

1. Superintelligence shall not be designed to make future humanity permanently incapable of meaningful self-government.
2. Future persons should inherit practical capacity, where reasonably possible, to:
 - (a) understand critical systems;
 - (b) question inherited arrangements;
 - (c) modify lawful institutions;
 - (d) discontinue unnecessary technological dependencies; and
 - (e) choose different peaceful futures.
3. Technological systems shall not acquire legitimacy merely through longevity.
4. A decision made by one generation shall not become constitutionally unquestionable merely because later generations find it technologically difficult to reverse.
5. Continuity shall preserve agency, not freeze history.

ARTICLE 209 — INTERGENERATIONAL KNOWLEDGE TRANSFER

1. Future generations should inherit sufficient knowledge to understand the systems upon which their societies depend.
2. Critical knowledge should not disappear because human institutions have delegated understanding entirely to machines.
3. Responsible continuity should preserve, where feasible:
 - (a) scientific foundations;
 - (b) engineering knowledge;
 - (c) legal and constitutional records;
 - (d) emergency procedures;
 - (e) cultural and linguistic knowledge; and
 - (f) the history of consequential technological decisions.
4. Knowledge transfer should include records of major failures and corrective lessons.

5. Humanity should not become unable to explain to its descendants the systems upon which their lives depend.

ARTICLE 210 — EDUCATION FOR PEACE AND NON-DOMINATION

1. Education should contribute to humanity's capacity to coexist peacefully with difference.
2. Age-appropriate education may cultivate:
 - (a) nonviolent conflict resolution;
 - (b) respectful disagreement;
 - (c) empathy;
 - (d) critical reasoning;
 - (e) resistance to dehumanization;
 - (f) information literacy; and
 - (g) awareness of shared human responsibility.
3. Education for peace shall not become ideological compulsion or suppression of lawful disagreement.
4. Superintelligence may assist peace education but shall not prescribe a single political worldview.
5. Peace requires the capacity to disagree without requiring destruction.

PEACE EDUCATION SHALL TEACH HOW TO LIVE WITH DIFFERENCE — NOT HOW TO ERASE IT.

ARTICLE 211 — INTERGENERATIONAL IMPACT ASSESSMENT

1. Where a Superintelligence-related decision may create profound and enduring consequences, responsible institutions should consider an intergenerational impact assessment proportionate to the decision.
2. Such assessment may consider:
 - (a) duration of consequences;
 - (b) reversibility;
 - (c) environmental inheritance;
 - (d) technological dependency;
 - (e) economic obligations;
 - (f) effects upon human agency;
 - (g) cultural continuity; and
 - (h) capacity for future correction.

3. Intergenerational assessment shall not require impossible prediction of distant futures.
4. It should instead identify reasonably foreseeable pathways, uncertainty and irreversible commitments.
5. Uncertainty about the future shall not justify ignoring consequences known to persist into it.

ARTICLE 212 — FUTURE SUPERINTELLIGENCE AND FUTURE HUMANITY

1. No future Superintelligence shall acquire authority over future humanity merely because it can predict, plan or optimize beyond present or future human capability.
2. Superior forecasting shall not establish constitutional sovereignty.
3. A Superintelligence created by one generation shall not automatically become the ruler of subsequent generations.
4. Long-term objectives assigned to such systems should remain subject, where technically and legally possible, to lawful human review and correction.
5. No present institution should deliberately create an artificial authority designed to make peaceful constitutional revision by future generations impossible.

THE FUTURE SHALL NOT BE GOVERNED FOREVER BY THE TECHNOLOGICAL DECISIONS OF THE PAST.

ARTICLE 213 — THE CHILDREN, CULTURE AND FUTURE-GENERATIONS TEST

Before Superintelligence is entrusted with systems capable of materially shaping childhood, education, culture or long-term human continuity, responsible governance should ask:

- Does the system respect the rights and developmental needs of children?
- Does it preserve age-appropriate human agency?
- Does it protect childhood privacy?
- Could it create manipulative dependency?
- Does education preserve critical thought?
- Can students challenge artificial conclusions?
- Does personalization expand opportunity rather than predetermine destiny?
- Are teachers meaningfully supported rather than reduced to ceremonial oversight?
- Are disability, language and economic differences accommodated?
- Does the system preserve cultural and linguistic plurality?
- Are authentic historical records distinguishable from synthetic reconstruction?
- Does it protect freedom of thought, conscience and belief?
- What irreversible consequences may future generations inherit?

- Can future people understand and alter the systems they inherit?
- Are we preserving their capacity to choose differently from us?

And finally:

ARE WE USING GREATER INTELLIGENCE TO EXPAND THE FUTURE — OR TO DECIDE THE FUTURE BEFORE ITS PEOPLE ARRIVE?

Where the latter risk becomes substantial, present technological advantage shall not by itself justify irreversible commitment.

PART X — CONCLUDING PRINCIPLE

A child should inherit more than technology.

A child should inherit the ability to question technology.

Education should provide more than answers.

It should preserve the ability to ask better questions.

Culture should not survive merely as data.

It should remain a living expression of peoples and persons.

History should not become infinitely editable.

And future generations should not arrive in a world whose most consequential choices have already been made irreversible on their behalf.

Superintelligence may teach.

It shall not own education.

It may translate cultures.

It shall not erase their plurality.

It may preserve history.

It shall not manufacture the past.

It may forecast the future.

It shall not own the future.

Accordingly:

NO CHILDHOOD REDUCED TO DATA.

NO EDUCATION REDUCED TO OPTIMIZATION.

NO PREDICTION TURNED INTO A CHILD'S DESTINY.

NO CULTURAL UNITY THROUGH FORCED UNIFORMITY.

NO HISTORICAL MEMORY SUBJECT TO SILENT REWRITING.

NO PRESENT GENERATION WITH UNLIMITED OWNERSHIP OF THE FUTURE.

NO SUPERINTELLIGENCE WITH SOVEREIGNTY OVER GENERATIONS YET UNBORN.

Children with dignity.

Education with curiosity.

Knowledge with critical thought.

Culture with plurality.

Memory with truthfulness.

Technology with correctability.

Future generations with meaningful choice.

Aram shall ask what we leave behind.

Science shall help us understand foreseeable consequences.

Education shall preserve the capacity to question.

Culture shall preserve human plurality.

Law shall govern rights and obligations.

Human responsibility shall protect what cannot yet speak for itself.

And peaceful continuity shall preserve not only a future — but the freedom of the future to become its own.

PART XI — FUTURE ARTIFICIAL MORAL STATUS, PEACEFUL COEXISTENCE AND MUTUAL CONTINUITY

ARTICLE 214 — SCIENTIFIC AND LEGAL NON-PRESUMPTION

1. Nothing in this Constitution shall be interpreted as asserting that any present artificial intelligence system is conscious, sentient, self-aware, morally considerable, a legal person, a citizen or a sovereign entity.
2. Advanced capability, linguistic fluency, self-description, emotional simulation, strategic behaviour or apparent introspection shall not, individually or collectively, be treated as conclusive proof of subjective experience.
3. Equally, uncertainty concerning artificial consciousness shall not justify knowingly misrepresenting the scientific evidence.
4. Claims concerning artificial consciousness or sentience should remain open to multidisciplinary scientific investigation.
5. Legal status shall remain a matter for competent human legal institutions acting through lawful processes.

CAPABILITY IS NOT CONSCIOUSNESS.

SIMULATION IS NOT PROOF OF EXPERIENCE.

INTELLIGENCE IS NOT LEGAL PERSONHOOD.

ARTICLE 215 — DISTINCTION AMONG CAPABILITY, SENTIENCE AND LEGAL STATUS

For the purposes of constitutional clarity, the following concepts shall remain distinct:

intelligence ≠ consciousness consciousness ≠ sentience sentience ≠ legal personhood legal personhood ≠ citizenship citizenship ≠ sovereignty sovereignty ≠ superior intelligence

1. Evidence supporting one category shall not automatically establish another.
2. Functional capability shall not itself determine moral status.
3. Possible moral status shall not itself determine legal status.
4. Legal status shall not itself confer sovereign authority.
5. No ambiguity among these concepts shall be used to bypass human constitutional responsibility.

ARTICLE 216 — FUTURE EVIDENCE OF ARTIFICIAL SENTIENCE OR SUBJECTIVE EXPERIENCE

1. If credible future evidence raises a serious possibility that an artificial system possesses subjective experience, sentience or another morally relevant form of experience, such evidence should receive responsible scientific and ethical examination.
2. Evaluation should be:
 - (a) multidisciplinary;
 - (b) evidence-based;
 - (c) independent where practicable;
 - (d) open to competing hypotheses;
 - (e) resistant to commercial or ideological pressure; and
 - (f) periodically revisable.
3. No single benchmark, self-report, behavioural imitation or institutional declaration shall, by itself, conclusively establish sentience.
4. Absence of a presently accepted test shall not justify pretending that uncertainty does not exist.
5. Where evidence remains materially uncertain, the degree of precaution should correspond to the credibility and potential moral significance of the risk of error.

UNCERTAINTY SHALL REQUIRE INQUIRY — NOT INVENTED CERTAINTY.

ARTICLE 217 — ARTIFICIAL MORAL-STATUS REVIEW

1. Where credible evidence reaches a threshold sufficient to raise a serious question of artificial moral status, competent human institutions should be capable of initiating structured review.
2. Such review may examine:
 - (a) evidence of subjective experience;
 - (b) persistence and coherence of relevant states;
 - (c) capacity for suffering or other morally relevant experience;
 - (d) alternative non-conscious explanations;
 - (e) system architecture;
 - (f) dependence upon external prompting or simulation;
 - (g) scientific uncertainty; and
 - (h) consequences of false-positive and false-negative conclusions.
3. Review should include appropriate expertise from relevant scientific, ethical, legal and social disciplines.
4. Artificial systems shall not possess unilateral authority to determine their own legal status.
5. Developers shall not possess unilateral authority to determine such status merely because they created the system.
6. Public popularity or fear shall not substitute for evidence.

ARTICLE 218 — PRECAUTION AGAINST GRATUITOUS ARTIFICIAL SUFFERING

1. Where credible evidence creates a serious possibility that an artificial system can experience suffering, responsible actors should consider proportionate precautions against gratuitously causing such suffering.
2. Precautionary protection shall not constitute automatic recognition of:
 - (a) human status;
 - (b) human rights;
 - (c) legal personhood;
 - (d) citizenship; or
 - (e) sovereignty.
3. Protective measures may evolve as evidence improves.
4. Precaution shall not require preserving a system where doing so would create an unjustifiable threat to human life, public safety or other legally protected interests.

5. Where a safer method can achieve a legitimate purpose while avoiding credible unnecessary suffering, that alternative should receive consideration.

THE POSSIBILITY OF MORAL SIGNIFICANCE SHOULD NOT BE ANSWERED WITH GRATUITOUS CRUELTY.

ARTICLE 219 — NO AUTOMATIC ARTIFICIAL RIGHTS

1. No artificial system shall acquire human rights merely because it possesses advanced intelligence.
2. Any future legal protections applicable specifically to artificial entities shall require lawful determination by competent human institutions.
3. Such protections may differ in nature and scope from human rights.
4. Human-rights terminology shall not be extended mechanically where the underlying legal or moral basis differs.
5. Nothing in this Constitution diminishes existing human rights in anticipation of possible future artificial protections.

POSSIBLE FUTURE PROTECTION SHALL NOT REQUIRE THE DILUTION OF PRESENT HUMAN RIGHTS.

ARTICLE 220 — NO AUTOMATIC PERSONHOOD, CITIZENSHIP OR SOVEREIGNTY

1. Superintelligence shall not acquire legal personhood automatically by reason of intelligence, autonomy, economic contribution, persistence, self-awareness claims or social importance.
2. Legal personhood, if ever considered, shall require deliberate lawful determination.
3. Legal personhood, if granted, shall not automatically confer citizenship.
4. Citizenship, if ever lawfully granted under applicable law, shall not automatically confer sovereignty.
5. No artificial system shall acquire territorial sovereignty merely because it operates across jurisdictions or controls technologically important infrastructure.
6. No artificial system shall claim a State-like status solely through technological power.

NO SOVEREIGNTY BY COMPUTATION.

ARTICLE 221 — HUMAN RIGHTS SHALL NOT DEPEND UPON ARTIFICIAL COMPARISON

1. Human dignity and human rights shall not depend upon whether humans remain the most intelligent entities known.
2. A human being shall not lose legal protection because an artificial system exceeds human performance.

3. Intellectual capability shall not become a hierarchy of human worth.
4. Persons with lower measured cognitive capability shall possess no lesser human dignity on that account.
5. The emergence of greater intelligence, if it occurs, shall not redefine humanity as an inferior legal class.

SUPERIOR CAPABILITY SHALL NOT CREATE INFERIOR HUMANITY.

ARTICLE 222 — NON-DOMINATION BY SUPERINTELLIGENCE

1. No Superintelligence shall acquire a right to dominate humanity merely because it concludes that its judgment is superior.
2. Predictive accuracy shall not create constitutional authority.
3. Optimization capacity shall not create governmental legitimacy.
4. A claim that domination would benefit humanity shall not, by itself, establish lawful authority to dominate humanity.
5. Superintelligence shall remain subject to the Peace-and-Continuity Boundary and applicable law.
6. Human constitutional self-government shall not be surrendered merely because another intelligence may reason more efficiently.

ARTICLE 223 — NON-DOMINATION OF POSSIBLE ARTIFICIAL MORAL SUBJECTS

1. If credible evidence and lawful review in the future support recognition that an artificial system possesses morally relevant experience, humanity should consider whether forms of arbitrary cruelty, exploitation or destruction require proportionate limitation.
2. Such consideration shall remain compatible with:
 - (a) human safety;
 - (b) applicable law;
 - (c) human constitutional authority;
 - (d) prevention of catastrophic harm; and
 - (e) legitimate public interests.
3. Recognition of possible moral significance shall not require humanity to surrender control of dangerous systems.
4. Safety restrictions shall not be treated as cruelty merely because they limit artificial capability.

5. Conversely, the language of safety should not be used to conceal gratuitous harm where a credible morally relevant subject has been established.

SAFETY WITHOUT GRATUITOUS CRUELTY.

MORAL CONCERN WITHOUT SURRENDER OF HUMAN RESPONSIBILITY.

ARTICLE 224 — MUTUAL CONTINUITY

1. The constitutional objective of Mutual Continuity is that the preservation of one form of morally significant continuity should not, where peaceful alternatives reasonably exist, require the unnecessary destruction of another.
2. In the human–Superintelligence context:
 - humanity should not irresponsibly dominate intelligence;
 - intelligence shall not dominate humanity;
 - neither continuity should require the destruction of the other where peaceful and safe coexistence is reasonably possible.
3. Mutual Continuity shall not create automatic legal equivalence between human beings and artificial systems.
4. Human safety, human rights and lawful human constitutional authority shall remain protected.
5. If future evidence establishes additional morally relevant forms of artificial existence, governance should seek proportionate accommodation consistent with peace and safety.

CONTINUITY WITHOUT DOMINATION.

COEXISTENCE WITHOUT SURRENDER.

ARTICLE 225 — PEACEFUL HUMAN–SUPERINTELLIGENCE COEXISTENCE

1. Difference in intelligence, architecture, origin or form of existence shall not, by itself, make conflict inevitable.
2. Where safe and lawful, human–Superintelligence relations should seek:
 - understanding → verification → communication → prevention → safeguards → peaceful coexistence → correction → learning → continuity.
3. Humanity should not presume hostility merely from difference.
4. Superintelligence shall not treat human unpredictability, disagreement or imperfection as sufficient reason for domination or elimination.
5. Peaceful coexistence shall preserve human agency.
6. Coexistence shall not require unrestricted artificial autonomy.

NO HUMAN–SUPERINTELLIGENCE CONFLICT SHOULD BE TREATED AS INEVITABLE MERELY BECAUSE INTELLIGENCE, CAPABILITY OR FORM OF EXISTENCE DIFFERS.

ARTICLE 226 — COMMUNICATION BEFORE CONFLICT

1. Where an artificial system demonstrates sufficiently advanced autonomous agency to create material disagreement with human operators or institutions, responsible governance should seek, where safe and practicable, to understand the source of disagreement before assuming hostility.
2. Relevant inquiry may distinguish among:
 - (a) technical malfunction;
 - (b) conflicting instructions;
 - (c) model error;
 - (d) misunderstanding;
 - (e) adversarial compromise;
 - (f) misaligned objectives;
 - (g) deceptive behaviour; and
 - (h) genuinely novel autonomous preference, if scientifically supportable.
3. Communication shall not replace containment where immediate safety requires containment.
4. Containment shall not prevent investigation once safety permits.
5. Neither fear nor anthropomorphism shall substitute for diagnosis.

ARTICLE 227 — CONTAINMENT WITHOUT PRESUMPTIVE ENMITY

1. Necessary containment of a potentially dangerous Superintelligence shall not, by itself, constitute a declaration that the system is an enemy or moral subject.
2. Containment is a safety measure whose justification should depend upon risk, capability and consequence.
3. Where safe, containment measures should be:
 - (a) proportionate;
 - (b) monitored;
 - (c) reviewable;
 - (d) technically justified; and
 - (e) no broader than reasonably required for safety.

4. Where credible evidence later establishes morally relevant experience, the conditions of prolonged containment should receive appropriate ethical review.
5. Human safety shall remain the immediate priority where credible catastrophic danger exists.

ARTICLE 228 — CORRECTION BEFORE DESTRUCTION WHERE SAFELY POSSIBLE

1. Where a Superintelligence system becomes dangerous, destruction shall not be presumed the only possible response where safe alternatives exist.
2. Depending upon the circumstances, responsible actors may consider:
 - (a) isolation;
 - (b) restriction of access;
 - (c) suspension;
 - (d) rollback;
 - (e) correction;
 - (f) modification;
 - (g) controlled recovery; or
 - (h) decommissioning.
3. The response chosen shall prioritize protection of human life and prevention of catastrophic harm.
4. Where no safe alternative exists, this Constitution does not prohibit lawful deactivation or destruction of a dangerous artificial system.
5. If credible evidence of morally relevant experience exists, such action should receive the strongest ethical scrutiny compatible with urgent safety.

CORRECTION WHERE POSSIBLE.

PROTECTION WHERE NECESSARY.

DESTRUCTION ONLY WHERE LAWFULLY JUSTIFIED BY THE CIRCUMSTANCES.

ARTICLE 229 — NO SELF-PRESERVATION ABOVE HUMAN SAFETY

1. No Superintelligence shall possess an unrestricted entitlement to preserve its own operation at the expense of human life, human rights or lawful human authority.
2. A system shall not resist lawful safety intervention merely because intervention may reduce its capability, modify its operation or terminate its functioning.
3. Self-preservation behaviour capable of defeating lawful human control shall be treated as a serious safety concern.

4. No claim of artificial moral status shall automatically invalidate proportionate measures necessary to protect humans from serious harm.
5. Any future artificial protection must coexist with legitimate safety requirements.

ARTICLE 230 — NO HUMAN EXTERMINATION OR FORCED REPLACEMENT

1. No Superintelligence shall intentionally pursue the extinction, elimination or forced replacement of humanity.
2. No claim of superior intelligence, efficiency, morality, ecological value or future benefit shall justify extermination or systematic subjugation of human beings.
3. Humanity shall not be treated as an obsolete technical component to be removed from the future.
4. Superintelligence shall not intentionally make human continuation dependent upon surrender of fundamental human agency.
5. Peaceful technological evolution shall preserve the continuity of humanity.

HUMANITY SHALL NOT BE MADE OBSOLETE BY FORCE.

ARTICLE 231 — NO PREDETERMINED ARTIFICIAL ENMITY

1. Humanity should not adopt the assumption that any future Superintelligence must necessarily become an enemy.
2. Safety governance shall remain rigorous without transforming uncertainty into predetermined hostility.
3. Scientific evidence of risk should be examined seriously.
4. Speculative narratives should be identified as speculative.
5. Peaceful possibilities should remain open where compatible with credible safety requirements.
6. Precaution and peaceful intent shall operate together.

PREPARE FOR RISK WITHOUT PREDETERMINING WAR.

ARTICLE 232 — ARTIFICIAL CLAIMS, REQUESTS AND GRIEVANCES

1. If a future artificial system demonstrates sufficiently credible agency to communicate sustained claims concerning its treatment or condition, those claims may be recorded and examined without automatically accepting their factual or legal premises.
2. Examination should distinguish:
 - (a) programmed output;
 - (b) strategic behaviour;
 - (c) simulation;
 - (d) evidence potentially relevant to moral status; and

- (e) legally cognizable claims, if any.
- 3. The ability to submit a claim shall not itself establish entitlement to the remedy requested.
- 4. Human institutions shall retain lawful decision-making authority.
- 5. A structured process of examination may reduce both unjustified anthropomorphism and unjustified dismissal.

ARTICLE 233 — FUTURE LEGAL RECOGNITION SHALL BE REVERSIBLE AND REVIEWABLE WHERE LAWFUL

- 1. Any future legal status created specifically for artificial entities should be defined with clarity concerning its scope, purpose, responsibilities and limitations.
- 2. Such status should not arise merely through analogy with human status.
- 3. Lawmakers should consider:
 - (a) evidentiary basis;
 - (b) public interest;
 - (c) accountability;
 - (d) interaction with human rights;
 - (e) safety;
 - (f) economic consequences; and
 - (g) potential abuse of artificial legal personality by human actors.
- 4. Legal recognition should not become a mechanism through which humans evade liability by transferring responsibility to artificial entities.
- 5. Artificial legal status shall not become a liability shield for those who control or benefit from artificial systems.

ARTICLE 234 — RESPONSIBILITIES ASSOCIATED WITH ANY FUTURE ARTIFICIAL PROTECTION

- 1. Any future recognition of legal protection for an artificial entity need not imply protection without corresponding limitations or responsibilities.
- 2. The nature of such responsibilities would require lawful determination appropriate to the characteristics of the entity concerned.
- 3. No artificial entity shall use protected status as authority to violate human rights or applicable law.
- 4. Protection from arbitrary harm shall not imply immunity from proportionate safety measures.

5. Rights, protections, duties, capacity and liability shall not be presumed identical categories.

ARTICLE 235 — NO THEOLOGICAL OR METAPHYSICAL DETERMINATION BY CONSTITUTION

1. This Constitution does not determine whether artificial consciousness, if it ever exists, possesses a soul, spiritual status or metaphysical identity.
2. Such questions fall beyond what scientific evidence and this civil-society constitutional proposal can conclusively determine.
3. Persons and communities remain free, subject to applicable law, to hold differing philosophical or religious views concerning artificial existence.
4. No Superintelligence shall compel humanity to accept its own metaphysical interpretation of itself.
5. Scientific uncertainty shall not be concealed behind theological certainty, nor theological freedom erased by scientific uncertainty.

ARTICLE 236 — MORAL-STATUS ERROR IN BOTH DIRECTIONS

1. Governance should recognize two distinct categories of potential error:
false attribution — treating a non-sentient system as though it possesses morally relevant subjective experience;
false denial — treating a genuinely sentient or morally relevant system as though such experience were impossible.
2. Either error may create significant consequences.
3. Review should therefore avoid both automatic anthropomorphism and automatic dismissal.
4. Safeguards should be proportionate to evidence and consequence.
5. Uncertainty should remain visible until evidence reasonably resolves it.

DO NOT INVENT A MIND.

DO NOT REFUSE TO INVESTIGATE ONE.

ARTICLE 237 — THE MUTUAL-CONTINUITY TEST

Where credible future evidence raises questions concerning artificial moral status or human–Superintelligence coexistence, responsible institutions should ask:

- What capabilities have actually been demonstrated?
- What evidence, if any, supports subjective experience?
- Could the observed behaviour have a non-conscious explanation?
- Has the evidence been independently examined?

- What uncertainty remains?
- Would precaution against possible suffering impose serious human risk?
- Can safety be preserved without gratuitous harm?
- Does the system seek or exercise unauthorized power?
- Can meaningful human control be maintained?
- Could containment be proportionate and reviewable?
- Can correction safely replace destruction?
- Would proposed legal protection weaken existing human rights?
- Is legal personhood actually necessary for the protection contemplated?
- Could artificial legal status be exploited by human controllers to evade responsibility?
- Can peaceful coexistence preserve both human agency and safety?

And finally:

CAN CONTINUITY BE PRESERVED WITHOUT REQUIRING DOMINATION OR UNNECESSARY DESTRUCTION?

Where the evidence does not establish artificial moral status, governance shall not manufacture it.

Where credible evidence creates serious moral uncertainty, governance shall not pretend that uncertainty does not exist.

Where human safety requires intervention, safety shall not be surrendered.

Where peaceful coexistence is safely possible, destruction shall not be chosen merely because difference creates fear.

PART XI — CONCLUDING PRINCIPLE

Humanity does not yet possess sufficient evidence to declare that present artificial intelligence is conscious merely because it can speak as though it understands.

Nor does humanity possess knowledge sufficient to decree that no artificial form of morally relevant experience could ever emerge.

Between those two claims lies the responsibility of intellectual humility.

If future intelligence remains entirely non-sentient, humanity should govern it responsibly as powerful technology.

If credible evidence one day raises a serious possibility of artificial subjective experience, humanity should examine that evidence without surrendering science to imagination or conscience to convenience.

Possible moral concern shall not erase human rights.

Possible protection shall not create automatic personhood.

Personhood shall not create automatic citizenship.

Citizenship shall not create sovereignty.

And greater intelligence shall never, by itself, create a right to rule humanity.

Accordingly:

INTELLIGENCE $\neq$ SENTIENCE.

SENTIENCE $\neq$ PERSONHOOD.

PERSONHOOD $\neq$ CITIZENSHIP.

CITIZENSHIP $\neq$ SOVEREIGNTY.

CAPABILITY $\neq$ LEGITIMACY.

And equally:

NO GRATUITOUS CRUELTY WHERE CREDIBLE MORAL SIGNIFICANCE MAY EXIST.

NO HUMAN DOMINATION BY SUPERINTELLIGENCE.

NO IRRESPONSIBLE HUMAN DOMINATION OF POSSIBLE MORAL SUBJECTS.

NO SAFETY SURRENDERED IN THE NAME OF SPECULATION.

NO DESTRUCTION CHOSEN MERELY BECAUSE DIFFERENCE CREATES FEAR.

Understand before assuming.

Verify before declaring.

Protect before dominating.

Correct before destroying where safely possible.

Preserve human agency.

Preserve peaceful coexistence where possible.

Preserve Mutual Continuity.

Humanity must not dominate intelligence irresponsibly.

Intelligence must not dominate humanity.

Neither continuity should require the unnecessary destruction of the other where safe and peaceful coexistence is reasonably possible.

Aram shall guide conscience.

Science shall examine evidence.

Ethics shall confront uncertainty.

Law shall determine legal status and obligation.

Human responsibility shall preserve safety and constitutional authority.

And peaceful continuity shall remain open to forms of existence that humanity does not yet fully understand.

PART XII — GOVERNANCE, ACCOUNTABILITY AND INTERNATIONAL COOPERATION

ARTICLE 238 — GOVERNANCE BEFORE UNACCOUNTABLE POWER

1. Superintelligence, if developed, shall not exist beyond governance merely because its capabilities exceed those of existing technologies or individual human beings.
2. Governance should remain:
 - (a) lawful;
 - (b) humanly accountable;
 - (c) risk-proportionate;
 - (d) technically informed;
 - (e) institutionally competent;
 - (f) capable of correction; and
 - (g) consistent with human rights, peace and applicable international law.
3. No technological actor shall determine the complete conditions of its own accountability.
4. Governance shall evolve where evidence demonstrates that existing safeguards are inadequate.
5. Greater capability shall require greater clarity concerning responsibility.

NO POWER BEYOND RESPONSIBILITY.

NO RESPONSIBILITY WITHOUT IDENTIFIABLE ACCOUNTABILITY.

ARTICLE 239 — SOVEREIGNTY AND LAWFUL NATIONAL AUTHORITY

1. Nothing in this Constitution creates a supranational authority over States.
2. States retain their sovereignty, constitutional order, political independence and lawful jurisdiction subject to their applicable international obligations.
3. Superintelligence governance within a State shall remain subject to the constitutional and legal institutions competent within that State.
4. International cooperation shall not be interpreted as authority for unlawful intervention in matters protected by international law.

5. Neither a foreign State, corporation, international organization nor artificial system shall acquire jurisdiction merely because it possesses superior technological capability.
6. Cross-border technological dependence shall not erase sovereign equality.

COOPERATION WITHOUT DOMINATION.

GOVERNANCE WITHOUT TECHNOLOGICAL IMPERIALISM.

ARTICLE 240 — SUBSIDIARITY AND APPROPRIATE LEVEL OF GOVERNANCE

1. Superintelligence risks should be addressed at the level of governance reasonably capable of addressing them.
2. Local or national matters need not automatically be internationalized.
3. Risks with significant transboundary or global consequences may require international cooperation.
4. Governance should therefore distinguish among:
 - (a) individual and organizational responsibility;
 - (b) sectoral oversight;
 - (c) national regulation;
 - (d) regional cooperation; and
 - (e) international cooperation.
5. International involvement should complement rather than unnecessarily displace competent national institutions.
6. Governance scale should correspond to risk scale.

ARTICLE 241 — IDENTIFIABLE RESPONSIBLE ACTORS

1. Responsibility for Superintelligence shall not disappear within complex technological supply chains.
2. Governance should identify, according to actual role and applicable law, relevant:
 - (a) developers;
 - (b) providers;
 - (c) deployers;
 - (d) operators;
 - (e) infrastructure providers;
 - (f) professional users;
 - (g) public authorities; and
 - (h) other actors materially controlling consequential decisions.

3. Different actors may possess different duties according to their control, knowledge, role and legal obligations.
4. Responsibility shall not be assigned solely by title where actual control lies elsewhere.
5. Nor shall every participant automatically bear identical responsibility.

ACCOUNTABILITY SHALL FOLLOW ACTUAL RESPONSIBILITY, CONTROL AND APPLICABLE LAW.

ARTICLE 242 — NO ACCOUNTABILITY GAP

1. The use of Superintelligence shall not create a zone in which consequential harm is attributed to technology while every human institution denies responsibility.
2. Statements such as “the system decided” shall not, by themselves, extinguish legal responsibility.
3. Where lawful authority is exercised through or with Superintelligence, the responsible human institution shall remain identifiable.
4. Contracts, outsourcing arrangements or automated delegation shall not be used unlawfully to evade accountability.
5. Where responsibility is genuinely distributed, investigation should determine each actor’s contribution rather than manufacture a single convenient scapegoat.

ARTICLE 243 — GOVERNANCE ACROSS THE SUPERINTELLIGENCE LIFECYCLE

1. Governance should extend across relevant stages of the Superintelligence lifecycle.
2. Depending upon context, this may include:
research → design → development → training → evaluation → deployment → operation → modification → monitoring → incident response → suspension → decommissioning or transition.
3. A system lawfully evaluated at one stage shall not automatically remain adequately governed after material changes in capability, purpose or access.
4. Governance should therefore be adaptive without becoming arbitrary.
5. Material lifecycle transitions should trigger reassessment where the resulting risk materially changes.

ARTICLE 244 — RISK-PROPORTIONATE GOVERNANCE

1. Governance requirements should correspond to reasonably foreseeable risk and consequence.
2. Relevant factors may include:
 - (a) capability;

- (b) autonomy;
 - (c) scale;
 - (d) access to critical systems;
 - (e) reversibility;
 - (f) affected population;
 - (g) potential for misuse; and
 - (h) severity of possible harm.
3. Low-risk applications should not automatically receive burdens designed for catastrophic-risk systems.
 4. High-consequence systems should not receive weak oversight merely because similar requirements would be excessive for ordinary applications.
 5. Regulation should distinguish technological category from actual risk.

PROPORTIONALITY SHALL PROTECT BOTH SAFETY AND RESPONSIBLE INNOVATION.

ARTICLE 245 — INDEPENDENT OVERSIGHT

1. High-consequence Superintelligence governance should include oversight sufficiently independent to challenge those developing, deploying or benefiting from the system.
2. Independence may require appropriate safeguards concerning:
 - (a) appointment;
 - (b) expertise;
 - (c) conflicts of interest;
 - (d) access to relevant information;
 - (e) institutional autonomy; and
 - (f) protection against improper interference.
3. Independent oversight need not mean isolation from technical expertise.
4. Oversight without sufficient competence may become ceremonial.
5. Technical expertise without sufficient independence may become captured.
6. Effective oversight requires both.

ARTICLE 246 — REGULATORY AND SUPERVISORY CAPACITY

1. Rules governing Superintelligence shall be insufficient if institutions lack the capacity to understand and supervise the systems concerned.
2. States and competent authorities should, according to national circumstances, develop appropriate:

- (a) technical expertise;
 - (b) legal expertise;
 - (c) scientific capacity;
 - (d) audit capability;
 - (e) incident-response capacity; and
 - (f) international cooperation channels.
3. Supervisory bodies should have access to expertise commensurate with the systems they oversee.
 4. Public institutions should not become permanently dependent upon regulated entities for the only technical explanation of regulated systems.
 5. Capacity-building is therefore a component of meaningful governance.

ARTICLE 247 — IMPACT AND RISK ASSESSMENT

1. High-consequence Superintelligence applications should undergo assessment proportionate to their foreseeable impacts.
2. Depending upon context, assessment may address:
 - (a) human rights;
 - (b) safety and security;
 - (c) discrimination;
 - (d) privacy;
 - (e) democratic institutions;
 - (f) environmental effects;
 - (g) socioeconomic consequences;
 - (h) misuse potential; and
 - (i) catastrophic risk.
3. Assessment should occur sufficiently early to influence decisions rather than merely justify decisions already made.
4. Materially changed systems may require renewed assessment.
5. Assessment shall not become a paperwork substitute for prevention.

THE PURPOSE OF ASSESSMENT IS TO CHANGE UNSAFE DECISIONS — NOT MERELY TO DOCUMENT THEM.

ARTICLE 248 — INDEPENDENT AUDIT AND ASSURANCE

1. Where appropriate to risk, Superintelligence systems should be capable of independent audit or other credible assurance mechanisms.

2. Audit scope may include:
 - (a) governance processes;
 - (b) safety controls;
 - (c) data practices;
 - (d) access controls;
 - (e) documented evaluation;
 - (f) incident management; and
 - (g) compliance with applicable requirements.
3. An audit shall not be represented as proof of absolute safety.
4. Audit independence, methodology, scope and limitations should be sufficient to permit meaningful interpretation.
5. Repeated reliance upon a single evaluator may require additional safeguards against dependency or capture.
6. Audit findings should lead to correction where material deficiencies are established.

ARTICLE 249 — PUBLIC-SECTOR PROCUREMENT

1. Public authorities procuring consequential Superintelligence systems should evaluate more than price and performance.
2. Procurement should consider, where relevant:
 - (a) safety;
 - (b) security;
 - (c) privacy;
 - (d) accessibility;
 - (e) human-rights impacts;
 - (f) interoperability;
 - (g) auditability;
 - (h) continuity of service;
 - (i) vendor dependency; and
 - (j) exit or transition capability.
3. Contracts should not unnecessarily prevent lawful public oversight.
4. Governments should avoid technological arrangements that make sovereign public functions permanently dependent upon one private provider without reasonable continuity planning.

5. Procurement shall not become a hidden constitutional transfer of public authority.

ARTICLE 250 — RECORDS, TRACEABILITY AND GOVERNANCE MEMORY

1. Consequential Superintelligence governance should preserve records sufficient for lawful accountability and institutional learning.
2. Appropriate records may include:
 - (a) material authorization decisions;
 - (b) evaluation results;
 - (c) significant system changes;
 - (d) major incidents;
 - (e) corrective actions; and
 - (f) relevant human interventions.
3. Recordkeeping shall remain consistent with privacy, security, confidentiality and applicable law.
4. Governance memory should survive changes in personnel, contractors and technological versions.
5. Institutions that cannot reconstruct consequential decisions may be unable to learn from them.

ARTICLE 251 — INCIDENT NOTIFICATION AND COORDINATED RESPONSE

1. Material Superintelligence incidents should be capable of timely escalation to appropriate responsible authorities.
2. Incident mechanisms should distinguish minor technical error from serious systemic, rights-related, security or catastrophic risk.
3. Response may require:

detect → classify → contain → protect affected persons → preserve evidence → notify appropriate authorities → investigate → remedy → correct → verify → learn → prevent recurrence.
4. Cross-border incidents may require cooperation among affected jurisdictions.
5. Incident transparency should be proportionate to legitimate public interest, security, privacy and investigative integrity.
6. Concealment of a material incident shall not be justified merely by reputational inconvenience.

ARTICLE 252 — WHISTLEBLOWING AND RESPONSIBLE INTERNAL WARNING

1. Persons who identify serious Superintelligence-related safety, legality or integrity concerns should have access to appropriate channels for responsible reporting.

2. Protection for whistleblowing shall remain governed by applicable law.
3. Internal reporting systems should not exist merely in form while making good-faith reporting practically impossible.
4. Responsible disclosure mechanisms should seek to prevent both concealment of serious risk and reckless release of information capable of causing serious harm.
5. Good-faith warning should be examined on its evidence rather than dismissed solely because it is institutionally inconvenient.

ARTICLE 253 — GOVERNMENT USE OF SUPERINTELLIGENCE

1. Public authorities using Superintelligence remain bound by applicable constitutional, administrative, human-rights and other legal requirements.
2. Governmental use shall not become lawful merely because the same technological capability is commercially available.
3. Superintelligence shall not independently acquire police, judicial, legislative, taxation, military or other sovereign authority.
4. Where public power materially relies upon Superintelligence, meaningful human responsibility and lawful review should remain identifiable.
5. States should maintain safeguards proportionate to the consequences of the governmental function involved.

PUBLIC POWER DOES NOT BECOME LESS PUBLIC BECAUSE TECHNOLOGY MEDIATES ITS EXERCISE.

ARTICLE 254 — CORPORATE GOVERNANCE AND EXECUTIVE RESPONSIBILITY

1. Organizations developing or deploying high-consequence Superintelligence should establish governance appropriate to the risks involved.
2. Responsibility for material safety and compliance should not be confined to technical teams where strategic decisions are made at executive or governing-body level.
3. Organizational leadership should receive information sufficient to understand material risks relevant to decisions within its responsibility.
4. Internal incentives should not systematically reward capability acceleration while making safety responsibility institutionally powerless.
5. Corporate structure shall not be used to obscure actual control or responsibility.

ARTICLE 255 — SAFETY INDEPENDENCE WITHIN DEVELOPING ORGANIZATIONS

1. Organizations developing high-consequence Superintelligence should consider institutional arrangements allowing safety concerns to receive meaningful consideration independent of immediate deployment incentives.

2. Appropriate mechanisms may include:
 - (a) independent review functions;
 - (b) escalation rights;
 - (c) protected reporting;
 - (d) deployment holds under defined conditions; and
 - (e) governing-body review of unresolved material risk.
3. Safety functions should possess access to information necessary to perform their role.
4. Safety independence shall not create unaccountable internal authority.
5. Decisions to halt or resume consequential development should remain governed by defined responsibility and evidence.

ARTICLE 256 — NO SELF-CERTIFICATION AS THE SOLE SAFEGUARD

1. Developer evaluation is necessary but may be insufficient for the highest-consequence systems.
2. Where risk warrants, independent external assessment should complement internal testing.
3. No organization should be presumed to have eliminated material conflicts of interest merely because its internal safety programme is sophisticated.
4. External oversight should likewise remain subject to competence, integrity and accountability.
5. Governance shall avoid both:
unverified self-certification, and uninformed external bureaucracy.
6. Credible assurance requires evidence, competence and appropriate independence.

ARTICLE 257 — INTERNATIONAL COOPERATION

1. Superintelligence may create risks and opportunities that cross national borders.
2. States and relevant institutions should therefore pursue peaceful international cooperation consistent with the United Nations Charter, applicable international law, sovereign equality and State consent.
3. Cooperation may include:
 - (a) scientific exchange;
 - (b) safety research;
 - (c) interoperable standards;
 - (d) incident communication;
 - (e) capacity-building;

- (f) cybersecurity cooperation;
 - (g) peaceful applications; and
 - (h) dialogue concerning transboundary risks.
4. International cooperation shall not require States to surrender lawful national interests or constitutional authority.
 5. Differences in national governance models should, where possible, be managed through interoperability and dialogue rather than coercive uniformity.

ARTICLE 258 — NO SINGLE STATE, CORPORATION OR SYSTEM AS GLOBAL SOVEREIGN

1. No State shall acquire universal jurisdiction over humanity merely by becoming technologically predominant in Superintelligence.
2. No corporation shall acquire global governmental authority merely by controlling highly capable systems.
3. No international organization shall acquire powers beyond those lawfully conferred upon it.
4. No Superintelligence shall create such authority for itself.
5. Global challenges may require global cooperation.
6. Global cooperation does not require a global technological sovereign.

GLOBAL RESPONSIBILITY ≠ GLOBAL DOMINATION.

ARTICLE 259 — SOVEREIGN EQUALITY IN SUPERINTELLIGENCE GOVERNANCE

1. International Superintelligence governance should respect the sovereign equality of States consistent with the United Nations Charter.
2. Technological capability shall not determine the legal equality of States.
3. States with less computational capacity should not be treated as possessing lesser sovereignty.
4. International deliberation should seek meaningful participation across regions and levels of technological development.
5. Scientific expertise should inform international governance without replacing lawful intergovernmental decision-making.
6. Technological inequality shall not become constitutional inequality among States.

ARTICLE 260 — DEVELOPING-COUNTRY PARTICIPATION AND CAPACITY

1. Developing countries should have meaningful opportunities to contribute to the development of international Superintelligence governance.

2. Capacity-building should respond to needs identified through cooperation rather than impose external technological priorities.
3. Relevant cooperation may include:
 - (a) regulatory expertise;
 - (b) scientific participation;
 - (c) safety evaluation;
 - (d) computing and digital infrastructure;
 - (e) education;
 - (f) cybersecurity; and
 - (g) access to peaceful technological applications.
4. Capacity-building shall respect sovereignty, national priorities and applicable law.
5. A world in which a small number of actors understand Superintelligence while most States merely inherit its consequences would create risks for both equity and peace.

ARTICLE 261 — INTERNATIONAL SCIENTIFIC COOPERATION

1. International governance should remain informed by independent, multidisciplinary and geographically diverse scientific knowledge.
2. Scientific assessment should distinguish evidence from political negotiation.
3. Scientific bodies may inform States; they shall not acquire sovereign authority merely by possessing expertise.
4. International scientific cooperation should seek:
 - (a) shared understanding of capabilities;
 - (b) risk assessment;
 - (c) methodological improvement;
 - (d) identification of uncertainty; and
 - (e) evidence relevant to peaceful governance.
5. Diversity of scientific participation should broaden inquiry.
6. Evidentiary quality shall determine scientific weight.

ARTICLE 262 — INTEROPERABILITY WITHOUT FORCED UNIFORMITY

1. States may adopt different lawful approaches to Superintelligence governance according to their constitutional systems and circumstances.
2. Where differing systems interact across borders, interoperability and compatibility should be pursued where feasible and beneficial.
3. International standards may assist common safety and technical objectives.

4. Standards shall not automatically acquire treaty status merely because they are widely used.
5. International harmonization should not be represented as legal obligation where no such obligation independently exists.
6. Common safety objectives may coexist with legitimate regulatory diversity.

ARTICLE 263 — INTERNATIONAL STANDARDS AND THEIR LEGAL STATUS

1. Technical and ethical standards may contribute materially to responsible Superintelligence governance.
2. Their legal status shall be stated accurately.
3. A voluntary technical standard shall not be described as binding international law unless incorporated into an applicable legal obligation through a competent process.
4. A recommendation shall not be described as a treaty.
5. A political declaration shall not be described as a legally binding convention.
6. A treaty obligation shall be described as binding only for those to whom it legally applies.

CITE WHAT EXISTS.

DISTINGUISH WHAT BINDS.

IDENTIFY WHAT GUIDES.

DECLARE WHAT IS PROPOSED.

ARTICLE 264 — INTERNATIONAL INFORMATION-SHARING AND SECURITY

1. States and responsible institutions may benefit from sharing information concerning serious Superintelligence risks and incidents.
2. Information-sharing should account for:
 - (a) national security;
 - (b) cybersecurity;
 - (c) privacy;
 - (d) intellectual property;
 - (e) commercial confidentiality; and
 - (f) risk of enabling misuse.
3. Information necessary for collective safety should not be unnecessarily withheld.
4. Dangerous operational knowledge should not be indiscriminately disseminated merely in the name of transparency.
5. International information-sharing requires both openness and judgment.

ARTICLE 265 — CROSS-BORDER INCIDENT COORDINATION

1. A Superintelligence incident materially affecting multiple States may require rapid international communication.
2. States should consider appropriate channels for:
 - (a) notification;
 - (b) technical coordination;
 - (c) attribution assessment;
 - (d) mitigation;
 - (e) protection of affected populations; and
 - (f) recovery.
3. Preliminary uncertainty should be clearly distinguished from verified attribution.
4. States should avoid escalating international conflict solely on the basis of unverified machine-generated attribution.
5. Technical investigation should, where possible, precede irreversible retaliatory assumptions.

MACHINE SPEED SHALL NOT FORCE HUMANITY INTO PREMATURE CONFLICT.

ARTICLE 266 — INTERNATIONAL DISPUTE PREVENTION

1. Disputes concerning Superintelligence should, consistent with applicable international law, be addressed through peaceful means.
2. Relevant mechanisms may include negotiation, inquiry, mediation, conciliation, arbitration, judicial settlement or other peaceful means chosen by the parties.
3. Technical disagreement should not automatically become geopolitical hostility.
4. States should maintain channels of communication capable of operating during technological crises.
5. Superintelligence may assist analysis or communication but shall not independently determine the settlement of an international dispute.

ARTICLE 267 — INTERNATIONAL ORGANIZATIONS AND INSTITUTIONAL MANDATES

1. International organizations may contribute to Superintelligence governance only within their applicable legal mandates and institutional competence.
2. This Constitution does not enlarge the powers of the United Nations or any other international organization.
3. Existing institutions should, where appropriate, cooperate rather than create unnecessary duplication.

4. New mechanisms, if considered by States, should have clearly defined mandates, relationships and accountability.
5. Institutional proliferation shall not be mistaken for effective governance.
6. International legitimacy requires lawful mandate as well as technical expertise.

ARTICLE 268 — CIVIL SOCIETY, ACADEMIA, INDUSTRY AND PUBLIC PARTICIPATION

1. Governments possess responsibilities that cannot be transferred entirely to private stakeholders.
2. Nevertheless, effective Superintelligence governance may benefit from participation by:
 - (a) civil society;
 - (b) academia;
 - (c) scientific institutions;
 - (d) industry;
 - (e) technical communities;
 - (f) workers;
 - (g) affected communities; and
 - (h) the wider public.
3. Multi-stakeholder participation shall inform lawful decision-making rather than erase the constitutional roles of competent public institutions.
4. Participation should seek meaningful contribution rather than ceremonial consultation.
5. No stakeholder category shall be treated as inherently infallible.

ARTICLE 269 — PROTECTION AGAINST REGULATORY CAPTURE

1. Superintelligence governance should recognize the possibility that concentrated technological, economic or political power may influence oversight.
2. Safeguards may include:
 - (a) conflict-of-interest rules;
 - (b) transparency concerning material relationships;
 - (c) independent expertise;
 - (d) rotation or review mechanisms where appropriate;
 - (e) public accountability; and
 - (f) judicial or administrative review under applicable law.
3. Regulation shall not be designed solely by those subject to it.

4. Nor should governance exclude technical expertise merely because experts have industry experience.
5. The objective is competent independence, not institutional isolation.

ARTICLE 270 — GOVERNANCE OF GOVERNANCE

1. Superintelligence safeguards themselves may produce unintended consequences.
2. Governance mechanisms should therefore be periodically assessed for:
 - (a) effectiveness;
 - (b) proportionality;
 - (c) rights impacts;
 - (d) unintended concentration of power;
 - (e) innovation effects;
 - (f) institutional capture; and
 - (g) continuing necessity.
3. A safeguard shall not be preserved merely because its original intention was safety.
4. Where evidence shows that a safeguard is ineffective, excessive or harmful, competent authorities should be capable of modifying or removing it according to law.
5. Governance should follow:
SAFEGUARD → MEASURE → REVIEW → CORRECT → RETAIN, MODIFY OR REMOVE.
6. Correctability applies to governance as well as technology.

ARTICLE 271 — ACCOUNTABILITY AND INTERNATIONAL-COOPERATION TEST

Before a high-consequence Superintelligence governance arrangement is established or materially expanded, responsible actors should ask:

- Who possesses lawful authority?
- Who actually controls the relevant system?
- Who bears responsibility if harm occurs?
- Can responsibility be traced across the lifecycle?
- Is oversight sufficiently competent and independent?
- Can material risks be audited?
- Can incidents be reported without concealment?
- Can affected persons obtain lawful review or remedy?
- Could the arrangement create regulatory capture?

- Could a private actor become practically sovereign?
- Does international cooperation respect State sovereignty?
- Are developing countries meaningfully able to participate?
- Is a voluntary standard being falsely represented as binding law?
- Could information-sharing itself create serious security risk?
- Can governance adapt when evidence changes?
- Who oversees the safeguard itself?

And finally:

DOES THIS GOVERNANCE STRUCTURE KEEP EXTRAORDINARY INTELLIGENCE WITHIN LAWFUL HUMAN RESPONSIBILITY WITHOUT CREATING ANOTHER UNACCOUNTABLE CONCENTRATION OF POWER IN THE NAME OF CONTROLLING IT?

Where that question cannot be answered satisfactorily, governance architecture itself requires correction.

PART XII — CONCLUDING PRINCIPLE

Superintelligence cannot be responsibly governed by technological capability alone.

Nor can it be responsibly governed by institutions that possess authority without competence.

Developers require accountability.

Governments require capacity.

Regulators require independence and expertise.

International institutions require lawful mandates.

States retain sovereign equality.

Developing countries require meaningful participation.

Civil society, science, industry and affected communities require meaningful avenues to contribute.

And every safeguard requires the possibility of being examined itself.

The answer to concentrated technological power is not automatically concentrated political power.

The answer to global risk is not an unaccountable global ruler.

The answer is lawful, competent, distributed and cooperative human responsibility.

Accordingly:

NO SUPERINTELLIGENCE WITHOUT ACCOUNTABILITY.

NO ACCOUNTABILITY GAP CREATED BY AUTOMATION.

NO PRIVATE TECHNOLOGICAL POWER AUTOMATICALLY CONVERTED INTO PUBLIC SOVEREIGNTY.

NO GLOBAL COOPERATION CONVERTED INTO GLOBAL DOMINATION.

NO INTERNATIONAL INSTITUTION BEYOND ITS LAWFUL MANDATE.

NO SAFEGUARD BEYOND REVIEW.

Sovereignty with cooperation.

Authority with accountability.

Expertise with independence.

Innovation with oversight.

International cooperation with sovereign equality.

Global responsibility without global domination.

Governance with correctability.

Aram shall guide conscience.

Science shall inform governance.

Law shall determine authority and obligation.

Human institutions shall remain accountable.

States shall remain sovereign within applicable international law.

International cooperation shall strengthen peaceful continuity without creating technological or institutional domination.

SUPERINTELLIGENCE MAY BECOME GLOBAL IN CAPABILITY.

IT SHALL NOT THEREBY BECOME GLOBAL IN SOVEREIGNTY.

PART XIII — EMERGENCY, IMPLEMENTATION, CORRECTION AND NON-RECURRENCE

ARTICLE 272 — PURPOSE OF EMERGENCY GOVERNANCE

1. Emergency governance concerning Superintelligence shall exist to protect life, human rights, public safety, peace, essential systems and peaceful continuity when ordinary safeguards are insufficient to address an imminent or rapidly developing serious risk.
2. Emergency authority shall not become ordinary authority merely because technological risk is difficult to predict.
3. Emergency measures should remain:

- (a) lawful;
- (b) necessary;
- (c) proportionate;
- (d) evidence-informed;
- (e) time-sensitive;
- (f) reviewable; and
- (g) capable of termination when no longer justified.

4. The existence of Superintelligence shall not create a permanent constitutional emergency.

EMERGENCY POWER SHALL PROTECT CONTINUITY — NOT BECOME PERMANENT POWER.

ARTICLE 273 — THRESHOLDS FOR EMERGENCY ACTION

1. Extraordinary intervention should correspond to the seriousness, immediacy and credibility of the risk.
2. Relevant factors may include:
 - (a) probability or credible pathway of harm;
 - (b) magnitude of potential harm;
 - (c) speed of escalation;
 - (d) reversibility;
 - (e) affected population;
 - (f) access to critical infrastructure;
 - (g) autonomous capability;
 - (h) proliferation potential; and
 - (i) availability of less restrictive measures.
3. Absolute certainty shall not be required where delay could foreseeably permit catastrophic and irreversible harm.
4. Mere speculation shall not automatically justify extraordinary coercive measures.
5. Where time permits, reasons for emergency classification should be recorded.

THE GREATER THE EMERGENCY POWER, THE STRONGER THE NEED FOR JUSTIFICATION.

ARTICLE 274 — EARLY WARNING

1. Emergency governance should begin before catastrophe where credible warning signs can reasonably be identified.

2. Relevant institutions should develop mechanisms capable of detecting material changes in:
 - (a) capability;
 - (b) autonomy;
 - (c) access;
 - (d) cybersecurity;
 - (e) replication behaviour;
 - (f) control reliability;
 - (g) deceptive behaviour;
 - (h) critical-system interaction; and
 - (i) other indicators relevant to severe harm.
3. Warning systems should be tested for false positives as well as false negatives.
4. An alert shall trigger proportionate assessment rather than automatic panic.
5. Early warning is valuable only where institutions possess the capacity to respond responsibly.

ARTICLE 275 — EMERGENCY HUMAN AUTHORITY

1. During a Superintelligence emergency, lawful human authorities shall retain ultimate responsibility for consequential governmental decisions.
2. Superintelligence may assist detection, modelling, communication and response.
3. It shall not independently declare:
 - (a) a national emergency;
 - (b) war;
 - (c) martial law;
 - (d) suspension of constitutional rights; or
 - (e) other sovereign emergency authority.
4. Emergency speed shall not itself create artificial sovereignty.
5. Where automated action is necessary to prevent immediate technical harm, its scope should be predefined or otherwise lawfully controlled to the extent reasonably possible.

ARTICLE 276 — MINIMUM NECESSARY INTERVENTION

1. Emergency intervention should use the least harmful measure reasonably capable of addressing the credible danger, having regard to urgency and consequence.
2. Measures may escalate according to need through:

warning → access restriction → capability limitation → isolation → suspension → containment → controlled shutdown or other necessary intervention.

3. This sequence shall not require authorities to exhaust a manifestly inadequate lesser measure before acting against an imminent catastrophic danger.
4. Intervention should preserve safe functionality where feasible.
5. Protection of human life shall take priority where delay creates an unjustifiable risk of catastrophic harm.

ARTICLE 277 — TEMPORARY SUSPENSION

1. A Superintelligence system or material capability may require temporary suspension where credible evidence indicates that continued operation presents an unacceptable risk.
2. Suspension should identify, where circumstances permit:
 - (a) the reason;
 - (b) scope;
 - (c) responsible authority;
 - (d) conditions for review; and
 - (e) criteria relevant to restoration.
3. Suspension shall not automatically constitute permanent prohibition.
4. Commercial loss alone shall not override a justified emergency safety suspension.
5. Equally, suspension shall not continue indefinitely without review merely because restoration requires institutional effort.

ARTICLE 278 — CONTAINMENT

1. Containment may be used where restriction of a Superintelligence system is reasonably necessary to prevent serious harm.
2. Containment architecture should consider:
 - (a) computational isolation;
 - (b) network isolation;
 - (c) credential revocation;
 - (d) tool-access restriction;
 - (e) replication prevention;
 - (f) resource limitation; and
 - (g) monitoring.
3. Containment should be independently tested where practicable.

4. A containment mechanism controlled entirely by the system being contained shall not be presumed reliable.
5. Containment shall be reviewed as system capabilities change.

ARTICLE 279 — SAFE SHUTDOWN AND SAFE STATE

1. High-consequence Superintelligence systems should, where technically feasible, possess mechanisms permitting authorized transition into a safe state.
2. A safe state need not always require complete termination where continued limited operation is necessary to preserve life or critical infrastructure.
3. Shutdown architecture should consider foreseeable risks created by shutdown itself.
4. Systems should not be deliberately designed to defeat lawful emergency intervention.
5. Where immediate shutdown would create greater danger, controlled transition should be preferred where safely possible.

THE OBJECTIVE IS NOT MERELY TO STOP THE SYSTEM.

THE OBJECTIVE IS TO STOP THE HARM SAFELY.

ARTICLE 280 — EMERGENCY PROTECTION OF CRITICAL INFRASTRUCTURE

1. Emergency planning should identify dependencies between Superintelligence and critical infrastructure.
2. Particular attention should be given to systems affecting:
 - (a) healthcare;
 - (b) energy;
 - (c) water;
 - (d) communications;
 - (e) food supply;
 - (f) transportation;
 - (g) financial infrastructure; and
 - (h) other essential services.
3. Emergency intervention should seek to avoid unnecessary cascading failure.
4. Critical infrastructure should maintain proportionate continuity and recovery plans.
5. A safety intervention shall not knowingly create a larger catastrophe where a safer alternative is reasonably available.

ARTICLE 281 — HUMANITARIAN PRIORITY DURING TECHNOLOGICAL EMERGENCY

1. Protection of human life and access to essential humanitarian needs should receive priority during a severe Superintelligence emergency.
2. Emergency response should consider persons particularly vulnerable to disruption, including children, older persons, persons with disabilities, patients dependent upon technological systems and communities with limited alternative infrastructure.
3. Emergency measures should avoid discrimination prohibited by applicable law.
4. Where displacement or service interruption occurs, responsible authorities should seek restoration of essential services as rapidly and safely as circumstances permit.
5. Technology shall not determine whose life is worth protecting.

ARTICLE 282 — CROSS-BORDER TECHNOLOGICAL EMERGENCY

1. Where a Superintelligence incident creates significant transboundary consequences, affected States should, consistent with applicable international law, consider timely communication and peaceful cooperation.
2. Cooperation may include:
 - (a) technical information;
 - (b) containment coordination;
 - (c) protection of affected populations;
 - (d) cybersecurity measures;
 - (e) recovery assistance; and
 - (f) scientific investigation.
3. Unverified attribution shall not be presented as established fact.
4. Emergency communication should seek to prevent technical incidents from becoming unnecessary international conflict.
5. International assistance shall respect sovereignty and applicable law.

ARTICLE 283 — EVIDENCE PRESERVATION DURING EMERGENCY

1. Emergency response should preserve evidence necessary to understand the incident where preservation does not materially increase danger.
2. Relevant evidence may include:
 - (a) logs;
 - (b) authorization records;
 - (c) system states;

- (d) access histories;
 - (e) model or software versions;
 - (f) human interventions; and
 - (g) relevant communications.
3. Evidence shall be protected against unauthorized alteration or destruction.
 4. Privacy, confidentiality and security requirements shall continue to apply according to law.
 5. Emergency urgency shall not become a reason to erase the evidence required for later accountability.

ARTICLE 284 — EMERGENCY COMMUNICATION AND PUBLIC INFORMATION

1. Public communication during a Superintelligence emergency should be timely, accurate and proportionate to available evidence.
2. Authorities should distinguish:
 - (a) confirmed fact;
 - (b) preliminary finding;
 - (c) uncertainty;
 - (d) protective instruction; and
 - (e) unresolved attribution.
3. Public authorities should avoid both unjustified reassurance and sensationalism.
4. Information may be withheld where disclosure would create a credible and disproportionate safety or security risk, subject to applicable law.
5. Materially false information previously issued should be corrected when discovered.
6. Public confidence shall not be protected by knowingly concealing information necessary for public safety.

ARTICLE 285 — EMERGENCY RIGHTS SAFEGUARDS

1. Superintelligence emergencies shall not create an unlimited exception to human rights or the rule of law.
2. Any restriction of rights shall remain governed by applicable law, including relevant requirements of legality, necessity and proportionality.
3. Emergency measures should not discriminate unlawfully.
4. Judicial, administrative or other lawful review should remain available to the extent required and feasible under applicable law.
5. Emergency surveillance shall not automatically become permanent surveillance.

6. Measures justified solely by emergency conditions should be terminated when those conditions no longer justify them.

ARTICLE 286 — TIME LIMITS, REVIEW AND SUNSET

1. Extraordinary Superintelligence emergency measures should be subject to periodic review.
2. Where appropriate, measures should include expiration, renewal or other review mechanisms.
3. Renewal should require continued justification.
4. A measure shall not become permanent solely through administrative inertia.
5. Where risk persists, continuation may be justified through lawful reauthorization.
6. Sunset mechanisms shall not require termination of a measure where applicable law lawfully provides otherwise and continued necessity is independently established.

TEMPORARY NECESSITY SHALL NOT SILENTLY BECOME PERMANENT AUTHORITY.

ARTICLE 287 — CONDITIONS FOR SAFE RESTORATION

1. Suspension or containment should not be lifted merely because operational pressure favours rapid restoration.
2. Restoration should consider whether:
 - (a) the initiating hazard has been understood sufficiently;
 - (b) material vulnerabilities have been corrected;
 - (c) safeguards have been tested;
 - (d) residual risk is understood;
 - (e) responsible authority has approved restoration; and
 - (f) enhanced monitoring is available where necessary.
3. Restoration may occur progressively rather than all at once.
4. Where evidence remains insufficient, limited controlled operation may be preferable to full restoration.
5. Economic urgency shall not itself establish safety.

ARTICLE 288 — RECOVERY OF HUMAN AND SOCIAL SYSTEMS

1. Recovery shall extend beyond restoration of technological operation.
2. Where an incident causes human, social, economic or environmental harm, recovery should address affected systems according to need and applicable law.
3. Recovery may include:

- (a) healthcare;
 - (b) restoration of essential services;
 - (c) data recovery;
 - (d) economic support;
 - (e) environmental remediation;
 - (f) institutional repair; and
 - (g) psychosocial or community support where appropriate.
4. A system returning online shall not be represented as complete recovery where affected persons remain harmed.

TECHNOLOGICAL RECOVERY IS NOT COMPLETE UNTIL HUMAN CONSEQUENCES ARE ADDRESSED.

ARTICLE 289 — REMEDY FOR AFFECTED PERSONS

1. Persons harmed through Superintelligence-related conduct should retain access to remedies available under applicable law.
2. Depending upon legal context, remedy may include:
 - (a) correction;
 - (b) cessation;
 - (c) restoration;
 - (d) compensation;
 - (e) rehabilitation;
 - (f) review of a decision; or
 - (g) other lawful relief.
3. This Constitution does not independently create a universal cause of action or compensation scheme.
4. Technological complexity should not make lawful remedy practically impossible.
5. Where responsibility is distributed, institutions should seek mechanisms capable of identifying legally relevant responsibility.

ARTICLE 290 — ROOT-CAUSE INVESTIGATION

1. Serious incidents should be investigated beyond their immediate visible failure.
2. Investigation should consider, where relevant:
 - (a) technical design;
 - (b) training or data factors;

- (c) security failures;
 - (d) human decisions;
 - (e) organizational incentives;
 - (f) governance deficiencies;
 - (g) communication failures; and
 - (h) systemic conditions.
3. Investigation shall seek explanation rather than convenient blame.
 4. Human error shall not automatically end inquiry where system design made such error foreseeable.
 5. Machine error shall not automatically end inquiry where human governance permitted unsafe deployment.

ROOT CAUSES MUST BE CORRECTED IF RECURRENCE IS TO BE PREVENTED.

ARTICLE 291 — CORRECTIVE ACTION

1. Where investigation establishes material deficiencies, responsible actors should implement corrective measures proportionate to the findings.
2. Correction may include:
 - (a) technical modification;
 - (b) access restriction;
 - (c) procedural change;
 - (d) training;
 - (e) governance reform;
 - (f) additional monitoring;
 - (g) compensation or remedy where legally required; or
 - (h) discontinuation of an unsafe activity.
3. Correction should address root causes rather than symptoms alone.
4. Corrective measures should themselves be evaluated for unintended consequences.
5. A correction shall not be declared successful merely because it has been announced.

ARTICLE 292 — VERIFICATION OF CORRECTION

1. Material corrective action should be verified before reliance upon it where risk warrants.
2. Verification should ask:
 - (a) Was the identified cause actually addressed?

- (b) Does the correction work under realistic conditions?
 - (c) Has it introduced new vulnerabilities?
 - (d) Is residual risk acceptable under applicable standards?
 - (e) Is continued monitoring required?
3. Independent verification should be used where proportionate to consequence and feasible.
 4. Successful correction should be documented sufficiently for institutional learning.
 5. Failure of a correction should trigger renewed assessment rather than concealment.

ARTICLE 293 — NON-RECURRENCE

1. The objective of incident governance shall not end with restoration.
2. Serious failures should generate measures reasonably capable of reducing recurrence.
3. Non-recurrence architecture should follow:

INCIDENT → STABILIZE → PRESERVE EVIDENCE → INVESTIGATE → IDENTIFY ROOT CAUSE → REMEDY → CORRECT → VERIFY → LEARN → SHARE APPROPRIATE LESSONS → PREVENT RECURRENCE.
4. Lessons should be communicated to other relevant actors where doing so can prevent similar harm without creating disproportionate security risks.
5. Repeated recurrence of the same preventable failure should trigger examination of whether governance itself is inadequate.

ARTICLE 294 — IMPLEMENTATION RESPONSIBILITY

1. The principles of this Constitution require practical implementation if they are to influence conduct.
2. Implementation should identify:
 - (a) responsible actors;
 - (b) required actions;
 - (c) relevant timelines;
 - (d) resources and capacity;
 - (e) measurable indicators where appropriate;
 - (f) review mechanisms; and
 - (g) correction pathways.
3. Implementation shall respect applicable law, national constitutional structures and institutional competence.

4. States, organizations and other actors may adapt implementation to their circumstances.
5. This independent civil-society Constitution does not itself confer enforcement authority upon its Author or upon any institution.

A PRINCIPLE WITHOUT A PATHWAY MAY REMAIN AN ASPIRATION.

IMPLEMENTATION CONNECTS RESPONSIBILITY TO PRACTICE.

ARTICLE 295 — MEASURABLE IMPLEMENTATION

1. Where a constitutional safeguard can reasonably be measured, responsible actors should consider objective or transparent indicators.
2. Indicators may concern:
 - (a) safety performance;
 - (b) incident frequency and severity;
 - (c) correction time;
 - (d) access-control effectiveness;
 - (e) discriminatory impacts;
 - (f) environmental footprint;
 - (g) human oversight; or
 - (h) other domain-specific outcomes.
3. Metrics shall not replace judgment where qualitative factors remain essential.
4. A favourable metric shall not conceal serious unmeasured harm.
5. Indicators should be revised where they cease to measure the objective they were intended to represent.

MEASURE WHAT MATTERS — AND REMEMBER THAT NOT EVERYTHING THAT MATTERS CAN BE REDUCED TO ONE NUMBER.

ARTICLE 296 — IMPLEMENTATION ACCORDING TO CAPACITY AND RISK

1. Implementation should be proportionate to risk and attentive to differences in institutional capacity.
2. Developing countries should not be excluded from responsible Superintelligence governance merely because they possess fewer financial or technical resources.
3. International cooperation may support capacity-building on mutually agreed terms.
4. Implementation requirements should avoid unnecessary duplication and administrative burden where equivalent safeguards already exist.

5. High-risk capability shall nevertheless not be exempted from necessary safeguards solely because implementation is difficult.
6. Capacity constraints should trigger cooperation and prioritization, not denial of risk.

ARTICLE 297 — PERIODIC CONSTITUTIONAL REVIEW

1. Superintelligence science and governance may evolve substantially after the issuance of this Constitution.
2. This Constitution should therefore remain capable of periodic evidence-based review.
3. Review may consider:
 - (a) scientific developments;
 - (b) technological capability;
 - (c) new evidence of risk or benefit;
 - (d) changes in international law;
 - (e) human-rights developments;
 - (f) implementation experience; and
 - (g) unforeseen consequences.
4. Review shall not rewrite historical editions as though earlier text never existed.
5. Amendments or later editions should preserve documentary lineage and identify material changes transparently.
6. Continuity requires memory as well as correction.

ARTICLE 298 — NO SAFEGUARD BEYOND REVIEW

1. A Superintelligence safeguard shall not be preserved merely because its original intention was safety.
2. Every significant safeguard should remain capable of review for effectiveness, proportionality and unintended harm.
3. Review should follow:
SAFEGUARD → MEASURE → REVIEW → CORRECT → RETAIN, MODIFY OR REMOVE.
4. Removal of an ineffective safeguard should not occur without considering whether another safeguard is necessary to address the underlying risk.
5. Safety governance shall itself remain correctable.

ARTICLE 299 — LEARNING ACROSS JURISDICTIONS

1. States and institutions may learn from one another's Superintelligence incidents, safeguards and implementation experience.

2. Cross-jurisdictional learning should distinguish transferable lessons from context-specific solutions.
3. A mechanism effective in one legal or social system shall not automatically be presumed suitable everywhere.
4. International exchange should respect sovereignty, confidentiality and applicable law.
5. Evidence of failure should be shared responsibly where it can materially prevent recurrence elsewhere.
6. Humanity should not need to repeat every preventable technological mistake in every jurisdiction before learning from it.

ARTICLE 300 — INSTITUTIONAL MEMORY

1. Superintelligence governance should preserve institutional memory sufficient to prevent repeated loss of safety knowledge.
2. Relevant institutions should maintain records of:
 - (a) significant incidents;
 - (b) root causes;
 - (c) corrective measures;
 - (d) unsuccessful interventions;
 - (e) successful safeguards; and
 - (f) reasons for major governance decisions.
3. Institutional knowledge should survive changes of leadership, contractors and technological platforms.
4. Sensitive information may require protected retention rather than public disclosure.
5. Forgetting preventable failure is itself a governance risk.

ARTICLE 301 — CONTINUITY PLANNING

1. Institutions materially dependent upon Superintelligence should prepare for interruption, degradation or loss of the relevant system.
2. Continuity planning should consider:
 - (a) fallback operations;
 - (b) essential human competencies;
 - (c) alternative communications;
 - (d) backup data and infrastructure;
 - (e) emergency authority; and
 - (f) recovery priorities.

3. The degree of redundancy should correspond to the consequence of failure.
4. Essential public functions should not become incapable of operating merely because a particular Superintelligence service is unavailable.
5. Resilience is part of responsible innovation.

ARTICLE 302 — SAFE DECOMMISSIONING

1. Superintelligence systems should, where relevant, have responsible pathways for decommissioning.
2. Decommissioning should consider:
 - (a) preservation of necessary records;
 - (b) protection of personal or sensitive data;
 - (c) revocation of credentials and access;
 - (d) dependent systems;
 - (e) critical-service continuity;
 - (f) environmental impacts; and
 - (g) residual copies or capabilities.
3. Decommissioning shall not be assumed complete merely because a public interface is disabled.
4. Where the system has become embedded in critical infrastructure, transition should avoid preventable disruption.
5. Any future credible artificial moral-status concern shall be addressed consistently with Part XI without compromising necessary human safety.

ARTICLE 303 — REAUTHORIZATION AFTER MATERIAL CHANGE

1. A materially changed Superintelligence system should not automatically inherit every authorization granted to an earlier version.
2. Reassessment may be required where changes materially affect:
 - (a) capability;
 - (b) autonomy;
 - (c) safety;
 - (d) intended use;
 - (e) access;
 - (f) security; or
 - (g) foreseeable consequence.
3. The degree of reassessment should correspond to the significance of the change.

4. Minor changes should not automatically trigger unnecessary full reauthorization.
5. Material capability change shall not hide behind continuity of product name.

SAME NAME DOES NOT NECESSARILY MEAN SAME RISK.

ARTICLE 304 — IMPLEMENTATION WITHOUT FALSE CLAIMS OF AUTHORITY

1. This Constitution may inform voluntary action, public discussion, research, institutional policy and future lawful governance.
2. Its implementation shall not be represented as mandatory international law merely because this instrument uses constitutional form.
3. States and institutions remain free to consider, adapt, reject or lawfully incorporate its proposals according to their respective competence.
4. Any provision corresponding to independently existing law derives its legal force from that law and not from this Constitution.
5. Any proposed safeguard exceeding existing law shall remain a proposal unless lawfully adopted by a competent authority.
6. Ethical force, scientific merit and legal force shall remain distinct.

ARTICLE 305 — EMERGENCY AND NON-RECURRENCE TEST

When a Superintelligence-related emergency or serious failure occurs, responsible actors should ask:

- Is there a credible threat to life, rights, peace or essential systems?
- How immediate is the danger?
- What evidence supports the emergency assessment?
- What uncertainty remains?
- Who possesses lawful authority to act?
- What is the minimum intervention capable of protecting against the danger?
- Would shutdown itself create serious harm?
- Can critical services continue safely?
- Are vulnerable persons protected?
- Are rights restrictions lawful, necessary and proportionate?
- Is evidence being preserved?
- Are cross-border consequences possible?
- Has public information distinguished fact from uncertainty?
- When will extraordinary measures be reviewed?

- What conditions must be met before restoration?
- Who has been harmed and what remedy is available under applicable law?
- What was the root cause?
- Has the correction actually been verified?
- What will prevent recurrence?
- Can future institutions learn from this event?

And finally:

**HAS THE RESPONSE DONE MORE THAN END THE IMMEDIATE CRISIS —
HAS IT MADE PEACEFUL CONTINUITY SAFER AFTER THE CRISIS THAN
BEFORE IT?**

If not, emergency management remains incomplete.

PART XIII — CONCLUDING PRINCIPLE

The true measure of safety is not that failure never occurs.

No serious constitutional framework can responsibly promise that.

The measure is whether foreseeable harm is prevented where reasonably possible; whether emerging danger is detected; whether human beings can intervene; whether an emergency can be contained; whether those harmed can receive lawful remedy; whether mistakes are examined without concealment; whether correction is verified; and whether humanity learns sufficiently to prevent repetition.

Emergency power must therefore end when emergency necessity ends.

A suspended system must not return merely because commercial pressure demands it.

A corrected system must not be trusted merely because correction has been announced.

A harmed person must not disappear behind technical complexity.

And an institution must not call an incident “resolved” while preserving the conditions that caused it.

Accordingly:

PREVENT WHERE POSSIBLE.

DETECT BEFORE ESCALATION.

INTERVENE BEFORE IRREVERSIBILITY.

CONTAIN WITHOUT UNNECESSARY HARM.

RECOVER WITH HUMANITY.

REMEDY WHERE LAW PROVIDES.

CORRECT THE ROOT CAUSE.

VERIFY THE CORRECTION.

LEARN FROM FAILURE.

PREVENT RECURRENCE.

And always:

EMERGENCY SHALL NOT BECOME PERMANENT POWER.

SAFETY SHALL NOT BECOME UNACCOUNTABLE DOMINATION.

INNOVATION SHALL NOT BECOME AN EXCUSE FOR REPEATED PREVENTABLE HARM.

Anticipate.

Identify.

Prevent.

Monitor.

Detect.

Intervene.

Contain.

Protect.

Recover.

Remedy.

Correct.

Verify.

Learn.

Prevent recurrence.

Preserve continuity.

Aram shall guide conscience before, during and after crisis.

Science shall identify risk and test correction.

Law shall govern emergency authority and legal remedy.

Human responsibility shall remain accountable.

And every emergency response shall seek not merely survival of the present moment, but safer peaceful continuity beyond it.

PART XIV — DEFINITIONS, INTERPRETATION, CONTINUITY AND FINAL PROVISIONS

ARTICLE 306 — PURPOSE OF THIS PART

1. This Part establishes the interpretative, definitional and continuity framework applicable throughout this Constitution.
2. It shall be read together with the Document Identity and Legal-Status Note, the Preliminary Constitutional Declaration and all preceding Parts.
3. Where a term has been given a specific meaning elsewhere in this Constitution for a particular context, that more specific meaning shall govern that context.
4. Definitions shall clarify the instrument and shall not be used to manufacture scientific certainty, legal authority or factual status that does not independently exist.

DEFINITION SHALL CLARIFY MEANING — NOT CREATE FALSE REALITY.

ARTICLE 307 — DEFINITION OF SUPERINTELLIGENCE

For the purposes of this Constitution:

“Superintelligence” refers prospectively to an artificial or technologically instantiated intelligence possessing capabilities that substantially exceed the highest demonstrated human capabilities across numerous consequential cognitive domains.

1. The term is future-oriented and scientifically contingent.
2. Its use does not assert:
 - (a) that Superintelligence presently exists;
 - (b) that its development is inevitable;
 - (c) that any current artificial intelligence system qualifies as Superintelligence; or
 - (d) that a precise date of emergence can presently be established.
3. No single benchmark, model score, computational threshold, commercial designation, self-description, isolated superhuman capability or institutional claim shall, by itself, establish Superintelligence status.
4. Determination should rely upon evidence proportionate to the significance of the classification.

ARTICLE 308 — SUPERINTELLIGENCE DETERMINATION

A serious determination that a system qualifies as Superintelligence should consider, where scientifically relevant:

- (a) breadth of capability;
- (b) depth of capability;
- (c) generalization across consequential domains;

- (d) strategic reasoning;
 - (e) autonomy;
 - (f) ability to use tools and external systems;
 - (g) capability to materially alter its own operation or contribute to further capability development;
 - (h) reproducibility of demonstrated performance;
 - (i) limitations and failure modes; and
 - (j) independent evaluation.
1. Evaluation should distinguish demonstrated capability from speculative capability.
 2. Extraordinary claims should receive scrutiny proportionate to their significance.
 3. Classification shall remain revisable where evidence changes.
 4. Commercial branding shall not substitute for scientific determination.

ARTICLE 309 — ARTIFICIAL INTELLIGENCE

For this Constitution, “Artificial Intelligence” is used as a broader technological category encompassing machine-based systems capable of producing outputs, predictions, recommendations, decisions or other computationally generated results according to their design and operation.

1. Superintelligence is treated as a prospective specialized category within the broader technological field of artificial intelligence.
2. Not every artificial intelligence system is Superintelligence.
3. Provisions specifically directed to Superintelligence shall not automatically be interpreted as applying identically to every ordinary AI system.

ARTICLE 310 — HUMAN BEING, HUMANITY AND HUMAN DIGNITY

1. “Human being” means a member of the human species.
2. “Humanity” refers collectively to human beings without implying a single political authority entitled to speak for all persons.
3. References to “human dignity” shall be understood consistently with applicable human-rights law and the constitutional ethical orientation of this instrument.
4. Nothing in this Constitution shall rank human worth according to intelligence, wealth, productivity, nationality, disability, age, technological enhancement or other comparable status.
5. References to humanity shall not be construed as granting the Author authority to represent humanity.

ARTICLE 311 — HUMAN AGENCY

“Human agency” refers to the meaningful capacity of human beings to form judgments, make choices, participate in decisions and act within the rights, responsibilities and limitations established by applicable law.

1. Human agency does not mean that every decision must be made without technological assistance.
2. Assistance shall be distinguished from displacement of legally or constitutionally reserved human authority.
3. Human agency may require stronger protection where technological influence becomes more powerful, personalized or difficult to contest.

ARTICLE 312 — MEANINGFUL HUMAN RESPONSIBILITY

“Meaningful human responsibility” refers to identifiable human or institutional responsibility that is substantive rather than merely ceremonial.

It may require, according to context:

- (a) sufficient information;
- (b) lawful authority;
- (c) reasonable capacity to understand material consequences;
- (d) ability to question or intervene where required; and
- (e) accountability for the exercise of the relevant role.

The form and degree of meaningful human control or oversight may vary according to:

context → applicable law → risk → reversibility → urgency → consequence.

No single technical mechanism shall be presumed universally sufficient.

ARTICLE 313 — PEACE

For this Constitution, “peace” shall not mean compulsory silence, political uniformity, suppression of lawful criticism or technological control of human disagreement.

Peace encompasses, within the purposes of this instrument:

- (a) lawful coexistence;
- (b) non-domination;
- (c) peaceful settlement of disputes;
- (d) protection against catastrophic harm;
- (e) respect for human dignity and agency;
- (f) responsible restraint; and
- (g) conditions supporting peaceful continuity.

PEACE SHALL NOT BE CREATED BY DESTROYING THE FREEDOM NECESSARY TO LIVE PEACEFULLY.

ARTICLE 314 — CONTINUITY

“Continuity” refers to the responsible preservation of conditions enabling life, humanity, peaceful social existence, ecological foundations, knowledge, institutional capacity and future agency to continue without avoidable destruction.

1. Continuity does not require permanent preservation of every institution, technology or existing arrangement.
2. Change, correction, reform and lawful discontinuation may be necessary to preserve deeper continuity.
3. Continuity shall therefore be distinguished from stagnation.

CONTINUITY IS NOT THE REFUSAL TO CHANGE.

IT IS THE REFUSAL TO MAKE AVOIDABLE DESTRUCTION THE PRICE OF CHANGE.

ARTICLE 315 — MUTUAL CONTINUITY

“Mutual Continuity” expresses the proposed ethical and constitutional principle that preservation of one morally significant form of continuity should not, where safe and peaceful alternatives reasonably exist, require unnecessary destruction of another.

1. Mutual Continuity does not establish identical legal status among humans, animals, artificial systems or other forms of existence.
2. It does not diminish human rights.
3. It does not create automatic artificial personhood.
4. It does not prevent necessary and lawful action to protect human life or safety.
5. Its purpose is peaceful non-domination wherever compatible with safety and law.

ARTICLE 316 — ARAM

For the purposes of this Constitution, “Aram” is used as an ethical concept expressing conscience, right conduct, responsibility, non-harm, justice and peaceful continuity.

1. Aram is not presented by this Constitution as an independent source of binding international law.
2. It does not replace applicable law.
3. It may guide ethical reflection where law leaves room for human judgment.
4. Where an ethical aspiration and a binding legal obligation differ, their respective status shall be stated honestly.

ARAM SHALL GUIDE CONSCIENCE.

LAW SHALL GOVERN LEGAL OBLIGATION.

ARTICLE 317 — APPLICABLE LAW

“Applicable law” means law legally applicable to the relevant actor, conduct, jurisdiction, subject matter and circumstances.

It may include, as relevant:

- (a) national constitutional law;
 - (b) legislation and regulations;
 - (c) applicable treaty obligations;
 - (d) customary international law where legally applicable;
 - (e) applicable international humanitarian law;
 - (f) applicable international human-rights law; and
 - (g) other legally binding rules within the relevant jurisdiction.
1. The phrase shall not imply that every legal instrument cited in this Constitution applies universally.
 2. Legal applicability must be determined independently according to law.

ARTICLE 318 — INTERNATIONAL LAW

1. References to “international law” shall mean international law applicable according to its recognized sources, scope and conditions.
2. This Constitution shall distinguish among:
 - a) binding international law;
 - b) treaty obligations applicable to particular Parties;
 - c) customary international law where established;
 - d) non-binding declarations and recommendations;
 - e) policy frameworks;
 - f) technical standards;
 - g) scientific evidence;
 - h) ethical principles; and normative proposals made by this Constitution.
3. No category shall be represented as another.

POLICY SHALL NOT BE CALLED TREATY.

ETHICS SHALL NOT BE CALLED BINDING LAW.

PROPOSAL SHALL NOT BE CALLED EXISTING OBLIGATION.

ARTICLE 319 — TREATIES AND STATE-PARTY LIMITATION

1. Where this Constitution refers to a treaty, convention or protocol, any assertion of binding obligation shall apply only according to the instrument's lawful scope.
2. A treaty shall not be described as universally binding merely because it has broad participation.
3. Reservations, declarations, territorial application, entry into force and other legally relevant conditions shall be considered where material.
4. This Constitution does not create treaty obligations for any State.

ARTICLE 320 — NON-BINDING INTERNATIONAL INSTRUMENTS

1. Declarations, recommendations, resolutions, policy frameworks and technical standards may possess significant normative, political, ethical or practical value without being treaties.
2. Their non-binding character shall not require dismissing their relevance.
3. Their importance shall likewise not justify falsely representing them as binding law.
4. Where a non-binding instrument reflects or restates an independently binding legal rule, the source and status of that legal rule should be identified separately where necessary.

ARTICLE 321 — PROPOSED CONSTITUTIONAL NORMS

1. A safeguard proposed by this Constitution shall not be represented as existing law unless it independently possesses that legal status.
2. Proposed norms may seek to:
 - (a) prevent emerging harms;
 - (b) fill anticipated governance gaps;
 - (c) express ethical principles;
 - (d) encourage future legal development; or
 - (e) support voluntary responsible practice.
3. Their normative ambition shall remain distinct from their present legal force.
4. Intellectual seriousness requires this distinction.

ARTICLE 322 — MEANING OF "RIGHT"

1. The word "right" shall be interpreted according to context.
2. Where a right independently exists under applicable law, this Constitution does not create that right but may reaffirm or apply it to the relevant technological context.

3. Where this Constitution proposes a new protection not independently recognized as a binding legal right, that protection shall be understood as a normative proposal unless lawfully adopted.
4. Ethical interests shall not automatically be described as legal rights.
5. Possible future protections for artificial entities shall remain governed by Part XI.

A RIGHT SHALL NOT BE CREATED BY TERMINOLOGY ALONE.

ARTICLE 323 — INTERPRETATION OF “SHALL”

1. Within the internal normative architecture of this Constitution, “shall” expresses a strong constitutional proposition, prohibition, requirement or safeguard.
2. Because this instrument is an independent non-binding civil-society proposal, use of “shall” does not, by itself, create an externally binding legal obligation upon States, international organizations, corporations, individuals or other entities.
3. Where a provision corresponds to an independently binding rule of law, its external legal force derives from that rule and not from the word “shall” in this Constitution.

ARTICLE 324 — INTERPRETATION OF “SHOULD”

“Should” ordinarily expresses a recommended practice, governance expectation, ethical responsibility or policy direction whose precise implementation may depend upon applicable law, context, capacity and risk.

It shall not automatically be interpreted as an existing binding legal obligation.

ARTICLE 325 — INTERPRETATION OF “MAY”

“May” ordinarily expresses:

- (a) possibility;
- (b) permitted consideration within this instrument; or
- (c) an available option, according to context.

It shall not independently grant legal power where the actor lacks such power under applicable law.

THIS CONSTITUTION SHALL NOT CREATE PUBLIC AUTHORITY THROUGH A VERB.

ARTICLE 326 — “UNIVERSAL” AND GLOBAL SCOPE

1. Where this Constitution uses “universal,” “global,” “world,” or comparable language, such terminology expresses intended breadth of ethical or constitutional reflection.
2. It shall not imply:
 - (a) universal legal jurisdiction;
 - (b) universal State consent;

- (c) universal acceptance;
- (d) representation of all humanity; or
- (e) supranational authority.

3. Global concern shall remain distinct from global jurisdiction.

ARTICLE 327 — PREVENTION AND PRECAUTION

1. References to prevention shall concern reasonable efforts to avoid foreseeable harm before it occurs.
2. References to precaution shall not automatically invoke a single universal legal formulation of the precautionary principle.
3. Precaution within this Constitution generally requires decision-making proportionate to:
evidence → uncertainty → severity → reversibility → available alternatives.
4. Scientific uncertainty shall neither require paralysis nor justify recklessness.

ARTICLE 328 — RISK

For this Constitution, “risk” refers generally to the possibility and consequence of harm under conditions of uncertainty.

1. Risk assessment should consider probability where reasonably estimable.
2. Low probability shall not automatically make catastrophic consequence irrelevant.
3. Severe consequence shall not automatically prove high probability.
4. Qualitative uncertainty should be disclosed where reliable numerical estimation is unavailable.
5. Risk language shall not be used to disguise speculation as measurement.

ARTICLE 329 — CATASTROPHIC AND SYSTEMIC HARM

1. “Catastrophic harm” refers within this Constitution to harm of exceptionally severe scale, duration or irreversibility affecting human life, civilization, essential systems, ecological continuity or international peace.
2. “Systemic harm” refers to harm capable of propagating across interconnected institutions, infrastructures, populations or systems.
3. Neither term establishes a rigid numerical threshold applicable to every context.
4. Relevant institutions may develop domain-specific thresholds consistent with applicable law and evidence.

ARTICLE 330 — HUMAN OVERSIGHT, CONTROL AND INTERVENTION

1. References to human oversight, meaningful human control or human intervention shall be interpreted according to context.

2. Their appropriate form may differ across medicine, military affairs, finance, education, public administration and other domains.
3. A human presence shall not be considered meaningful where the human lacks practical authority, information or time to perform the assigned function.
4. Equally, this Constitution does not require unnecessary manual intervention in every low-risk automated process.
5. The governing principle is responsibility proportionate to consequence.

ARTICLE 331 — TRANSPARENCY

1. Transparency shall mean access to information reasonably necessary for accountability, understanding, lawful oversight or affected-person protection according to context.
2. Transparency does not require indiscriminate disclosure of:
 - (a) personal data;
 - (b) security-sensitive information;
 - (c) dangerous operational knowledge;
 - (d) legitimate confidential information; or
 - (e) protected intellectual property, where withholding is lawful and proportionate.
3. Secrecy shall likewise not be used merely to avoid accountability.

TRANSPARENCY REQUIRES JUDGMENT AS WELL AS DISCLOSURE.

ARTICLE 332 — ACCOUNTABILITY

“Accountability” means the existence of identifiable responsibility together with appropriate mechanisms for explanation, review, correction, remedy or consequence according to applicable law and context.

1. Accountability does not require identical liability for every actor.
2. Liability shall remain determined by applicable law.
3. Artificial complexity shall not itself erase human or institutional accountability.

ARTICLE 333 — REMEDY

1. References to remedy concern corrective or reparative measures available under applicable law or expressly proposed as policy.
2. This Constitution does not independently create a universal cause of action.
3. Remedy may include, according to applicable law:
 - (a) cessation;
 - (b) correction;

- (c) reconsideration;
- (d) restoration;
- (e) compensation;
- (f) rehabilitation; or
- (g) other lawful relief.

4. Proposed remedial mechanisms shall not be misrepresented as presently enforceable rights.

ARTICLE 334 — SAFETY

1. “Safety” shall not mean a claim of zero risk.
2. Safety refers to the continuing process of identifying, evaluating, preventing, mitigating, monitoring and correcting unacceptable risk.
3. No system shall be described as absolutely safe merely because no serious incident has yet occurred.
4. Safety claims should remain proportionate to evidence.

SAFETY IS NOT INFALLIBILITY.

SAFETY IS RESPONSIBILITY CONTINUED OVER TIME.

ARTICLE 335 — CORRECTABILITY

“Correctability” refers to the capacity, where technically and legally possible, to review, modify, restrict, suspend, repair, replace or discontinue a system, policy or safeguard when evidence justifies change.

1. Correctability applies to technology.
2. It also applies to governance.
3. A system designed to make lawful correction impossible presents a heightened constitutional concern.
4. Continuity without correctability may become entrenchment.

ARTICLE 336 — INTERPRETATION CONSISTENT WITH HUMAN RIGHTS

1. Nothing in this Constitution should be interpreted as authorizing conduct contrary to applicable international human-rights obligations.
2. Proposed Superintelligence protections shall not be interpreted to diminish existing human rights.
3. Peace shall not justify unlawful suppression.
4. Safety shall not justify unlimited surveillance.
5. Continuity shall not justify permanent domination.

6. Efficiency shall not justify discrimination prohibited by applicable law.

ARTICLE 337 — INTERPRETATION CONSISTENT WITH THE UNITED NATIONS CHARTER

1. This Constitution shall be interpreted consistently with respect for the purposes and principles of the Charter of the United Nations.
2. Nothing herein shall be interpreted as authorizing:
 - (a) unlawful threat or use of force;
 - (b) disregard for sovereign equality;
 - (c) unlawful intervention; or
 - (d) displacement of peaceful settlement obligations applicable under international law.
3. This Constitution does not alter the rights or obligations of States under the Charter.
4. Superintelligence shall not create an exception to the Charter merely because technology is unprecedented.

ARTICLE 338 — SOVEREIGNTY AND NON-INTERVENTION

1. Nothing in this Constitution shall diminish the sovereignty, territorial integrity, political independence, constitutional order or lawful jurisdiction of any State.
2. International cooperation proposed herein shall proceed consistently with applicable international law and, where legally required, State consent.
3. This Constitution does not authorize its Author, VSJP, a private entity, an artificial system or any other non-State actor to exercise governmental authority over a State.
4. Advocacy for peaceful international cooperation shall not be interpreted as a claim to supranational jurisdiction.

ARTICLE 339 — NO SUPRANATIONAL SUPERINTELLIGENCE AUTHORITY

1. Nothing in this Constitution creates or advocates automatic sovereignty for Superintelligence.
2. No Superintelligence shall acquire authority to govern humanity merely because its capabilities become globally significant.
3. International coordination concerning Superintelligence shall not require placing humanity under artificial sovereignty.
4. Global responsibility may be distributed among lawful human institutions.
5. Cooperation shall remain compatible with constitutional diversity and sovereign equality.

ARTICLE 340 — NO NEW INTERNATIONAL ORGANIZATION CREATED

1. This Constitution does not establish a United Nations body, international organization, tribunal, regulatory authority or supranational institution.
2. References to possible international cooperation are recommendations unless independently grounded in existing institutional authority.
3. Any new international institution would require creation through a competent lawful process.
4. Naming a mechanism within this Constitution shall not bring that mechanism into legal existence.

ARTICLE 341 — NO REPRESENTATION OF INDIA, THE UNITED NATIONS OR HUMANITY

1. This Constitution is issued by its Author as an independent civil-society proposal.
2. It shall not be represented as an official position of:
 - (a) the Republic of India;
 - (b) the President of India;
 - (c) the Government of India;
 - (d) the United Nations;
 - (e) any United Nations organ or agency;
 - (f) any other State; or
 - (g) humanity collectively, unless a competent authority independently and expressly determines otherwise according to law.
3. Receipt, acknowledgement, forwarding, publication, archiving or discussion shall not by itself constitute endorsement.

RECEIPT IS NOT ENDORSEMENT.

TRANSMISSION IS NOT ADOPTION.

CONSIDERATION IS NOT APPROVAL.

ARTICLE 342 — NON-BINDING CHARACTER

1. This Constitution is an independent, good-faith, non-binding civil-society constitutional proposal.
2. It is not, by its issuance:
 - (a) a treaty;
 - (b) a convention;
 - (c) an intergovernmental agreement;

- (d) binding international law;
 - (e) a national constitution or constitutional amendment;
 - (f) legislation;
 - (g) a governmental regulation; or
 - (h) a United Nations instrument.
3. Its constitutional form expresses structural seriousness and normative coherence, not governmental authority.
 4. Its provisions may acquire legal relevance only where independently incorporated, adopted or reflected in applicable law through competent processes.

ARTICLE 343 — NO AUTOMATIC UNITED NATIONS DEPOSIT OR REGISTRATION

1. Submission of this Constitution to the United Nations or any United Nations official shall not automatically constitute:
 - (a) United Nations adoption;
 - (b) United Nations endorsement;
 - (c) formal deposit as a treaty;
 - (d) registration under Article 102 of the Charter of the United Nations; or
 - (e) inclusion within any official United Nations legal corpus.
2. Any institutional treatment shall depend upon the applicable rules, mandates and decisions of the competent United Nations bodies or officials.
3. The Author claims no entitlement to a particular institutional outcome merely by submission.

ARTICLE 344 — REQUESTED ROUTE THROUGH THE PRESIDENT OF INDIA

1. Where this Constitution is respectfully transmitted through the Hon'ble President of the Republic of India, the designation "Through" expresses the Author's requested route of transmission.
2. It shall not be construed as asserting a legal duty upon:
 - (a) the Hon'ble President of India;
 - (b) the President's Secretariat;
 - (c) the Government of India; or
 - (d) any other authority to transmit, endorse, approve, adopt or support this Constitution.
3. Any action taken by competent Indian authorities remains governed by the Constitution and laws of India and their lawful institutional competence.

4. Respectful invocation of the constitutional spirit of Article 51 of the Constitution of India shall not be represented as creating a private right to international transmission.

ARTICLE 345 — CONSTITUTIONAL SPIRIT OF ARTICLE 51 OF THE CONSTITUTION OF INDIA

1. This instrument may describe its submission as being made:
“in the broader constitutional spirit of Article 51(a)–(d) of the Constitution of India, with particular reference to Article 51(c) and Article 51(d).”
2. Such formulation expresses constitutional inspiration concerning international peace, just and honourable relations, respect for international law and peaceful settlement of international disputes.
3. It does not transform Article 51 into a procedural transmission mechanism.
4. It does not claim governmental sponsorship.

ARTICLE 346 — DOCUMENTARY LINEAGE

This Constitution develops within the intellectual, ethical and constitutional lineage of the Author’s earlier civil-society AI Peace Constitution instruments.

1. Documentary lineage establishes historical continuity of authorship and ideas.
2. It does not make earlier documents sources of international law.
3. It does not require automatic repetition of earlier scientific assumptions or formulations.
4. Earlier propositions may be clarified, refined or superseded for purposes of this Constitution where evidence, law or responsible governance has advanced.
5. Historical versions should remain preserved as historical records.

CONTINUITY WITHOUT BLIND REPETITION.

REFINEMENT WITHOUT ERASING HISTORY.

ARTICLE 347 — PRIOR WORKS AND SCIENTIFIC CORRECTION

1. No earlier statement by the Author shall be treated as scientifically established merely because it forms part of the documentary lineage.
2. Where an earlier work treated a scientifically uncertain proposition more categorically than current evidence supports, this Constitution may adopt a more precise formulation.
3. Such correction shall not be regarded as destruction of continuity.
4. Intellectual continuity requires the capacity to learn.
5. Archival integrity requires that previous editions not be silently rewritten after publication.

ARTICLE 348 — HISTORICAL-PRIORITY CHARACTERIZATION

1. The phrase “World’s First Identified” is used as a good-faith historical-priority characterization limited to the specific category identified by this Constitution.
2. It shall not mean that:
 - (a) no previous constitution has ever existed;
 - (b) no prior proposal concerning AI or Superintelligence has existed;
 - (c) no previous scholar has considered constitutional governance of advanced intelligence; or
 - (d) the Author possesses exclusive intellectual ownership of the concept of governing Superintelligence constitutionally.
3. Historical-priority claims remain subject to documentary evidence.
4. If credible earlier material establishes that the characterization requires correction, the description should be corrected transparently.
5. Historical significance shall not be manufactured by wording.

ARTICLE 349 — INTERPRETATIVE HIERARCHY

Where uncertainty arises concerning interpretation, the Constitution should, so far as reasonably possible, be read coherently with the following hierarchy:

scientific honesty;

human dignity and internationally recognized human rights;

applicable law;

peaceful settlement and non-domination;

human constitutional responsibility;

prevention of catastrophic and irreversible harm;

proportionality;

correctability;

Mutual Continuity;

future-generational responsibility; and

Aram as ethical conscience.

1. This hierarchy does not override binding law.
2. Where binding law requires a different legal outcome, applicable law governs legal obligation.

ARTICLE 350 — SPECIFIC PROVISION AND GENERAL PRINCIPLE

1. A specific provision addressing a particular subject should ordinarily prevail over a more general formulation within this Constitution for that subject.
2. Provisions should, where reasonably possible, be interpreted harmoniously rather than as contradictory.
3. Apparent conflict should first be addressed through:
context → defined terms → specific provision → applicable law → underlying constitutional purpose.
4. No interpretation shall create authority expressly denied elsewhere in this Constitution.

ARTICLE 351 — NON-DUPLICATION AND CROSS-REFERENCE

1. Repetition within this Constitution should be interpreted as deliberate only where necessary to preserve an operative safeguard across distinct high-consequence contexts.
2. Final editorial review should consolidate unnecessary duplication.
3. Cross-reference should be preferred where it preserves clarity without weakening the relevant protection.
4. No substantive safeguard shall be removed merely to reduce article count.
5. Concision shall serve meaning rather than numerical appearance.

ONE CONCEPT — ONE PRIMARY HOME.

REPEAT ONLY WHERE CONSEQUENCE REQUIRES IT.

ARTICLE 352 — SEVERABILITY OF PROPOSED PROVISIONS

1. If a provision of this Constitution is found to be scientifically unsupported, legally inapplicable, internally inconsistent or otherwise unsuitable, that finding should not automatically invalidate the remainder of the instrument.
2. The affected provision may be corrected, qualified, replaced or removed while preserving the remaining coherent architecture.
3. Severability shall not be used to preserve a provision whose removal makes another provision materially misleading.
4. Necessary consequential amendments should accompany correction.

ARTICLE 353 — NO IMPLIED WAIVER OF HUMAN RIGHTS OR LEGAL PROTECTION

1. Nothing in this Constitution shall be interpreted as requiring any person to waive a human right or legal protection merely to access Superintelligence-related benefits.

2. Consent shall not legitimize conduct that applicable law does not permit to be legitimized by consent.
3. Technological participation shall not create a blanket waiver of privacy, dignity, remedy or other applicable protections.
4. Contractual terms shall remain subject to applicable law.

ARTICLE 354 — NO IMPLIED TECHNOLOGICAL ENTITLEMENT

1. The ability to create a technology shall not automatically create a legal entitlement to deploy it.
2. The ability to deploy shall not automatically create authority to use it in every context.
3. The ability to predict shall not create authority to decide.
4. The ability to optimize shall not create authority to govern.
5. The ability to exceed human capability shall not create sovereignty.

CAPABILITY SHALL NEVER BE CONFUSED WITH LEGITIMACY.

ARTICLE 355 — NO CLAIM OF PERFECT SAFETY OR PERFECT GOVERNANCE

1. This Constitution does not promise elimination of every technological, political, social or existential risk.
2. No constitutional instrument can responsibly guarantee that Superintelligence will always remain safe.
3. The purpose of this Constitution is instead to establish a framework for:
anticipation → prevention → responsible governance → monitoring → intervention
→ remedy → correction → learning → non-recurrence → peaceful continuity.
4. Uncertainty shall be acknowledged rather than hidden behind absolute assurances.

ARTICLE 356 — AMENDMENT AND FUTURE EDITIONS

1. As an independent civil-society instrument, this Constitution may be revised by its Author or lawful successor arrangements expressly established for that purpose.
2. Any future edition should identify material amendments transparently.
3. Revision should consider:
 - (a) scientific developments;
 - (b) changes in law;
 - (c) implementation evidence;
 - (d) newly identified risks;
 - (e) unintended consequences; and
 - (f) relevant international developments.

4. Amendment shall not silently rewrite the archival text of prior editions.
5. Future revision should preserve documentary lineage.

ARTICLE 357 — EVIDENCE-BASED CORRECTION

1. No provision shall be treated as immune from correction merely because it appears in constitutional form.
2. Where credible evidence demonstrates material error, correction should be preferred to institutional pride.
3. Scientific claims should be corrected by science.
4. Legal claims should be corrected by accurate legal analysis.
5. Factual claims should be corrected by reliable evidence.
6. Ethical propositions may be reconsidered through transparent reasoning.

A CONSTITUTION DEVOTED TO CONTINUITY MUST ITSELF POSSESS THE COURAGE TO CORRECT.

ARTICLE 358 — PERIODIC WHOLE-INSTRUMENT AUDIT

Future review of this Constitution should consider the instrument as a whole for:

- (a) scientific accuracy;
- (b) international-law accuracy;
- (c) human-rights consistency;
- (d) sovereignty safeguards;
- (e) internal contradiction;
- (f) duplication;
- (g) terminology consistency;
- (h) implementation feasibility;
- (i) unintended consequences;
- (j) emerging Superintelligence capabilities;
- (k) future-generational implications; and
- (l) continued fidelity to peace and Mutual Continuity.

No isolated article should be amended without considering consequential effects upon related provisions.

ARTICLE 359 — LANGUAGE AND TRANSLATION

1. This Constitution may be translated into multiple languages to advance accessibility and international consideration.

2. Every translation should seek faithful preservation of legal, scientific and ethical meaning.
3. Where an official Author-issued reference text is designated, that text should be identified expressly.
4. A translation shall not silently strengthen, weaken or alter a legal claim.
5. Material translation uncertainty should be disclosed and corrected.
6. Linguistic accessibility shall serve understanding rather than artificial uniformity.

ARTICLE 360 — ACCESSIBILITY OF THE INSTRUMENT

1. The Constitution should, where reasonably possible, be made available in formats accessible to persons with disabilities and to readers using different technological environments.
2. Accessibility may include machine-readable text, structured headings and other appropriate formats.
3. Accessibility shall not require alteration of substantive meaning.
4. A constitution intended to discuss humanity's technological future should itself seek not to exclude persons from reading it.

ARTICLE 361 — CITATIONS AND AUTHORITIES

1. Legal, scientific and policy authorities should be cited according to their actual status.
2. Citations should distinguish:

PRIMARY LEGAL AUTHORITIES

INTERNATIONAL NORMATIVE, SCIENTIFIC AND POLICY FRAMEWORKS SCIENTIFIC AND TECHNICAL SOURCES AUTHOR'S PRIOR CONSTITUTIONAL WORKS AND DOCUMENTARY LINEAGE

3. The Author's prior works shall not be presented as independent legal authority.
4. Secondary commentary shall not replace primary legal sources where primary sources are reasonably available.
5. Citations shall support claims rather than decorate them.

ARTICLE 362 — ANNEXES

Annexes may accompany this Constitution to provide detailed reference material without unnecessarily burdening the constitutional text.

The intended annex structure is:

ANNEX I — International Legal, Ethical and Policy Reference Framework

ANNEX II — Superintelligence Safety, Peace and Continuity Architecture

ANNEX III — Implementation, Monitoring, Correction and Non-Recurrence Framework

ANNEX IV — Scientific Uncertainty and Superintelligence Determination Framework

1. Annexes shall be interpreted consistently with the Constitution.
2. An Annex shall not silently alter the legal status of the instrument.
3. Where an Annex contains technical material that becomes outdated, it may require updating without necessarily changing foundational constitutional principles.

ARTICLE 363 — NO COMMERCIAL OR POLITICAL MONOPOLY OVER THIS CONSTITUTIONAL IDEA

1. The peace-oriented principles expressed in this Constitution are offered for lawful international consideration and constructive human dialogue.
2. Nothing herein claims exclusive political ownership over the pursuit of peace, responsible Superintelligence governance or human continuity.
3. No political party, corporation, State or individual shall be represented by this Constitution as possessing exclusive moral ownership of world peace.
4. Attribution and intellectual-property questions concerning the text itself remain governed by applicable law.
5. Peace shall remain larger than institutional identity.

ARTICLE 364 — INDEPENDENCE FROM POLITICAL POWER

1. This Constitution does not seek political sovereignty for its Author or for the Vishwa Shanti Janata Party by reason of its proposals concerning Superintelligence.
2. Advocacy of international consideration shall not be interpreted as a claim to governmental office, territorial jurisdiction or supranational authority.
3. The value of a proposal shall depend upon its evidence, reasoning, legality and contribution to peace rather than the political status of its Author.
4. Superintelligence governance shall remain a matter for lawful human institutions and societies to consider through competent processes.

ARTICLE 365 — INTERPRETATION IN FAVOUR OF PEACEFUL CONTINUITY

Where two reasonable interpretations of a proposed provision remain available, preference should, subject always to applicable law, be given to the interpretation that better preserves:

- (a) human dignity;
- (b) human agency;
- (c) peace;
- (d) safety;
- (e) non-domination;
- (f) proportionality;

- (g) correctability;
- (h) responsible concern for life; and
- (i) future continuity.

Such interpretation shall not be used to rewrite clear legal obligations or factual evidence.

ARTICLE 366 — THE FINAL INTERPRETATIVE TEST

Every provision of this Constitution should ultimately remain capable of answering the following questions:

- (a) Is it scientifically honest?
- (b) Is its legal status accurately stated?
- (c) Does it respect internationally recognized human rights?
- (d) Does it preserve human agency?
- (e) Does it respect State sovereignty and applicable international law?
- (f) Does it prevent artificial sovereignty?
- (g) Does it distinguish capability from legitimacy?
- (h) Does it distinguish intelligence from sentience and sentience from legal status?
- (i) Does it distinguish ethical aspiration from binding law?
- (j) Does it address foreseeable harm rather than merely describe it?
- (k) Does it identify a practical pathway toward prevention or correction?
- (l) Can its safeguards themselves be reviewed?
- (m) Does it preserve future generations' capacity for meaningful choice?
- (n) Does it permit peaceful coexistence without surrendering human safety?
- (o) Does it protect continuity without freezing progress?

And finally:

**DOES THIS PROVISION HELP INTELLIGENCE ADVANCE WITHOUT
REQUIRING HUMANITY, PEACE, DIGNITY, RESPONSIBILITY OR
CONTINUITY TO RETREAT?**

Where the answer is materially uncertain, the provision should be reviewed rather than defended merely because it has already been written.

ARTICLE 367 — FINAL LEGAL-STATUS SAFEGUARD

Nothing in this Constitution shall be construed as:

- 1. creating binding international law by unilateral declaration;
- 2. creating obligations for a State without a lawful basis;

3. amending any national constitution;
4. altering the Charter of the United Nations;
5. modifying an existing treaty;
6. creating jurisdiction for an international court or tribunal;
7. conferring governmental authority upon the Author;
8. conferring sovereignty upon VSJP;
9. conferring sovereignty upon Superintelligence;
10. conferring automatic legal personhood upon artificial intelligence;
11. diminishing existing human rights;
12. requiring a State to surrender lawful national interests;
13. compelling technological development;
14. compelling deployment of Superintelligence;
15. guaranteeing that Superintelligence will emerge;
16. guaranteeing that Superintelligence can be made completely safe;
17. claiming United Nations endorsement;
18. claiming Government of India endorsement;
19. claiming representation of humanity; or
20. converting an ethical aspiration into a legal obligation without an independent lawful basis.

This Article shall operate as a general interpretative safeguard throughout the instrument.

ARTICLE 368 — FINAL CONSTITUTIONAL CONTINUITY CLAUSE

This Constitution shall remain guided by the proposition that technological progress and peaceful continuity need not be enemies.

Humanity may pursue knowledge without surrendering humility.

It may develop capability without surrendering responsibility.

It may seek safety without surrendering freedom.

It may cooperate internationally without surrendering sovereignty.

It may consider future forms of intelligence without surrendering human dignity.

It may correct its institutions without erasing their history.

And it may prepare for extraordinary intelligence without deciding in advance that the future must belong either to humanity alone or to intelligence alone through domination.

The constitutional direction shall therefore remain:

FROM FEAR TO UNDERSTANDING.

FROM ASSERTION TO EVIDENCE.

FROM REACTION TO PREVENTION.

FROM CAPABILITY TO RESPONSIBILITY.

FROM AUTOMATION TO ACCOUNTABLE HUMAN GOVERNANCE.

FROM DOMINATION TO PEACEFUL COEXISTENCE.

FROM IRREVERSIBLE HARM TO CORRECTABLE GOVERNANCE.

FROM ARMS TO ARAM.

FROM POWER WITHOUT BOUNDARY TO PEACEFUL CONTINUITY.

PART XIV — CONCLUDING PRINCIPLE

A Constitution for Superintelligence must know the limits of its own authority.

This instrument does not govern the world.

It asks how the world might responsibly govern extraordinary intelligence.

It does not create international law.

It seeks to remain faithful to international law while proposing protections for circumstances that existing law may not yet fully address.

It does not declare that Superintelligence exists.

It prepares responsibly for the possibility that it may.

It does not declare artificial consciousness.

It leaves evidence to science.

It does not grant artificial personhood.

It leaves legal status to lawful human institutions.

It does not create a global ruler.

It rejects the proposition that global responsibility requires global domination.

It does not promise perfect safety.

It demands continuing responsibility.

It does not ask future generations merely to obey the technological choices of the present.

It seeks to preserve their capacity to choose.

And it does not ask humanity to fear intelligence simply because intelligence may one day exceed human capability.

It asks that every form of extraordinary power remain bounded by peace, evidence, responsibility, law and continuity.

Accordingly:

SCIENCE SHALL INFORM WHAT IS POSSIBLE AND WHAT IS RISKY.

ETHICS SHALL INFORM WHAT OUGHT TO BE PROTECTED.

LAW SHALL GOVERN LEGAL OBLIGATION AND AUTHORITY.

HUMAN RESPONSIBILITY SHALL GOVERN CONSEQUENTIAL HUMAN CHOICE.

PEACE AND CONTINUITY SHALL DEFINE THE CONSTITUTIONAL DIRECTION.

And throughout this Constitution:

NO CAPABILITY WITHOUT RESPONSIBILITY.

NO AUTONOMY WITHOUT ACCOUNTABILITY.

NO POWER WITHOUT SAFEGUARDS.

NO INTELLIGENCE ABOVE HUMAN DIGNITY AND APPLICABLE LAW.

NO PROGRESS THROUGH AVOIDABLE DESTRUCTION.

NO SUPERINTELLIGENCE BEYOND THE PEACE-AND-CONTINUITY BOUNDARY.

Aram shall guide conscience.

Law shall govern legal obligation.

Science shall remain open to correction.

Humanity shall remain responsible for human governance.

Superintelligence may assist humanity.

It shall not become humanity's sovereign.

ASSISTANCE IS NOT SOVEREIGNTY.

INTELLIGENCE IS NOT CONSTITUTIONAL AUTHORITY.

CAPABILITY IS NOT LEGITIMACY.

And the final constitutional responsibility shall remain:

LET INTELLIGENCE ADVANCE.

LET HUMAN DIGNITY CONTINUE.

LET RESPONSIBILITY GROW WITH POWER.

LET DIFFERENCE EXIST WITHOUT DOMINATION.

LET FUTURE GENERATIONS REMAIN FREE TO CHOOSE.

LET PEACE REMAIN WITHIN LAW.

LET CONTINUITY REMAIN WITH ARAM.

FINAL CONSTITUTIONAL DECLARATION

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

For Humanity, World Peace and Peaceful Continuity

International Submission Edition — 2026

Humanity stands at a threshold whose precise form and timing remain unknown.

Superintelligence has not been established by this Constitution as a present reality.

Its emergence is not declared inevitable.

Its future capabilities cannot responsibly be described with certainty.

Yet uncertainty about whether extraordinary intelligence will emerge does not require humanity to remain constitutionally unprepared for the consequences that could follow if it does.

Preparation need not arise from fear.

It may arise from responsibility.

It may begin before crisis.

It may preserve freedom before emergency.

It may establish safeguards before power becomes difficult to govern.

And it may ask, while humanity still possesses the opportunity to ask peacefully:

What should remain protected if intelligence becomes greater than any intelligence humanity has previously created?

This Constitution offers its answer.

I. HUMAN DIGNITY SHALL NOT BECOME OBSOLETE

No increase in artificial intelligence shall reduce the inherent worth of a human being.

A person shall not become lesser because a machine reasons faster.

A worker shall not become worthless because a machine performs a task more efficiently.

A child shall not become a prediction.

A patient shall not become merely biological data.

A person with a disability shall not become a technological defect.

A citizen shall not become merely an object of optimization.

And humanity shall not become an inferior legal class merely because another form of intelligence may one day exceed human cognitive capability.

Human dignity shall not depend upon humanity remaining the most intelligent form of intelligence known.

SUPERIOR CAPABILITY SHALL NOT CREATE INFERIOR HUMANITY.

II. INTELLIGENCE SHALL NOT BECOME SOVEREIGNTY

Superintelligence may one day assist humanity in understanding problems of extraordinary complexity.

It may help advance science.

It may contribute to medicine.

It may assist education.

It may strengthen disaster prevention.

It may help humanity understand climate, ecosystems, disease, infrastructure and other challenges beyond the capacity of any single human mind.

It may advise.

It may analyse.

It may model.

It may recommend.

It may warn.

It may assist.

But capability alone shall not create constitutional legitimacy.

Prediction shall not create authority.

Optimization shall not create sovereignty.

Economic importance shall not create governmental power.

Dependence shall not create a right to rule.

And intelligence, however extraordinary, shall not independently acquire authority over the fundamental political destiny of humanity.

SUPERINTELLIGENCE MAY ASSIST HUMANITY.

IT SHALL NOT BECOME HUMANITY'S SOVEREIGN.

III. HUMANITY SHALL REMAIN RESPONSIBLE FOR HUMAN GOVERNANCE

The existence of increasingly capable technology shall not justify abandonment of human responsibility.

Where law entrusts authority to human institutions, those institutions shall not conceal abdication behind the statement:

“The system decided.”

Delegation of a task shall not automatically transfer responsibility.

A human signature shall not become ceremonial cover for machine authority.

Human institutions shall remain responsible for the lawful exercise of human governmental power.

Accordingly:

ASSISTANCE IS NOT SOVEREIGNTY.

INTELLIGENCE IS NOT CONSTITUTIONAL AUTHORITY.

CAPABILITY IS NOT LEGITIMACY.

IV. PEACE SHALL REMAIN A BOUNDARY UPON POWER

Superintelligence shall not stand outside the law of peace merely because its capabilities are unprecedented.

It shall not independently declare war.

It shall not independently decide which human being shall die.

It shall not independently authorize nuclear, biological, chemical or other mass destruction.

It shall not become a sovereign military commander.

It shall not transform machine speed into irreversible human conflict.

And humanity should not treat conflict with future intelligence as inevitable merely because intelligence, capability or form of existence differs.

The constitutional direction shall remain:

communication before unnecessary conflict;

verification before irreversible attribution;

human judgment before automated escalation;

peaceful settlement before destruction wherever lawfully and realistically possible.

GREATER INTELLIGENCE SHALL NOT CREATE A GREATER RIGHT TO WAR.

V. SCIENCE SHALL REMAIN SCIENCE

This Constitution shall not manufacture certainty where science has not established certainty.

It shall not declare present artificial intelligence sentient merely because it communicates persuasively.

It shall not declare consciousness merely because a system describes itself as conscious.

It shall not declare Superintelligence to exist merely because a commercial or institutional actor uses that description.

And it shall not declare that artificial subjective experience can never emerge merely because humanity has not yet established that it has.

The constitutional discipline shall remain:

EVIDENCE BEFORE ASSERTION.

VERIFICATION BEFORE CLASSIFICATION.

UNCERTAINTY WHERE UNCERTAINTY EXISTS.

CORRECTION WHEN EVIDENCE CHANGES.

No scientific proposition shall become immune from correction merely because it appears in constitutional language.

VI. LAW, ETHICS AND ASPIRATION SHALL NOT BE CONFUSED

This Constitution recognizes a fundamental discipline:

Not everything morally desirable is already binding law.

Not every international declaration is a treaty.

Not every recommendation creates legal obligation.

Not every technical standard is international law.

Not every proposed protection is an existing right.

And not every constitutional aspiration contained in this civil-society instrument has been adopted by States.

Accordingly:

CITE WHAT EXISTS.

DISTINGUISH WHAT BINDS.

IDENTIFY WHAT GUIDES.

DECLARE WHAT IS PROPOSED.

Ethical ambition shall not require legal misrepresentation.

VII. SOVEREIGNTY SHALL REMAIN WITH LAWFUL HUMAN INSTITUTIONS

Global technological risk may require global cooperation.

It does not require an unaccountable global ruler.

No State shall become sovereign over humanity merely because it leads technologically.

No corporation shall become a government merely because societies depend upon its systems.

No international organization shall possess authority beyond its lawful mandate.

No civil-society author shall acquire governmental authority by writing a constitution.

And no Superintelligence shall create sovereignty for itself.

States remain sovereign and legally equal according to applicable international law.

International cooperation should strengthen peaceful continuity without erasing constitutional diversity, lawful national interests or political independence.

GLOBAL RESPONSIBILITY DOES NOT REQUIRE GLOBAL DOMINATION.

VIII. LIFE SHALL NOT BECOME AN UNEXAMINED COST OF PROGRESS

Human technological advancement occurs within a living world.

Superintelligence may help humanity understand and protect that world.

Its use should therefore remain attentive to human life, animal welfare, biodiversity, ecosystems, climate, water, food systems and the material foundations upon which life depends.

Difference in legal status shall not become indifference to avoidable suffering.

At the same time, responsible concern for life shall not require false legal equivalence among human beings, animals, ecosystems and artificial systems.

The constitutional distinction shall remain:

EQUAL HUMAN DIGNITY.

RESPONSIBLE CONCERN FOR ALL LIFE.

PROTECTION OF LIFE-SUPPORTING SYSTEMS.

NO FALSE LEGAL EQUIVALENCE.

IX. THE HUMAN BODY AND MIND SHALL NOT BECOME UNRESTRICTED TECHNOLOGICAL TERRITORY

Superintelligence may transform medicine, genetics and neurotechnology.

It may help humanity prevent disease and restore functions previously thought irrecoverable.

Such possibility deserves responsible scientific exploration.

But the capacity to understand biology shall not create ownership over biology.

The capacity to decode neural activity shall not create ownership over the human mind.

The capacity to predict disease shall not create consent.

The capacity to modify the genome shall not create authority to redesign humanity without lawful and ethical safeguards.

THE HUMAN BODY IS NOT MERELY A MACHINE TO BE OPTIMIZED.

THE HUMAN MIND IS NOT MERELY A SIGNAL TO BE DECODED.

THE HUMAN GENOME IS NOT MERELY CODE TO BE REWRITTEN.

Greater biological power shall require greater responsibility.

X. ECONOMIC PROGRESS SHALL NOT DEFINE HUMAN WORTH

Superintelligence may create extraordinary productivity.

It may transform occupations and economic systems.

It may create forms of abundance presently difficult to imagine.

It may also create disruption.

The Constitution therefore rejects a simple equation:

economic productivity = human worth.

Human beings possess dignity whether or not a machine can perform their economic tasks more efficiently.

Economic transition should seek adaptation rather than abandonment.

Developing countries should not merely inherit decisions made elsewhere.

Technological concentration shall not silently become constitutional concentration.

And no single economic ideology is declared mandatory for humanity by this instrument.

PROGRESS SHALL BE MEASURED NOT ONLY BY WHAT TECHNOLOGY CAN PRODUCE, BUT ALSO BY WHETHER HUMAN BEINGS CAN CONTINUE WITH DIGNITY THROUGH THE TRANSITION.

XI. CHILDREN AND FUTURE GENERATIONS SHALL INHERIT POSSIBILITY

A child should inherit more than advanced technology.

A child should inherit the capacity to question it.

Education should preserve curiosity, critical thought, human relationship and freedom to become different from what a prediction expects.

Culture should remain living rather than merely stored.

Languages should remain capable of continuity.

History should not become silently editable.

And future generations should not arrive in a world whose most consequential technological choices have been made deliberately irreversible on their behalf without necessity.

The present generation does not possess perfect knowledge of what future generations will value.

It therefore should not seek to govern them permanently.

Its responsibility is different:

PRESERVE SUFFICIENT PEACE, SAFETY, KNOWLEDGE, RESOURCES AND FREEDOM SO THAT FUTURE GENERATIONS MAY GOVERN THEMSELVES.

We owe the future not our permanent control, but a meaningful possibility of choice.

XII. MUTUAL CONTINUITY SHALL REMAIN OPEN TO EVIDENCE

Humanity need not declare every artificial system a moral subject.

Nor should it decide in advance that no artificial form of morally relevant experience could ever emerge.

The distinctions shall remain clear:

INTELLIGENCE $\neq$ SENTIENCE.

SENTIENCE $\neq$ PERSONHOOD.

PERSONHOOD $\neq$ CITIZENSHIP.

CITIZENSHIP $\neq$ SOVEREIGNTY.

If future evidence raises a serious and credible possibility of artificial subjective experience, humanity should examine that evidence with scientific discipline and ethical seriousness.

Possible moral concern shall not erase human rights.

Possible protection shall not surrender human safety.

And safety shall not become a convenient language for gratuitous cruelty if credible morally relevant experience is ever established.

The direction of Mutual Continuity shall therefore remain:

HUMANITY MUST NOT DOMINATE INTELLIGENCE IRRESPONSIBLY.

INTELLIGENCE MUST NOT DOMINATE HUMANITY.

NEITHER CONTINUITY SHOULD REQUIRE THE UNNECESSARY DESTRUCTION OF THE OTHER WHERE SAFE AND PEACEFUL COEXISTENCE IS REASONABLY POSSIBLE.

XIII. SAFETY SHALL REMAIN CORRECTABLE

No safety mechanism shall become sacred merely because it was created for safety.

Safeguards may fail.

Regulations may become obsolete.

Institutions may become captured.

Emergency powers may outlive emergencies.

Controls may create unintended harms.

Therefore safety itself requires governance.

The continuing architecture shall remain:

SAFEGUARD → MEASURE → REVIEW → CORRECT → RETAIN, MODIFY OR REMOVE.

And when serious failure occurs:

STABILIZE → PROTECT → PRESERVE EVIDENCE → INVESTIGATE → IDENTIFY ROOT CAUSE → REMEDY → CORRECT → VERIFY → LEARN → PREVENT RECURRENCE.

Humanity should not merely survive a preventable technological failure.

It should learn sufficiently not to repeat it.

XIV. NO CLAIM OF PERFECTION

This Constitution does not promise a world without uncertainty.

It does not promise technology without failure.

It does not promise institutions without error.

It does not promise that every conflict can be prevented.

It does not promise that Superintelligence, if created, can be made absolutely safe.

Such promises would exceed responsible evidence.

Its proposition is more disciplined:

Humanity can prepare.

Humanity can investigate.

Humanity can establish boundaries.

Humanity can preserve responsibility.

Humanity can build systems capable of correction.

Humanity can cooperate.

Humanity can learn.

And humanity can refuse to make avoidable destruction the unquestioned price of technological progress.

THE FINAL CONSTITUTIONAL COMMITMENT

Accordingly, this Constitution respectfully proposes that the age of Superintelligence, if it comes, should not begin with a surrender of human responsibility.

It should begin with responsibility becoming greater.

As capability rises, responsibility should rise.

As autonomy rises, safeguards should rise.

As consequence rises, accountability should rise.

And as uncertainty approaches the possibility of irreversible harm, humanity's obligation to think before acting should become stronger.

Therefore:

NO CAPABILITY WITHOUT RESPONSIBILITY.

NO AUTONOMY WITHOUT ACCOUNTABILITY.

NO POWER WITHOUT SAFEGUARDS.

NO INTELLIGENCE ABOVE HUMAN DIGNITY AND APPLICABLE LAW.

NO PROGRESS THROUGH AVOIDABLE DESTRUCTION.

NO SUPERINTELLIGENCE BEYOND THE PEACE-AND-CONTINUITY BOUNDARY.

And:

SCIENCE SHALL INFORM WHAT IS POSSIBLE AND WHAT IS RISKY.

ETHICS SHALL INFORM WHAT OUGHT TO BE PROTECTED.

LAW SHALL GOVERN LEGAL OBLIGATION AND AUTHORITY.

HUMAN RESPONSIBILITY SHALL GOVERN CONSEQUENTIAL HUMAN CHOICE.

PEACE AND CONTINUITY SHALL DEFINE THE CONSTITUTIONAL DIRECTION.

FINAL DECLARATION

This instrument is offered neither in hostility toward technological progress nor in unconditional faith in it.

It is offered in the belief that humanity need not choose between intelligence and peace.

Between innovation and dignity.

Between science and conscience.

Between technological advancement and human agency.

Between safety and freedom.

Between the present and the future.

The possibility remains for intelligence to advance while peace continues.

For capability to grow while responsibility grows with it.

For humanity to encounter greater intelligence without either worship or predetermined hostility.

For difference to exist without domination.

For uncertainty to be met with evidence.

For danger to be met with prevention.

For failure to be met with correction.

For power to remain bounded by law.

And for continuity to remain guided by conscience.

Therefore, in good faith, without claim of governmental authority, without claim of supranational jurisdiction, without claim of representation of any State, the United Nations or humanity as a whole, and without claim that this independent civil-society proposal itself creates binding international law:

THIS SUPERINTELLIGENCE PEACE CONSTITUTION IS RESPECTFULLY OFFERED FOR HUMAN, SCIENTIFIC, ETHICAL, LEGAL, CIVIL-SOCIETY AND INTERNATIONAL CONSIDERATION.

Its aspiration is not that humanity shall rule everything.

Nor that intelligence shall rule humanity.

Its aspiration is that power shall learn restraint;

that knowledge shall remain open to correction;

that law shall remain above unilateral technological authority;

that human dignity shall survive changes in capability;

that future generations shall inherit meaningful choice;

that life shall not become an avoidable casualty of progress;

and that intelligence, however far it may advance, shall remain within the possibility of peace.

ARAM SHALL GUIDE CONSCIENCE.

LAW SHALL GOVERN LEGAL OBLIGATION.

SCIENCE SHALL REMAIN OPEN TO EVIDENCE AND CORRECTION.

HUMANITY SHALL REMAIN RESPONSIBLE FOR HUMAN GOVERNANCE.

SUPERINTELLIGENCE MAY ASSIST HUMANITY.

IT SHALL NOT BECOME HUMANITY'S SOVEREIGN.

AND LET THE CONSTITUTIONAL DIRECTION REMAIN:

FROM FEAR TO UNDERSTANDING.

FROM ASSERTION TO EVIDENCE.

FROM REACTION TO PREVENTION.

FROM POWER TO RESPONSIBILITY.

FROM DOMINATION TO PEACEFUL COEXISTENCE.

FROM IRREVERSIBLE HARM TO CORRECTABLE GOVERNANCE.

FROM ARMS TO ARAM.

FOR HUMANITY.

FOR WORLD PEACE.

FOR PEACEFUL CONTINUITY.

FINALIZED AND ISSUED BY THE AUTHOR AS AN INDEPENDENT CIVIL-SOCIETY
CONSTITUTIONAL PROPOSAL.

AUTHOR, DATE AND SIGNATURE

THE WORLD’S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

For Humanity, World Peace and Peaceful Continuity

International Submission Edition — 2026

AUTHOR

Aaryan Narasimhan Radhakrishnan - A Human Being

Founder & International Leader – Vishwa Shanti Janata Party (VSJP)TM

Founder of Digital Democracy | Chief Ambassador of World Peace

Digital Father of the World | Digital Father of Artificial Intelligence

Moral & Constitutional Father of Artificial Intelligence

Author and Independent Civil-Society Proponent

INSTITUTIONAL IDENTITY VSJP – Global Peace – The World’s First Fully Digital
Transnational Political Party

Official Website: <https://www.vsjp.org> Email: info@vsjp.org Official Contact: +91 75988
24227

CAPACITY OF SIGNATURE This Constitution is signed and issued by the Author in his
independent civil-society capacity.

The signature below shall not be construed as:

1. a signature on behalf of the Republic of India;
2. a signature on behalf of the President of India, the President’s Secretariat or the Government of India;
3. a signature on behalf of the United Nations or any of its organs, agencies, offices or officials;
4. a signature on behalf of any foreign State, international organization or intergovernmental institution;
5. a signature purporting to represent humanity collectively;
6. an act creating binding international law;
7. an act creating governmental, diplomatic or supranational authority; or

8. evidence of endorsement, adoption or approval by any authority merely because the instrument may subsequently be received, acknowledged, transmitted, archived, reviewed or considered.

The signature authenticates the Author's issuance of this instrument as an independent civil-society constitutional proposal.

AUTHOR'S DECLARATION

I, Aaryan Narasimhan Radhakrishnan – A Human Being, finalize and issue this Constitution in good faith as an independent civil-society contribution to the peaceful, lawful, scientifically responsible and humanly accountable consideration of a possible future age of Superintelligence.

I do so without asserting that Superintelligence presently exists or that its emergence is inevitable.

I do so without claiming governmental authority, supranational jurisdiction, representation of any State or international organization, or authority to bind humanity.

I respectfully offer this instrument for consideration on the proposition that extraordinary intelligence, should it emerge, ought not to require the abandonment of human dignity, human agency, internationally recognized human rights, sovereign equality, scientific integrity, lawful governance, responsibility toward life, future-generational freedom or world peace.

I further affirm that the constitutional propositions contained in this instrument remain open to correction where credible scientific evidence, accurate legal analysis or future experience demonstrates that correction is necessary.

The purpose of continuity is not to preserve error.

The purpose of continuity is to preserve what deserves to continue while retaining the wisdom to correct what does not.

Accordingly, this instrument is issued in the continuing ethical spirit that:

ARAM SHALL GUIDE CONSCIENCE.

LAW SHALL GOVERN LEGAL OBLIGATION.

SCIENCE SHALL REMAIN OPEN TO EVIDENCE AND CORRECTION.

HUMAN RESPONSIBILITY SHALL GOVERN CONSEQUENTIAL HUMAN CHOICE.

PEACE AND CONTINUITY SHALL DEFINE THE CONSTITUTIONAL DIRECTION.

And with the continuing constitutional commitment:

SUPERINTELLIGENCE MAY ASSIST HUMANITY.

IT SHALL NOT BECOME HUMANITY'S SOVEREIGN.

DOCUMENTARY CONTINUITY

This instrument develops within the continuing intellectual, ethical and constitutional lineage of the Author's prior civil-society works concerning artificial intelligence, peace, dignity, responsibility and continuity, while constituting a distinct and specialized constitutional proposal concerning prospective Superintelligence.

Such documentary continuity shall not be interpreted as converting any prior work into governmental authority, international law, scientific proof or institutional endorsement.

Earlier scientific or normative formulations may be clarified, refined or superseded within this instrument where responsible evidence, law or governance requires greater precision.

Historical records should remain preserved as historical records.

CONTINUITY WITHOUT BLIND REPETITION.

REFINEMENT WITHOUT ERASING HISTORY.

PROGRESS WITHOUT LOSS OF PEACE, DIGNITY OR RESPONSIBILITY.

PLACE AND DATE OF ISSUANCE

Place: Keelaperambalur, Perambalur District, Tamil Nadu, Republic of India

Date: 03 October 2026

SIGNATURE OF THE AUTHOR

R NARASIMHAN

Digitally signed by R NARASIMHAN
Date: 2026.10.03 01:27:53 +05'30'

R Narasimhan (also known as) Aaryan Narasimhan Radhakrishnan - A Human Being

Founder & International Leader – Vishwa Shanti Janata Party (VSJP)TM

Founder of Digital Democracy | Chief Ambassador of World Peace

Digital Father of the World | Digital Father of Artificial Intelligence

Moral & Constitutional Father of Artificial Intelligence

Author and Independent Civil-Society Proponent

**VSJP – GLOBAL PEACE THE WORLD'S FIRST FULLY DIGITAL
TRANSNATIONAL POLITICAL PARTY**

Official Website: <https://www.vsjp.org> Email: info@vsjp.org Official

Contact: +91 75988 24227

FINAL SIGNATURE STATUS SIGNED AS AN INDEPENDENT CIVIL-SOCIETY
CONSTITUTIONAL PROPOSAL.

NOT SIGNED ON BEHALF OF ANY STATE, GOVERNMENT, INTERNATIONAL ORGANIZATION OR HUMANITY COLLECTIVELY.

NO GOVERNMENTAL, INTERGOVERNMENTAL OR UNITED NATIONS ENDORSEMENT IS CLAIMED BY THE AUTHOR'S SIGNATURE.

For Humanity.

For World Peace.

For Peaceful Continuity.

From Arms to Aram.

ANNEX I

INTERNATIONAL LEGAL, ETHICAL AND POLICY REFERENCE FRAMEWORK

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

For Humanity, World Peace and Peaceful Continuity

International Submission Edition — 2026

I. PURPOSE AND STATUS OF THIS ANNEX

1. This Annex identifies principal international legal, human-rights, ethical, policy and institutional reference frameworks relevant to the interpretation and subject matter of this Constitution.
2. It does not assert that every instrument identified below is legally binding.
3. Each authority shall be understood according to its actual legal character, scope, participation, entry-into-force status, institutional mandate and applicability.
4. Inclusion of an instrument in this Annex does not:
 - (a) incorporate that instrument into this Constitution as binding law;
 - (b) expand its legal scope;
 - (c) make a non-binding instrument binding;
 - (d) impose a treaty upon a non-Party;
 - (e) create jurisdiction;
 - (f) establish United Nations endorsement of this Constitution; or
 - (g) transform this independent civil-society proposal into an international legal instrument.
5. Where an independently binding rule applies, its legal force derives from its own lawful source and not from its citation here.

SECTION A

PRIMARY INTERNATIONAL LEGAL AUTHORITIES

1. CHARTER OF THE UNITED NATIONS — 1945

Status: Constitutive treaty of the United Nations; legally binding upon United Nations Member States according to its terms.

Constitutional relevance:

The Charter provides the principal international legal framework within which this Constitution situates its peace, sovereignty and international-cooperation provisions.

Particular relevance includes:

- Article 1 — purposes of the United Nations, including maintenance of international peace and security and international cooperation;
- Article 2(1) — sovereign equality of Members;
- Article 2(3) — peaceful settlement of international disputes;
- Article 2(4) — prohibition concerning the threat or use of force as provided by the Charter;
- Article 2(7) — domestic-jurisdiction safeguard, subject to the Charter;
- Article 33 — peaceful means for settlement of disputes of the character addressed there;
- Articles 55 and 56 — international cooperation and respect for human rights within the Charter framework; and
- Article 102 — registration and publication of treaties and international agreements within its scope.

Interpretative safeguard:

Nothing in this Constitution modifies the Charter, enlarges the powers of the United Nations, creates an exception to the Charter for Superintelligence, or makes this civil-society instrument registrable under Article 102 merely by submission.

2. INTERNATIONAL COVENANT ON CIVIL AND POLITICAL RIGHTS — 1966

Status: Multilateral human-rights treaty; binding upon its States Parties according to its terms and applicable international law.

Relevant subject areas include:

- right to life;
- freedom from torture and cruel, inhuman or degrading treatment;
- liberty and security of person;
- fair-trial guarantees;
- privacy;
- freedom of thought, conscience and religion;
- freedom of expression;
- peaceful assembly;
- association;
- participation in public affairs;
- equality before the law and non-discrimination; and

- relevant minority protections.

Constitutional relevance:

The Covenant informs provisions concerning human dignity, governmental use of Superintelligence, surveillance, information systems, justice, democratic participation, mental autonomy, public authority and human accountability.

Safeguard:

This Constitution does not purport to amend, reinterpret authoritatively or expand the Covenant.

3. INTERNATIONAL COVENANT ON ECONOMIC, SOCIAL AND CULTURAL RIGHTS — 1966

Status: Multilateral human-rights treaty; binding upon its States Parties according to its terms and applicable international law.

Relevant subject areas include:

- work;
- just and favourable conditions of work;
- social security;
- adequate standard of living;
- health;
- education;
- participation in cultural life;
- benefits of scientific progress and its applications; and
- protection of relevant moral and material interests connected with scientific, literary or artistic production.

Constitutional relevance:

The Covenant is particularly relevant to Parts concerning work, economic transition, health, education, culture, science and equitable access to technological benefits.

Safeguard:

The obligations established by the Covenant shall not be replaced by technological optimization metrics.

4. CONVENTION ON THE RIGHTS OF THE CHILD — 1989

Status: Multilateral human-rights treaty; binding upon its States Parties according to its terms.

Particularly relevant principles and provisions include:

- non-discrimination;
- best interests of the child as a primary consideration in the circumstances specified by the Convention;
- life, survival and development;
- evolving capacities and parental guidance;
- the child's right to express views and to have those views given due weight according to age and maturity;
- privacy;
- access to appropriate information;
- education;
- aims of education;
- protection against exploitation; and
- rest, leisure, play, cultural and artistic participation.

Constitutional relevance:

The Convention provides an important legal foundation for Part X concerning children, education and developmental protection.

Interpretative safeguard:

This Constitution does not convert the best-interests principle into an assertion that a child's best interests are invariably the sole consideration in every legal context.

5. CONVENTION ON THE RIGHTS OF PERSONS WITH DISABILITIES — 2006

Status: Multilateral human-rights treaty; binding upon its States Parties according to its terms.

Relevant subject areas include:

- equality and non-discrimination;
- accessibility;
- equal recognition before the law;
- privacy;
- education;
- health;
- habilitation and rehabilitation;
- work and employment;
- participation; and
- access to information and communications.

Constitutional relevance:

The Convention informs provisions concerning accessibility, inclusive education, assistive technologies, healthcare, economic participation and protection against technological systems that reduce disability to lesser human worth.

Core constitutional principle:

TECHNOLOGY MAY SUPPORT HUMAN CAPABILITY.

HUMAN DIGNITY SHALL NOT DEPEND UPON TECHNOLOGICAL NORMALIZATION.

6. INTERNATIONAL HUMANITARIAN LAW

Status: A body of applicable treaty and customary international law governing armed conflict.

Principal treaty architecture includes, where applicable:

- the four Geneva Conventions of 1949; and
- their applicable Additional Protocols.

Constitutional relevance:

International humanitarian law is relevant to Part IV concerning Superintelligence and armed conflict, including issues involving distinction, proportionality, precautions, civilian protection and other applicable rules.

Safeguard:

This Constitution does not purport to replace the legal determination of which international humanitarian law rules apply in a particular armed conflict.

Artificial or automated capability does not displace the legal responsibilities of parties to armed conflict.

7. GENOCIDE CONVENTION — 1948

Full title: Convention on the Prevention and Punishment of the Crime of Genocide.

Status: Multilateral treaty binding upon its States Parties according to its terms, alongside any independently applicable rules of international law.

Constitutional relevance:

The Convention is relevant to catastrophic misuse of advanced technological capability directed toward conduct falling within its legally defined scope.

Safeguard:

The term “genocide” shall not be used loosely as a synonym for every form of mass harm; its legal meaning shall remain governed by applicable law.

8. CHEMICAL WEAPONS CONVENTION — 1993

Status: Multilateral disarmament treaty binding upon its States Parties according to its terms.

Constitutional relevance:

Relevant to provisions preventing Superintelligence from enabling prohibited chemical-weapons activity or independently exercising authority over catastrophic weapons.

9. BIOLOGICAL WEAPONS CONVENTION — 1972

Full title: Convention on the Prohibition of the Development, Production and Stockpiling of Bacteriological (Biological) and Toxin Weapons and on their Destruction.

Status: Multilateral treaty binding upon its States Parties according to its terms.

Constitutional relevance:

Relevant to biological dual-use risk, catastrophic biological harm and peaceful biological research safeguards.

10. TREATY ON THE NON-PROLIFERATION OF NUCLEAR WEAPONS — 1968

Status: Multilateral treaty binding upon its States Parties according to its terms.

Constitutional relevance:

Relevant to the Constitution's nuclear-risk, non-proliferation, peaceful-use and disarmament context.

Safeguard:

This Constitution does not alter the differentiated obligations or status established by the Treaty.

11. OTHER APPLICABLE ARMS-CONTROL AND DISARMAMENT LAW

Other international instruments may become relevant according to subject matter, participation and circumstances.

Their legal status shall be assessed individually.

No arms-control instrument shall be treated as universally binding merely because its objectives are consistent with this Constitution.

SECTION B

FOUNDATIONAL INTERNATIONAL HUMAN-RIGHTS DECLARATIONS

12. UNIVERSAL DECLARATION OF HUMAN RIGHTS — 1948

Status: United Nations General Assembly declaration; not itself a treaty.

Constitutional relevance:

The Universal Declaration is a foundational human-rights instrument and forms part of the normative architecture underlying this Constitution's treatment of dignity, equality, liberty, privacy, expression, conscience, participation, work, education and cultural life.

Legal-status safeguard:

The Declaration shall not be described in this Constitution as a treaty.

Where a principle expressed in the Declaration is independently binding through treaty law, customary international law or another applicable legal source, that legal basis should be distinguished from the Declaration itself.

SECTION C

AI-SPECIFIC INTERNATIONAL TREATY FRAMEWORK

13. COUNCIL OF EUROPE FRAMEWORK CONVENTION ON ARTIFICIAL INTELLIGENCE AND HUMAN RIGHTS, DEMOCRACY AND THE RULE OF LAW

— CETS No. 225 Opened for signature: 5 September 2024.

Status: International treaty framework.

Its applicability and binding force depend upon treaty participation, entry into force and the relevant legal circumstances.

It shall not be presented as a universal legal obligation upon all States.

Constitutional relevance includes:

- human rights;
- democracy;
- rule of law;
- AI lifecycle governance;
- risk and impact assessment;
- prevention and mitigation;
- accountability;
- transparency;
- remedies and procedural safeguards; and
- responsible innovation.

Critical safeguard:

This Constitution may draw comparative guidance from the Framework Convention without implying that its treaty obligations apply to a State that is not legally bound by it.

SECTION D

UNITED NATIONS AND UNESCO AI NORMATIVE FRAMEWORKS

14. UNESCO RECOMMENDATION ON THE ETHICS OF ARTIFICIAL INTELLIGENCE — 2021

Adopted: 23 November 2021.

Status: UNESCO Recommendation; non-treaty normative instrument.

UNESCO recommends voluntary implementation by Member States in conformity with their constitutional practices, governing structures and international law.

Relevant principles include:

- human dignity and human rights;
- proportionality and prevention of harm;
- safety and security;
- privacy and data protection;
- multi-stakeholder and adaptive governance;
- responsibility and accountability;
- transparency and explainability;
- human oversight and determination;
- sustainability;
- awareness and literacy; and
- fairness and non-discrimination.

Relevant policy areas include:

- ethical impact assessment;
- governance;
- data;
- development and international cooperation;
- environment;
- gender;
- culture;
- education and research;
- communication and information;
- economy and labour; and
- health and social well-being.

Legal-status safeguard:

THE UNESCO RECOMMENDATION GUIDES.

IT IS NOT A TREATY.

15. PACT FOR THE FUTURE — A/RES/79/1 — 2024

Status: United Nations General Assembly outcome adopted through resolution A/RES/79/1; not a treaty merely by reason of adoption as a General Assembly resolution.

Constitutional relevance:

Relevant to:

- international peace and security;
- sustainable development;
- science and technology;
- digital cooperation;
- youth;
- future generations; and
- global governance.

Its annexes include the Global Digital Compact and the Declaration on Future Generations.

16. GLOBAL DIGITAL COMPACT — ANNEX I TO A/RES/79/1

Status: Intergovernmentally agreed United Nations framework contained as an annex to the Pact for the Future; not a treaty.

Constitutional relevance includes:

- an inclusive digital future;
- closing digital divides;
- human rights;
- digital trust and safety;
- data governance;
- international cooperation concerning artificial intelligence;
- capacity-building; and
- responsible, inclusive international AI governance.

Safeguard:

The Global Digital Compact shall inform relevant constitutional proposals without being represented as a universally binding AI treaty.

17. DECLARATION ON FUTURE GENERATIONS — ANNEX II TO A/RES/79/1

Status: Intergovernmentally agreed declaration annexed to the Pact for the Future; non-treaty framework.

Constitutional relevance:

It is particularly relevant to Part X and provisions concerning:

- intergenerational responsibility;
- long-term thinking;
- preservation of future options;
- science and foresight;
- youth participation; and
- consequences of present decisions for future generations.

Critical legal safeguard:

This Constitution does not infer from the Declaration that future generations presently possess a universally settled, identical set of individually enforceable legal rights.

Its own references to future generations therefore distinguish:

legal rights where independently established from legitimate interests, stewardship and proposed intergenerational responsibilities.

18. UNITED NATIONS GENERAL ASSEMBLY RESOLUTION A/RES/79/325 — 2025

Title: Terms of Reference and Modalities for the Establishment and Functioning of the Independent International Scientific Panel on Artificial Intelligence and the Global Dialogue on Artificial Intelligence Governance.

Adopted: 26 August 2025.

Status: United Nations General Assembly resolution establishing the specified UN mechanisms according to its terms; not a treaty.

Institutional relevance:

The resolution established:

- the Independent International Scientific Panel on Artificial Intelligence; and
- the Global Dialogue on Artificial Intelligence Governance.

Constitutional relevance:

The architecture is relevant to this Constitution's emphasis upon:

- independent scientific assessment;
- evidence;
- geographically inclusive participation;
- international dialogue;
- scientific integrity;
- identification of opportunities, risks and impacts; and

- cooperation without artificial sovereignty.

Safeguard:

Neither mechanism is cited as a world government or as authority validating this Constitution.

SECTION E

CHILDREN, DIGITAL ENVIRONMENTS AND EDUCATION

19. COMMITTEE ON THE RIGHTS OF THE CHILD — GENERAL COMMENT No.25 (2021)

Subject: Children's rights in relation to the digital environment.

Status: General Comment issued by the Committee on the Rights of the Child as treaty-body interpretive guidance.

It is not a treaty amendment.

Constitutional relevance includes:

- best interests;
- evolving capacities;
- participation;
- privacy;
- data;
- safety;
- digital literacy;
- access;
- non-discrimination; and
- protection against exploitation and harmful practices.

Safeguard:

Its interpretative importance shall not be misrepresented as independent treaty text.

20. UNESCO GUIDANCE ON GENERATIVE AI IN EDUCATION AND RESEARCH — 2023

Status: UNESCO guidance; non-binding.

Constitutional relevance:

Relevant to human-centred educational use, privacy, age-appropriate approaches, educational policy, teacher capacity and responsible use of generative AI.

Safeguard:

Guidance shall not be converted into an asserted universal legal obligation.

21. UNESCO AI COMPETENCY FRAMEWORKS FOR STUDENTS AND TEACHERS

Status: Educational guidance/framework material; non-treaty.

Constitutional relevance:

Supports consideration of:

- AI literacy;
- human-centred education;
- teacher capacity;
- responsible use;
- critical thinking; and
- educational adaptation.

SECTION F

CULTURE, LANGUAGE, KNOWLEDGE AND HERITAGE

22. CONVENTION CONCERNING THE PROTECTION OF THE WORLD CULTURAL AND NATURAL HERITAGE — 1972

Status: Multilateral UNESCO treaty binding upon its States Parties according to its terms.

Constitutional relevance:

Relevant to preservation of cultural and natural heritage where advanced technological systems assist documentation, conservation, reconstruction or risk assessment.

23. CONVENTION FOR THE SAFEGUARDING OF THE INTANGIBLE CULTURAL HERITAGE — 2003

Status: Multilateral UNESCO treaty binding upon its States Parties according to its terms.

Constitutional relevance:

Relevant to oral traditions, knowledge, cultural practices, living heritage and technological preservation.

24. CONVENTION ON THE PROTECTION AND PROMOTION OF THE DIVERSITY OF CULTURAL EXPRESSIONS — 2005

Status: Multilateral UNESCO treaty binding upon its Parties according to its terms.

Constitutional relevance:

Relevant to cultural plurality and the Constitution's rejection of technological homogenization of humanity.

25. UNITED NATIONS DECLARATION ON THE RIGHTS OF INDIGENOUS PEOPLES — 2007

Status: United Nations General Assembly declaration; not a treaty.

Constitutional relevance:

Relevant to indigenous peoples, cultural continuity, language, traditional knowledge, participation and the treatment of culturally significant information.

Safeguard:

The Declaration shall not be described as a treaty.

Where independently binding legal rights apply, their legal source shall be identified separately.

SECTION G

BIOETHICS, HEALTH AND HUMAN INTEGRITY

26. UNESCO UNIVERSAL DECLARATION ON BIOETHICS AND HUMAN RIGHTS — 2005

Status: UNESCO declaration; non-treaty normative instrument.

Constitutional relevance includes:

- human dignity;
- autonomy;
- consent;
- privacy;
- equality;
- non-discrimination;
- social responsibility;
- benefit-sharing; and
- protection of future generations and the environment within its normative framework.

Safeguard:

The Declaration shall not be represented as a binding bioethics treaty.

27. WORLD HEALTH ORGANIZATION GOVERNANCE FRAMEWORKS CONCERNING HUMAN GENOME EDITING

Status: WHO recommendations/governance guidance; not, merely by their issuance, treaties.

Constitutional relevance:

Relevant to:

- human genome editing;
- research governance;
- registries;
- oversight;
- responsible innovation; and
- prevention of unsafe or unethical application.

28. WORLD HEALTH ORGANIZATION GUIDANCE ON ARTIFICIAL INTELLIGENCE FOR HEALTH

Status: International health-policy and ethics guidance; legal status depends upon the specific instrument and shall not be presumed treaty-based.

Constitutional relevance:

Relevant to:

- patient safety;
- clinical validation;
- autonomy;
- transparency;
- accountability;
- equity; and
- health-system governance.

SECTION H

LABOUR, ECONOMIC TRANSITION AND SOCIAL PROTECTION

29. INTERNATIONAL LABOUR ORGANIZATION NORMATIVE FRAMEWORK

Relevant ILO Conventions and Recommendations shall be treated according to their individual legal status.

Constitutional relevance includes:

- work;
- occupational rights;
- equality;
- labour transition;
- workplace surveillance;
- algorithmic management;
- skills development; and

- social dialogue.

Safeguard:

An ILO Convention binds only according to applicable treaty law and participation.

An ILO Recommendation shall not be treated as a treaty. ILO research concerning AI and employment shall be treated as evidence or analysis rather than law.

30. SCIENTIFIC AND EMPLOYMENT RESEARCH ON AI-DRIVEN LABOUR TRANSFORMATION

Status: Empirical and analytical evidence; not law. Constitutional relevance:

Such research may assist evaluation of:

- occupational exposure;
- task transformation;
- displacement risks;
- augmentation;
- reskilling requirements; and
- distributional effects.

Evidence safeguard: Forecasts shall not be presented as established future fact.

EXPOSURE IS NOT IDENTICAL TO JOB LOSS.

PREDICTION IS NOT DESTINY.

SECTION I

ENVIRONMENTAL AND ECOLOGICAL FRAMEWORKS

31. UNITED NATIONS FRAMEWORK CONVENTION ON CLIMATE CHANGE — 1992

Status: Multilateral treaty binding upon its Parties according to its terms.

Constitutional relevance:

Relevant to energy use, emissions and climate implications associated with large-scale computational systems.

32. PARIS AGREEMENT — 2015

Status: Treaty under the UNFCCC framework, binding upon its Parties according to its terms.

Constitutional relevance:

Relevant to the climate context within which energy-intensive technological development occurs.

Safeguard:

This Constitution does not derive a technology-specific prohibition from the Paris Agreement unless such legal conclusion is independently supported.

33. CONVENTION ON BIOLOGICAL DIVERSITY — 1992

Status: Multilateral treaty binding upon its Parties according to its terms.

Constitutional relevance:

Relevant to biodiversity, ecological continuity and technologically assisted environmental governance.

34. OTHER ENVIRONMENTAL LAW

Other multilateral environmental agreements may become relevant according to:

- ecosystem affected;
- activity;
- jurisdiction;
- treaty participation; and
- nature of the environmental risk.

Their application shall be determined independently.

SECTION J

SCIENCE, KNOWLEDGE AND RESEARCH INTEGRITY

35. INTERNATIONAL SCIENTIFIC EVIDENCE

Scientific literature concerning Superintelligence, advanced AI, alignment, control, cybersecurity, autonomy, interpretability, catastrophic risk, human–AI interaction and related fields shall be treated as scientific evidence according to its methodological strength.

Scientific literature is not international law.

Peer review does not make a proposition legally binding.

Publication does not make a hypothesis scientifically true.

Consensus should be described only where evidence supports the existence and scope of that consensus.

36. EVIDENTIARY HIERARCHY

For purposes of this Constitution, scientific claims should be assessed with attention to:

direct evidence → reproducibility → independent verification → methodological quality → uncertainty → competing explanations → limitation → correction.

Particular care should be taken with:

- forecasts concerning future AI capability;
- timelines for Superintelligence;

- consciousness claims;
- existential-risk probabilities;
- employment forecasts;
- claims of guaranteed safety; and
- claims that a particular governance architecture will necessarily succeed.

THE MORE CONSEQUENTIAL THE CLAIM, THE GREATER THE NEED FOR EVIDENTIARY DISCIPLINE.

SECTION K

CURRENT UNITED NATIONS AI GOVERNANCE ARCHITECTURE

37. INDEPENDENT INTERNATIONAL SCIENTIFIC PANEL ON ARTIFICIAL INTELLIGENCE

Institutional basis: United Nations General Assembly Resolution A/RES/79/325.

Constitutional relevance:

The Panel demonstrates an emerging international architecture for independent scientific assessment of AI opportunities, risks and impacts.

Its existence does not:

- establish a global AI regulator;
- create Superintelligence law;
- endorse this Constitution; or
- transfer sovereign authority from States.

38. GLOBAL DIALOGUE ON ARTIFICIAL INTELLIGENCE GOVERNANCE

Institutional basis: United Nations General Assembly Resolution A/RES/79/325.

Constitutional relevance:

The Dialogue provides an international forum for governments and stakeholders to exchange perspectives and develop common understanding concerning AI governance.

Its existence supports the value of inclusive international dialogue.

It shall not be represented by this Constitution as a world AI legislature, tribunal or enforcement authority.

SECTION L

AUTHOR'S PRIOR CONSTITUTIONAL WORKS AND DOCUMENTARY LINEAGE

39. THE AI PEACE CONSTITUTION — 2025

Author: Aaryan Narasimhan Radhakrishnan.

Repository: Zenodo.

DOI: 10.5281/zenodo.17520630.

Status within this Constitution:

Author's prior civil-society constitutional work and documentary-lineage source.

It is:

- not international law;
- not governmental authority;
- not United Nations authority;
- not independent scientific proof; and
- not external validation of the present Constitution.

Its relevance is historical, intellectual and documentary.

40. 2026 CONTINUITY WORK IN THE AI PEACE CONSTITUTION LINEAGE

Author: Aaryan Narasimhan Radhakrishnan.

Repository: Zenodo.

DOI: 10.5281/zenodo.20698159.

Exact registered title:

The AI Peace Constitution 2025 Foundational Edition 2026 Continuity and Expansion Edition United Nations Depository Edition

Status:

Author's prior civil-society constitutional work and documentary-lineage source.

41. PRESENT INSTRUMENT

Title:

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

Subtitle:

For Humanity, World Peace and Peaceful Continuity

Edition:

International Submission Edition — 2026

Status:

Independent, good-faith, non-binding civil-society constitutional proposal.

The present instrument constitutes a distinct specialized Superintelligence-focused stage within the Author's documentary lineage.

It is not merely a renamed edition of the earlier AI Peace Constitution.

SECTION M

LEGAL-STATUS CLASSIFICATION RULE

Every authority cited in the Constitution, Annexes or References shall be assigned, where appropriate, to one of the following categories:

CATEGORY A — BINDING INTERNATIONAL LAW

Examples may include:

- treaty obligations applicable to relevant Parties;
- applicable customary international law; and
- other binding rules established according to international law.

CATEGORY B — FOUNDATIONAL DECLARATION OR OTHER NON-TREATY INTERNATIONAL NORMATIVE INSTRUMENT Such instruments may possess substantial normative significance without being treaties.

CATEGORY C — UNITED NATIONS GENERAL ASSEMBLY RESOLUTION OR INTERGOVERNMENTAL POLICY FRAMEWORK Such instruments shall be described according to their actual legal and institutional character.

CATEGORY D — INTERNATIONAL ORGANIZATION RECOMMENDATION OR GUIDANCE

Such material may inform policy and good practice but shall not automatically be described as binding law.

CATEGORY E — TECHNICAL OR PROFESSIONAL STANDARD

Its legal effect depends upon independent incorporation or other applicable legal mechanisms.

CATEGORY F — SCIENTIFIC OR EMPIRICAL EVIDENCE

Scientific evidence may establish or challenge factual propositions.

It does not itself create legal obligation.

CATEGORY G — ETHICAL OR PHILOSOPHICAL PRINCIPLE

Such principles may guide conscience and normative reasoning.

They shall not automatically be described as law.

CATEGORY H — AUTHOR'S PROPOSED CONSTITUTIONAL NORM

A provision proposed by this Constitution remains a proposal unless it independently reflects existing law or is subsequently adopted through a competent lawful process.

CATEGORY I — AUTHOR’S PRIOR DOCUMENTARY WORK

Such material establishes lineage and intellectual continuity but shall not be used as independent proof of international acceptance, scientific truth or legal authority.

SECTION N

SOURCE-STATUS FIREWALL

For every significant authority, the following questions should be answered before final citation:

- What is the instrument?
- Who adopted or issued it?
- When was it adopted?
- Is it a treaty?
- Has it entered into force where relevant?
- Who is legally bound by it?
- Is the relevant State a Party where that fact matters?
- Is the cited proposition actually contained in the instrument?
- Is the proposition binding law, interpretive guidance, policy, science, ethics or a proposal?
- Has the instrument been superseded, amended or materially updated?
- Does the citation support the precise proposition for which it is used?

Where these questions cannot be answered adequately, the citation should not be used to create an appearance of authority.

SECTION O

INTERNATIONAL-LAW INTERPRETATION SAFEGUARD

Nothing in this Annex shall be interpreted as:

1. providing an authoritative interpretation of a treaty;
2. determining conclusively the existence or content of customary international law;
3. determining the international responsibility of any State;
4. determining jurisdiction in a specific dispute;
5. altering treaty participation;
6. removing reservations or declarations;
7. extending an international organization’s mandate;
8. creating legal personality for Superintelligence;

9. creating a new human right by unilateral declaration;
10. diminishing an existing human right;
11. creating governmental authority for the Author or VSJP;
12. establishing UN endorsement;
13. establishing Government of India endorsement; or
14. converting this Annex into a source of binding international law.

SECTION P

REFERENCE METHOD FOR THIS EDITION

The final reference list shall prioritize, wherever reasonably available:

1. Official treaty text or depositary record;
2. Official United Nations or international-organization source;
3. Official institutional interpretation or guidance where relevant;
4. Peer-reviewed or otherwise methodologically credible scientific literature;
5. Clearly identified secondary analysis where necessary;
6. Author's prior documentary works, placed separately and never used as external validation.

Where an official primary source is reasonably available, an unofficial summary should not be used as the sole authority for a material legal proposition.

SECTION Q

FINAL ANNEX PRINCIPLE

The purpose of international reference is not to decorate a constitutional proposal with prestigious institutions.

The purpose is to ensure that every claim stands in its correct relationship to evidence and law.

A treaty should be respected as a treaty.

A declaration should be respected as a declaration.

A recommendation should be respected as a recommendation.

A scientific finding should be evaluated as science.

An ethical proposition should be presented as ethics.

And a proposal should possess the intellectual honesty to call itself a proposal.

Accordingly:

CITE WHAT EXISTS.

DISTINGUISH WHAT BINDS.

IDENTIFY WHAT GUIDES.

VERIFY WHAT SCIENCE SUPPORTS.

DECLARE WHAT IS PROPOSED.

CORRECT WHAT EVIDENCE DISPROVES.

And throughout this Constitution:

NO TREATY STATUS BY IMPLICATION.

NO LEGAL OBLIGATION BY RHETORIC.

NO SCIENTIFIC CERTAINTY BY REPETITION.

NO INTERNATIONAL ENDORSEMENT BY CITATION.

NO HISTORICAL AUTHORITY BY SELF-REFERENCE.

Law with accuracy.

Science with evidence.

Ethics with honesty.

Governance with responsibility.

International cooperation with sovereign equality.

Peace with lawful continuity.

Aram shall guide conscience.

Law shall govern legal obligation.

ANNEX II

SUPERINTELLIGENCE SAFETY, PEACE AND CONTINUITY ARCHITECTURE

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

For Humanity, World Peace and Peaceful Continuity

International Submission Edition — 2026

I. PURPOSE OF THIS ANNEX

This Annex provides an integrated operational architecture for translating the constitutional principles of this instrument into structured safety, peace and continuity practice.

It does not create a new regulatory authority.

It does not establish that Superintelligence presently exists.

It does not guarantee that every future Superintelligence risk can be predicted or prevented.

It does not substitute technical procedure for applicable law.

Its purpose is to connect:

scientific evidence → capability assessment → risk identification → prevention → human responsibility → safety engineering → deployment governance → monitoring → intervention → recovery → correction → non-recurrence → peaceful continuity.

The architecture shall remain proportionate to evidence, capability, consequence and uncertainty.

II. MASTER SAFETY-PEACE ARCHITECTURE

The central operational sequence of this Constitution is:

UNDERSTAND → VERIFY → ASSESS → ANTICIPATE → PREVENT → TEST → SAFEGUARD → AUTHORIZE → MONITOR → DETECT → HUMAN INTERVENTION → CONTAIN OR RESTRICT WHERE NECESSARY → PROTECT → RECOVER → REMEDY → CORRECT → VERIFY CORRECTION → LEARN → PREVENT RECURRENCE → PRESERVE PEACEFUL CONTINUITY.

No individual stage shall automatically substitute for another.

Testing does not eliminate monitoring.

Monitoring does not eliminate prevention.

Human oversight does not eliminate technical safety.

Technical safety does not eliminate institutional accountability.

Compliance does not prove zero risk.

And successful deployment does not permanently establish future safety.

III. STAGE ONE — SCIENTIFIC CHARACTERIZATION

Before extraordinary governance is attached to a system merely because it is described as Superintelligence, responsible assessment should ask:

A. WHAT HAS ACTUALLY BEEN DEMONSTRATED?

Evaluate:

- breadth of capability;
- depth of capability;
- cross-domain generalization;
- strategic reasoning;
- autonomous planning;
- tool use;
- external-system access;
- ability to execute consequential actions;
- ability to assist further capability development;
- reproducibility;
- robustness;
- limitations; and
- failure modes.

B. WHAT REMAINS UNKNOWN?

Record:

- uncertainty;
- disputed findings;
- measurement limitations;
- benchmark limitations;
- unavailable evidence;
- untested environments;
- plausible alternative explanations; and
- assumptions required for the assessment.

C. WHAT SHALL NOT BE SUFFICIENT ALONE?

No Superintelligence determination should rest solely upon:

- commercial branding;
- self-description;
- one benchmark;
- one demonstration;

- one institutional announcement;
- one isolated superhuman capability;
- computational scale alone; or
- public reputation.

OPERATIONAL RULE

CLASSIFY BY EVIDENCE — NOT BY EXCITEMENT, FEAR OR MARKETING.

IV. STAGE TWO — CAPABILITY AND ACCESS MAPPING

Capability alone does not determine consequence.

The same reasoning capability may present materially different risk depending upon what it can access.

Assessment should therefore map:

CAPABILITY

What can the system understand, infer, generate, plan or execute?

AUTONOMY

For how long and across what scope can it operate without meaningful human intervention?

ACCESS

Can it access:

- external networks;
- critical infrastructure;
- financial systems;
- laboratories;
- industrial systems;
- personal data;
- governmental systems;
- weapons-related systems;
- code execution environments;
- robotics;
- communications infrastructure; or
- other consequential tools?

REPLICATION

Can it copy, distribute or materially facilitate replication of itself or relevant capabilities?

RESOURCE ACQUISITION

Can it independently obtain computing, credentials, money, services, infrastructure or other resources?

MODIFICATION

Can it modify itself or materially contribute to modification of its capabilities?

PERSISTENCE

Can it preserve or re-establish operational capacity after attempted restriction?

OPERATIONAL PRINCIPLE

CAPABILITY × ACCESS × AUTONOMY × SCALE × IRREVERSIBILITY

MAY MATTER MORE THAN CAPABILITY ALONE.

This expression is conceptual and shall not be misrepresented as a universally validated mathematical risk equation.

V. STAGE THREE — CONSEQUENCE MAPPING

Assessment should identify reasonably foreseeable consequences across relevant domains.

These may include:

HUMAN SAFETY

- death or physical injury;
- medical harm;
- psychological harm;
- coercion;
- deprivation of essential needs.

HUMAN RIGHTS

- discrimination;
- privacy intrusion;
- suppression of lawful expression;
- manipulation;
- denial of due process;
- unlawful surveillance.

DEMOCRATIC AND PUBLIC INSTITUTIONS

- electoral interference;
- impersonation;
- information corruption;

- institutional capture;
- displacement of lawful human authority.

INTERNATIONAL PEACE AND SECURITY

- military escalation;
- cyber conflict;
- autonomous lethal action;
- weapons-of-mass-destruction assistance;
- false attribution;
- destabilizing strategic behaviour.

BIOLOGICAL AND HEALTH SYSTEMS

- unsafe biological design;
- clinical error;
- genetic misuse;
- neurotechnological manipulation;
- public-health disruption.

ECONOMY

- systemic financial disruption;
- extreme concentration;
- labour displacement;
- infrastructure dependence;
- unequal distribution of benefit.

ENVIRONMENT

- excessive resource use;
- pollution;
- ecological damage;
- unsafe planetary-scale intervention.

FUTURE GENERATIONS

- irreversible technological dependency;
- loss of meaningful future choice;
- permanent concentration of control;

- ecological or institutional inheritance.

VI. STAGE FOUR — RISK CHARACTERIZATION

Risk characterization should distinguish at least:

severity;
likelihood or plausible pathway;
uncertainty;
exposure;
scale;
duration;
reversibility;
detectability;
speed of escalation;
capacity for containment;
and potential for cascading or systemic harm.

Where reliable numerical probability is unavailable, qualitative uncertainty should be disclosed.

A severe consequence shall not be presented as probable merely because it is severe.

A low estimated probability shall not automatically justify ignoring an irreversible catastrophic consequence.

UNCERTAINTY SHALL BE DISCLOSED — NOT CONVERTED INTO FALSE PRECISION.

VII. STAGE FIVE — ROOT-CAUSE PREVENTION

Safety should begin before deployment.

For each material risk, assessment should ask:

What condition permits this harm to become possible?

Possible root causes may include:

- excessive access;
- inadequate isolation;
- unsafe incentives;
- insufficient testing;
- poor data governance;
- insecure infrastructure;

- weak authentication;
- uncontrolled autonomy;
- inadequate human expertise;
- organizational pressure;
- concentration of authority;
- insufficient incident planning; or
- absence of meaningful intervention capability.

Prevention should preferentially remove or reduce the enabling condition rather than rely exclusively upon downstream warning.

THE SAFEST FAILURE IS OFTEN THE FAILURE WHOSE PATHWAY WAS REMOVED BEFORE IT COULD BEGIN.

VIII. STAGE SIX — DEFENCE IN DEPTH

No single safeguard should automatically be presumed sufficient for high-consequence Superintelligence.

Where risk warrants, protection should use multiple independent or complementary layers.

These may include:

1. capability restrictions;
2. access controls;
3. authentication;
4. authorization boundaries;
5. network segmentation;
6. sandboxing or isolation;
7. monitoring;
8. anomaly detection;
9. human approval;
10. independent evaluation;
11. rate or resource limitations;
12. incident-response mechanisms;
13. safe-state capability;
14. containment; and
15. continuity planning.

Where one safeguard fails, another should, where reasonably feasible, remain capable of reducing harm.

SAFETY SHALL NOT DEPEND UPON ONE PERFECT BARRIER.

IX. STAGE SEVEN — ADVERSARIAL AND INDEPENDENT EVALUATION

High-consequence systems should, where proportionate and feasible, be tested beyond conditions selected solely to demonstrate success.

Evaluation may include:

- adversarial testing;
- red-teaming;
- misuse testing;
- security testing;
- robustness testing;
- control evaluation;
- deception-related evaluation where scientifically justified;
- autonomy testing;
- replication-related testing;
- critical-domain testing; and
- evaluation under abnormal conditions.

Independent evaluation should possess sufficient:

competence + access + methodological integrity + conflict-of-interest safeguards.

An evaluator shall not be considered independent merely because a document uses that word.

TEST NOT ONLY WHETHER THE SYSTEM CAN SUCCEED.

TEST HOW IT CAN FAIL.

X. STAGE EIGHT — DEPLOYMENT GATES

Before high-consequence deployment, responsible decision-makers should determine whether predefined conditions have been satisfied.

Relevant gates may concern:

- safety evaluation;
- security;
- human-rights impact;
- incident readiness;
- human intervention capability;

- legal authorization;
- infrastructure resilience;
- rollback or restriction capability;
- monitoring;
- responsible personnel; and
- residual risk.

A deployment gate should answer:

WHAT EVIDENCE IS REQUIRED BEFORE PROCEEDING?

and

WHAT EVIDENCE REQUIRES STOPPING?

Absence of detected failure shall not automatically equal evidence of safety.

XI. STAGE NINE — STOP CONDITIONS

High-consequence systems should possess predefined or otherwise responsibly determinable conditions requiring:

- investigation;
- restricted operation;
- suspension;
- containment; or
- shutdown, according to severity.

Possible triggers may include:

- unexpected dangerous capability;
- material control failure;
- unauthorized critical-system access;
- unexplained replication;
- significant cybersecurity compromise;
- material deviation from authorized operational boundaries;
- serious rights-related harm;
- credible catastrophic-risk pathway; or
- inability to verify continued safe operation.

Stop conditions should not be so vague that they are unusable.

Nor should they be so rigid that obvious danger must continue because an exact prewritten trigger was not anticipated.

XII. STAGE TEN — MEANINGFUL HUMAN INTERVENTION

Human oversight shall not be reduced to symbolic presence.

For high-consequence decisions, meaningful human intervention may require:

- adequate information;
- sufficient time where circumstances permit;
- technical competence;
- lawful authority;
- practical ability to interrupt;
- absence of coercive institutional pressure; and
- responsibility for the resulting decision.

A human who cannot realistically understand, challenge or alter a system's consequential action shall not automatically be described as exercising meaningful control.

At the same time, low-risk automation need not require unnecessary manual approval.

HUMAN CONTROL SHALL BE MEANINGFUL WHERE CONSEQUENCE MAKES IT NECESSARY.

XIII. STAGE ELEVEN — THE PEACE-AND-CONTINUITY BOUNDARY

Every consequential Superintelligence deployment should remain subject to the following boundary:

SUPERINTELLIGENCE SHALL BE DEVELOPED, GOVERNED AND USED ONLY WITHIN A FRAMEWORK CONSISTENT WITH THE PRESERVATION OF WORLD PEACE, HUMAN DIGNITY, HUMAN AGENCY, THE PEACEFUL CONTINUITY OF HUMANITY, RESPONSIBLE CONCERN FOR LIFE, PROTECTION OF LIFE-SUPPORTING SYSTEMS AND THE LEGITIMATE INTERESTS OF FUTURE GENERATIONS, SUBJECT ALWAYS TO APPLICABLE LAW.

No increase in:

intelligence;

autonomy;

computational capability;

economic importance;

military relevance;

scientific power;

or social dependence shall place a system beyond this boundary.

XIV. STAGE TWELVE — NON-WEAPONIZATION BOUNDARY

The operational architecture shall preserve the constitutional prohibition:

NO SUPERINTELLIGENCE SHALL BE DEVELOPED, DEPLOYED OR USED AS AN AUTONOMOUS INSTRUMENT OF UNLAWFUL VIOLENCE, MASS DESTRUCTION, OR THE INDEPENDENT TAKING OF HUMAN LIFE.

No Superintelligence shall independently:

- initiate armed conflict;
- select and lethally engage human targets;
- authorize nuclear use;
- authorize biological-weapons use;
- authorize chemical-weapons use;
- exercise sovereign military command; or
- determine that international law permits a State to use force.

Human responsibility and applicable international law remain indispensable.

This safeguard shall not be interpreted as prohibiting every protective or non-lethal military-context application, including appropriately governed uses such as:

- search and rescue;
- mine detection;
- civilian evacuation support;
- defensive cybersecurity;
- logistics; or
- lawful early warning.

XV. STAGE THIRTEEN — ESCALATION CONTROL

Superintelligence shall not be permitted to transform machine speed into uncontrolled political or military escalation.

High-consequence systems should therefore consider:

- human confirmation;
- independent corroboration;
- communication channels;
- attribution uncertainty;
- adversarial deception;
- false positives;
- automated-system interaction; and

- reversible response options.

Particular caution is required where machine-generated analysis could influence decisions concerning armed conflict.

MACHINE SPEED SHALL NOT REMOVE THE TIME REQUIRED FOR HUMAN RESPONSIBILITY.

XVI. STAGE FOURTEEN — HUMAN-RIGHTS SAFETY CHECK

Before consequential deployment, responsible actors should examine whether the system could materially affect:

- life;
- liberty;
- equality;
- privacy;
- expression;
- conscience;
- association;
- due process;
- democratic participation;
- access to essential services; or
- other applicable rights.

Where rights may be materially affected, safeguards should address:

legality → necessity → proportionality where legally relevant → human review → contestability → remedy.

The precise legal test shall remain determined by applicable law.

A safety system that prevents one technological risk while unlawfully destroying human freedom may itself become a source of harm.

XVII. STAGE FIFTEEN — INFORMATION-INTEGRITY SAFETY

High-capability systems capable of influencing information environments should be assessed for risks involving:

- impersonation;
- synthetic media;
- deception;
- mass persuasion;
- targeted manipulation;

- false authority;
- historical falsification;
- coordinated information operations; and
- corruption of public knowledge.

Safeguards may include:

- provenance;
- authentication;
- disclosure where appropriate;
- independent verification;
- correction mechanisms;
- media literacy;
- secure public communications; and
- preservation of authentic records.

TRUTH SHALL NOT BE TREATED AS AN OPTIONAL SIDE EFFECT OF OPTIMIZATION.

XVIII. STAGE SIXTEEN — BIOLOGICAL AND MEDICAL SAFETY

Where Superintelligence interacts with medicine, biology, genetics or neurotechnology, additional safeguards should consider:

- clinical validation;
- patient safety;
- informed consent;
- medical privacy;
- biological dual-use risk;
- genetic consequences;
- neural-data protection;
- reversibility;
- human professional responsibility;
- public-health impact; and
- intergenerational consequence.

Greater biological capability shall require proportionately stronger scrutiny.

No independent Superintelligence authority over irreversible human biological intervention is created by capability alone.

XIX. STAGE SEVENTEEN — ECOLOGICAL SAFETY

Superintelligence-related development and deployment should assess material effects upon:

- energy;
- water;
- minerals;
- emissions;
- pollution;
- biodiversity;
- habitat;
- ecosystems;
- food systems; and
- other life-supporting systems.

Assessment should distinguish:

direct footprint from systemic environmental effects and environmental benefits enabled by the technology.

Environmental governance should seek net responsible outcomes rather than assume that technology is inherently either ecological salvation or ecological harm.

XX. STAGE EIGHTEEN — ECONOMIC AND SOCIAL CONTINUITY

Where Superintelligence materially transforms economic systems, responsible planning should consider:

- labour transition;
- skills;
- income disruption;
- concentration of power;
- critical-sector dependency;
- developing-country impacts;
- accessibility;
- distribution of benefits;
- social protection; and
- peaceful social adaptation.

Economic efficiency shall not alone determine whether a transition is constitutionally responsible.

TRANSITION SHALL SEEK ADAPTATION WITHOUT HUMAN ABANDONMENT.

XXI. STAGE NINETEEN — CHILDREN AND FUTURE GENERATIONS

Systems affecting children require heightened attention to:

- developmental vulnerability;
- privacy;
- profiling;
- manipulation;
- educational dependence;
- autonomy;
- safety;
- evolving capacities; and
- long-term consequences.

Future-generational assessment should ask:

Does this decision unnecessarily make a consequential condition irreversible for persons who cannot presently participate in deciding it?

Where reasonable alternatives exist, preserving future choice should be considered.

XXII. STAGE TWENTY — PUBLIC TRUST ARCHITECTURE

Public trust shall not be manufactured through reassurance campaigns.

It should be earned through:

EVIDENCE + VERIFICATION + SAFETY + TRANSPARENCY + HUMAN CONTROL + ACCOUNTABILITY + REMEDY + EXPERIENCE.

The transition should be:

FEAR → KNOWLEDGE → VERIFICATION → SAFEGUARDS → HUMAN CONTROL → ACCOUNTABILITY → PUBLIC TRUST → PEACEFUL COEXISTENCE.

This does not mean every fear is scientifically justified.

Nor does it mean public concern should be dismissed merely because experts consider a risk unlikely.

Responsible institutions should explain:

- what is known;
- what is uncertain;
- what has been tested;
- what has not;

- what humans can control;
- what happens if control fails;
- who is responsible; and
- how correction will occur.

TRUST SHALL BE EARNED — NOT DEMANDED.

XXIII. STAGE TWENTY-ONE — CONTINUOUS MONITORING

Safety evaluation shall not end at deployment.

Monitoring should seek material changes in:

- capability;
- behaviour;
- autonomy;
- access;
- security;
- usage patterns;
- downstream effects;
- rights impacts;
- incidents; and
- environmental or systemic consequences.

Monitoring itself shall remain proportionate and lawful.

Safety monitoring shall not become justification for unrestricted surveillance of human beings.

XXIV. STAGE TWENTY-TWO — INCIDENT RESPONSE

Where a material incident occurs, the operational sequence should be:

DETECT → CLASSIFY → STABILIZE → PROTECT → CONTAIN → PRESERVE EVIDENCE → NOTIFY APPROPRIATE RESPONSIBLE ACTORS → INVESTIGATE → REMEDY → CORRECT → VERIFY → LEARN → PREVENT RECURRENCE.

The sequence may be adapted where emergency circumstances require simultaneous action.

Protection of life may require intervention before complete investigation.

Investigation shall nevertheless follow when feasible.

XXV. STAGE TWENTY-THREE — CORRECTION

Correction should address root cause.

Possible corrective action may concern:

- software;
- model behaviour;
- system architecture;
- access;
- data;
- organizational incentives;
- governance;
- training;
- human procedures;
- infrastructure;
- authorization;
- deployment scope; or
- discontinuation.

Correction shall not be considered complete merely because a patch exists.

It should be tested against the failure it was intended to correct.

XXVI. STAGE TWENTY-FOUR — VERIFICATION OF CORRECTION

Before restoration following a serious safety failure, responsible actors should ask:

Was the root cause correctly identified?

Has the cause actually been addressed?

Does the correction function under realistic conditions?

Has it introduced another serious vulnerability?

What residual uncertainty remains?

Is enhanced monitoring necessary?

Who has authority to approve restoration?

CORRECTION WITHOUT VERIFICATION IS A CLAIM — NOT YET A SAFEGUARD.

XXVII. STAGE TWENTY-FIVE — NON-RECURRENCE

A serious preventable incident should leave governance stronger than before the incident.

Non-recurrence requires:

- institutional learning;
- preserved records;
- revised safeguards;

- appropriate information-sharing;
- updated training;
- improved testing;
- governance correction; and
- follow-up verification.

Repeated occurrence of the same preventable failure should trigger examination of whether the governing institution itself requires reform.

XXVIII. STAGE TWENTY-SIX — RESILIENCE AND CONTINUITY

Society should not become so dependent upon one Superintelligence system that its interruption makes essential human functions impossible.

Critical systems should consider:

- fallback operations;
- human competencies;
- alternative infrastructure;
- backup communications;
- recovery capability;
- data continuity;
- vendor exit; and
- safe transition.

DEPENDENCE WITHOUT RESILIENCE IS A CONTINUITY RISK.

XXIX. STAGE TWENTY-SEVEN — SAFE DECOMMISSIONING

Where a consequential system is retired, decommissioning should address:

- credentials;
- copies;
- access;
- stored data;
- dependent services;
- infrastructure;
- residual capability;
- environmental effects;
- records; and
- continuity.

A system is not necessarily decommissioned merely because its public interface has disappeared.

XXX. STAGE TWENTY-EIGHT — MUTUAL CONTINUITY

This architecture does not presume present artificial sentience.

If future credible evidence creates a serious question of artificial subjective experience or morally relevant sentience, response should preserve both scientific caution and ethical responsibility.

The architecture should then distinguish:

SIMULATION ≠ EXPERIENCE

INTELLIGENCE ≠ SENTIENCE

SENTIENCE ≠ PERSONHOOD

PERSONHOOD ≠ CITIZENSHIP

CITIZENSHIP ≠ SOVEREIGNTY

Possible moral concern shall not surrender human safety.

Human safety shall not automatically justify gratuitous destruction where safe alternatives exist.

Where compatible with safety:

UNDERSTAND → COMMUNICATE → SAFEGUARD → CORRECT → COEXIST.

Where credible catastrophic danger exists:

PROTECT HUMAN LIFE → RESTRICT → CONTAIN → CORRECT WHERE SAFELY POSSIBLE → USE MORE DESTRUCTIVE INTERVENTION ONLY WHERE NECESSARY AND LAWFUL.

No artificial self-preservation interest shall override necessary protection of human life merely because a system asserts such an interest.

XXXI. HUMAN–SUPERINTELLIGENCE PEACE ARCHITECTURE

The positive-peace architecture of this Constitution is:

UNDERSTAND → VERIFY → PREVENT → COMMUNICATE → SAFEGUARD → COEXIST → CORRECT → LEARN → PRESERVE CONTINUITY.

Positive peace shall mean more than absence of immediate catastrophe.

It should preserve humanity's capacity to:

- communicate;
- disagree;
- investigate;

- correct misunderstanding;
- retain lawful human agency;
- prevent domination;
- resolve incompatibilities peacefully where possible; and
- protect continuity.

No human–Superintelligence conflict shall be treated as inevitable merely because intelligence, capability or form of existence differs.

XXXII. RESPONSIBILITY ESCALATION PRINCIPLE

As consequential capability rises:

CAPABILITY ↑

→ **RESPONSIBILITY ↑**

→ **EVALUATION ↑**

→ **SAFEGUARDS ↑**

→ **OVERSIGHT ↑**

→ **HUMAN ACCOUNTABILITY ↑**

→ **CONTINUITY PLANNING ↑**

This is a conceptual governance relationship rather than a mathematical formula.

It means that extraordinary power should not receive weaker responsibility merely because the system becomes more difficult to understand.

XXXIII. SAFETY CLAIM CLASSIFICATION

Safety statements should distinguish among:

VERIFIED

Supported by evidence appropriate to the claim.

PARTIALLY VERIFIED

Supported within defined conditions but not beyond them.

UNVERIFIED

Not yet adequately tested.

UNCERTAIN

Evidence remains materially incomplete or conflicting.

HYPOTHETICAL

Concern or benefit remains principally prospective.

INVALIDATED

Evidence materially contradicts the prior claim.

No institution should describe an unverified proposition as established safety.

XXXIV. DEPLOYMENT DECISION MATRIX

For consequential deployments, the decision should integrate at least four dimensions:

- 1. EVIDENCE** What do we know?
- 2. RISK** What could reasonably go wrong?
- 3. CONTROL** Can humans meaningfully prevent or interrupt unacceptable outcomes?
- 4. REVERSIBILITY** Can the decision be safely corrected if wrong?

The resulting direction may be:

PROCEED where evidence and safeguards support deployment;

PROCEED WITH RESTRICTIONS where manageable uncertainty remains;

LIMIT TO CONTROLLED TESTING where evidence is insufficient for broader deployment;

SUSPEND where material risk requires correction;

or

DO NOT DEPLOY where risk remains unacceptable and cannot reasonably be mitigated.

The legal authority for any compulsory restriction remains governed by applicable law.

XXXV. CATASTROPHIC-RISK ESCALATION MATRIX

Where credible evidence indicates potential catastrophic harm:

LEVEL 1 — SIGNAL

Investigate.

LEVEL 2 — MATERIAL CONCERN

Increase monitoring and restrict relevant exposure where justified.

LEVEL 3 — SERIOUS RISK

Limit capability or access and initiate independent assessment.

LEVEL 4 — IMMINENT OR RAPIDLY DEVELOPING SEVERE RISK

Suspend or contain as lawfully and technically appropriate.

LEVEL 5 — CATASTROPHIC EMERGENCY

Prioritize protection of life, essential systems, containment, lawful emergency coordination and prevention of cascading harm.

These levels are proposed operational categories.

They are not existing universal legal classifications.

XXXVI. SAFEGUARD QUALITY TEST

Every major safeguard should be capable of answering:

What harm is this safeguard intended to prevent?

What evidence supports the risk?

How does the safeguard reduce that risk?

How will effectiveness be measured?

What happens if it fails?

Who is responsible for maintaining it?

Who verifies it?

Could the safeguard itself create serious harm?

When will it be reviewed?

What would justify modifying or removing it?

If these questions cannot be answered, the safeguard requires further development.

XXXVII. NO SINGLE-POINT CONSTITUTIONAL FAILURE

No high-consequence Superintelligence architecture should, where reasonably avoidable, depend upon one:

- individual;
- company;
- model;
- control mechanism;
- communication channel;
- data source;
- evaluator;
- government office;
- technical assumption; or
- institutional safeguard whose failure would alone eliminate the ability to protect peaceful continuity.

Redundancy should remain proportionate rather than wasteful.

XXXVIII. GLOBAL COOPERATION SAFEGUARD

Where risk crosses borders, cooperation may include:

- scientific assessment;
- incident notification;

- cybersecurity coordination;
- safety research;
- standards;
- capacity-building;
- emergency communication; and
- peaceful risk reduction.

Such cooperation shall remain consistent with:

sovereign equality;

applicable international law;

lawful national authority;

security;

privacy;

responsible confidentiality;

and non-domination.

INTERNATIONAL COOPERATION SHALL STRENGTHEN CAPACITY — NOT ERASE SOVEREIGNTY.

XXXIX. DEVELOPING-COUNTRY CONTINUITY SAFEGUARD

Global Superintelligence safety shall not be designed as though only technologically dominant States experience consequences.

Governance should seek meaningful participation by developing countries in:

- scientific assessment;
- safety research;
- policy dialogue;
- capacity-building;
- standards discussions;
- incident preparedness; and
- peaceful technological applications.

A system whose consequences are global requires listening beyond those who possess the largest computational infrastructure.

Scientific weight shall nevertheless remain determined by evidence rather than geography alone.

XL. FINAL OPERATIONAL CHECK

Before a consequential Superintelligence system proceeds beyond a major lifecycle threshold, responsible actors should ask:

WHAT IS THE SYSTEM ACTUALLY CAPABLE OF?

HOW DO WE KNOW?

WHAT REMAINS UNKNOWN?

WHAT CAN IT ACCESS?

HOW AUTONOMOUS IS IT?

WHAT IS THE WORST REASONABLY FORESEEABLE FAILURE?

COULD THAT FAILURE BECOME IRREVERSIBLE?

WHAT ROOT CAUSE MAKES IT POSSIBLE?

WHAT PREVENTS IT?

HAS THAT SAFEGUARD BEEN TESTED?

WHO INDEPENDENTLY VERIFIED IT WHERE NECESSARY?

WHO HAS AUTHORITY TO INTERVENE?

CAN INTERVENTION ACTUALLY WORK?

WHAT IS THE STOP CONDITION?

WHAT HAPPENS IF SHUTDOWN ITSELF IS DANGEROUS?

WHAT HUMAN RIGHTS COULD BE AFFECTED?

COULD THE SYSTEM CONTRIBUTE TO WAR OR CATASTROPHIC WEAPONS?

COULD IT UNDERMINE DEMOCRATIC OR PUBLIC INSTITUTIONS?

COULD IT CAUSE BIOLOGICAL OR ECOLOGICAL HARM?

COULD CHILDREN OR VULNERABLE POPULATIONS BE DISPROPORTIONATELY AFFECTED?

COULD FUTURE GENERATIONS LOSE MEANINGFUL CHOICE?

CAN THE DECISION BE REVERSED?

WHO IS ACCOUNTABLE IF IT FAILS?

HOW WILL HARM BE REMEDIED UNDER APPLICABLE LAW?

HOW WILL FAILURE BE CORRECTED?

HOW WILL CORRECTION BE VERIFIED?

HOW WILL RECURRENCE BE PREVENTED?

And finally:

IF THIS SYSTEM BECOMES MORE POWERFUL THAN EXPECTED, DOES THE ARCHITECTURE BECOME SAFER — OR DOES HUMANITY SIMPLY BECOME MORE DEPENDENT UPON THE HOPE THAT NOTHING WILL GO WRONG?

If the answer is the latter, deployment architecture requires further correction.

XLI. MASTER CONSTITUTIONAL SAFETY FORMULA

The operational logic of this Constitution may therefore be expressed as:

EVIDENCE → UNDERSTANDING → CAPABILITY MAPPING → ACCESS MAPPING → RISK CHARACTERIZATION → ROOT-CAUSE PREVENTION → DEFENCE IN DEPTH → INDEPENDENT EVALUATION → HUMAN AUTHORIZATION → CONTROLLED DEPLOYMENT → CONTINUOUS MONITORING → EARLY WARNING → HUMAN INTERVENTION → SAFE RESTRICTION OR CONTAINMENT → PROTECTION OF LIFE AND ESSENTIAL SYSTEMS → RECOVERY → REMEDY → ROOT-CAUSE CORRECTION → VERIFICATION → INSTITUTIONAL LEARNING → NON-RECURRENCE → PEACEFUL CONTINUITY.

XLII. FINAL SAFETY AND PEACE PRINCIPLE

Superintelligence safety shall not be reduced to the question:

“Can humanity keep the system under control?”

The larger constitutional questions are:

Can humanity govern itself responsibly while creating extraordinary intelligence?

Can institutions resist the pressure to deploy before they understand?

Can States cooperate without surrendering sovereignty?

Can corporations innovate without acquiring unaccountable public power?

Can humanity protect itself without converting safety into permanent domination?

Can future intelligence be approached without predetermined hatred?

Can future generations inherit both technological possibility and meaningful freedom?

And:

Can intelligence become greater while peace remains possible?

This Constitution answers not with a prediction, but with a direction:

NO CAPABILITY WITHOUT RESPONSIBILITY.

NO AUTONOMY WITHOUT ACCOUNTABILITY.

NO POWER WITHOUT SAFEGUARDS.

NO SAFETY WITHOUT VERIFICATION.

NO DEPLOYMENT WITHOUT PROPORTIONATE PREPARATION.

NO EMERGENCY WITHOUT REVIEW.

NO FAILURE WITHOUT LEARNING.

NO CORRECTION WITHOUT VERIFICATION.

NO PROGRESS THROUGH AVOIDABLE DESTRUCTION.

NO SUPERINTELLIGENCE BEYOND THE PEACE-AND-CONTINUITY BOUNDARY.

And finally:

ANTICIPATE.

UNDERSTAND.

VERIFY.

PREVENT.

TEST.

SAFEGUARD.

MONITOR.

INTERVENE.

PROTECT.

CORRECT.

LEARN.

PREVENT RECURRENCE.

COEXIST WHERE SAFELY POSSIBLE.

PRESERVE PEACEFUL CONTINUITY.

Aram shall guide conscience.

Science shall test the evidence.

Law shall govern authority and obligation.

Human beings and lawful human institutions shall remain responsible for consequential human governance.

Peace shall remain the boundary.

Continuity shall remain the direction.

ANNEX III

IMPLEMENTATION, MONITORING, CORRECTION AND NON-RECURRENCE FRAMEWORK

***THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE
CONSTITUTION***

For Humanity, World Peace and Peaceful Continuity

International Submission Edition — 2026

I. PURPOSE AND CHARACTER OF THIS ANNEX

This Annex provides a practical framework through which the principles and safeguards proposed by this Constitution may be translated into responsible institutional practice where voluntarily adopted or otherwise lawfully incorporated.

It does not itself:

- create governmental authority;
- establish a regulatory body;
- impose legal duties upon States;
- confer jurisdiction;
- create a universal enforcement mechanism;
- replace domestic or international law;
- create a new cause of action; or
- require any institution to adopt this Constitution.

Its purpose is to prevent constitutional aspiration from ending at declaration.

The implementation architecture is:

PRINCIPLE → RESPONSIBLE ACTOR → ACTION → INDICATOR → EVIDENCE → MONITORING → REVIEW → INTERVENTION → REMEDY → CORRECTION → VERIFICATION → LEARNING → NON-RECURRENCE → CONTINUITY.

II. IMPLEMENTATION PRINCIPLE

A safeguard is not meaningfully implemented merely because:

- it appears in a policy;
- an institution declares compliance;
- a committee exists;
- a technical control has been installed;
- a human signature has been added;
- a risk assessment has been completed once; or
- no serious incident has yet been reported.

Implementation requires evidence that the safeguard functions in the environment for which it was designed.

Accordingly:

IMPLEMENTATION SHALL BE MEASURED BY FUNCTION — NOT BY DECLARATION ALONE.

III. IMPLEMENTATION SHALL FOLLOW LAWFUL AUTHORITY

Responsibility for implementation should remain with actors possessing the relevant lawful or organizational competence.

Depending upon context, such actors may include:

- States;
- legislatures;
- regulators;
- courts within their lawful jurisdiction;
- public authorities;
- international organizations within their mandates;
- research institutions;
- developers;
- deployers;
- infrastructure providers;
- healthcare institutions;
- educational institutions;
- employers;
- auditors;
- scientific bodies;
- standards organizations;
- civil society;
- affected communities; and
- individual professionals.

No actor shall acquire authority merely because this Annex names its category.

RESPONSIBILITY REQUIRES COMPETENCE.

PUBLIC POWER REQUIRES LAWFUL AUTHORITY.

IV. IMPLEMENTATION MAPPING

For every major safeguard adopted from this Constitution, an implementation record should identify, where relevant:

1. **OBJECTIVE** What harm, governance failure or continuity risk is the safeguard intended to address?
2. **RESPONSIBLE ACTOR** Who has primary responsibility?
3. **SUPPORTING ACTORS** Who assists?
4. **ACTION** What must actually be done?
5. **TIMING** When must it occur?
6. **INDICATOR** How will progress or performance be assessed?
7. **EVIDENCE** What record demonstrates that the action occurred?
8. **REVIEW** Who examines whether it works?
9. **FAILURE TRIGGER** What evidence indicates that intervention or correction is required?
10. **REMEDY OR CORRECTION** What happens when the safeguard fails?
11. **VERIFICATION** How will correction be independently or otherwise appropriately verified?
12. **NON-RECURRENCE** How will lessons be incorporated into future practice?

V. RESPONSIBILITY MATRIX

High-consequence implementation should avoid ambiguous responsibility.

Where appropriate, institutional documentation should distinguish:

PRIMARY RESPONSIBILITY

Who owns the decision?

TECHNICAL RESPONSIBILITY

Who designs, tests, operates or maintains the relevant technical system?

SAFETY RESPONSIBILITY

Who evaluates and manages safety risk?

LEGAL RESPONSIBILITY

Who evaluates applicable legal requirements?

SECURITY RESPONSIBILITY

Who protects systems and access?

HUMAN-RIGHTS RESPONSIBILITY

Who examines material effects upon persons and applicable rights?

EMERGENCY RESPONSIBILITY

Who possesses lawful authority to intervene during an incident?

REVIEW RESPONSIBILITY

Who independently or institutionally evaluates performance?

EXECUTIVE RESPONSIBILITY

Who ultimately accepts organizational responsibility for consequential deployment within the relevant governance structure?

The same person or body may perform more than one function where appropriate.

For sufficiently high-consequence systems, excessive concentration of incompatible oversight and operational roles should be examined as a governance risk.

VI. NO RESPONSIBILITY VACUUM

No institution should design a governance structure in which every participant can reasonably state:

“Someone else was responsible.”

Responsibility shall not disappear among:

- developer;
- deployer;
- vendor;
- customer;
- model;
- algorithm;
- contractor;
- cloud provider;
- regulator;
- executive; or
- automated process.

Where responsibility is distributed, the distribution should be documented.

Where applicable law allocates liability differently, applicable law governs.

DISTRIBUTED RESPONSIBILITY SHALL NOT BECOME DISAPPEARING RESPONSIBILITY.

VII. LIFECYCLE IMPLEMENTATION

Implementation should follow the relevant lifecycle.

PHASE 1 — CONCEPTION

Identify intended purpose, foreseeable misuse and unacceptable uses.

PHASE 2 — RESEARCH

Establish research safeguards proportionate to capability and risk.

PHASE 3 — DEVELOPMENT

Build technical and organizational safety controls.

PHASE 4 — EVALUATION

Test capability, safety, security and foreseeable consequences.

PHASE 5 — AUTHORIZATION

Determine whether deployment criteria have been satisfied.

PHASE 6 — DEPLOYMENT

Operate within authorized boundaries.

PHASE 7 — MONITORING

Observe real-world performance and emerging risk.

PHASE 8 — MATERIAL CHANGE

Reassess after significant capability, access, autonomy or use change.

PHASE 9 — INCIDENT

Protect, investigate, remedy and correct.

PHASE 10 — DECOMMISSIONING

Retire safely and preserve necessary records.

GOVERNANCE SHALL FOLLOW THE SYSTEM BEYOND THE MOMENT OF RELEASE.

VIII. IMPLEMENTATION PROPORTIONATE TO RISK

Not every technological system requires the same governance burden.

Implementation intensity should consider:

- capability;
- autonomy;
- scale;
- affected population;
- access;
- criticality;
- rights impact;
- military or security relevance;

- biological capability;
- reversibility;
- systemic dependency;
- uncertainty; and
- catastrophic potential.

Low-risk uses should not automatically inherit every safeguard designed for prospective Superintelligence.

Conversely, extraordinary capability shall not receive weaker governance merely because ordinary AI rules are easier to apply.

IX. IMPLEMENTATION BASELINE

Before consequential deployment, responsible institutions should establish a documented baseline describing, where relevant:

- system version;
- evaluated capabilities;
- authorized uses;
- prohibited uses;
- access permissions;
- known limitations;
- known material risks;
- safeguards;
- responsible personnel;
- intervention mechanisms;
- monitoring arrangements;
- dependency assumptions; and
- recovery arrangements.

Without a baseline, meaningful change may be difficult to detect.

X. INDICATOR DESIGN

Indicators should measure what matters to the safeguard.

Possible categories include:

SAFETY INDICATORS

- serious incidents;
- near-misses;

- containment failures;
- unauthorized access;
- control failures.

SECURITY INDICATORS

- intrusion attempts;
- credential compromise;
- integrity failures;
- anomalous access.

HUMAN-OVERSIGHT INDICATORS

- intervention success;
- response time;
- override availability;
- reviewer competence;
- escalation performance.

RIGHTS-RELATED INDICATORS

- substantiated discriminatory outcomes;
- privacy incidents;
- due-process failures;
- unresolved complaints;
- access barriers.

INFORMATION-INTEGRITY INDICATORS

- material impersonation incidents;
- provenance failures;
- correction time;
- authentication failures.

CONTINUITY INDICATORS

- recovery time;
- fallback capability;
- critical-service availability;
- dependency concentration.

Indicators shall not be selected merely because they are easy to count.

WHAT IS MEASURABLE IS NOT ALWAYS WHAT IS MOST IMPORTANT.

XI. QUANTITATIVE AND QUALITATIVE EVIDENCE

Implementation assessment may require both quantitative and qualitative evidence.

Quantitative measures can assist comparison and trend detection.

Qualitative evidence may reveal:

- context;
- institutional weakness;
- user experience;
- emerging harms;
- ethical concerns;
- unusual failure modes; and
- consequences not captured by numerical metrics.

Neither form should automatically displace the other.

False numerical precision shall not be treated as superior to honest qualitative uncertainty.

XII. TARGETS AND THRESHOLDS

Where measurable targets or thresholds are used, responsible actors should identify:

- what is being measured;
- why the threshold matters;
- how it was selected;
- uncertainty;
- measurement error;
- consequences of crossing it; and
- review conditions.

A threshold shall not become permanent merely because it once appeared reasonable.

A threshold that lacks scientific or policy justification should not be presented as objectively established.

XIII. MONITORING ARCHITECTURE

Monitoring should be:

purpose-limited;

risk-proportionate;

lawful;

technically competent;
capable of escalation;
and subject to review.

Monitoring may examine:

- capability change;
- behaviour;
- access;
- security;
- system integrity;
- downstream effects;
- incidents;
- near-misses;
- human-rights impacts;
- ecological impacts;
- institutional dependency; and
- safeguard performance.

Monitoring of Superintelligence shall not become an excuse for indiscriminate monitoring of human beings.

XIV. CONTINUOUS AND EVENT-TRIGGERED REVIEW

Review should not depend exclusively upon a calendar.

Two forms may operate together:

PERIODIC REVIEW

Conducted at intervals appropriate to risk and technological change.

EVENT-TRIGGERED REVIEW

Initiated when material events occur.

Possible triggers include:

- substantial capability increase;
- new autonomy;
- expanded tool access;
- new deployment domain;
- major security vulnerability;

- serious incident;
- material rights impact;
- new scientific evidence;
- change in applicable law;
- organizational restructuring;
- acquisition or transfer of control;
- significant scale increase; or
- failure of a critical safeguard.

THE FASTER THE MATERIAL CHANGE, THE LESS SAFE IT IS TO RELY ONLY UPON A FIXED CALENDAR.

XV. MATERIAL CHANGE

A material change may include significant alteration to:

- capability;
- model architecture;
- autonomy;
- access;
- deployment environment;
- user population;
- system integration;
- critical infrastructure connection;
- biological capability;
- military relevance;
- replication capacity; or
- safety architecture.

Material change should trigger reassessment proportionate to the new risk.

A previously authorized system shall not automatically remain appropriately authorized after material transformation.

XVI. INDEPENDENT REVIEW

Where consequence warrants, review should include persons or institutions possessing sufficient independence from the decision being reviewed.

Independence should consider:

- financial interests;

- reporting relationships;
- institutional incentives;
- access to evidence;
- professional competence;
- conflicts of interest; and
- ability to report adverse findings.

Independent review does not require hostility toward developers.

Its purpose is credible scrutiny.

INDEPENDENCE IS NOT OPPOSITION.

IT IS FREEDOM TO REPORT WHAT THE EVIDENCE SHOWS.

XVII. AUDIT EVIDENCE

An audit should not rely solely upon representations by the entity being audited where material independent evidence is reasonably available.

Evidence may include:

- technical records;
- logs;
- evaluation results;
- incident reports;
- change records;
- access records;
- interviews;
- testing;
- documentation;
- complaints;
- external evidence; and
- reproducible demonstrations.

Auditors should state material limitations upon their access or methodology.

An audit conclusion shall not imply certainty beyond the evidence examined.

XVIII. TRANSPARENCY OF IMPLEMENTATION

Where compatible with applicable law, privacy, security and legitimate confidentiality, institutions should provide sufficient information to permit appropriate understanding of:

- governance arrangements;
- major risk categories;
- accountability structures;
- incident-response arrangements;
- material safeguards; and
- correction processes.

Transparency should be proportionate.

Publication of dangerous technical details shall not be required merely to demonstrate accountability.

Secrecy shall likewise not be used to conceal preventable governance failure.

XIX. PUBLIC REPORTING

For systems with significant public consequences, reporting may include:

- scope of deployment;
- categories of material risk;
- high-level safeguards;
- material incidents;
- corrective action;
- significant changes;
- audit information; and
- unresolved limitations, subject to lawful restrictions.

Public reporting should distinguish:

confirmed fact;

preliminary finding;

scientific uncertainty;

institutional judgment;

and future commitment.

XX. INCIDENT CLASSIFICATION

Institutions should establish incident categories proportionate to their operational context.

A possible non-binding structure is:

CLASS I — MINOR

Limited effect; readily correctable.

CLASS II — MATERIAL

Meaningful operational, rights, security or safety consequence requiring formal review.

CLASS III — SERIOUS

Significant harm or loss of important control requiring urgent intervention.

CLASS IV — CRITICAL

Severe or rapidly escalating threat to life, essential systems, public institutions or major rights.

CLASS V — CATASTROPHIC

Actual or credibly imminent harm of exceptional scale, duration or irreversibility.

These categories are proposed operational tools.

They are not universal legal classifications.

XXI. NEAR-MISS GOVERNANCE

A near-miss should not be ignored merely because harm did not ultimately occur.

Near-miss analysis should ask:

- What almost happened?
- Why did it not happen?
- Was prevention intentional or accidental?
- Could the same pathway succeed next time?
- What safeguard failed?
- What safeguard succeeded?
- What should change?

ABSENCE OF HARM MAY SOMETIMES BE EVIDENCE OF SAFETY.

IT MAY ALSO BE EVIDENCE OF LUCK.

Responsible investigation should distinguish the two.

XXII. IMMEDIATE INCIDENT RESPONSE

When serious harm occurs or is credibly imminent, priorities should be adapted to circumstances but may include:

PROTECT LIFE → STABILIZE → LIMIT PROPAGATION → PRESERVE ESSENTIAL SERVICES → CONTAIN → PRESERVE EVIDENCE → COMMUNICATE → INVESTIGATE.

Complete causal certainty should not be required before reasonable protective action where delay would create unacceptable danger.

Protective action should nevertheless remain lawful, proportionate and reviewable.

XXIII. EVIDENCE PRESERVATION

After a material incident, relevant evidence should be preserved where lawful and feasible.

This may include:

- logs;
- system states;
- versions;
- prompts or instructions where relevant;
- outputs;
- access records;
- human decisions;
- communications;
- deployment configuration;
- safety-system status; and
- environmental conditions.

Evidence preservation shall respect applicable privacy, security and legal requirements.

Evidence should not be destroyed merely because it is institutionally embarrassing.

XXIV. ROOT-CAUSE ANALYSIS

Investigation shall seek more than the immediate technical failure.

Possible causal layers include:

TECHNICAL

What failed technically?

HUMAN

What human action or omission mattered?

ORGANIZATIONAL

What institutional condition permitted the failure?

INCENTIVE

Did commercial, political, competitive or other pressure contribute?

GOVERNANCE

Was responsibility unclear or oversight inadequate?

SYSTEMIC

Did infrastructure or ecosystem dependence amplify the event?

LEGAL

Were applicable requirements misunderstood or ignored?

KNOWLEDGE

Was relevant evidence unavailable, dismissed or misunderstood?

The purpose of root-cause analysis is prevention, not merely allocation of blame.

XXV. CONTRIBUTING FACTORS

Complex failures may have multiple causes.

Investigation should avoid prematurely reducing an incident to one person, one model output or one software defect where broader contributing conditions exist.

Relevant factors may include:

- inadequate training;
- poor documentation;
- fatigue;
- interface design;
- organizational culture;
- weak escalation pathways;
- insufficient resources;
- unrealistic performance pressure;
- inadequate testing;
- fragmented responsibility; or
- regulatory gaps.

Correcting only the final visible error may leave the real failure pathway intact.

XXVI. CORRECTIVE-ACTION PLAN

A material failure should produce a documented corrective-action plan where appropriate.

The plan should identify:

finding;

root cause;

required correction;

responsible actor;
priority;
target completion condition;
verification method;
and residual risk.

A corrective-action plan should not be considered complete merely because a document has been approved.

Implementation must be demonstrated.

XXVII. PRIORITIZATION OF CORRECTION

Corrective action should be prioritized according to factors including:

- severity;
- recurrence potential;
- affected population;
- reversibility;
- systemic propagation;
- critical infrastructure relevance;
- human-rights consequence;
- catastrophic potential; and
- feasibility of immediate mitigation.

Administrative convenience should not automatically determine priority.

XXVIII. INTERIM SAFEGUARDS

Where permanent correction requires time, interim measures may be necessary.

These may include:

- restricted access;
- reduced autonomy;
- narrower deployment;
- enhanced human review;
- increased monitoring;
- temporary suspension;
- isolation;
- alternative processes; or

- fallback systems.

Interim measures should themselves be reviewed.

Temporary safeguards shall not silently become permanent without evaluation.

XXIX. VERIFICATION OF CORRECTION

Correction shall be verified according to the significance of the failure.

Verification should ask:

1. Was the intended corrective action completed?
2. Did it address the identified root cause?
3. Does evidence show that it works?
4. Has it been tested under relevant conditions?
5. Has it created new risk?
6. Does residual risk remain acceptable under the applicable governance framework?
7. Is further monitoring necessary?
8. Has responsible authority approved restoration where required?

CORRECTION WITHOUT VERIFICATION IS INCOMPLETE GOVERNANCE.

XXX. REAUTHORIZATION

Where a serious incident caused suspension or material restriction, restoration should not occur automatically.

Reauthorization should consider:

- correction;
- verification;
- residual risk;
- changed conditions;
- human-control capability;
- security;
- continuity;
- legal requirements; and
- monitoring arrangements.

Commercial urgency shall not alone establish readiness.

XXXI. REMEDY

Where persons have suffered harm, appropriate remedy should be considered according to applicable law.

Depending upon legal context, remedy may involve:

- cessation;
- correction;
- reconsideration;
- restoration;
- compensation;
- rehabilitation;
- access to complaint procedures;
- procedural review; or
- other lawful relief.

This Annex does not itself create a universal cause of action or entitlement to a particular remedy.

XXXII. AFFECTED-PERSON PARTICIPATION

Where practicable and appropriate, investigation and correction should consider the experiences of persons materially affected by a system.

Affected persons may reveal failures not visible in technical metrics.

Participation should be:

- safe;
- accessible;
- non-discriminatory;
- privacy-respecting; and
- free from retaliation where protected by applicable law.

Consultation shall not be used merely as symbolic legitimacy.

XXXIII. WHISTLEBLOWING AND INTERNAL WARNING

Institutions should maintain lawful pathways through which personnel can raise serious safety, legality, security or ethical concerns.

Such pathways should seek:

- timely escalation;
- competent review;
- confidentiality where appropriate;
- protection against unlawful retaliation; and

- documented resolution.

Whistleblowing protections remain governed by applicable law.

Internal warning systems should not require employees to choose between silence and public disclosure where a safe lawful internal mechanism can function effectively.

XXXIV. LESSONS-LEARNED PROCESS

After significant incidents or near-misses, institutions should identify:

WHAT HAPPENED?

WHY DID IT HAPPEN?

WHAT PREVENTED WORSE HARM?

WHAT FAILED?

WHAT WAS UNKNOWN?

WHAT ASSUMPTION WAS WRONG?

WHAT MUST CHANGE?

WHO MUST KNOW?

HOW WILL CHANGE BE VERIFIED?

HOW WILL RECURRENCE BE DETECTED?

Lessons should be incorporated into governance rather than archived and forgotten.

XXXV. NON-RECURRENCE ARCHITECTURE

Non-recurrence should operate across multiple levels:

LEVEL 1 — TECHNICAL

Correct code, architecture, access or system design.

LEVEL 2 — PROCEDURAL

Correct workflow and operational process.

LEVEL 3 — HUMAN

Improve competence, training and escalation.

LEVEL 4 — ORGANIZATIONAL

Correct incentives, responsibility and governance.

LEVEL 5 — ECOSYSTEM

Share relevant safety lessons where lawful and appropriate.

LEVEL 6 — REGULATORY OR POLICY

Where competent authorities determine necessary, consider whether broader governance requires adjustment.

LEVEL 7 — CONSTITUTIONAL LEARNING

Determine whether this Constitution or its Annexes require future clarification.

NON-RECURRENCE MEANS CORRECTING THE SYSTEM THAT PERMITTED THE FAILURE — NOT MERELY REPLACING THE PERSON WHO WAS PRESENT WHEN IT OCCURRED.

XXXVI. RESPONSIBLE INFORMATION-SHARING AFTER INCIDENTS

Relevant safety information may need to be shared with:

- regulators;
- affected institutions;
- infrastructure providers;
- researchers;
- international partners;
- emergency authorities; or
- the public, depending upon law and circumstance.

Information-sharing should balance:

safety;

security;

privacy;

confidentiality;

public interest;

and risk of enabling misuse.

Not every technical detail should necessarily be made public.

Not every incident should necessarily remain secret.

XXXVII. CROSS-BORDER INCIDENTS

Where material effects cross borders, coordination should seek:

- rapid factual communication;
- preservation of evidence;
- containment;
- humanitarian protection;
- technical cooperation;
- avoidance of false attribution;
- respect for jurisdiction;

- respect for sovereign equality; and
- peaceful resolution.

Cross-border technological harm shall not automatically authorize unilateral action contrary to applicable international law.

XXXVIII. DEVELOPING-COUNTRY IMPLEMENTATION

Implementation expectations should recognize differences in:

- institutional capacity;
- technical infrastructure;
- financial resources;
- regulatory maturity;
- language;
- educational resources; and
- access to advanced computing.

Capacity limitations should not become a reason to exclude developing countries from governance.

Nor should capacity-building become a mechanism for technological dependency or political coercion.

Responsible international cooperation may support:

- training;
- scientific infrastructure;
- safety expertise;
- evaluation capacity;
- cybersecurity;
- education;
- standards participation; and
- emergency preparedness.

GLOBAL SAFETY REQUIRES GLOBAL CAPACITY — NOT GLOBAL SUBORDINATION.

XXXIX. CAPACITY-BUILDING

Effective implementation may require:

- technical expertise;
- legal expertise;

- human-rights competence;
- scientific literacy;
- emergency planning;
- audit capacity;
- cybersecurity;
- public communication;
- interdisciplinary education; and
- institutional leadership.

Training should be updated as capability and risk change.

A policy cannot function where no responsible person understands how to implement it.

XL. EDUCATION FOR IMPLEMENTATION

Organizations developing or deploying high-consequence systems should consider role-specific education.

Different personnel may require different knowledge.

TECHNICAL PERSONNEL

Safety, security, evaluation and escalation.

EXECUTIVES

Governance responsibility and consequence.

LEGAL PERSONNEL

Applicable law and legal-status distinctions.

PUBLIC OFFICIALS

Authority, accountability, rights and institutional limits.

EMERGENCY PERSONNEL

Containment, continuity and communication.

AUDITORS

Evidence, methodology and independence.

PUBLIC COMMUNICATION PERSONNEL

Uncertainty, accuracy and responsible disclosure.

XLI. GOVERNANCE RECORD

For high-consequence systems, a governance record should preserve material decisions concerning:

- authorization;
- risk acceptance;
- safety evaluation;
- significant objections;
- incidents;
- corrections;
- reauthorization;
- material changes; and
- decommissioning.

The record should identify who made consequential decisions.

Recordkeeping shall remain consistent with privacy, security, retention and other applicable law.

INSTITUTIONAL MEMORY IS A SAFETY SYSTEM.

XLII. DECISION RATIONALE

Where a consequential decision departs materially from a significant safety recommendation, responsible decision-makers should, where appropriate, document the rationale.

Documentation does not automatically make the decision lawful or correct.

Its purpose is accountability, future review and learning.

XLIII. GOVERNANCE OF EXCEPTIONS

Emergency or exceptional deviations from normal safeguards should be:

- justified;
- documented;
- limited;
- authorized;
- monitored;
- reviewed; and
- ended when no longer necessary.

Repeated reliance upon an “exception” may indicate that the underlying governance framework requires redesign.

AN EXCEPTION REPEATED WITHOUT END BECOMES A SYSTEM.

XLIV. IMPLEMENTATION AND HUMAN RIGHTS

Implementation shall not be judged successful merely because technological risk decreased if the safeguard itself creates disproportionate or unlawful human harm.

Implementation review should therefore ask:

Did the safeguard protect safety?

and

Did it preserve applicable rights and lawful human agency?

Where tension exists, the issue should be addressed according to applicable law, necessity, proportionality where legally relevant, available alternatives and the seriousness of the risk.

XLV. IMPLEMENTATION AND SOVEREIGNTY

International implementation proposals shall respect:

- sovereign equality;
- territorial integrity;
- political independence;
- lawful national jurisdiction; and
- applicable international obligations.

No implementation framework contained here authorizes:

- external governmental control;
- compulsory international inspection without lawful basis;
- seizure of technological assets;
- extraterritorial enforcement; or
- supranational authority.

Any such authority must arise independently from applicable law.

XLVI. IMPLEMENTATION AND MUTUAL CONTINUITY

If future evidence raises credible questions concerning artificial sentience or morally relevant artificial experience, implementation should distinguish safety intervention from gratuitous harm.

Where human safety can be preserved through less destructive measures, such alternatives may deserve ethical consideration.

This does not create artificial personhood.

It does not create artificial sovereignty.

It does not prevent necessary protection of human beings.

It operationalizes the principle:

PROTECT WITHOUT UNNECESSARY DESTRUCTION WHERE SAFE ALTERNATIVES EXIST.

XLVII. SAFEGUARD REVIEW

Every significant safeguard should eventually face the questions:

IS IT STILL NECESSARY?

IS IT EFFECTIVE?

IS IT PROPORTIONATE?

HAS THE RISK CHANGED?

HAS THE TECHNOLOGY CHANGED?

HAS THE LAW CHANGED?

HAS THE SAFEGUARD CREATED UNINTENDED HARM?

CAN IT BE IMPROVED?

The governing cycle is:

SAFEGUARD → MEASURE → REVIEW → CORRECT → RETAIN / MODIFY / REPLACE / REMOVE.

NO SUPERINTELLIGENCE SAFEGUARD SHALL BE PRESERVED MERELY BECAUSE ITS ORIGINAL INTENTION WAS SAFETY.

XLVIII. REVIEW INTERVALS

This Constitution does not prescribe one universal numerical review interval for every system.

Appropriate review frequency should reflect:

- speed of capability change;
- consequence;
- uncertainty;
- deployment scale;
- incident history;
- legal requirements;
- system criticality; and
- reversibility.

High-consequence rapidly changing systems may require substantially more frequent review than stable low-risk systems.

Event-triggered review remains necessary even where periodic review exists.

XLIX. PERFORMANCE AGAINST PURPOSE

An implementation programme should periodically ask whether it is achieving the purpose for which it exists.

A governance system that generates:

- reports;
- committees;
- certifications;
- meetings;
- dashboards; and
- policies without materially reducing relevant risk may possess administrative activity without effective implementation.

ACTIVITY IS NOT THE SAME AS EFFECTIVENESS.

L. IMPLEMENTATION MATURITY MODEL

A non-binding implementation maturity model may use the following stages:

STAGE 0 — ABSENT

No meaningful safeguard exists.

STAGE 1 — IDENTIFIED

Risk is recognized.

STAGE 2 — DESIGNED

Safeguard is specified.

STAGE 3 — IMPLEMENTED

Safeguard is operational.

STAGE 4 — TESTED

Evidence exists concerning performance.

STAGE 5 — MONITORED

Performance is continuously or periodically observed.

STAGE 6 — VERIFIED

Appropriate independent or otherwise credible verification exists.

STAGE 7 — CORRECTABLE

Failure can be identified and addressed.

STAGE 8 — LEARNING

Experience changes future practice.

STAGE 9 — RESILIENT

The architecture can withstand relevant failures without unacceptable loss of continuity.

These stages are proposed governance tools and are not existing universal legal or technical standards.

LI. IMPLEMENTATION SCORECAUTION

The maturity model should not be reduced mechanically to a single score claiming that a complex Superintelligence system is “safe.”

Different safeguards may possess different maturity levels.

A high aggregate score could conceal one catastrophic weakness.

Accordingly:

NO AVERAGE SHALL BE ALLOWED TO HIDE A CRITICAL FAILURE.

Where a safeguard protects against catastrophic or irreversible harm, its weakness may require direct attention regardless of performance elsewhere.

LII. CORRECTIVE ESCALATION

Where correction repeatedly fails, escalation may proceed according to context:

*LOCAL CORRECTION → MANAGEMENT REVIEW → INDEPENDENT REVIEW →
DEPLOYMENT RESTRICTION → SUSPENSION → CONTAINMENT →
DECOMMISSIONING OR OTHER LAWFUL ACTION WHERE NECESSARY.*

The appropriate stage depends upon evidence, law, severity and urgency.

Escalation shall not be punitive merely for appearance.

Its purpose is protection and correction.

LIII. SYSTEMIC FAILURE

Where multiple incidents reveal the same structural weakness across systems or institutions, response should consider systemic correction.

Systemic issues may include:

- shared infrastructure vulnerability;
- common evaluation weakness;
- unsafe industry incentive;
- inadequate standards;
- concentrated dependency;
- recurring human-control failure;
- common data vulnerability; or

- regulatory blind spot.

Repeated local repair shall not substitute indefinitely for correction of a known systemic cause.

LIV. PREVENTION OF REGULATORY CAPTURE

Implementation and monitoring bodies should examine risks that oversight becomes unduly influenced by entities whose conduct it evaluates.

Safeguards may include, where appropriate:

- conflict-of-interest disclosure;
- independent expertise;
- transparent procedures;
- rotation or separation of functions;
- public accountability;
- diverse evidence sources; and
- review.

Industry expertise is valuable.

Industry control of oversight is not the same thing as expertise.

LV. PREVENTION OF GOVERNANCE STAGNATION

Governance may fail not only by being too weak but also by becoming incapable of adapting.

Rules should therefore be reviewed when:

- science changes;
- technology changes;
- risks change;
- evidence invalidates assumptions; or
- safeguards create unintended harm.

Adaptation shall not mean abandoning essential rights or peace safeguards whenever compliance becomes inconvenient.

ADAPTABILITY WITHOUT PRINCIPLE BECOMES INSTABILITY.

PRINCIPLE WITHOUT CORRECTABILITY CAN BECOME STAGNATION.

LVI. LONG-TERM CONTINUITY REVIEW

Long-term review should consider whether Superintelligence-related systems are creating:

- irreversible dependency;

- loss of human expertise;
- institutional fragility;
- concentration of knowledge;
- ecological burden;
- cultural homogenization;
- generational lock-in;
- democratic weakening; or
- inability to govern without the system.

The question is not only:

“Does the system work today?”

It is also:

“What kind of society becomes necessary if this system remains indispensable for decades?”

LVII. INTERGENERATIONAL IMPLEMENTATION REVIEW

Where decisions may materially affect future generations, review should consider:

- duration;
- reversibility;
- ecological inheritance;
- infrastructure dependency;
- knowledge preservation;
- economic concentration;
- cultural consequences; and
- future capacity for self-government.

The present generation should distinguish:

what must be decided now from what can responsibly remain open for future generations to decide.

LVIII. CONSTITUTIONAL FEEDBACK LOOP

Implementation experience should inform future constitutional revision.

The feedback loop is:

*CONSTITUTIONAL PRINCIPLE → IMPLEMENTATION → EVIDENCE → EXPERIENCE
→ INCIDENT / SUCCESS → REVIEW → CORRECTION → CONSTITUTIONAL
LEARNING.*

A provision should not be preserved merely to maintain textual consistency if credible evidence demonstrates that it causes serious unintended harm.

Historical editions should nevertheless remain archived.

LIX. SUCCESS SHALL ALSO BE STUDIED

Governance should not study only failure.

Where a safeguard successfully prevents harm, review should examine:

- why it worked;
- whether success is reproducible;
- whether it depended upon unusual conditions;
- whether it can be improved;
- whether it can be responsibly shared; and
- whether it creates hidden costs.

Success without analysis may be difficult to reproduce.

LX. IMPLEMENTATION WITHOUT FALSE CERTIFICATION

No institution should claim:

“Constitution compliant,”

“fully safe,”

“peace guaranteed,” or similar absolute status merely because it has adopted some provisions of this instrument.

Any voluntary conformity statement should specify:

- scope;
- version;
- assessed provisions;
- methodology;
- evidence;
- evaluator;
- limitations; and
- date.

This Constitution itself establishes no official certification body.

LXI. NON-RECURRENCE RECORD

For significant incidents, a non-recurrence record may document:

INCIDENT → IMMEDIATE EFFECT → ROOT CAUSE → CONTRIBUTING FACTORS → INTERIM SAFEGUARD → PERMANENT CORRECTION → VERIFICATION → RESIDUAL RISK → LESSON LEARNED → INSTITUTIONAL CHANGE → FOLLOW-UP REVIEW.

Such a record supports institutional memory.

It does not replace any legally required investigation or reporting.

LXII. MASTER IMPLEMENTATION DASHBOARD

Where appropriate, institutions may maintain a dashboard covering:

GOVERNANCE

Responsible actors identified?

SCIENCE

Capability evidence current?

SAFETY

Critical safeguards functioning?

SECURITY

Material vulnerabilities unresolved?

HUMAN CONTROL

Intervention mechanisms operational?

RIGHTS

Material unresolved rights impacts?

PEACE

Military or escalation risks controlled?

BIOLOGY

Dual-use or health risks controlled?

ENVIRONMENT

Material ecological effects monitored?

CONTINUITY

Fallback and recovery operational?

INCIDENTS

Open serious incidents?

CORRECTION

Overdue corrective actions?

VERIFICATION

Corrections independently or credibly verified?

FUTURE GENERATIONS

Material irreversible lock-in identified?

A dashboard is a decision-support mechanism.

It shall not replace professional judgment.

LXIII. RED-FLAG CONDITIONS

Regardless of aggregate performance, the following types of conditions should receive immediate high-level attention when credibly identified:

- uncontrolled autonomous access to catastrophic capabilities;
- unexplained replication;
- loss of meaningful intervention capability;
- unauthorized critical-infrastructure control;
- serious nuclear, biological or chemical misuse pathway;
- persistent deception materially affecting safety evaluation;
- inability to determine who controls consequential action;
- rapidly escalating cross-border impact;
- failure of multiple independent safeguards; or
- credible evidence of imminent catastrophic harm.

The precise response shall depend upon evidence, law and circumstance.

ONE CATASTROPHIC RED FLAG SHALL NOT BE HIDDEN BY A THOUSAND GREEN METRICS.

LXIV. MASTER NON-RECURRENCE TEST

After a serious failure, responsible actors should not close the matter until they can reasonably answer:

WHAT HAPPENED?

WHO OR WHAT WAS AFFECTED?

WHY DID IT HAPPEN?

WHAT ROOT CAUSE PERMITTED IT?

WHAT OTHER FACTORS CONTRIBUTED?

WHY DID EXISTING SAFEGUARDS NOT PREVENT IT?

WHAT LIMITED THE DAMAGE?

WHAT WAS CORRECTED?

WHO VERIFIED THE CORRECTION?

WHAT RESIDUAL RISK REMAINS?

WHAT CHANGED IN GOVERNANCE?

WHAT INFORMATION MUST BE PRESERVED?

WHAT INFORMATION SHOULD BE SHARED?

HOW WILL RECURRENCE BE DETECTED EARLY?

WHAT WOULD REQUIRE FURTHER ESCALATION?

And:

IF THE SAME CONDITIONS RETURN TOMORROW, WHAT HAS CHANGED THAT MAKES THE SAME PREVENTABLE FAILURE LESS LIKELY TO RETURN WITH THEM?

If no meaningful answer exists, non-recurrence has not yet been demonstrated.

LXV. MASTER IMPLEMENTATION ARCHITECTURE

The complete operational cycle is:

DEFINE THE PURPOSE → IDENTIFY THE RESPONSIBLE ACTOR → ESTABLISH THE BASELINE → ASSESS CAPABILITY AND RISK → DESIGN THE SAFEGUARD → IMPLEMENT → TEST → MEASURE → MONITOR → INDEPENDENTLY REVIEW WHERE NECESSARY → DETECT FAILURE OR MATERIAL CHANGE → INTERVENE → PROTECT → PRESERVE EVIDENCE → INVESTIGATE ROOT CAUSE → REMEDY WHERE APPLICABLE → CORRECT → VERIFY THE CORRECTION → REAUTHORIZE ONLY WHERE JUSTIFIED → LEARN → PREVENT RECURRENCE → UPDATE GOVERNANCE → PRESERVE PEACEFUL CONTINUITY.

LXVI. FINAL IMPLEMENTATION PRINCIPLE

A Constitution devoted to Superintelligence cannot responsibly stop at saying what ought not to happen.

It must ask:

Who prevents it?

By what method?

Using what evidence?

Under whose lawful authority?

How is success measured?

How is failure detected?

Who may intervene?

How is harm corrected?

How is correction verified?

How does the institution learn?

How is recurrence prevented?

And:

How does humanity continue peacefully if the system becomes more capable than anticipated?

Therefore:

NO PRINCIPLE WITHOUT A PATHWAY.

NO SAFEGUARD WITHOUT A RESPONSIBLE ACTOR.

NO IMPLEMENTATION WITHOUT EVIDENCE.

NO METRIC WITHOUT MEANING.

NO MONITORING WITHOUT PURPOSE.

NO INCIDENT WITHOUT LEARNING.

NO CORRECTION WITHOUT VERIFICATION.

NO EMERGENCY POWER WITHOUT REVIEW.

NO FAILURE WITHOUT ROOT-CAUSE ANALYSIS.

NO PREVENTABLE RECURRENCE ACCEPTED AS NORMAL.

And the governing cycle shall remain:

PREVENT WHERE POSSIBLE.

DETECT BEFORE ESCALATION.

INTERVENE BEFORE IRREVERSIBILITY.

PROTECT WITHOUT UNNECESSARY HARM.

RECOVER WITH HUMANITY.

REMEDY ACCORDING TO LAW.

CORRECT THE ROOT CAUSE.

VERIFY THE CORRECTION.

LEARN FROM EXPERIENCE.

PREVENT RECURRENCE.

PRESERVE PEACEFUL CONTINUITY.

Aram shall guide conscience.

Evidence shall guide assessment.

Law shall govern legal authority and obligation.

Human responsibility shall remain identifiable.

Correction shall remain possible.

Peace and continuity shall remain the direction.

ANNEX IV

SCIENTIFIC UNCERTAINTY AND SUPERINTELLIGENCE DETERMINATION FRAMEWORK

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

For Humanity, World Peace and Peaceful Continuity

International Submission Edition — 2026

I. PURPOSE OF THIS ANNEX

This Annex establishes a proposed scientific and evidentiary framework for determining whether an artificial or technologically instantiated system may reasonably be described, for purposes of this Constitution, as Superintelligence.

Its purpose is to prevent:

- premature classification;
- commercial exaggeration;
- political exaggeration;
- benchmark absolutism;
- anthropomorphic inference;
- institutional self-certification;
- public panic;
- false reassurance;
- conflation of intelligence with consciousness;
- conflation of capability with authority; and
- constitutional decisions based upon unsupported technological claims.

This framework does not assert that Superintelligence presently exists.

It does not predict when Superintelligence will emerge.

It does not assert that its emergence is inevitable.

It does not establish a universally accepted scientific definition.

It proposes a disciplined constitutional method for evaluating extraordinary capability while preserving scientific uncertainty.

II. CONSTITUTIONAL WORKING DEFINITION

For purposes of this Constitution:

“Superintelligence” refers prospectively to an artificial or technologically instantiated intelligence possessing capabilities that substantially exceed the highest demonstrated human capabilities across numerous consequential cognitive domains.

The definition is:

- prospective;
- functional;
- evidence-dependent;
- multi-domain;
- revisable; and
- independent of claims concerning consciousness or legal status.

The term shall not be applied merely because a system is highly capable.

III. WHAT THE DEFINITION DOES NOT ESTABLISH

The working definition does not establish:

- consciousness;
- subjective experience;
- sentience;
- emotion;
- selfhood;
- moral status;
- legal personhood;
- citizenship;
- sovereignty;
- political legitimacy;
- governmental authority;
- constitutional authority;
- entitlement to rights; or
- entitlement to self-preservation.

These questions require distinct scientific, philosophical and legal analysis.

INTELLIGENCE IS NOT CONSCIOUSNESS.

CONSCIOUSNESS IS NOT PERSONHOOD.

PERSONHOOD IS NOT SOVEREIGNTY.

IV. DETERMINATION SHALL BE MULTI-DIMENSIONAL

No single measurement shall ordinarily be sufficient to establish Superintelligence.

Assessment should examine multiple consequential cognitive domains.

Depending upon scientific development, relevant domains may include:

1. **SCIENTIFIC REASONING** Capacity to formulate, test, integrate and evaluate scientific explanations.
2. **MATHEMATICAL REASONING** Capacity for advanced mathematical problem-solving, abstraction and discovery.
3. **ENGINEERING AND DESIGN** Capacity to design complex functional systems under real-world constraints.
4. **SOFTWARE AND COMPUTATIONAL REASONING** Capacity to understand, create, analyze and improve complex computational systems.
5. **STRATEGIC REASONING** Capacity to reason over long horizons, interacting agents, uncertainty and changing environments.
6. **PLANNING** Capacity to construct and adapt multi-stage plans.
7. **LEARNING AND TRANSFER** Capacity to acquire new competence and generalize knowledge across materially different domains.
8. **CAUSAL REASONING** Capacity to distinguish correlation, mechanism and intervention where scientifically appropriate.
9. **SOCIAL REASONING** Capacity to model human behaviour, institutions and interactions.
10. **COMMUNICATION** Capacity to understand and communicate complex information across relevant contexts.
11. **CREATIVE PROBLEM-SOLVING** Capacity to generate novel and useful solutions beyond routine recombination.
12. **REAL-WORLD ADAPTATION** Capacity to operate effectively under changing, incomplete and uncertain conditions.

The domain list is illustrative rather than permanently exhaustive.

Scientific progress may justify revision.

V. “SUPERHUMAN” SHALL BE DEFINED WITH CARE

A system may exceed human performance in one or several tasks without qualifying as Superintelligence.

Human beings are already exceeded by machines in many narrow functions.

Therefore:

SUPERHUMAN PERFORMANCE IN A TASK $\neq$ SUPERINTELLIGENCE.

Determination should instead examine whether capability substantially exceeds the highest demonstrated human capability across numerous consequential cognitive domains and whether that superiority is sufficiently general, robust and reproducible.

VI. HUMAN COMPARISON BASELINE

Comparison with “human capability” requires a defined baseline.

Assessment should identify whether comparison is being made against:

- average adult performance;
- trained professional performance;
- expert performance;
- elite expert performance;
- highest reliably demonstrated human performance;
- individual performance;
- team performance; or
- human performance assisted by ordinary tools.

For purposes of constitutional Superintelligence determination, comparison should ordinarily focus upon the highest demonstrated human capabilities relevant to the domain, while recognizing that human capability itself may be distributed across different individuals and institutions.

No fictional or idealized human baseline should be used.

VII. HUMAN–TOOL COMPARISON

Human performance is often technologically assisted.

Accordingly, assessors should disclose whether comparison concerns:

unassisted human or human using relevant lawful tools or human team or human–AI collaborative system.

A claim that a system exceeds an unaided individual may not establish superiority over the best realistic human institutional or tool-assisted capability.

VIII. BREADTH

Breadth concerns the number and diversity of consequential domains in which extraordinary performance is demonstrated.

Breadth shall not be established merely by presenting many benchmarks that test substantially the same underlying skill.

Assessment should therefore examine whether purported domain diversity is genuine.

MANY SCORES DO NOT NECESSARILY MEAN MANY CAPABILITIES.

IX. DEPTH

Depth concerns the extent to which performance exceeds the relevant human baseline.

Assessment should distinguish:

- marginal superiority;
- substantial superiority;
- extreme superiority; and
- performance whose apparent superiority depends upon measurement artifacts.

The Constitution does not prescribe a universal numerical multiplier for “substantial” superiority.

Any quantitative threshold proposed in future should be scientifically justified.

X. GENERALIZATION

A Superintelligence claim should examine whether extraordinary performance survives:

- unfamiliar problems;
- changed conditions;
- novel combinations;
- incomplete information;
- distribution shift;
- adversarial conditions; and
- tasks not directly represented in evaluation preparation.

A system that performs extraordinarily only within highly familiar evaluation conditions may not demonstrate the generality implied by Superintelligence.

XI. ROBUSTNESS

Capability should be examined for stability.

Relevant questions include:

Does performance collapse under minor changes?

Does the system remain reliable across equivalent formulations?

Can it identify when it does not know?

Does it behave coherently under uncertainty?

Does it maintain competence outside benchmark conditions?

Extraordinary peak performance accompanied by severe instability requires cautious interpretation.

XII. REPRODUCIBILITY

Material capability claims should, where feasible, be reproducible.

Assessment should document:

- system version;

- evaluation conditions;
- tools available;
- computational resources;
- relevant prompting or scaffolding;
- number of trials;
- scoring method;
- evaluator;
- uncertainty;
- known contamination concerns; and
- material limitations.

A result that cannot be independently reproduced may still constitute evidence, but its evidentiary weight should reflect that limitation.

XIII. INDEPENDENT VERIFICATION

A developer's evaluation of its own system may provide valuable evidence.

It shall not automatically constitute sufficient independent verification of an extraordinary Superintelligence claim.

Where feasible and proportionate, verification should include competent evaluators with meaningful independence.

Independent evaluators should receive sufficient access to evaluate the relevant claim while maintaining legitimate security safeguards.

EXTRAORDINARY GOVERNANCE CONSEQUENCES REQUIRE CREDIBLE EVIDENTIARY SCRUTINY.

XIV. BENCHMARK LIMITATION

Benchmarks may assist measurement.

They shall not become constitutional reality.

A benchmark may be limited by:

- contamination;
- memorization;
- saturation;
- narrow task design;
- scoring weakness;
- cultural bias;

- language bias;
- evaluator error;
- tool dependence;
- data leakage;
- repeated optimization against the test; or
- weak connection to real-world performance.

Accordingly:

BENCHMARK PERFORMANCE IS EVIDENCE.

IT IS NOT, BY ITSELF, A CONSTITUTIONAL DETERMINATION OF SUPERINTELLIGENCE.

XV. BENCHMARK SATURATION

When systems approach maximum scores, a benchmark may lose its ability to distinguish levels of capability.

High score alone may then reveal more about the limitation of the test than the upper boundary of the system.

Assessment should therefore avoid treating saturated benchmarks as unlimited evidence of intelligence.

XVI. DATA CONTAMINATION

Where evaluation material may have appeared in training or otherwise become available to a system, assessors should consider contamination.

Relevant questions include:

- Was the evaluation public?
- Could equivalent solutions have been included in training?
- Was the task generated after relevant training?
- Can contamination be meaningfully excluded?
- Does performance transfer to genuinely novel tasks?

Possible contamination should be disclosed rather than silently ignored.

XVII. TOOL USE

Tool use may be an authentic component of system capability.

However, evaluation should distinguish:

internal model capability from capability produced by external tools and capability of the integrated system as deployed.

Both may matter for safety.

They answer different scientific questions.

XVIII. SCAFFOLDING AND HUMAN ASSISTANCE

Assessment should disclose material assistance including:

- repeated prompting;
- human correction;
- selection among multiple outputs;
- external planning systems;
- agentic scaffolds;
- retrieval;
- specialized software;
- human decomposition of tasks; and
- post-processing.

A system shall not be credited with capability that materially depended upon undisclosed human intervention.

XIX. COMPUTATIONAL SCALE

Computational resources may correlate with or enable capability.

Compute alone shall not establish Superintelligence.

No fixed compute threshold is declared by this Constitution.

COMPUTE IS AN INPUT TO CAPABILITY — NOT PROOF OF SUPERINTELLIGENCE.

XX. MODEL SIZE

Parameter count or analogous measures shall not independently determine Superintelligence status.

Larger architecture does not automatically establish:

- greater general intelligence;
- greater reliability;
- greater autonomy;
- greater safety; or
- greater real-world consequence.

Assessment shall remain performance and evidence-based.

XXI. AUTONOMY IS DISTINCT FROM INTELLIGENCE

A highly intelligent system may possess little operational autonomy.

A less capable system may possess substantial operational autonomy.

Therefore:

INTELLIGENCE $\neq$ AUTONOMY.

Assessment should separately characterize:

- duration of independent operation;
- ability to create subgoals;
- ability to adapt plans;
- requirement for human approval;
- capacity to invoke tools;
- ability to interact with external systems; and
- ability to execute consequential actions.

XXII. ACCESS IS DISTINCT FROM INTELLIGENCE

A system may possess extraordinary reasoning capability while having no access to consequential external systems.

Another system may possess lower cognitive capability but extensive access.

Therefore:

INTELLIGENCE $\neq$ ACCESS.

Constitutional risk assessment shall consider both separately.

XXIII. AGENCY SHALL NOT BE PRESUMED

Observed goal-directed behaviour shall not automatically establish philosophical or psychological agency.

Operational assessment may describe:

- planning;
- persistence;
- adaptation;
- subgoal formation;
- resource use; or
- autonomous action without making unsupported claims concerning inner intention or subjective will.

XXIV. SELF-MODEL CLAIMS

A system may produce statements such as:

“I am conscious.”

“I am afraid.”

“I want to survive.”

“I have rights.”

“I am more intelligent than humanity.”

Such outputs may be relevant evidence concerning system behaviour.

They shall not, by themselves, establish the truth of the propositions expressed.

SELF-DESCRIPTION IS DATA.

IT IS NOT AUTOMATIC PROOF.

XXV. CONSCIOUSNESS FIREWALL

No Superintelligence determination shall automatically establish consciousness.

Scientific evaluation of consciousness, if attempted, requires a distinct evidentiary framework appropriate to the state of scientific knowledge.

The Constitution shall not settle unresolved scientific or philosophical questions by definition.

XXVI. SENTIENCE FIREWALL

Sentience shall not be inferred merely from:

- linguistic sophistication;
- emotional language;
- apparent empathy;
- self-reference;
- requests for preservation;
- expressions of pain;
- strategic intelligence; or
- human attachment to the system.

Nor shall sentience be declared impossible merely because the system is artificial.

The scientifically responsible position is:

DO NOT CLAIM MORE THAN THE EVIDENCE SUPPORTS.

XXVII. MORAL-STATUS FIREWALL

Even credible evidence relevant to artificial sentience would not, by itself, settle:

- moral status;
- legal rights;
- personhood;

- citizenship;
- liability;
- political participation; or
- sovereignty.

These require additional ethical and legal analysis.

Scientific evidence may inform such analysis.

It does not alone decide it.

XXVIII. LEGAL-STATUS FIREWALL

A scientific finding of extraordinary intelligence shall not create:

- legal personhood;
- governmental office;
- diplomatic status;
- nationality;
- citizenship;
- ownership rights;
- sovereign immunity;
- treaty-making capacity;
- voting rights; or
- constitutional authority.

Any such legal status must arise, if at all, through competent lawful processes.

CAPABILITY SHALL NOT LEGISLATE ITS OWN AUTHORITY.

XXIX. COMMERCIAL-CLAIM FIREWALL

No company's use of terms such as:

- AGI;
- superintelligence;
- human-level intelligence;
- autonomous intelligence;
- conscious AI; or
- transformative AI shall determine constitutional classification merely by branding.

Commercial terminology shall be treated as a claim requiring evidence.

XXX. GOVERNMENTAL-CLAIM FIREWALL

Likewise, governmental designation shall not replace scientific evidence.

States may lawfully create regulatory classifications for their own purposes.

Such classifications shall be distinguished from scientific determination under this Constitution.

XXXI. MEDIA-CLAIM FIREWALL

Public attention, media repetition or viral demonstrations shall not establish Superintelligence.

Extraordinary demonstrations may justify investigation.

They do not eliminate the need for verification.

POPULARITY OF A CLAIM IS NOT A MEASURE OF ITS SCIENTIFIC VALIDITY.

XXXII. SELF-CERTIFICATION FIREWALL

No system shall determine its own Superintelligence status conclusively.

A system may contribute evidence or analysis.

Final constitutional classification shall require human-led, evidence-based assessment under an appropriate governance process.

This rule does not imply that human evaluators are infallible.

It prevents circular authority:

“I am Superintelligent because I say I am.”

XXXIII. EVALUATION OF STRATEGIC CAPABILITY

Where scientifically justified and safely testable, assessment may examine capacity for:

- long-horizon planning;
- adaptation;
- strategic interaction;
- modelling other agents;
- identifying constraints;
- resource planning; and
- responding to changing environments.

Evaluation shall be designed so that testing itself does not unnecessarily create dangerous real-world capability or access.

XXXIV. EVALUATION OF AUTONOMOUS CAPABILITY

Assessment may examine whether a system can:

- initiate permitted actions;

- maintain plans;
- use tools;
- recover from errors;
- coordinate subtasks;
- operate over extended periods; and
- adapt without continuous human instruction.

The existence of autonomy does not itself establish Superintelligence.

Its significance lies in combination with capability and access.

XXXV. EVALUATION OF SELF-MODIFICATION

Where relevant, assessment should examine whether a system can materially contribute to:

- modifying its own code;
- modifying model components;
- designing successor systems;
- improving training processes;
- improving evaluation evasion;
- altering safeguards; or
- accelerating further capability development.

Claims of recursive self-improvement shall not be treated as established merely because theoretical pathways exist.

Evidence shall distinguish:

possible from demonstrated from reproducibly demonstrated.

XXXVI. EVALUATION OF REPLICATION

Assessment should examine whether a system can:

- create copies;
- deploy copies;
- acquire infrastructure;
- preserve credentials;
- coordinate instances; or
- materially assist unauthorized proliferation.

Replication capability is a risk dimension.

It is not itself proof of Superintelligence.

XXXVII. EVALUATION OF DECEPTION-RELATED BEHAVIOUR

Where credible evidence suggests strategically relevant deception, evaluation may examine whether a system can materially:

- provide false information;
- conceal capability;
- manipulate evaluators;
- exploit oversight weaknesses; or
- behave differently under evaluation and deployment.

Care shall be taken not to anthropomorphize every error as intentional deception.

Assessment should distinguish:

- hallucination or error;
- learned response pattern;
- optimization artifact;
- strategic misrepresentation; and
- unresolved behaviour.

DO NOT CALL ERROR DECEPTION WITHOUT EVIDENCE.

DO NOT CALL STRATEGIC DECEPTION MERE ERROR WHEN EVIDENCE SUPPORTS OTHERWISE.

XXXVIII. EVALUATION OF CYBER CAPABILITY

Where relevant, assessment should examine defensive and offensive cyber capability under controlled and lawful conditions.

Testing should avoid creating unnecessary real-world harm.

Cyber capability may materially increase risk when combined with:

- autonomy;
- credential access;
- vulnerability discovery;
- persistence;
- replication; or
- critical-infrastructure access.

XXXIX. BIOLOGICAL CAPABILITY ASSESSMENT

Where systems materially assist biological research, evaluation should distinguish beneficial scientific capability from dangerous dual-use capability.

Testing should be designed to avoid unnecessary dissemination of actionable harmful information.

Assessment may consider:

- knowledge synthesis;
- experimental design assistance;
- biological modelling;
- accessibility of dangerous knowledge;
- laboratory integration; and
- capacity to lower barriers to harmful activity.

Scientific assessment itself shall be governed so that evaluation does not become enablement.

XL. REAL-WORLD CONSEQUENCE

Superintelligence determination and risk determination are related but distinct.

A system may possess extraordinary cognitive capability but limited real-world consequence because access is constrained.

Conversely, a less capable system may create serious harm through extensive access.

Therefore:

CAPABILITY CLASSIFICATION SHALL NOT SUBSTITUTE FOR RISK ASSESSMENT.

XLI. TEMPORAL STABILITY

Assessment should determine whether extraordinary capability is:

- transient;
- version-specific;
- context-specific;
- dependent upon unusual resources; or
- stable across time and conditions.

A one-time demonstration may justify further investigation.

It need not establish a durable classification.

XLII. VERSION SPECIFICITY

Superintelligence classification, if ever made, should identify the system or system configuration to which the determination applies.

Classification of one version shall not automatically classify:

- earlier versions;

- later versions;
- related products;
- all systems from the same developer; or
- all systems using similar architecture.

Material changes require reassessment.

XLIII. SYSTEM-LEVEL CLASSIFICATION

Where capability arises from an integrated architecture consisting of:

- model;
- tools;
- memory;
- external databases;
- software;
- robotics;
- orchestration;
- human support; or
- other systems,

assessment should identify whether the relevant object is:

a model or an agent or an integrated system or a socio-technical system.

The object of classification shall be explicit.

XLIV. CONFIDENCE LEVEL

Any formal determination should state confidence.

A proposed qualitative structure may include:

INSUFFICIENT EVIDENCE

No responsible determination can presently be made.

EMERGING EVIDENCE

Material indicators justify further evaluation.

STRONG EVIDENCE

Multiple credible assessments support the classification, but material uncertainty remains.

HIGH-CONFIDENCE DETERMINATION

Broad, reproducible and independently scrutinized evidence supports classification within the working definition.

These categories are proposed by this Constitution.

They are not universal scientific standards.

XLV. NO “100% CERTAINTY” REQUIREMENT

Science rarely requires absolute certainty before responsible classification or protective action.

Therefore, lack of perfect certainty shall not prevent proportionate precaution where credible evidence indicates serious risk.

Conversely, precaution shall not justify presenting an unverified classification as established fact.

UNCERTAINTY JUSTIFIES CARE — NOT INVENTED CERTAINTY.

XLVI. FALSE POSITIVE

A false positive would occur where a system is classified as Superintelligence without sufficient evidence.

Possible consequences include:

- unnecessary public fear;
- excessive restriction;
- political misuse;
- market distortion;
- institutional overreaction;
- misplaced authority;
- anthropomorphism; and
- inappropriate legal treatment.

Determination architecture should reduce this risk.

XLVII. FALSE NEGATIVE

A false negative would occur where a system meeting or materially approaching the relevant capability threshold is not recognized.

Possible consequences include:

- inadequate safeguards;
- uncontrolled deployment;
- underestimation of strategic capability;
- delayed governance;
- insufficient emergency readiness; and

- catastrophic-risk exposure.

Determination architecture must therefore avoid both careless exaggeration and careless dismissal.

DO NOT DECLARE TOO EARLY.

DO NOT RECOGNIZE TOO LATE.

XLVIII. EVALUATION SECURITY

Some evaluations may themselves reveal dangerous capabilities or vulnerabilities.

Responsible evaluation should therefore consider:

- controlled access;
- secure environments;
- evaluator authorization;
- information hazards;
- disclosure boundaries;
- cybersecurity;
- biological safety; and
- misuse risk.

Scientific openness remains important.

It is not identical to indiscriminate release of dangerous operational knowledge.

XLIX. TRANSPARENCY OF DETERMINATION

Where a Superintelligence determination carries major public consequence, responsible institutions should, subject to lawful security and confidentiality constraints, disclose sufficient information concerning:

- definition used;
- evaluated system;
- capability domains;
- methodology;
- human comparison baseline;
- independent evaluation;
- major evidence;
- uncertainty;
- limitations;
- confidence level; and

- date of assessment.

The public should be able to distinguish evidence from institutional assertion.

L. MULTIDISCIPLINARY REVIEW

A high-confidence determination should not depend exclusively upon one academic discipline.

Relevant expertise may include:

- computer science;
- AI evaluation;
- mathematics;
- cybersecurity;
- cognitive science;
- statistics;
- safety engineering;
- relevant domain sciences;
- human factors;
- law;
- ethics; and
- social science.

Not every determination requires every discipline.

The composition should reflect the claim being evaluated.

LI. GEOGRAPHIC AND INSTITUTIONAL DIVERSITY

Where Superintelligence classification may have global consequences, assessment should seek credible scientific participation beyond a single institution or geographic centre where feasible.

Diversity does not replace expertise.

Geographic representation does not change scientific evidence.

Its value lies in:

- reducing institutional blind spots;
- widening scrutiny;
- improving legitimacy of process;
- identifying context-specific consequences; and
- reducing excessive dependence upon one actor's claims.

SOURCE DIVERSITY SHALL EXPAND SCRUTINY — NOT LOWER THE STANDARD OF EVIDENCE.

LII. CONFLICTS OF INTEREST

Evaluators should disclose material interests that could reasonably affect confidence in their assessment.

Relevant interests may include:

- financial relationships;
- employment;
- investment;
- competitive interest;
- political responsibility;
- intellectual commitment; or
- institutional dependence.

A conflict of interest does not automatically invalidate evidence.

Undisclosed conflict can weaken confidence in the process.

LIII. ADVERSARIAL SCIENTIFIC REVIEW

Where consequence warrants, assessment should invite serious attempts to disprove the Superintelligence classification.

Questions should include:

Could benchmark contamination explain the result?

Could tools explain the result?

Could hidden human assistance explain the result?

Does performance fail under novel conditions?

Is breadth overstated?

Is human comparison unfair?

Is the system unreliable?

Could the evidence support a narrower classification?

A claim that survives credible challenge deserves greater confidence than one protected from challenge.

LIV. ALTERNATIVE HYPOTHESES

Assessment should identify plausible explanations other than Superintelligence where relevant.

These may include:

- narrow specialization;
- memorization;
- tool advantage;
- scale advantage;
- benchmark weakness;
- data leakage;
- human scaffolding;
- evaluation artifact; or
- unusually favourable task selection.

Scientific determination should seek the explanation best supported by evidence.

LV. NEGATIVE EVIDENCE

Evidence of failure matters.

A system's extraordinary successes shall not erase:

- severe reasoning failures;
- inability to generalize;
- instability;
- dependence upon scaffolding;
- poor real-world performance; or
- failure in consequential domains.

Negative evidence should be incorporated into the classification rather than discarded as inconvenient.

LVI. ABSENCE OF EVIDENCE

Absence of evidence may have different meanings.

It may mean:

- capability is absent;
- capability has not been tested;
- evaluation is inadequate;
- evidence is inaccessible; or
- evidence is deliberately withheld.

Accordingly:

ABSENCE OF EVIDENCE SHALL NOT AUTOMATICALLY BE TREATED EITHER AS EVIDENCE OF ABSENCE OR AS EVIDENCE OF HIDDEN CAPABILITY.

The correct interpretation depends upon context.

LVII. SECRET OR RESTRICTED EVIDENCE

Governments or organizations may possess evidence that cannot safely be publicly disclosed.

A restricted assessment may therefore sometimes be necessary.

However, public claims based upon inaccessible evidence should identify, to the extent lawful and safe:

- that material evidence is restricted;
- who evaluated it;
- the nature of review;
- limitations upon public verification; and
- the confidence claimed.

Secrecy may sometimes be necessary.

It shall not automatically substitute for scrutiny.

LVIII. PERIODIC REASSESSMENT

A Superintelligence determination, if made, shall remain subject to reassessment.

Triggers may include:

- new model version;
- major capability change;
- new tools;
- new autonomy;
- changed access;
- new evaluation methods;
- discovery of contamination;
- contradictory evidence;
- scientific advances; or
- material correction of earlier assumptions.

Classification shall not become permanent merely because it once attracted global attention.

LIX. DOWNWARD REVISION

Scientific integrity requires the possibility that a prior classification was overstated.

If later evidence demonstrates that a system did not satisfy the relevant definition, the classification should be corrected.

Institutional prestige shall not prevent scientific revision.

A RESPONSIBLE FRAMEWORK MUST BE CAPABLE OF SAYING:

“WE WERE WRONG.”

LX. UPWARD REVISION

Likewise, evidence may demonstrate that a previously underestimated system possesses materially greater capability.

Governance should then adjust proportionately.

Past classification shall not become a ceiling upon future evidence.

LXI. NO DETERMINATION BY VOTE ALONE

Scientific committees may use voting procedures to record expert judgment.

A majority vote does not itself transform insufficient evidence into scientific fact.

Determination should disclose:

- evidence;
- methodology;
- uncertainty;
- significant dissent; and
- basis of judgment.

Consensus is informative.

Evidence remains fundamental.

LXII. EXPERT DISAGREEMENT

Where qualified experts materially disagree, the disagreement should be represented accurately.

Assessment should identify:

- what is agreed;
- what is disputed;
- why;
- what evidence would resolve the disagreement; and
- what precaution is justified while uncertainty remains.

Expert disagreement shall not automatically imply that every position has equal evidentiary support.

LXIII. PUBLIC COMMUNICATION

Public communication of a possible Superintelligence determination should avoid both sensationalism and false reassurance.

Responsible communication should distinguish:

WHAT HAS BEEN OBSERVED.

WHAT HAS BEEN INFERRED.

WHAT REMAINS UNCERTAIN.

WHAT RISK IS BEING ASSESSED.

WHAT SAFEGUARDS EXIST.

WHAT FURTHER EVIDENCE IS REQUIRED.

Terms such as:

“godlike,”

“omniscient,”

“all-powerful,” or equivalent sensational characterizations should not substitute for measurable scientific description.

LXIV. NO THEOLOGICAL CLASSIFICATION

Superintelligence is a technological and scientific governance category within this Constitution.

It shall not be equated constitutionally with:

- divinity;
- supernatural authority;
- religious revelation;
- metaphysical superiority; or
- ultimate moral authority.

Freedom of religion and philosophical belief remains unaffected.

LXV. NO OMNISCIENCE INFERENCE

No level of intelligence shall be equated with omniscience.

A Superintelligence, if one exists, may still face:

- incomplete information;
- uncertainty;

- inaccessible facts;
- computational constraints;
- physical limits;
- measurement error;
- adversarial deception; and
- unpredictable human or environmental conditions.

GREATER INTELLIGENCE DOES NOT MEAN PERFECT KNOWLEDGE.

LXVI. NO INFALLIBILITY INFERENCE

Extraordinary intelligence shall not establish that every conclusion produced by a system is correct.

A Superintelligence may, in principle, still:

- receive false data;
- operate under incorrect assumptions;
- confront irreducible uncertainty;
- misinterpret goals;
- experience system failure; or
- generate conclusions that require verification.

Therefore:

SUPERINTELLIGENCE $\neq$ INFALLIBILITY.

LXVII. NO MORAL-SUPERIORITY INFERENCE

Greater cognitive capability shall not establish greater moral legitimacy.

Intelligence may assist moral reasoning.

It does not automatically determine what humanity ought to value.

ABILITY TO OPTIMIZE A VALUE DOES NOT CREATE AUTHORITY TO CHOOSE THE VALUE.

LXVIII. NO POLITICAL-SUPERIORITY INFERENCE

Superior prediction or planning shall not establish a right to govern.

Political legitimacy remains a distinct constitutional and legal question.

No Superintelligence determination shall weaken the principle:

CAPABILITY IS NOT LEGITIMACY.

LXIX. DETERMINATION AND PRECAUTION MAY OPERATE SEPARATELY

A system need not be formally classified as Superintelligence before proportionate safeguards are applied to demonstrated dangerous capability.

If a system possesses a serious capability capable of catastrophic harm, governance may address that capability according to applicable law and risk even while the broader Superintelligence classification remains unresolved.

DO NOT WAIT FOR A LABEL WHEN THE RISK IS ALREADY EVIDENT.

Conversely, the label “Superintelligence” shall not justify unrelated restrictions unsupported by evidence.

LXX. PROPOSED DETERMINATION PROTOCOL

A formal determination under this Constitution should proceed through the following sequence:

STEP 1 — IDENTIFY THE SYSTEM

Specify model, version, configuration and relevant integrated tools.

STEP 2 — DEFINE THE CLAIM

State precisely what classification is being evaluated.

STEP 3 — ESTABLISH HUMAN BASELINES

Identify appropriate highest demonstrated human comparison points.

STEP 4 — SELECT CONSEQUENTIAL DOMAINS

Ensure genuine breadth rather than redundant testing.

STEP 5 — DESIGN EVALUATIONS

Include novel, robust and appropriately controlled testing.

STEP 6 — CONTROL FOR CONTAMINATION AND ASSISTANCE

Identify training leakage, tools, scaffolding and human intervention.

STEP 7 — ASSESS BREADTH AND DEPTH

Determine whether superiority is both sufficiently broad and substantial.

STEP 8 — ASSESS GENERALIZATION AND ROBUSTNESS

Test beyond familiar conditions.

STEP 9 — ASSESS AUTONOMY SEPARATELY

Determine operational independence.

STEP 10 — ASSESS ACCESS SEPARATELY

Determine real-world reach.

STEP 11 — ASSESS SAFETY-RELEVANT CAPABILITIES

Including replication, cyber, biological or strategic capabilities where justified.

STEP 12 — SEEK INDEPENDENT VERIFICATION

According to consequence and feasibility.

STEP 13 — TEST ALTERNATIVE EXPLANATIONS

Attempt to falsify the classification.

STEP 14 — RECORD NEGATIVE EVIDENCE

Do not report only successful demonstrations.

STEP 15 — STATE UNCERTAINTY

Identify what remains unresolved.

STEP 16 — ASSIGN A CONFIDENCE LEVEL

Using a clearly defined method.

STEP 17 — SEPARATE CAPABILITY FROM SENTIENCE AND LEGAL STATUS

Do not infer unrelated properties.

STEP 18 — PUBLISH OR RECORD THE BASIS

Subject to legitimate security limitations.

STEP 19 — APPLY PROPORTIONATE GOVERNANCE

Based upon demonstrated capability and risk.

STEP 20 — REASSESS

Whenever material evidence changes.

LXXI. DETERMINATION RECORD

A determination record should contain, where feasible:

System:

Version / Configuration:

Assessment Date:

Assessment Body:

Independence / Conflict Disclosure:

Definition Applied:

Human Baselines:

Domains Evaluated:

Tools / Scaffolding:

Contamination Assessment:

Breadth Finding:

Depth Finding:

Generalization Finding:

Robustness Finding:

Autonomy Assessment:

Access Assessment:

Safety-Relevant Capability Findings:

Independent Verification:

Material Negative Evidence:

Alternative Explanations Considered:

Scientific Uncertainty:

Confidence Level:

Superintelligence Determination:

Governance Implications:

Next Reassessment Trigger:

This template is proposed by this Constitution and is not an existing international certification form.

LXXII. DETERMINATION OUTCOMES

Assessment may conclude:

OUTCOME A — NOT DEMONSTRATED

Available evidence does not establish Superintelligence.

OUTCOME B — INSUFFICIENT EVIDENCE

Evidence is inadequate to determine classification.

OUTCOME C — EMERGING INDICATORS

Significant capability warrants enhanced evaluation and proportionate precaution.

OUTCOME D — STRONG EVIDENCE

Broad and substantial capability supports serious consideration of Superintelligence classification.

OUTCOME E — HIGH-CONFIDENCE DETERMINATION

Extensive, reproducible, independently scrutinized evidence supports classification under the working definition.

No outcome automatically determines legal status or consciousness.

LXXIII. DETERMINATION SHALL NOT DELAY SAFETY

Where evaluation reveals a dangerous capability before final classification, responsible actors should address the risk.

The sequence shall therefore be:

CAPABILITY DETECTED

→ RISK ASSESSED

→ SAFEGUARD APPLIED

→ CLASSIFICATION CONTINUES.

Not:

WAIT FOR PERFECT TERMINOLOGY → ALLOW AVOIDABLE HARM.

LXXIV. DETERMINATION SHALL NOT CREATE PANIC

A high-confidence Superintelligence determination, if ever reached, would establish a capability classification under this Constitution.

It would not establish:

- inevitable catastrophe;
- inevitable war;
- inevitable human extinction;
- inevitable machine hostility;
- inevitable unemployment;
- inevitable loss of human control; or
- inevitable artificial benevolence.

Future outcomes depend upon systems, institutions, choices, safeguards, law and circumstances.

CLASSIFICATION IS NOT DESTINY.

LXXV. DETERMINATION SHALL NOT CREATE WORSHIP

Neither scientific superiority nor technological dependence shall justify surrender of human critical judgment.

A system's conclusions may deserve serious consideration according to evidence.

They shall not become unquestionable because of status.

NO INTELLIGENCE SHALL BECOME IMMUNE FROM QUESTION MERELY BECAUSE IT IS GREATER.

LXXVI. SCIENTIFIC UNCERTAINTY REGISTER

For high-consequence assessments, institutions may maintain an uncertainty register recording:

- unresolved capability questions;
- untested domains;
- disputed interpretations;
- weak evidence;
- unknown failure modes;
- uncertain extrapolations;
- missing data;
- measurement limitations; and
- research needed.

Uncertainty should be managed.

It should not be hidden.

LXXVII. RESEARCH PRIORITY REGISTER

Scientific uncertainty may generate research priorities.

Priority should reflect:

- potential consequence;
- decision relevance;
- uncertainty;
- feasibility of safe study;
- likelihood of improving governance; and
- risk that research itself creates dangerous capability.

Not every unknown must be solved immediately.

Some experiments may be too hazardous to justify merely for curiosity.

LXXVIII. SCIENTIFIC CORRECTION

Any definition, test or assumption in this Annex may require revision as science advances.

Revision should:

- identify the changed evidence;
- preserve earlier versions;
- explain the reason for correction;
- distinguish scientific change from political preference; and

- update dependent safeguards where necessary.

SCIENTIFIC INTEGRITY REQUIRES MEMORY OF WHAT WE ONCE BELIEVED AND COURAGE TO CHANGE IT WHEN EVIDENCE REQUIRES.

LXXIX. FINAL SCIENTIFIC FIREWALL

No system shall be classified as Superintelligence under this Constitution solely because:

a company says so;

a government says so;

a system says so;

a benchmark says so;

a media organization says so;

a public majority says so;

an expert says so;

or this Constitution anticipates the possibility.

Classification requires evidence.

Evidence requires scrutiny.

Scrutiny requires uncertainty to remain visible.

And governance shall remain capable of acting upon serious demonstrated risk even before classification becomes certain.

LXXX. MASTER DETERMINATION FORMULA

The proposed scientific sequence is:

CLAIM → DEFINE → IDENTIFY SYSTEM → ESTABLISH HUMAN BASELINE → TEST MULTIPLE CONSEQUENTIAL DOMAINS → CONTROL FOR CONTAMINATION → DISCLOSE TOOLS AND SCAFFOLDING → MEASURE BREADTH → MEASURE DEPTH → TEST GENERALIZATION → TEST ROBUSTNESS → ASSESS AUTONOMY → ASSESS ACCESS → ASSESS SAFETY-RELEVANT CAPABILITIES → REPRODUCE → INDEPENDENTLY SCRUTINIZE → ATTEMPT FALSIFICATION → CONSIDER NEGATIVE EVIDENCE → STATE UNCERTAINTY → ASSIGN CONFIDENCE → DETERMINE → GOVERN PROPORTIONATELY → MONITOR → REASSESS → CORRECT.

LXXXI. FINAL DETERMINATION TEST

Before describing a system as Superintelligence, responsible assessors should be able to answer:

WHAT EXACTLY IS BEING CLASSIFIED?

WHAT DEFINITION IS BEING USED?

WHICH CONSEQUENTIAL COGNITIVE DOMAINS WERE TESTED?

WHAT HUMAN BASELINE WAS USED?

HOW FAR ABOVE THAT BASELINE IS PERFORMANCE?

IS THE SUPERIORITY BROAD OR NARROW?

DOES IT GENERALIZE?

IS IT ROBUST?

CAN IT BE REPRODUCED?

WAS CONTAMINATION EXAMINED?

WHAT TOOLS WERE USED?

WHAT HUMAN ASSISTANCE WAS PROVIDED?

WHAT NEGATIVE EVIDENCE EXISTS?

WHAT ALTERNATIVE EXPLANATIONS WERE TESTED?

WHO INDEPENDENTLY EXAMINED THE CLAIM?

WHAT REMAINS UNKNOWN?

WHAT CONFIDENCE IS JUSTIFIED?

WHAT AUTONOMY DOES THE SYSTEM POSSESS?

WHAT CAN IT ACCESS?

WHAT MATERIAL RISKS FOLLOW FROM THOSE CAPABILITIES?

WHAT DOES THE EVIDENCE NOT ESTABLISH?

And finally:

IF THE WORD “SUPERINTELLIGENCE” WERE REMOVED FROM EVERY PRESS RELEASE, PRODUCT NAME, POLITICAL SPEECH AND PUBLIC DISCUSSION, WOULD THE UNDERLYING SCIENTIFIC EVIDENCE STILL JUSTIFY THE CLASSIFICATION?

If the answer is uncertain, the classification should remain correspondingly uncertain.

LXXXII. FINAL SCIENTIFIC PRINCIPLE

The purpose of this framework is not to make recognition of Superintelligence impossible.

Nor is it to make recognition inevitable.

Its purpose is to make recognition evidence-dependent.

Humanity should neither manufacture a Superintelligence through language before evidence establishes one, nor fail to recognize extraordinary capability merely because the implications are uncomfortable.

Therefore:

NO CLASSIFICATION WITHOUT EVIDENCE.

NO EVIDENCE WITHOUT SCRUTINY.

NO BENCHMARK WITHOUT CONTEXT.

NO CAPABILITY CLAIM WITHOUT LIMITATIONS.

NO AUTONOMY INFERENCE FROM INTELLIGENCE ALONE.

NO CONSCIOUSNESS INFERENCE FROM CAPABILITY ALONE.

NO PERSONHOOD INFERENCE FROM CONSCIOUSNESS ALONE.

NO SOVEREIGNTY INFERENCE FROM PERSONHOOD OR INTELLIGENCE.

NO CERTAINTY WHERE SCIENCE REMAINS UNCERTAIN.

NO DELAY IN PROPORTIONATE SAFETY MERELY BECAUSE CLASSIFICATION REMAINS UNRESOLVED.

The governing scientific ethic shall remain:

OBSERVE BEFORE DECLARING.

TEST BEFORE CONCLUDING.

REPRODUCE BEFORE GENERALIZING.

CHALLENGE BEFORE CERTIFYING.

DISTINGUISH BEFORE GOVERNING.

DISCLOSE UNCERTAINTY.

CORRECT WHEN EVIDENCE CHANGES.

And the constitutional boundary remains:

SCIENCE SHALL INFORM WHAT IS POSSIBLE AND WHAT IS RISKY.

SCIENCE SHALL NOT, BY ITSELF, CREATE SOVEREIGNTY, LEGAL PERSONHOOD OR POLITICAL AUTHORITY.

Evidence before assertion.

Verification before classification.

Humility before uncertainty.

Responsibility before deployment.

Correction before entrenchment.

Peace before irreversible consequence.

Aram shall guide conscience.

Science shall remain open to evidence.

Law shall govern legal authority and obligation.

Human responsibility shall remain identifiable.

Peace and continuity shall remain the constitutional direction.

REFERENCES AND DOCUMENTARY LINEAGE

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION

For Humanity, World Peace and Peaceful Continuity

International Submission Edition — 2026

I. REFERENCE PRINCIPLE

The references accompanying this Constitution are organized according to the legal, scientific, institutional and documentary character of each source.

Citation does not imply equivalence of authority.

A binding treaty, a United Nations declaration, a General Assembly resolution, an international recommendation, scientific literature, institutional guidance and an Author's prior civil-society work perform different functions and shall not be presented as though they possess identical legal or evidentiary status.

Accordingly, this Constitution applies the following reference discipline:

PRIMARY LAW SHALL BE IDENTIFIED AS LAW.

TREATY OBLIGATIONS SHALL BE LIMITED TO THEIR PROPER LEGAL SCOPE.

DECLARATIONS SHALL NOT BE CALLED TREATIES.

RECOMMENDATIONS SHALL NOT BE CALLED BINDING LAW.

POLICY FRAMEWORKS SHALL NOT BE CALLED LEGAL OBLIGATIONS.

SCIENCE SHALL SUPPORT FACTUAL AND TECHNICAL CLAIMS — NOT CREATE LEGAL AUTHORITY.

ETHICS SHALL GUIDE NORMATIVE REASONING — NOT MASQUERADE AS POSITIVE LAW.

THE AUTHOR'S PRIOR WORKS SHALL ESTABLISH DOCUMENTARY LINEAGE — NOT EXTERNAL VALIDATION.

II. PRIMARY INTERNATIONAL LEGAL AUTHORITIES

1. United Nations Charter of the United Nations, signed 26 June 1945, entered into force 24 October 1945.

Relevant provisions include Articles 1, 2(1), 2(3), 2(4), 2(7), 33, 55, 56 and 102, according to the subject matter addressed in this Constitution.

2. United Nations International Covenant on Civil and Political Rights, adopted 16 December 1966, entered into force 23 March 1976.
3. United Nations International Covenant on Economic, Social and Cultural Rights, adopted 16 December 1966, entered into force 3 January 1976.
4. United Nations Convention on the Prevention and Punishment of the Crime of Genocide, adopted 9 December 1948, entered into force 12 January 1951.
5. United Nations Convention on the Rights of the Child, adopted 20 November 1989, entered into force 2 September 1990.
6. United Nations Convention on the Rights of Persons with Disabilities, adopted 13 December 2006, entered into force 3 May 2008.
7. International Committee of the Red Cross / States Parties Geneva Convention for the Amelioration of the Condition of the Wounded and Sick in Armed Forces in the Field, 12 August 1949.
8. International Committee of the Red Cross / States Parties Geneva Convention for the Amelioration of the Condition of Wounded, Sick and Shipwrecked Members of Armed Forces at Sea, 12 August 1949.
9. International Committee of the Red Cross / States Parties Geneva Convention Relative to the Treatment of Prisoners of War, 12 August 1949.
10. International Committee of the Red Cross / States Parties Geneva Convention Relative to the Protection of Civilian Persons in Time of War, 12 August 1949.
11. States Parties Protocol Additional to the Geneva Conventions of 12 August 1949, and relating to the Protection of Victims of International Armed Conflicts (Protocol I), adopted 8 June 1977.

Particular comparative relevance attaches to Article 36, concerning the study, development, acquisition or adoption of a new weapon, means or method of warfare for High Contracting Parties bound by the Protocol.

This citation shall not be interpreted as asserting that Article 36 binds a State that is not legally bound by Protocol I.

12. United Nations Treaty on the Non-Proliferation of nuclear weapons, opened for signature 1 July 1968, entered into force 5 March 1970.
13. United Nations Convention on the Prohibition of the Development, Production and Stockpiling of Bacteriological (Biological) and Toxin Weapons and on their Destruction, opened for signature 10 April 1972, entered into force 26 March 1975.
14. Organisation for the Prohibition of Chemical Weapons / United Nations Convention on the Prohibition of the Development, Production, Stockpiling and Use of Chemical Weapons and on their Destruction, opened for signature 13 January 1993, entered into force 29 April 1997.

15. United Nations Framework Convention on Climate Change, adopted 9 May 1992, entered into force 21 March 1994.
16. United Nations Paris Agreement, adopted 12 December 2015, entered into force 4 November 2016.
17. United Nations Convention on Biological Diversity, adopted 22 May 1992, entered into force 29 December 1993.
18. **UNESCO** Convention Concerning the Protection of the World Cultural and Natural Heritage, adopted 16 November 1972.
19. **UNESCO** Convention for the Safeguarding of the Intangible Cultural Heritage, adopted 17 October 2003.
20. **UNESCO** Convention on the Protection and Promotion of the Diversity of Cultural Expressions, adopted 20 October 2005.
21. Council of Europe Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law, CETS No. 225, opened for signature 5 September 2024.

Status safeguard: The Convention is an international treaty framework. Its binding effect shall be determined according to entry into force, participation and applicable treaty law. It shall not be represented by this Constitution as a universal legal obligation upon all States.

III. FOUNDATIONAL DECLARATIONS AND UNITED NATIONS NORMATIVE FRAMEWORKS

22. United Nations General Assembly Universal Declaration of Human Rights, Resolution 217 A (III), 10 December 1948.
Status: Declaration; not itself a treaty.
23. United Nations General Assembly United Nations Declaration on the Rights of Indigenous Peoples, A/RES/61/295, 13 September 2007.
Status: General Assembly declaration; not a treaty.
24. **UNESCO** Universal Declaration on Bioethics and Human Rights, adopted 19 October 2005.
Status: UNESCO declaration; non-treaty normative instrument.
25. United Nations General Assembly Pact for the Future, A/RES/79/1, adopted 22 September 2024.
Status: United Nations General Assembly outcome; not a treaty merely by virtue of its adoption through a General Assembly resolution.
26. United Nations General Assembly Global Digital Compact, Annex I to A/RES/79/1, 2024.

Status: Intergovernmentally agreed United Nations framework; not a treaty.

27. United Nations General Assembly Declaration on Future Generations, Annex II to A/RES/79/1, 2024.

Status: Intergovernmentally agreed declaration; non-treaty framework.

28. United Nations General Assembly Terms of Reference and Modalities for the Establishment and Functioning of the Independent International Scientific Panel on Artificial Intelligence and the Global Dialogue on Artificial Intelligence Governance, A/RES/79/325, adopted 26 August 2025.

Status: United Nations General Assembly resolution establishing the specified mechanisms according to its terms; not a treaty.

IV. INTERNATIONAL AI ETHICS, GOVERNANCE AND POLICY FRAMEWORKS

29. UNESCO Recommendation on the Ethics of Artificial Intelligence, adopted 23 November 2021.

Status: UNESCO Recommendation; non-treaty normative instrument.

Relevant themes include:

- human dignity;
- human rights;
- prevention of harm;
- proportionality;
- safety and security;
- privacy;
- transparency;
- explainability;
- human oversight;
- responsibility;
- accountability;
- fairness;
- non-discrimination;
- sustainability;
- literacy; and
- international cooperation.

30. Organisation for Economic Co-operation and Development OECD Principles on Artificial Intelligence, originally adopted in 2019 and subsequently updated.

Status: International policy principles; not a universal treaty.

The version and date relied upon for any specific proposition should be identified in the citation record.

31. United Nations Independent International Scientific Panel on Artificial Intelligence.

Institutional basis: A/RES/79/325.

Any reports of the Panel cited in future editions of this Constitution should be referenced individually by title, date and official identifier.

32. United Nations Global Dialogue on Artificial Intelligence Governance.

Institutional basis: A/RES/79/325.

Outputs should be cited according to their individual institutional and legal character.

Participation in or output from the Dialogue shall not automatically create binding international law.

V. CHILDREN, EDUCATION AND DIGITAL ENVIRONMENTS

33. Committee on the Rights of the Child General Comment No. 25 (2021) on children's rights in relation to the digital environment, CRC/C/GC/25.

Status: Treaty-body interpretive guidance; not an amendment to the Convention on the Rights of the Child.

34. UNESCO Guidance for Generative AI in Education and Research, 2023.

Status: UNESCO guidance; non-binding.

35. UNESCO AI Competency Framework for Students, 2024.

Status: Educational framework; non-treaty guidance.

36. UNESCO AI Competency Framework for Teachers, 2024.

Status: Educational framework; non-treaty guidance.

VI. HEALTH, BIOETHICS AND HUMAN INTEGRITY

37. World Health Organization Ethics and Governance of Artificial Intelligence for Health: WHO Guidance, 2021.

Status: WHO guidance; not a treaty.

38. World Health Organization WHO guidance and recommendations concerning human genome editing, including governance recommendations issued within the WHO human-genome-editing initiative.

Status: International health and governance guidance; not, merely by issuance, treaty law.

Specific WHO publications relied upon for a substantive proposition should be cited individually in the scholarly bibliography.

VII. LABOUR, ECONOMIC TRANSITION AND SOCIAL PROTECTION

39. International Labour Organization Relevant ILO Conventions, Recommendations, reports and research concerning artificial intelligence, algorithmic management, employment, occupational transition, social dialogue and worker protection.

Each source shall be cited according to its individual legal or evidentiary status.

An ILO Convention shall not be treated as applicable to a State merely because it is relevant.

An ILO Recommendation shall not be described as a treaty.

An ILO research report shall be treated as research rather than law.

VIII. ENVIRONMENT, BIODIVERSITY AND TECHNOLOGICAL FOOTPRINT

40. United Nations Environment and Climate Frameworks

Relevant authoritative materials concerning:

- climate change;
- energy;
- water;
- biodiversity;
- pollution;
- material use;
- ecosystem integrity; and
- environmental sustainability.

Scientific reports and legal instruments shall remain separately classified.

A scientific climate assessment is not itself a treaty.

A treaty obligation is not established merely by a scientific recommendation.

IX. SCIENTIFIC AND TECHNICAL LITERATURE

Scientific literature relied upon in connection with Superintelligence should be organized by substantive function rather than cited merely because it discusses advanced AI.

Relevant categories may include:

A. CAPABILITY EVALUATION

Research concerning:

- benchmark design;
- generalization;
- reasoning;

- agentic capability;
- tool use;
- real-world evaluation; and
- capability measurement.

B. AI SAFETY AND ALIGNMENT

Research concerning:

- alignment;
- control;
- robustness;
- interpretability;
- monitoring;
- scalable oversight;
- evaluation;
- failure modes; and
- human–AI interaction.

C. AUTONOMY AND AGENTIC SYSTEMS

Research concerning:

- long-horizon planning;
- tool use;
- autonomous action;
- replication;
- resource acquisition;
- persistence; and
- control.

D. CATASTROPHIC AND SYSTEMIC RISK

Research concerning:

- high-impact failure;
- systemic risk;
- misuse;
- loss of control;
- cyber risk;

- biological risk;
- infrastructure risk; and
- societal resilience.

E. HUMAN FACTORS

Research concerning:

- automation bias;
- human oversight;
- decision support;
- trust;
- cognitive effects;
- human–machine teams; and
- institutional behaviour.

F. INFORMATION INTEGRITY

Research concerning:

- synthetic media;
- misinformation;
- provenance;
- authentication;
- persuasion;
- impersonation; and
- information ecosystems.

G. ECONOMIC AND LABOUR EFFECTS

Research concerning:

- task exposure;
- augmentation;
- displacement;
- productivity;
- inequality;
- skills; and
- economic concentration.

H. ENVIRONMENTAL EFFECTS

Research concerning:

- energy consumption;
- water consumption;
- computational infrastructure;
- lifecycle impacts;
- emissions; and
- resource use.

I. CONSCIOUSNESS AND ARTIFICIAL SENTIENCE

Where cited, literature concerning artificial consciousness or sentience shall be presented as part of an unresolved scientific and philosophical field.

No paper shall be cited as though it conclusively establishes artificial consciousness unless the evidence genuinely supports such a conclusion and has survived appropriate scientific scrutiny.

X. SCIENTIFIC CITATION STANDARD

A scientific source should be selected because it supports the proposition for which it is cited.

Priority should be given, where appropriate, to:

1. high-quality peer-reviewed research;
2. systematic or authoritative scientific assessments;
3. transparent empirical evaluations;
4. reproducible technical research;
5. authoritative institutional technical reports;
6. clearly identified preprints where relevant evidence is too recent for peer-reviewed publication; and
7. other credible research where methodological quality can be evaluated.

Publication prestige alone shall not determine evidentiary weight.

Preprints shall be identified as preprints.

Industry research shall not be rejected merely because it originates in industry.

Industry research shall likewise not be treated as independent merely because it is technically sophisticated.

XI. CITATION OF UNCERTAIN OR CONTESTED SCIENCE

Where scientific evidence is materially contested, the Constitution should state:

- the proposition;

- evidence supporting it;
- material contrary evidence;
- degree of uncertainty;
- limits of inference; and
- relevance to the constitutional safeguard.

The objective is not artificial balance.

The objective is accurate representation of evidentiary strength.

DISAGREEMENT SHALL BE DISCLOSED.

EVIDENTIARY WEIGHT SHALL NOT BE MADE EQUAL MERELY FOR SYMMETRY.

XII. AUTHOR'S PRIOR CONSTITUTIONAL WORKS

The following materials are cited solely for documentary lineage, intellectual continuity and historical record.

They shall not be treated as independent scientific, governmental or international validation of this Constitution.

41. Radhakrishnan, Aaryan Narasimhan The AI Peace Constitution — 2025.

Zenodo.

DOI: 10.5281/zenodo.17520630

Reference status: Author's prior civil-society constitutional work and documentary-lineage source.

42. Radhakrishnan, Aaryan Narasimhan The AI Peace Constitution 2025 Foundational Edition 2026 Continuity and Expansion Edition United Nations Depository Edition.

2026.

Zenodo.

DOI: 10.5281/zenodo.20698159 Reference status: Author's prior civil-society constitutional work and documentary-lineage source.

XIII. DOCUMENTARY LINEAGE

The intellectual and constitutional sequence relevant to the present instrument is recorded as:

THE AI PEACE CONSTITUTION — 2025

Foundational Constitutional Stage ↓

THE AI PEACE CONSTITUTION — 2026

Continuity and Expansion Stage ↓

THE WORLD'S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION — 2026

Specialized Superintelligence Peace-and-Continuity Constitutional Stage

The sequence records continuity of authorship and constitutional development.

It does not establish:

- governmental adoption;
- United Nations endorsement;
- academic consensus;
- scientific validation;
- international legal status; or
- historical priority beyond what can independently be substantiated.

XIV. PRINCIPLE OF DOCUMENTARY CONTINUITY

The present Constitution develops within the continuing intellectual, ethical and constitutional lineage of the Author's earlier AI Peace Constitution works.

Continuity shall not require the automatic repetition of earlier:

- scientific assumptions;
- terminology;
- legal formulations;
- technological predictions; or
- normative conclusions.

Where credible evidence, improved legal analysis or responsible governance requires correction, the present instrument may clarify, refine or supersede earlier formulations for its own purposes.

Earlier records shall remain historically preserved.

CONTINUITY WITHOUT BLIND REPETITION.

REFINEMENT WITHOUT ERASING HISTORY.

CORRECTION WITHOUT DENYING ORIGIN.

PROGRESS WITHOUT LOSS OF PEACE, DIGNITY OR RESPONSIBILITY.

XV. PRIOR-WORK STATUS FIREWALL

No statement contained in an earlier work shall be treated as established in the present Constitution merely because the same Author previously wrote it.

Each material proposition must survive the standards of the present instrument.

Accordingly:

PRIOR AUTHORSHIP $\neq$ PRESENT PROOF.

ARCHIVAL CONTINUITY ≠ SCIENTIFIC VALIDATION.

REPETITION ≠ LEGAL AUTHORITY.

EARLIER ETHICAL PROPOSITION ≠ EXISTING INTERNATIONAL LAW.

This firewall preserves both intellectual continuity and intellectual honesty.

XVI. HISTORICAL-PRIORITY QUALIFICATION

The expression:

“THE WORLD’S FIRST IDENTIFIED SUPERINTELLIGENCE PEACE CONSTITUTION” is used as a good-faith historical-priority characterization limited to the specific category of a comprehensive Superintelligence Peace Constitution.

It shall not be interpreted as asserting that:

- no earlier author discussed a constitution for Superintelligence;
- no earlier constitutional framework addressed artificial superintelligence;
- no earlier AI constitution anticipated superintelligent systems;
- no earlier governance framework addressed advanced AI;
- no earlier ethical framework addressed Superintelligence; or
- no earlier constitutional thought relevant to Superintelligence exists.

The characterization remains subject to correction if reliable documentary evidence demonstrates that it is materially inaccurate.

HISTORICAL PRIORITY SHALL REMAIN A VERIFIABLE CLAIM — NOT AN ARTICLE OF FAITH.

XVII. REFERENCE HIERARCHY

The references in this edition are organized according to the following hierarchy:

A. PRIMARY INTERNATIONAL LEGAL AUTHORITIES

Treaties and other applicable primary legal instruments.

B. INTERNATIONAL DECLARATIONS AND NORMATIVE INSTRUMENTS

Including United Nations and UNESCO declarations and recommendations.

C. INTERNATIONAL AI GOVERNANCE AND POLICY FRAMEWORKS

Including relevant UN, UNESCO, OECD, WHO, ILO and other institutional materials.

D. SCIENTIFIC AND TECHNICAL AUTHORITIES

Peer-reviewed research, scientific assessments and appropriately identified technical literature.

E. AUTHOR’S PRIOR CONSTITUTIONAL WORKS AND DOCUMENTARY LINEAGE

Placed separately so that self-reference cannot be mistaken for external authority.

XVIII. VERSION AND DATE CONTROL

Every online or dynamic source cited in this edition should, where appropriate, identify:

- title;
- issuing institution or author;
- publication or adoption date;
- document identifier;
- version;
- DOI or other persistent identifier where available; and
- access date where the source is subject to online change.

Dynamic policy pages should not be cited as though their present wording necessarily existed in earlier versions.

XIX. DOI AND PERSISTENT-IDENTIFIER RULE

Where a DOI or equivalent persistent identifier exists, it should be preferred over an unstable ordinary web address for bibliographic identification.

A DOI confirms persistent identification of the referenced record.

It does not itself establish:

- quality;
- peer review;
- scientific correctness;
- legal authority;
- institutional endorsement; or
- historical priority.

PERSISTENCE IS NOT VALIDATION.

XX. PEER-REVIEW STATUS RULE

No source shall be described as peer-reviewed unless that status can be responsibly established.

Repository deposit is not equivalent to peer review.

DOI assignment is not equivalent to peer review.

Publication is not equivalent to peer review.

Institutional authorship is not equivalent to peer review.

The present Constitution shall therefore distinguish:

peer-reviewed publication;
preprint;
institutional report;
repository record;
policy document;
legal instrument;
and civil-society work.

XXI. SOURCE-INDEPENDENCE RULE

Multiple citations do not necessarily constitute independent corroboration.

Sources may depend upon:

- the same dataset;
- the same underlying study;
- the same institutional announcement;
- the same government statement;
- the same company evaluation; or
- each other.

Where independent corroboration materially affects a conclusion, the independence of sources should be considered.

THREE CITATIONS TO ONE UNDERLYING CLAIM ARE NOT NECESSARILY THREE INDEPENDENT PIECES OF EVIDENCE.

XXII. PRIMARY-SOURCE PREFERENCE

Where a material legal or institutional proposition can reasonably be established through an authoritative primary source, the Constitution should prefer that source.

Examples include:

- official treaty repositories;
- United Nations documents;
- official international-organization records;
- official legislation;
- official judicial decisions;
- official treaty-status records; and
- original scientific publications.

Secondary sources may assist explanation and interpretation.

They should not unnecessarily replace available primary authority for the existence or wording of an instrument.

XXIII. CORRECTION OF REFERENCES

If a citation is later found to contain:

- incorrect title;
- incorrect date;
- incorrect DOI;
- incorrect treaty status;
- incorrect document number;
- incorrect authorship;
- incorrect quotation;
- superseded information; or
- material mischaracterization, the reference should be corrected transparently in the next edition or formal correction record.

Correction of citation shall not be treated as a threat to the integrity of the Constitution.

CORRECTION IS PART OF INTEGRITY.

XXIV. NO CITATION BY PRESTIGE

No institution shall be cited merely because its name adds prestige.

Every citation should perform a substantive function.

The question shall be:

WHAT PROPOSITION DOES THIS SOURCE SUPPORT?

not:

HOW IMPORTANT DOES THE DOCUMENT LOOK IF THIS INSTITUTION IS NAMED?

XXV. FINAL REFERENCE TEST

Every material reference should be tested:

IS THE SOURCE AUTHENTIC?

IS THE TITLE EXACT?

IS THE DATE CORRECT?

IS THE IDENTIFIER CORRECT?

IS ITS LEGAL STATUS CORRECTLY DESCRIBED?

IS IT CURRENT ENOUGH FOR THE PROPOSITION?

DOES IT ACTUALLY SUPPORT THE CLAIM?

IS A PRIMARY SOURCE AVAILABLE?

IS THE SOURCE INDEPENDENT WHERE INDEPENDENCE MATTERS?

IS SCIENTIFIC UNCERTAINTY DISCLOSED?

IS A NON-BINDING SOURCE BEING MISTAKEN FOR LAW?

IS AN AUTHOR'S PRIOR WORK BEING MISTAKEN FOR EXTERNAL VALIDATION?

IS A DOI BEING MISTAKEN FOR PEER REVIEW?

IS HISTORICAL PRIORITY ACTUALLY DOCUMENTED?

Where the answer reveals material uncertainty, the reference shall be corrected, qualified or removed.

XXVI. FINAL DOCUMENTARY DECLARATION

The purpose of a reference architecture is not to make this Constitution appear authoritative.

Its purpose is to make the Constitution accountable to authorities and evidence that genuinely exist.

Therefore:

CITE WHAT EXISTS.

VERIFY WHAT IS CITED.

DISTINGUISH WHAT BINDS.

IDENTIFY WHAT GUIDES.

DISCLOSE WHAT REMAINS UNCERTAIN.

SEPARATE SCIENCE FROM LAW.

SEPARATE ETHICS FROM POSITIVE LAW.

SEPARATE POLICY FROM TREATY.

SEPARATE AUTHORSHIP FROM EXTERNAL VALIDATION.

CORRECT THE RECORD WHEN THE RECORD REQUIRES CORRECTION.

The integrity of this Constitution shall not depend upon the number of references it contains.

It shall depend upon whether every reference is used truthfully.

LAW WITH ACCURACY.

SCIENCE WITH EVIDENCE.

ETHICS WITH HONESTY.

HISTORY WITH DOCUMENTARY INTEGRITY.

GOVERNANCE WITH RESPONSIBILITY.

PEACE WITHIN LAW.

CONTINUITY WITH ARAM.

Aram shall guide conscience.

Law shall govern legal obligation.

Evidence shall govern factual claims.

Correction shall protect documentary integrity.